

CSCE 435 Group 16 project

0. Group number: 16

1. Group members:

1. Anjali Kumar
2. Shawn Mathen
3. Ashrita Vadlapatla
4. Robert Eads

2. Project topic (e.g., parallel sorting algorithms)

Parallel Sorting Algorithms

2a. Communication Method

Our team's primary method of communication will be GroupMe with Slack as a secondary method.

2b. Brief project description (what algorithms will you be comparing and on what architectures)

Each of the selected sort algorithms, Bitonic, Merge, Selection, & Sample, will be run using MPI and CUDA separately.

2c. Pseudocode for each parallel algorithm

Algorithm 1: Bitonic Sort

MPI

1. Distribute Data: Use MPI_Scatter to distribute the data among processes. Each process will receive a portion of the dataset.
2. Local Bitonic Sort: Each process performs a local Bitonic Sort on its portion of the data. This involves performing Bitonic Sort on the local dataset using parallel threads or any parallel programming model within the process.
3. Bitonic Merge: Implement the Bitonic Merge operation for pairs of processes. This involves exchanging data between neighboring processes to create a bitonic sequence.
4. Global Bitonic Sort: Perform multiple iterations of the Bitonic Sort while exchanging and comparing adjacent elements between neighboring processes. In each iteration, processes exchange data and perform Bitonic Merge.
5. Repeat for Required Iterations: Repeat the Bitonic Sort and data exchange steps for the required number of iterations to ensure a fully sorted dataset. The number of iterations depends on the size of the dataset and the number of processes.
6. Verify Correctness: After the sorting is complete, use a verification step

to ensure the correctness of the sorted data. You can implement a global verification step where each process checks the correctness of its portion, or you can gather data to the root process for a centralized verification.

CUDA

1. Allocate Device Memory: Use `cudaMalloc` to allocate memory on the GPU for the dataset. Allocate memory for temporary buffers if needed.
2. Copy Data to GPU: Use `cudaMemcpy` to copy the data from the host to the GPU.
- Launch Kernel for Local Bitonic Sort:
3. Write a CUDA kernel to perform local Bitonic Sort on the GPU. Each thread will work on a portion of the dataset. Threads in the same block can cooperate to perform the sort using shared memory.
4. Synchronize Threads: Use `__syncthreads()` to synchronize threads within a block after the local sort.
5. Launch Bitonic Merge Kernel: Write a CUDA kernel to perform Bitonic Merge operation for pairs of blocks. This may involve exchanging data between neighboring blocks and creating bitonic sequences.
6. Synchronize Threads Again: Use `__syncthreads()` to synchronize threads after the merge operation.
7. Repeat for Multiple Iterations: Repeat steps 3-6 for the required number of iterations. The number of iterations depends on the size of the dataset and the structure of the bitonic sequence.
8. Copy Data Back to CPU: Use `cudaMemcpy` to copy the sorted data from the GPU back to the host.
9. Free Device Memory: Use `cudaFree` to release the allocated memory on the GPU.
10. Verify Correctness: After the sorting is complete, use a verification step to ensure the correctness of the sorted data. You can use a similar verification function as before.

General Pseudocode

```
void bitonicMerge(int a[], int low, int cnt, int dir)
{
    if (cnt>1)
    {
        int k = cnt/2;
        for (int i=low; i<low+k; i++)
            compAndSwap(a, i, i+k, dir);
        bitonicMerge(a, low, k, dir);
        bitonicMerge(a, low+k, k, dir);
    }
}

void bitonicSort(int a[], int low, int cnt, int dir)
{
    if (cnt>1)
```

```

{
    int k = cnt/2;

    // sort in ascending order since dir here is 1
    bitonicSort(a, low, k, 1);

    // sort in descending order since dir here is 0
    bitonicSort(a, low+k, k, 0);

    // Will merge whole sequence in ascending order
    // since dir=1.
    bitonicMerge(a,low, cnt, dir);
}
}

```

MPI Pseudocode

```

int main(int argc, char **argv) {
    // MPI Initialization

    // Scatter data among processes
    MPI_Scatter(/* ... */);

    // Local Bitonic Sort on each process
    localBitonicSort(/* ... */);

    // Global Bitonic Sort
    for (int k = 2; k <= numProcesses; k *= 2) {
        for (int j = k / 2; j > 0; j /= 2) {
            // Bitonic Merge for pairs of processes
            bitonicMerge(/* ... */);
        }
    }

    // Gather sorted data to root process
    MPI_Gather(/* ... */);

    // MPI Finalization

    return 0;
}

```

CUDA Pseudocode

```

__global__ void bitonicSortKernel(int *data, int dataSize) {
    // Calculate thread index
    int tid = blockIdx.x * blockDim.x + threadIdx.x;

```

```

// Bitonic Sort
for (int k = 2; k <= dataSize; k *= 2) {
    for (int j = k / 2; j > 0; j /= 2) {
        int ixj = tid ^ j;

        if ((ixj) > tid) {
            if ((tid & k) == 0) {
                // Sort ascending
                if (data[tid] > data[ixj]) {
                    // Swap elements
                    int temp = data[tid];
                    data[tid] = data[ixj];
                    data[ixj] = temp;
                }
            }
            if ((tid & k) != 0) {
                // Sort descending
                if (data[tid] < data[ixj]) {
                    // Swap elements
                    int temp = data[tid];
                    data[tid] = data[ixj];
                    data[ixj] = temp;
                }
            }
        }
    }
}

__syncthreads(); // Synchronize threads after each iteration
}
}
}

```

Algorithm 2: Merge Sort

MPI

1. Distribute data among processes
2. Each process performs a local sequential merge sort
3. Pairwise merging of sorted sublists using MPI_Send and MPI_Recv
4. Recursively perform merging until the data is fully sorted
5. Verify correctness

CUDA

1. Each CUDA thread loads a portion of the data into shared memory
2. Perform a local sequential merge sort within each thread's shared memory

3. Use CUDA parallel reduction to find the pivot elements
4. Broadcast to all threads
5. Each thread partitions its data around pivot
6. Calculate the offsets for each partition
7. CUDA scatter to move elements to their respective partitions
8. Recursively sort each partition using CUDA parallel sort
9. Perform parallel merge or merge in shared memory, depending on the size of the partitions

Pseudocode

Parallel:

The following parallel Merge Sort algorithms will be implemented using MPI.

```
function parallel_merge_sort(arr, num_threads)
    if length(arr) <= 1
        return arr

    mid = length(arr) // 2
    left = arr[0...mid-1]
    right = arr[mid...]

    if num_threads > 1
        left_thread = start_thread(parallel_merge_sort(left, num_threads/2))
        right_thread = start_thread(parallel_merge_sort(right, num_threads/2))
        join_thread(left_thread)
        join_thread(right_thread)

    else
        left = merge_sort(left)
        right = merge_sort(right)

    return parallel_merge(left, right)
```

The following is pseudocode for a CUDA implementation

```
__global__ void mergeCUDA(int* arr, int left, int mid, int right) {
    // Sequential merge
}

__global__ void mergeSortParallelCUDA(int* arr, int size) {
    if (size <= 1) return;

    int mid = size / 2;

    int* left = arr;
    int* right = arr + mid;
```

```

    mergeSortParallelCUDA<<<1, 1>>>(left, mid);
    mergeSortParallelCUDA<<<1, 1>>>(right, size - mid);

    mergeCUDA(arr, 0, mid - 1, size - 1);
}

```

Algorithm 3: Selection Sort

MPI

1. Distribute the data among processes using MPI_Scatter.
2. Each process applies the selection sort algorithm to its portion of the data.
3. The sort selects the minimum element and swaps it with the first unsorted element so that there is a local sorted segment.
3. Communicate with neighboring processes using MPI_Send and MPI_Recv to exchange boundary data so that elements are placed correctly.
4. There is no need for a scatter operation, as Selection Sort is an in-place sorting algorithm.
5. Repeat the Selection Sort for the required number of iterations, ensuring all elements are correctly sorted.
6. Verify the correctness of the sorted data.

CUDA

1. Each CUDA thread loads a portion of the data into shared memory
2. Implement selection sort within the shared memory to locally sort the data.
3. Utilize CUDA parallel reduction techniques to find the minimum element within the shared memory.
4. Share the minimum element among all threads.
5. Partition the data based on the broadcasted minimum element.
6. Determine offsets for each partition to facilitate data movement.
7. Utilize CUDA scatter operations to move elements to their respective partitions.
8. Continue the selection sort for a certain number of iterations.
9. As selection sort is an exchange-based sorting algorithm that works directly on the data in place, there is no merging.

```

void selectionSort(int arr[], int n)
{
    int i, j, min_idx;

    // One by one move boundary of

```

```

// unsorted subarray
for (i = 0; i < n - 1; i++) {

    // Find the minimum element in
    // unsorted array
    min_idx = i;
    for (j = i + 1; j < n; j++) {
        if (arr[j] < arr[min_idx])
            min_idx = j;
    }

    // Swap the found minimum element
    // with the first element
    if (min_idx != i)
        swap(arr[min_idx], arr[i]);
}
}

```

Algorithm 4: Sample Sort

MPI

```

FOR number of splitters selected by each process
    Push back randomly selected value from local data to sample vector
    Allocate memory for all sampled splitters
    MPI_Allgather sampled splitters to each process
    Sort sampled splitters and choose bucket bounds from throughout sampled
    splitters
    Evaluate local elements and place into send buffers
    Calculate required buffer size from each process
    FOR number of processes
        MPI_Gather the buffer sizes from each process on each process
    Allocate receive buffer for incoming data
    FOR number of processes
        MPI_Gatherv to send data to the correct process
    Sort with sequential quicksort

```

CUDA

All steps are performed with CUDA kernels

- Sample elements
- Sort samples
- Select pivots from sampled elements
- Gather required size and start position for each bucket
- Group elements into buckets within an array
- Have each threads sort the data within the bounds of its respective bucket

2d. Citations

Bitonic Sort

- <https://www.geeksforgeeks.org/bitonic-sort/>
- https://www.tools-of-computing.com/tc/CS/Sorts/bitonic_sort.htm

Merge Sort

- <https://computcademy.net/algorithmic-thinking-with-python-part-3-divide-and-conquer-strategy/#:~:text=There%20is%20a%20really%20clever%20trick%20that,the%20same%20type%20as%20the%20original%20problem.>
- <https://teivah.medium.com/parallel-merge-sort-in-java-e3213ae9fa2c>
- <https://pushkar2196.wordpress.com/2017/04/19/mergesort-cuda-implementation/>

Selection Sort

- <https://www.geeksforgeeks.org/selection-sort/>

Sample Sort

- <https://en.wikipedia.org/wiki/Samplesort>
- <https://cse.buffalo.edu/faculty/miller/Courses/CSE702/Nicolas-Barrios-Fall-2021.pdf>
- <https://www.geeksforgeeks.org/quicksort-using-random-pivoting/>

2e. Evaluation plan - what and how will you measure and compare

Each algorithm will be run with the input types of randomized, sorted, reverse sorted, and 1% perturbed. The input sizes are planned to be 2^{16} , 2^{18} , 2^{20} , 2^{22} , 2^{24} , 2^{26} and 2^{28} .

Graphs for strong and weak scaling, as well as strong speedup, will be created and used in the analysis and drawing of conclusions for each algorithm across the different input types and input sizes. A strong scaling comparison will also be run for the algorithms against each other to see if and where one will stand out from the rest.

The number of threads in a block on the GPU to be tested will be [64, 128, 256, 512, 1024].

3. Project Implementation

The listed algorithms were fully implemented using both MPI and CUDA separately.

4. Performance Evaluation

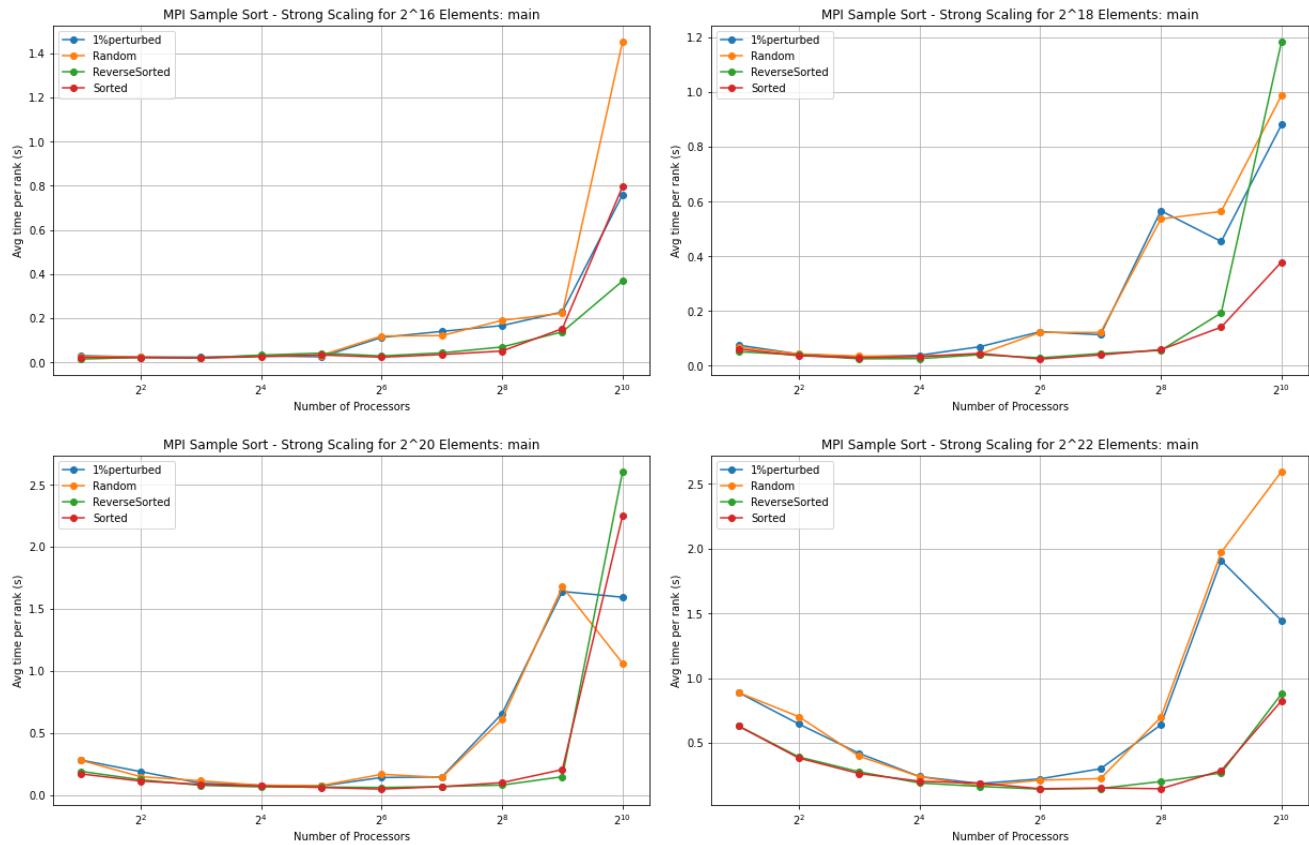
Sample Sort

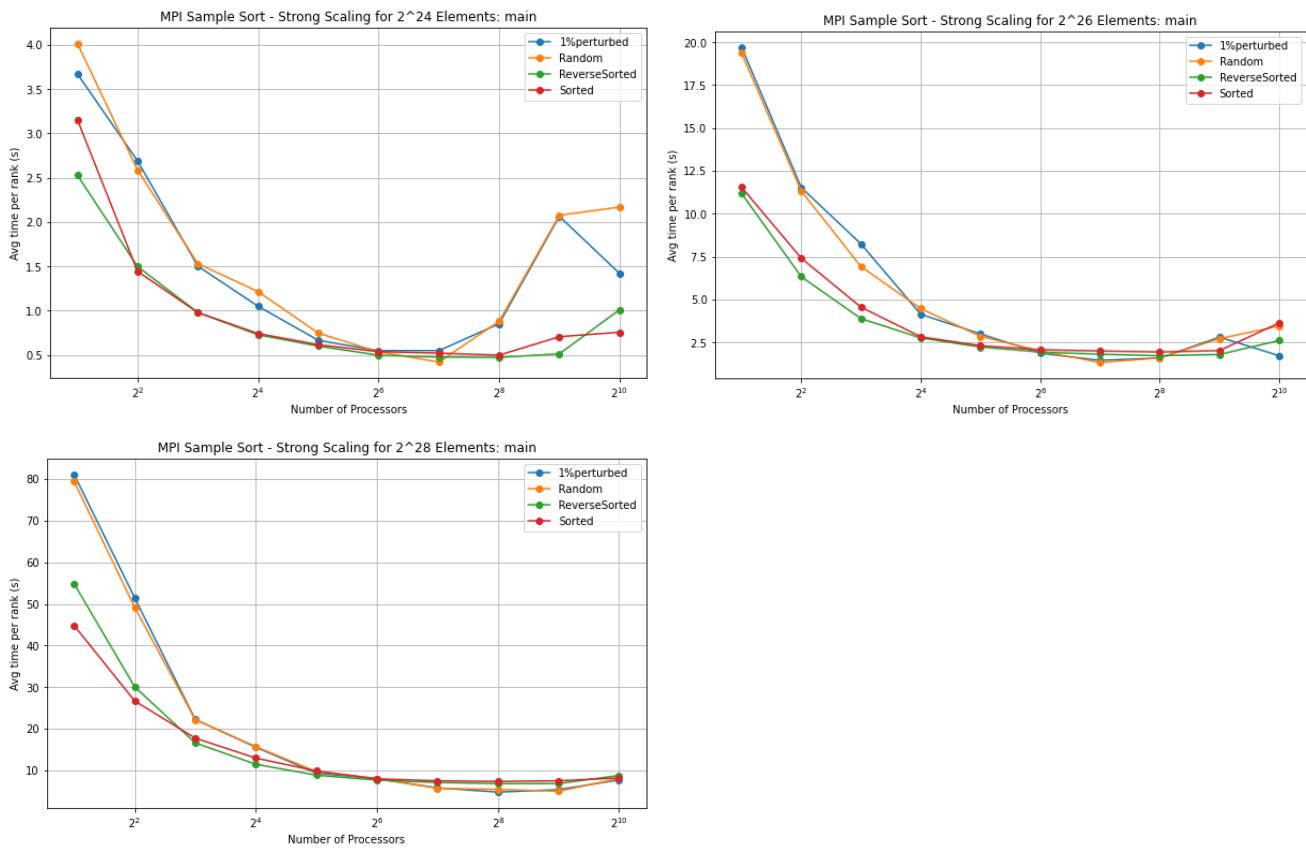
Note: The CUDA implementation of sample sort did not require any comm regions (cudaMemcpy), so the following section will have no comm graphs for the CUDA subsections.

Strong Scaling

main

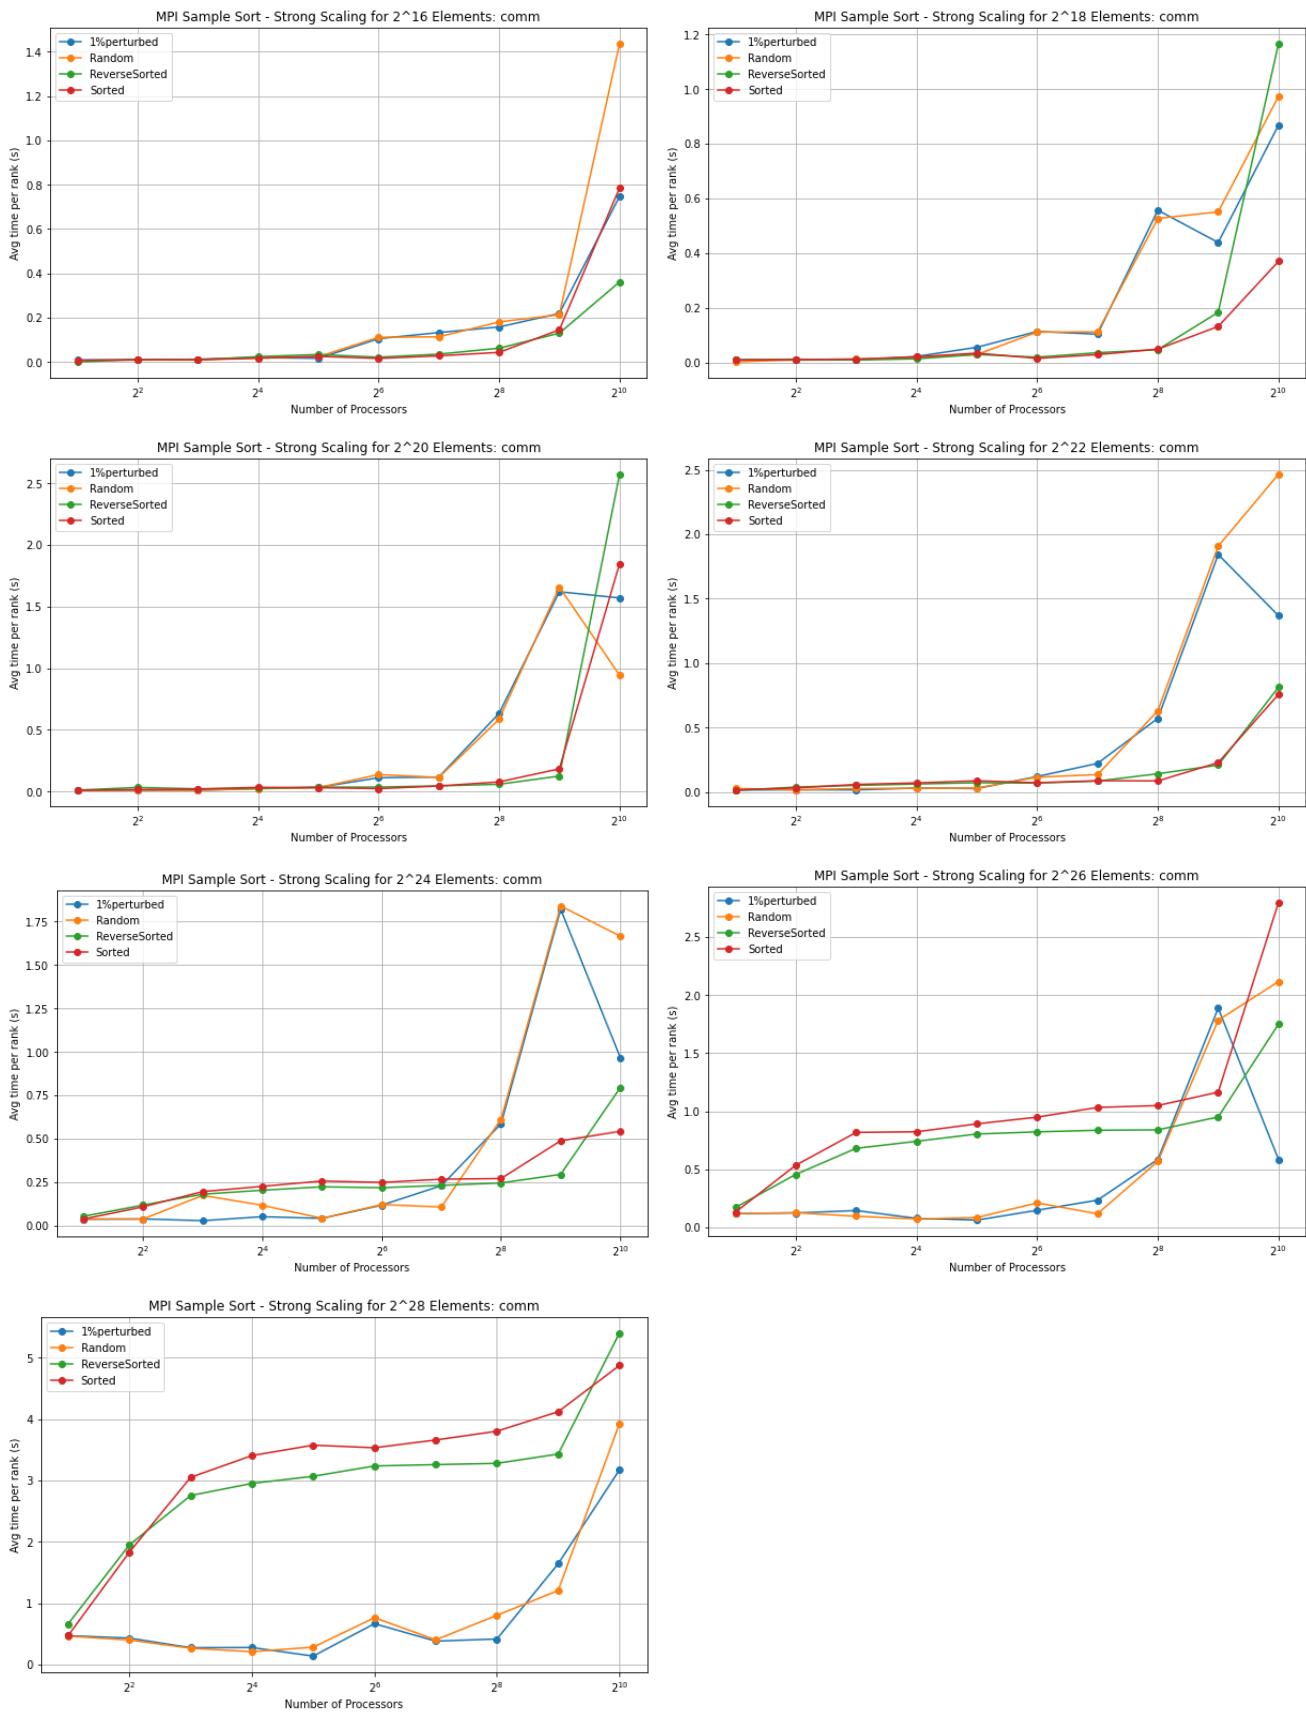
The strong scaling graphs for main tell conflicting stories. Looking at the lower input sizes, up to 2^{16} , the algorithm looks to have poor strong scaling performance as it is relatively flat for the beginning before spiking up as it gets closer to the end. However, looking at the larger input sizes, from 2^{18} and up, the graphs start to display good strong scaling performance as they slope downwards and flatten off once a point of diminishing returns is reached, around the 2^6 or 2^7 number of processes mark. This behavior is caused by the algorithm shifting from being communication-bound in the smaller input sizes to being computation-bound for the larger ones.





comm

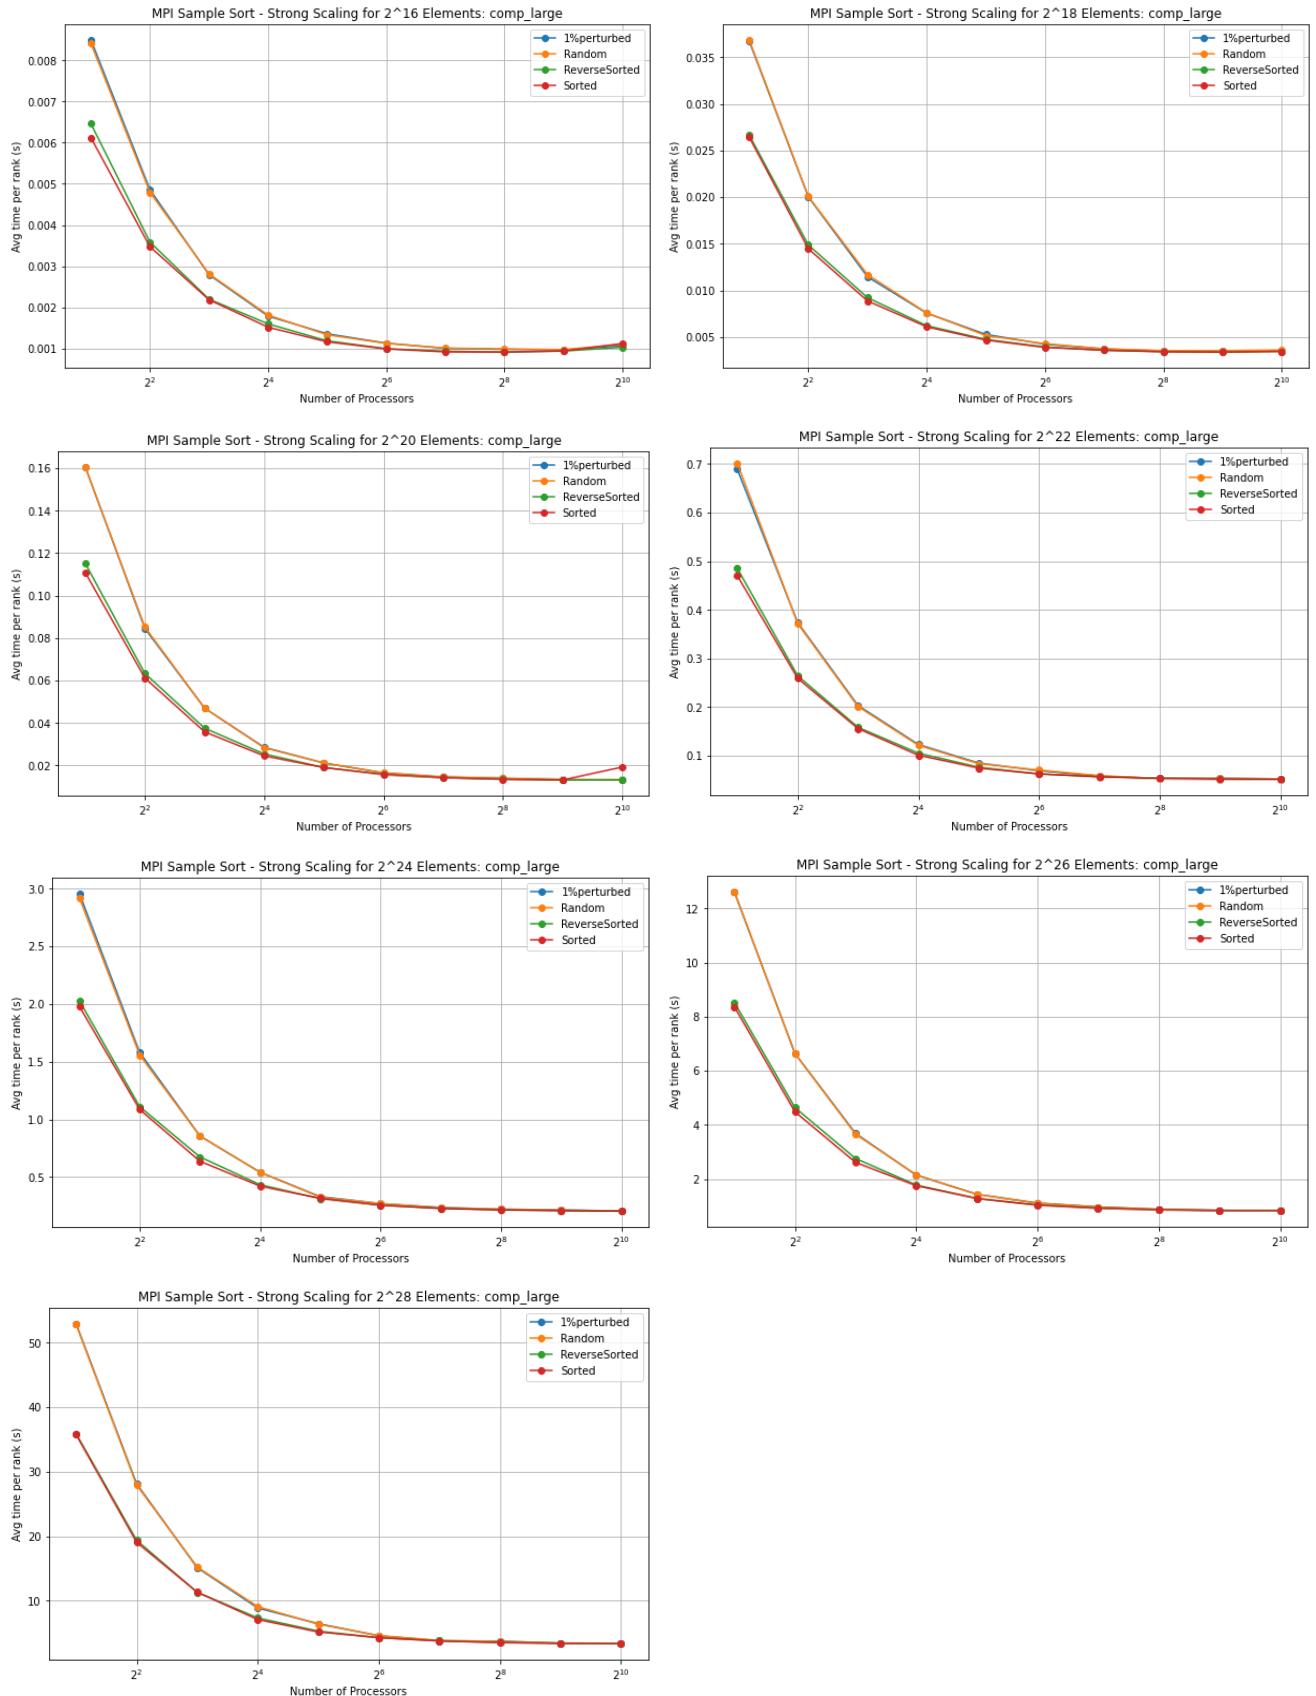
The communication part of this algorithm has poor strong scaling but does not make huge jumps in time as the computation does. Comm starts at less than 0.1 seconds for the smaller input sizes and takes up to a little over 5 seconds for the largest input size. This means comm will dominate the overall runtime in the beginning, but then become less important as input sizes get larger. This matches what is seen in the graphs for main.



comp_large

Comp_{large} has good strong scaling for all input sizes and hits a point of diminishing returns somewhere in the 2⁶ to 2⁷ processes range. As expected, the runtime grows from less than a second to around 50 seconds as larger input sizes are tested. This further explains why computation dominates the graphs of

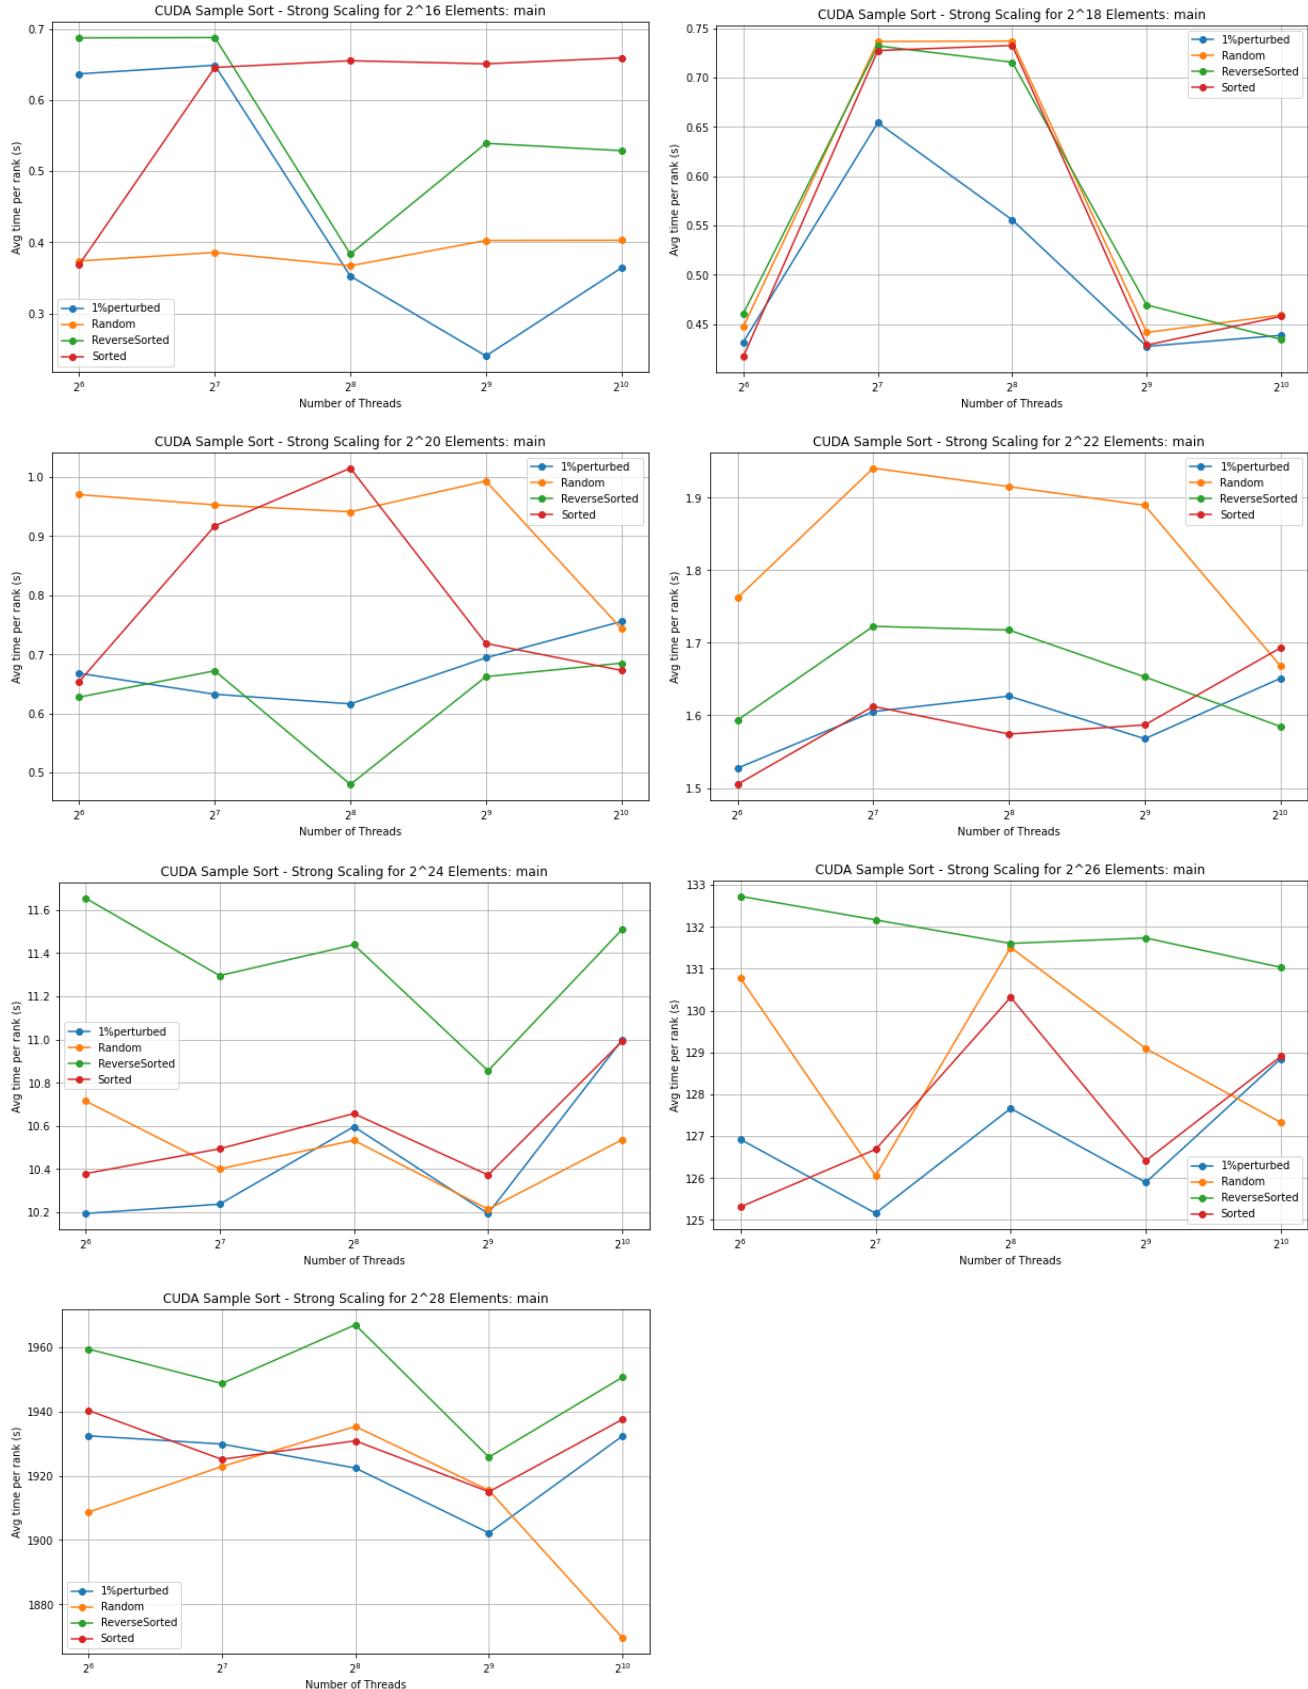
main only for the larger input sizes, as 2^{24} is where computation begins to consistently take more time than communication.



CUDA

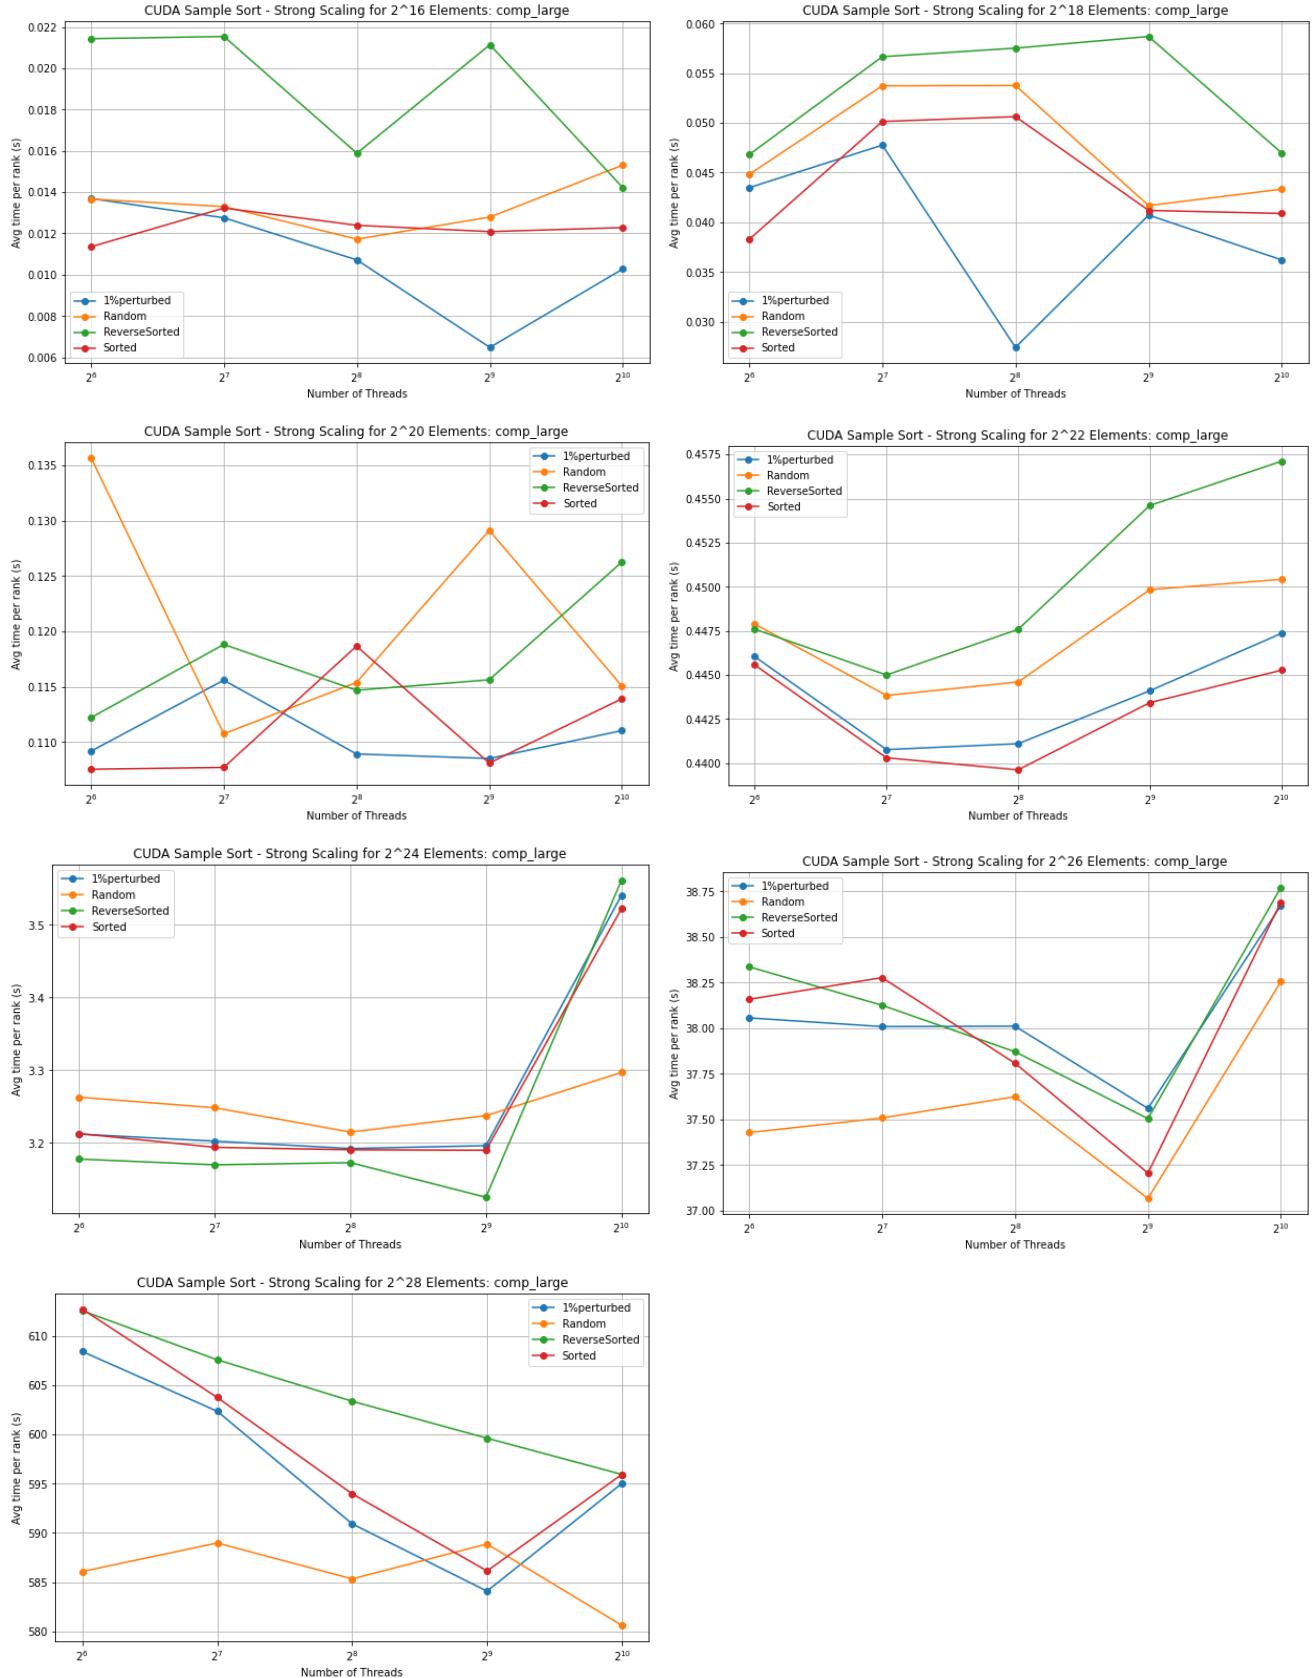
main

The strong scaling graphs for CUDA do not paint a very clear picture in either direction. The graphs are all over the place, but most tend to generally trend upwards overall, leading to a conclusion of poor strong scaling. This behavior is most likely due to the implementation not taking advantage of the GPU as much as it could. An interesting point to note when looking at these graphs, while the lines are very up and down, relative to the overall time taken, the range on the Y-axis is not that big.



comp_large

As CUDA did not have communication for the algorithm, comp_large looks fairly similar to main. While the 2^{28} graph shape seems to be an outlier with the lines trending downwards, the rest generally trend upwards implying poor strong scaling performance. The up-and-down nature of these graphs is again most likely due to the algorithm not taking proper advantage of the GPU.

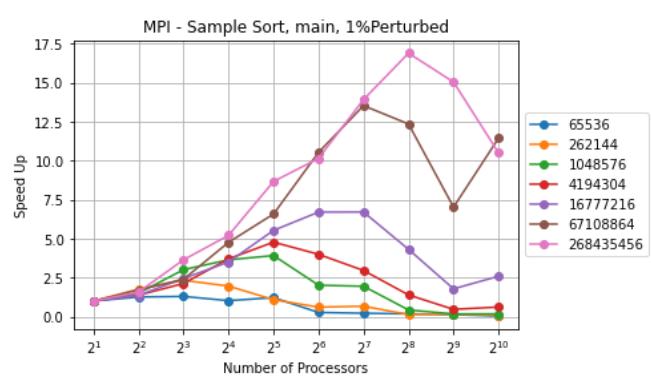
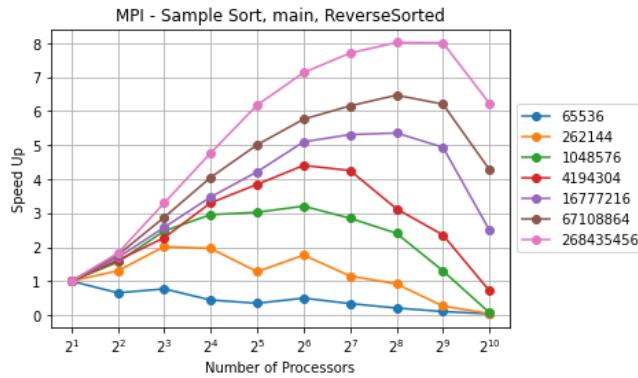
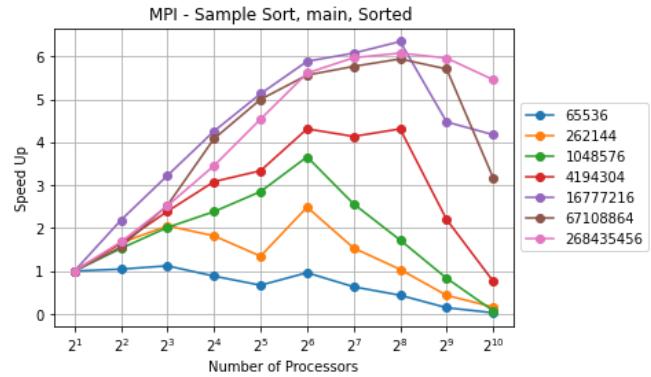
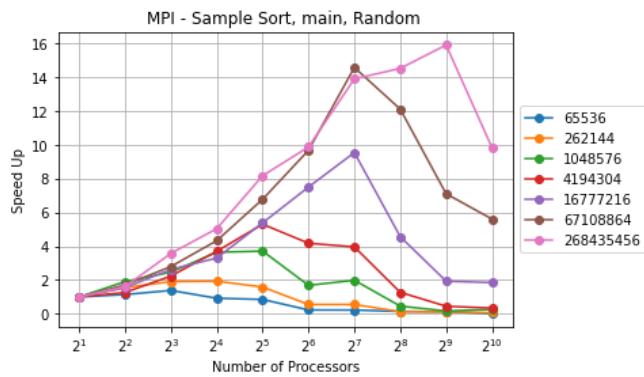


Speed up

MPI

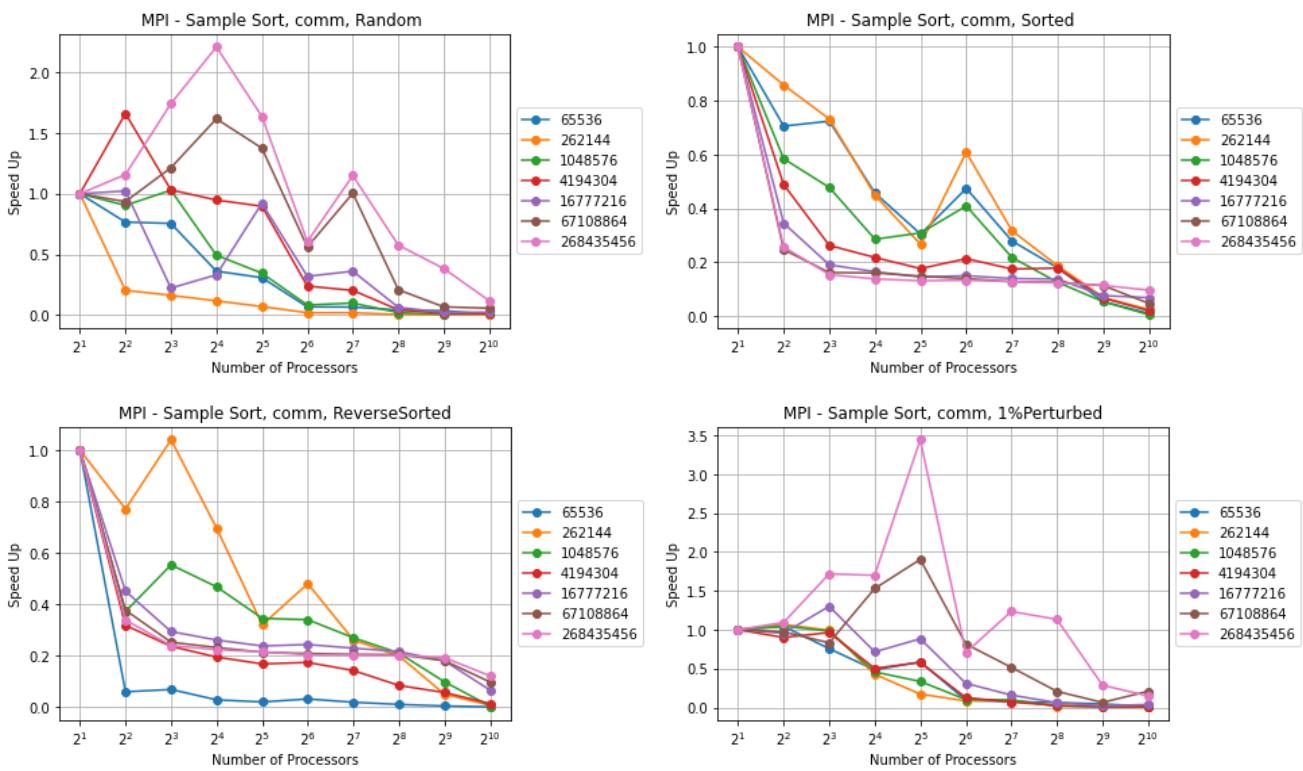
main

The MPI Speedup graphs for main have a similar pattern of switching what they are bounded by as seen in the strong scaling graphs, but reversed. Each one starts with a fairly good speedup shape as it is dominated by the speedup of computation in the early process counts, then starts to fall off as the communication overhead becomes greater and greater. As expected, the larger the input size, generally the better the speedup the algorithm sees.



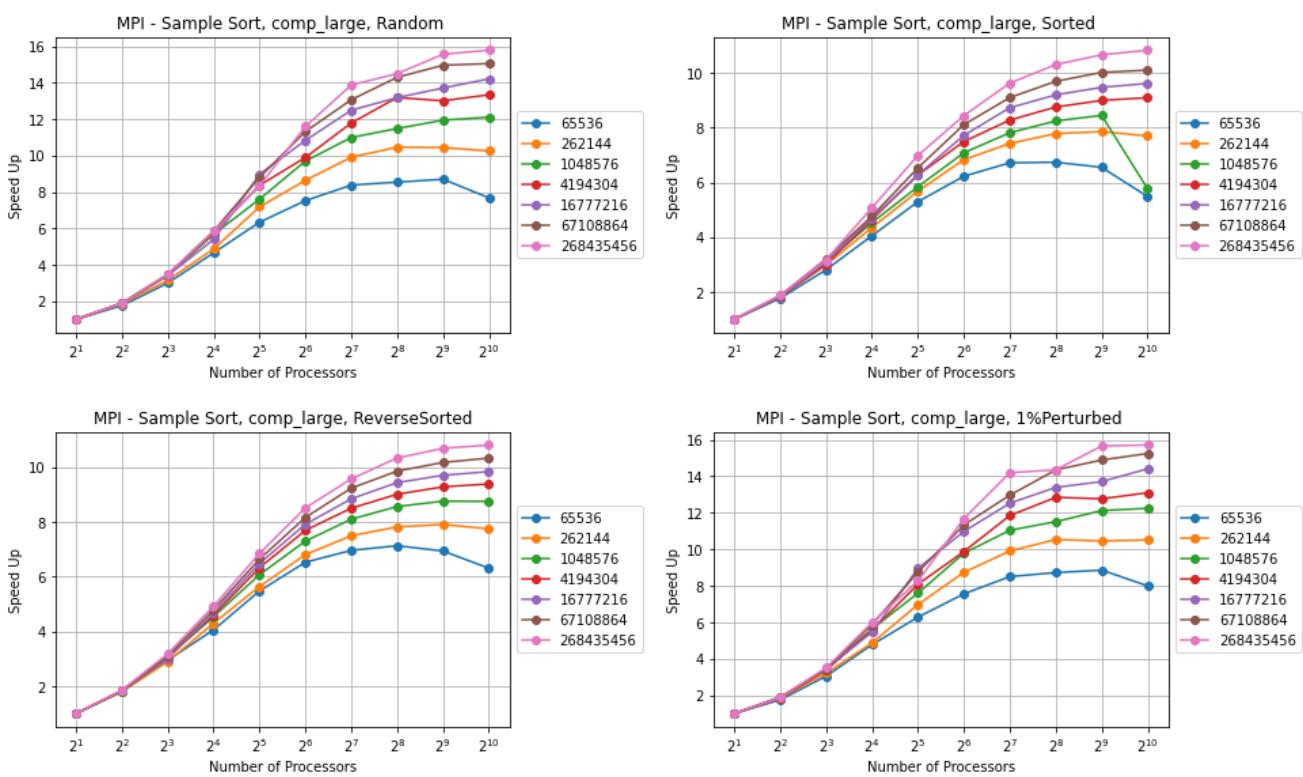
comm

Communication speedup graphs tell an expected story when considering the region has poor strong scaling performance. It starts off around 1.0 for the smaller input sizes but always ends up nearing 0 by the time the largest input size is reached.



comp_large

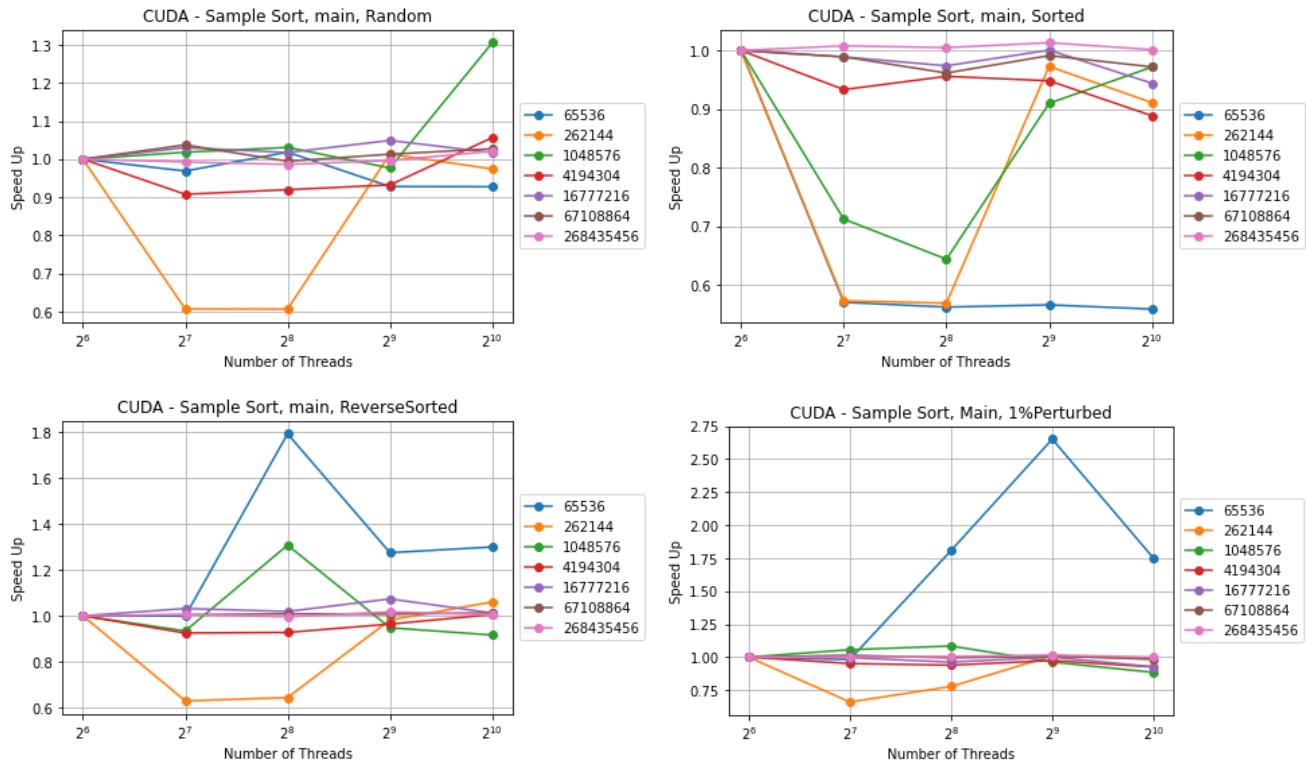
The comp_large speedup graphs are almost picture-perfect when talking about speedups found in the real world. All input types see large speedups in the beginning as more processes are added for small numbers, but then plateau to various degrees once it reaches the point of diminishing returns at around 2^7 processes. As expected, the larger input sizes tend to achieve better speedups as they allow the algorithm to better utilize the maximum allowed resources.



CUDA

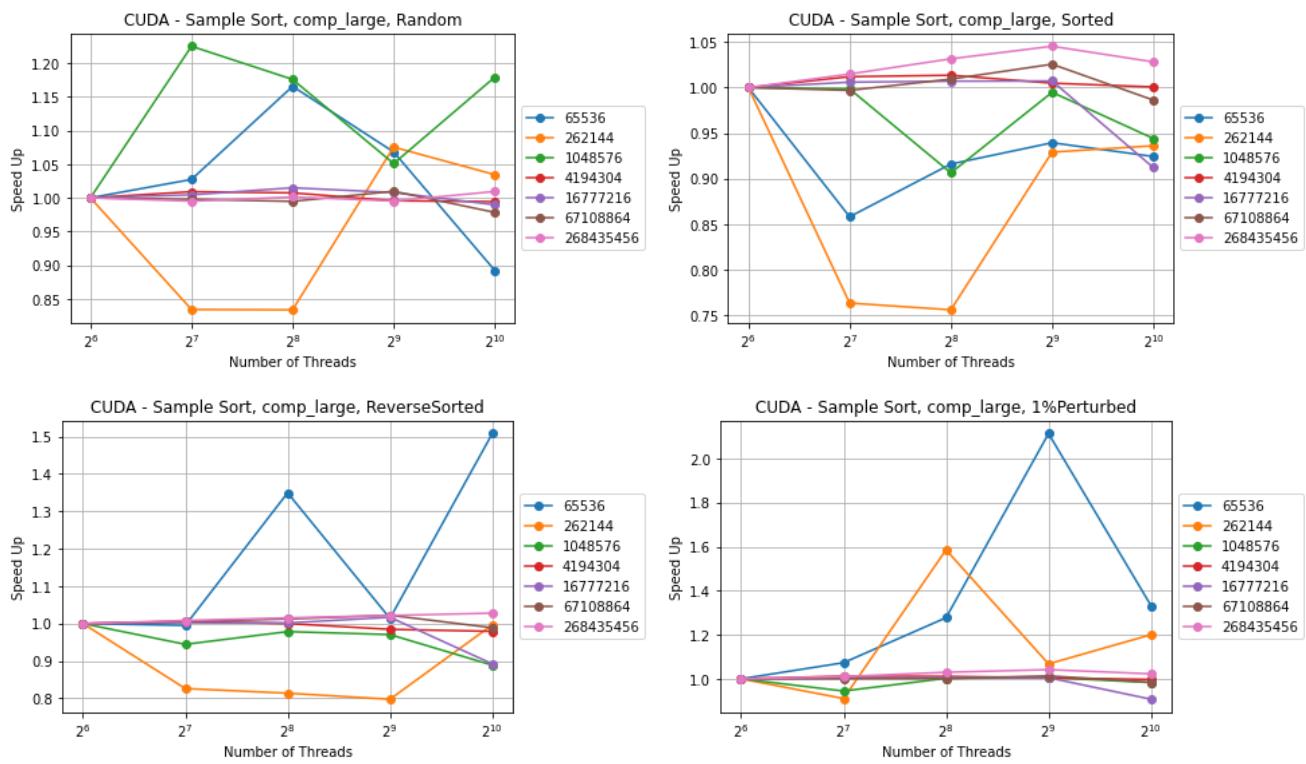
main

The CUDA speedup graphs for main show an interesting pattern of the larger input sizes hugging the 1.0 line and the smaller inputs bouncing around to a greater degree. This most likely means there are some small variations in the runtime no matter the input size and it is just more noticeable on the smaller sizes because their overall shorter times make them more susceptible to change.



comp_large

The comp_large graphs are similar to the main graphs, but there seems to be a little bit more randomness. This most likely means the data initialization and correctness check sections of the program also have unclear runtime trends, therefore stabilizing the overall times a little bit, leading to the calmer graphs for main shown above.



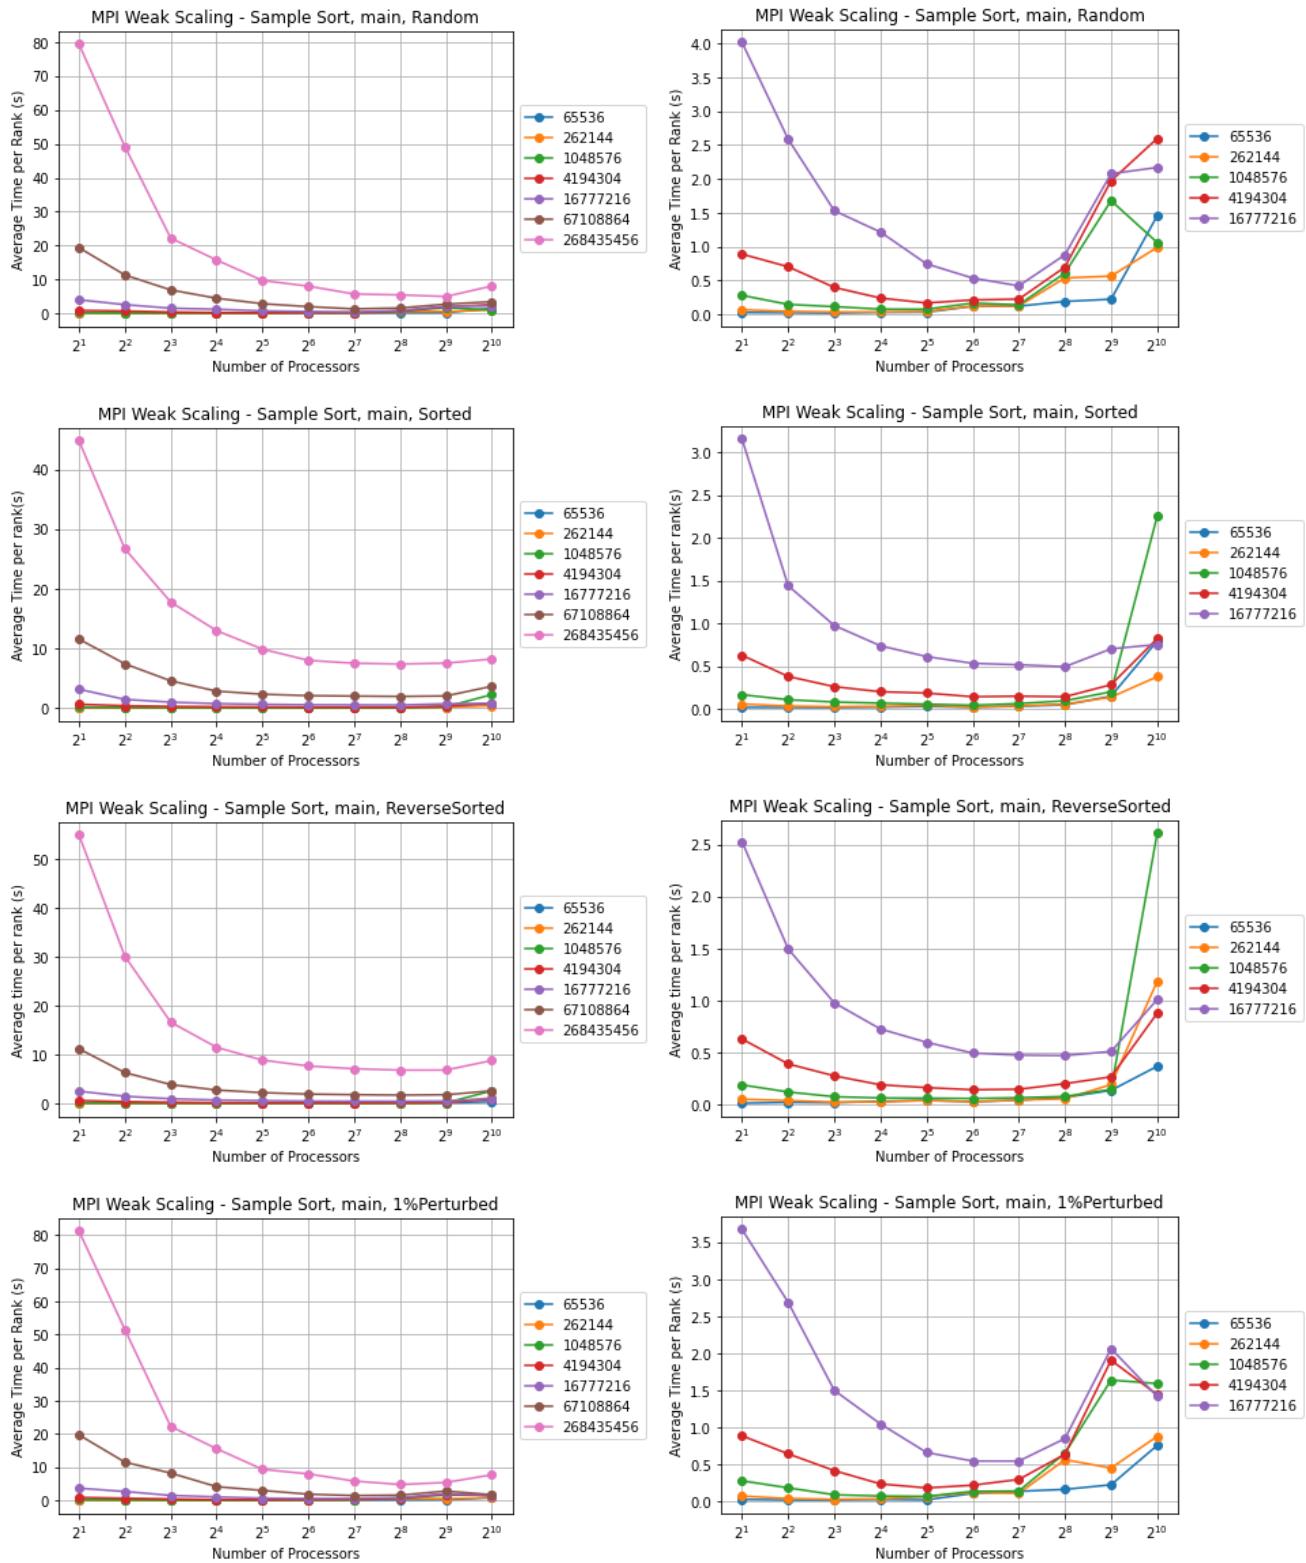
Weak Scaling

Note: Due to the longer runtimes of larger input sizes, the graph shape for the smaller inputs gets lost. To combat this, the left column of graphs has all input sizes and the right column has the same graph, but with the larger input sizes removed to enable analysis of smaller input sizes.

MPI

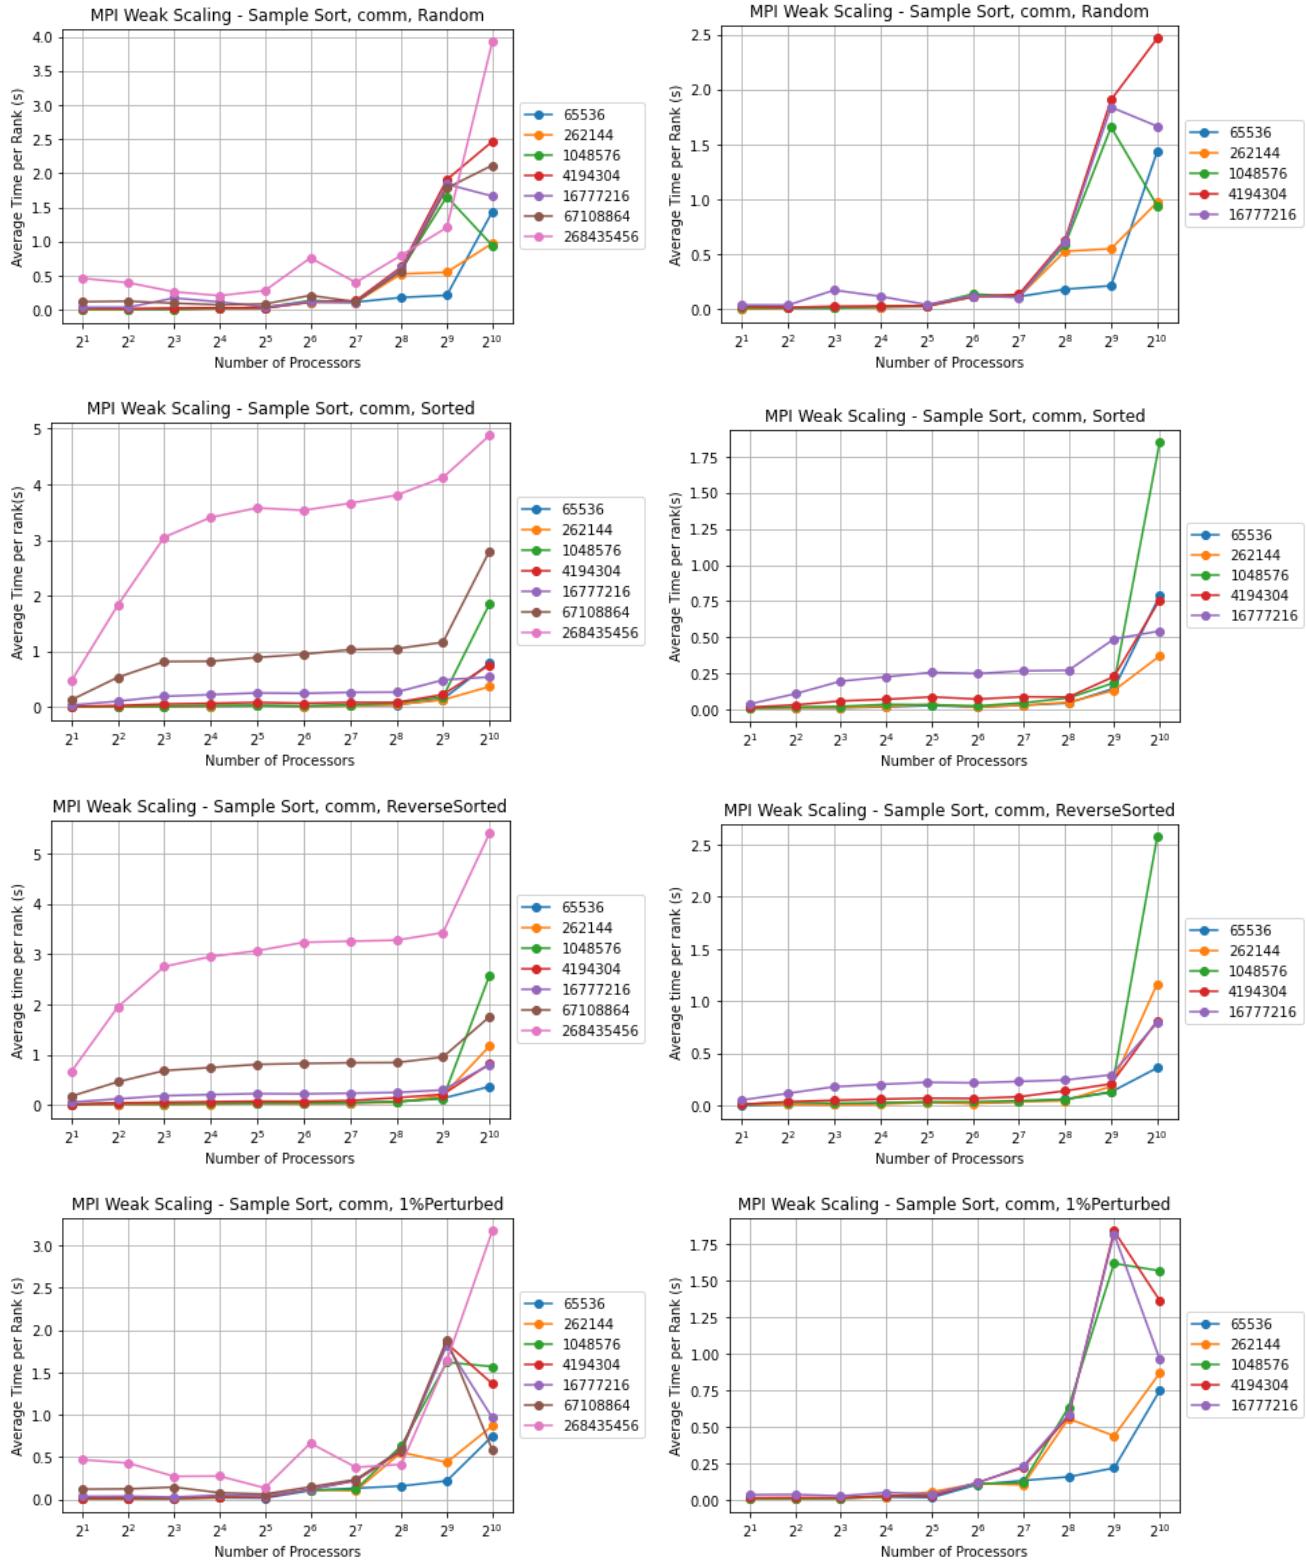
main

Weak scaling in main for the MPI implementation seems pretty good when looking at all input sizes. It is not until just the smaller input sizes are graphed that the upward spike near the largest number of processes can be seen. This is due to the communication overhead for larger number of processes that can be seen in the comm graphs below. Prior to that, the runtime is dominated by the good weak scaling of the comp_large region.



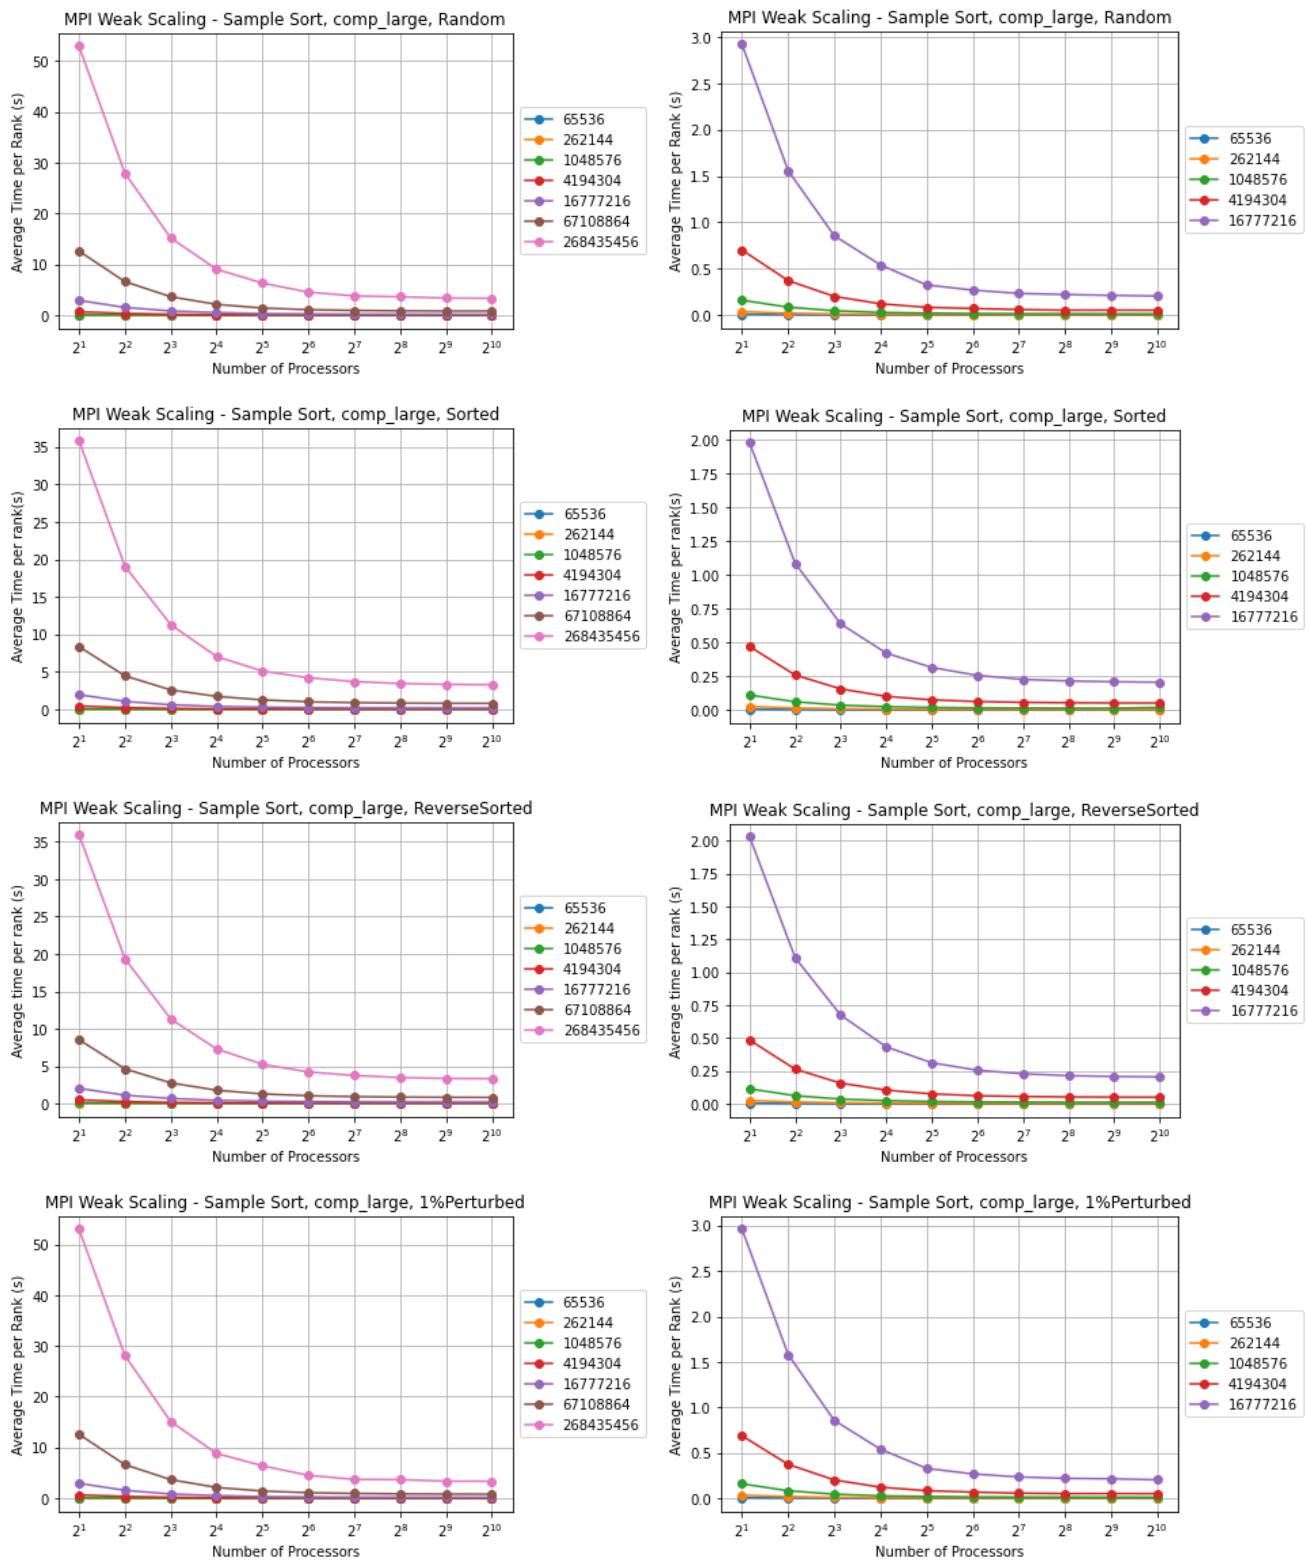
comm

Graphs for the comm region, with the exception of the largest input size, seem to show decent weak scaling for lower numbers of processes. It is not until 2^6 or 2^7 , when more than one node and therefore network communication is added, do the comm graphs start to show poor weak scaling and bleed into the main graphs.



comp_large

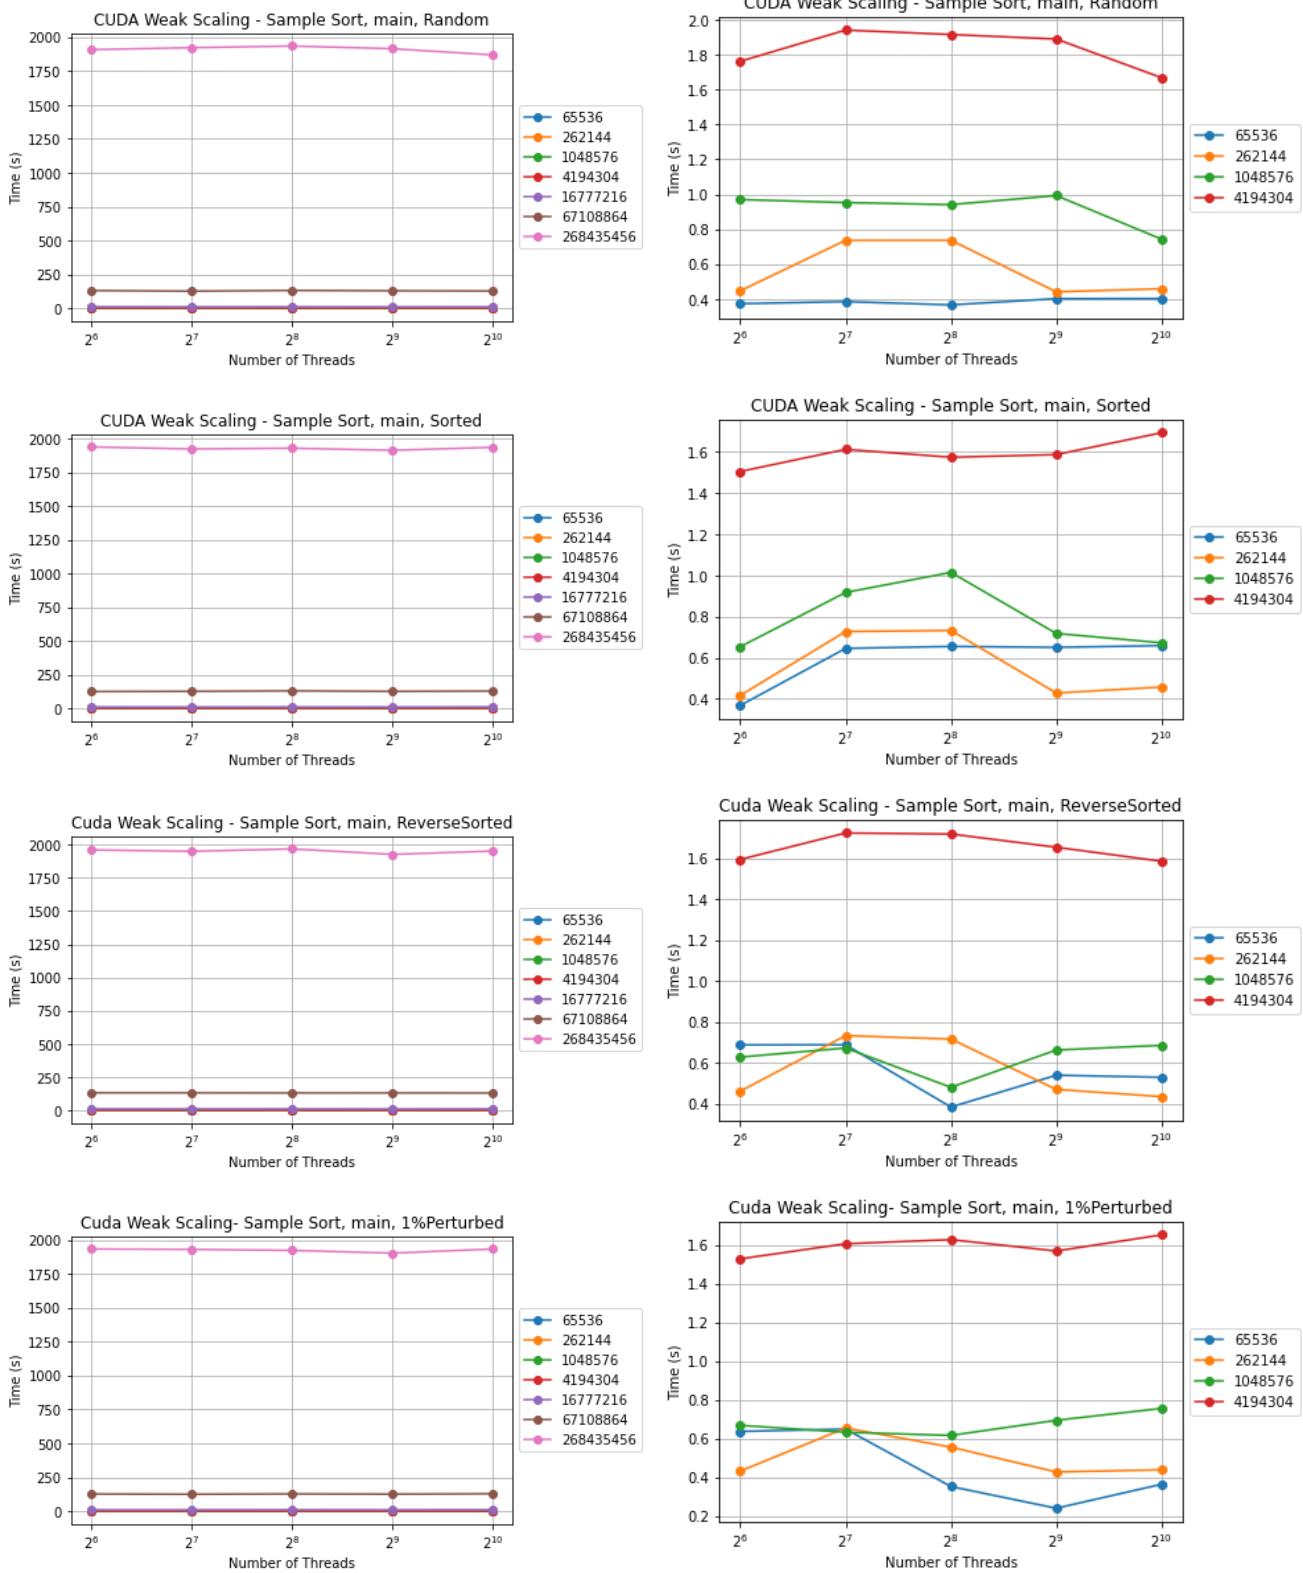
The graphs for comp_large show good weak scaling across the board for all input sizes as they have a steep drop at the beginning and then flatten out at the end. An interesting observation that can be seen on the graph of smaller input sizes is that sizes of 2^{22} and under do not see as steep of an initial dropoff as the larger sizes do.



CUDA

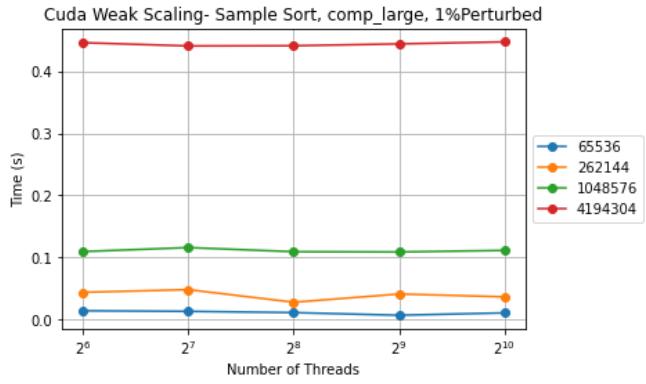
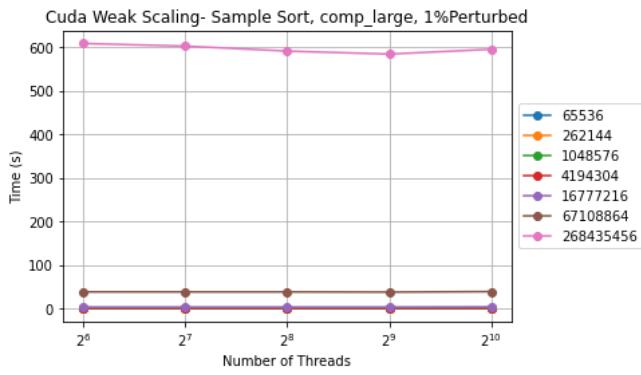
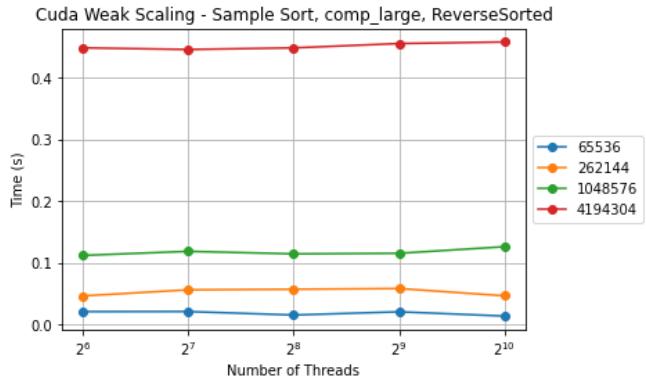
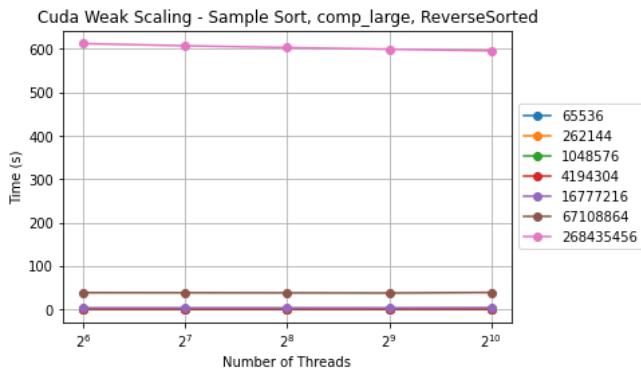
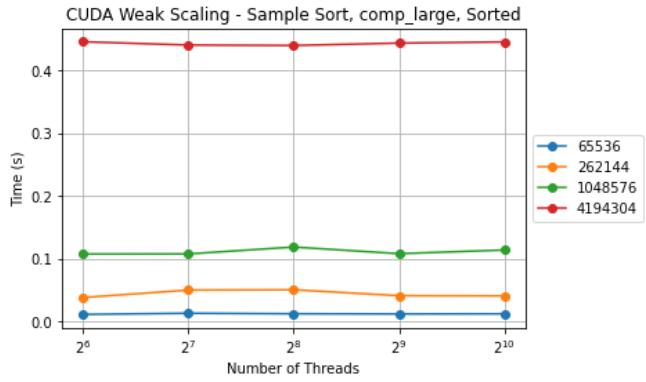
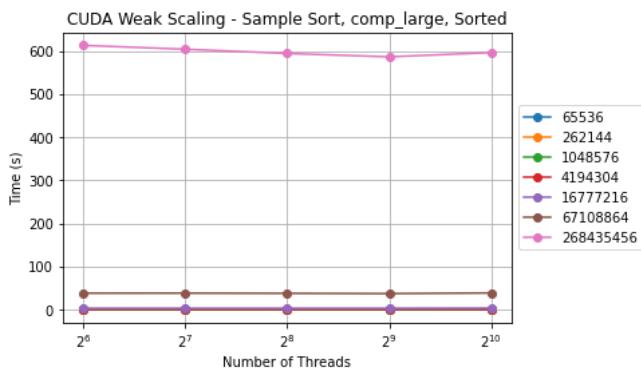
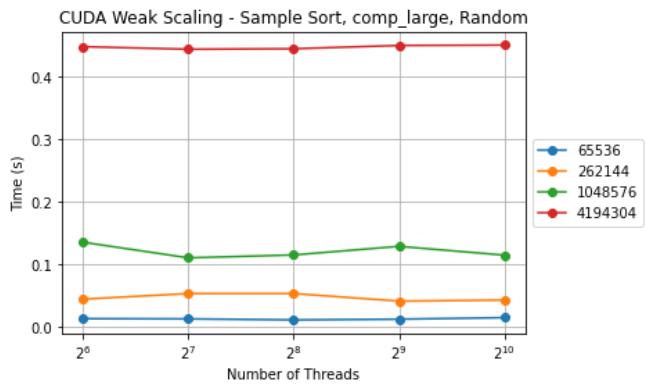
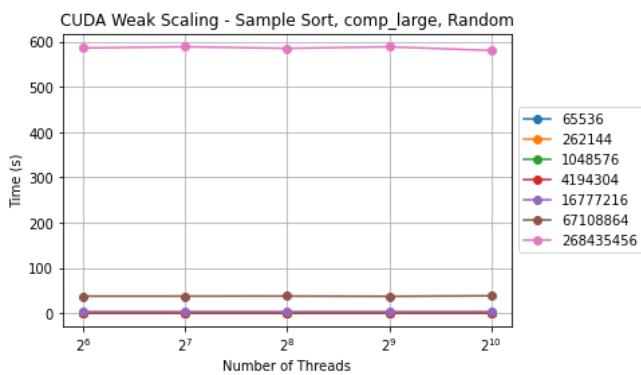
main

The CUDA weak scaling graphs for main show fairly flat lines overall on the graphs with larger input sizes implying good weak scaling performance. When looking at just the smaller input sizes, while a bit more variation can be seen, the Y-axis scale is much more narrow so small ups and downs are not surprising.



comp_large

Weak scaling comp_large graphs echo the story told by the main graphs with the exception of the smaller input sizes. These maintain much more of a flat line which means the variation found in the graphs above is due to either data initialization or correctness checking runtimes fluctuating, not the actual sorting computation time.



Merge Sort

Weak Scaling

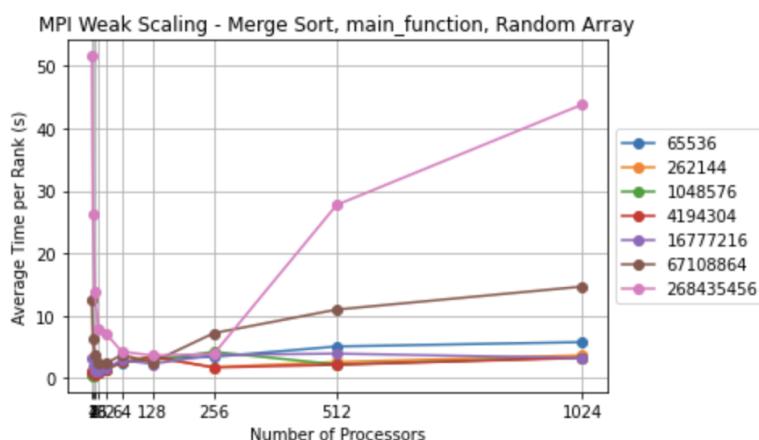
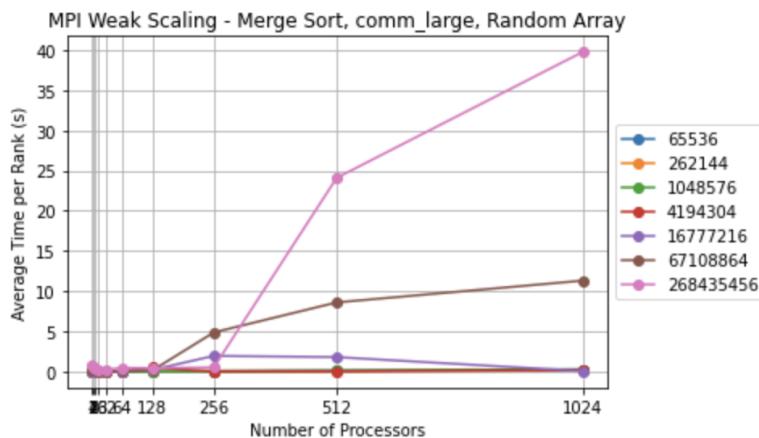
MPI

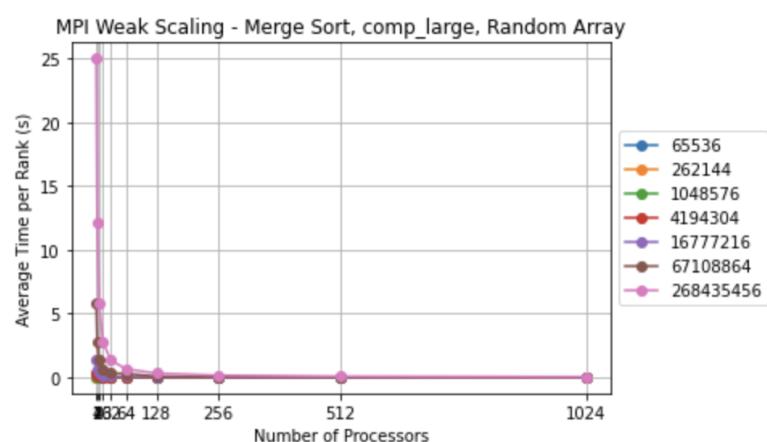
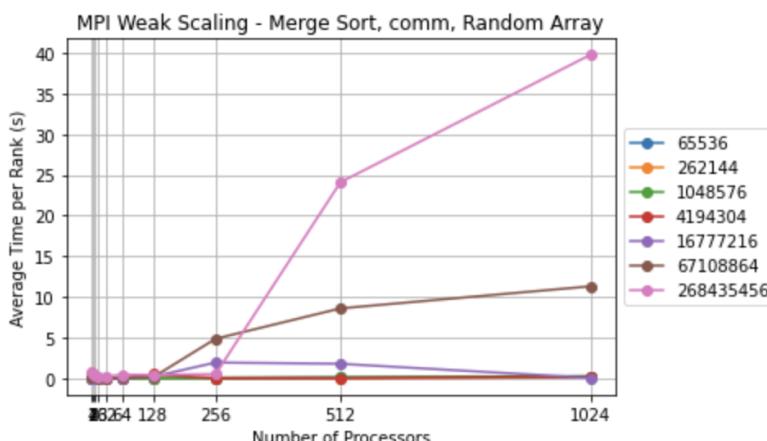
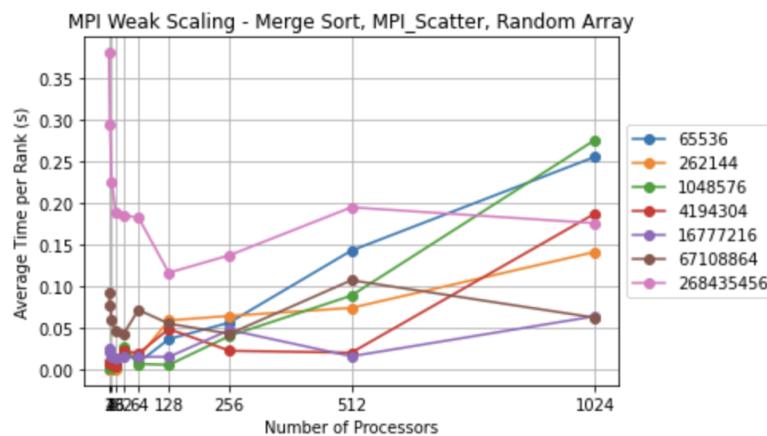
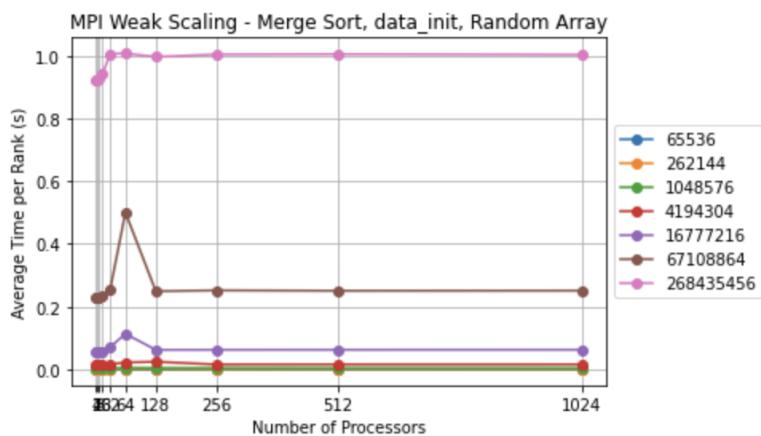
One thing to note when I did my weak scaling is that I chose to measure the average time per rank over number of processors. The reason I chose average time over total time is that discussing with the TA about

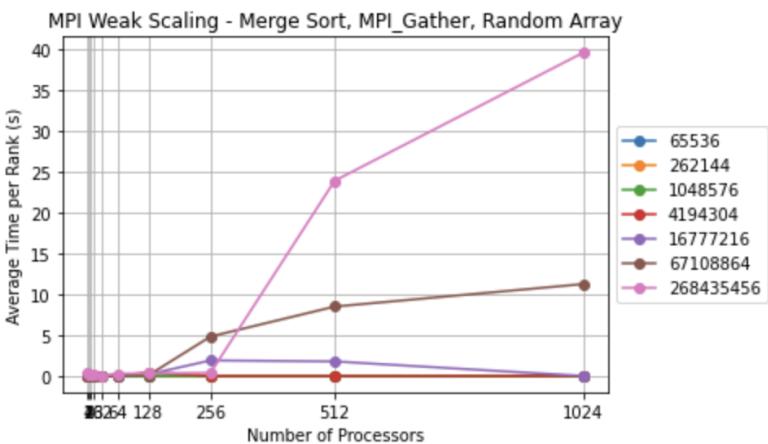
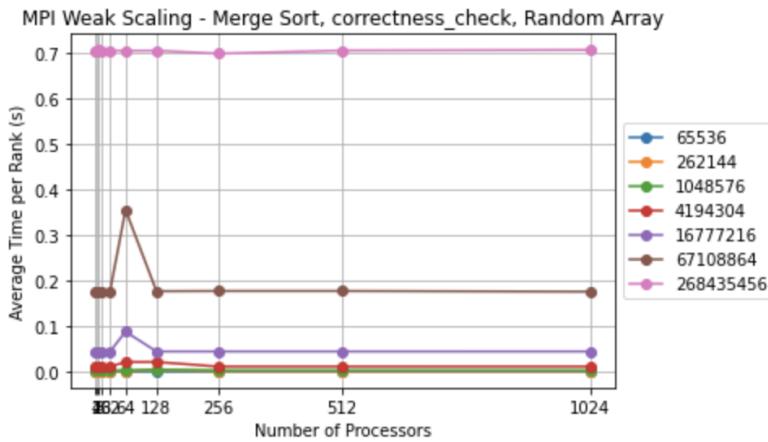
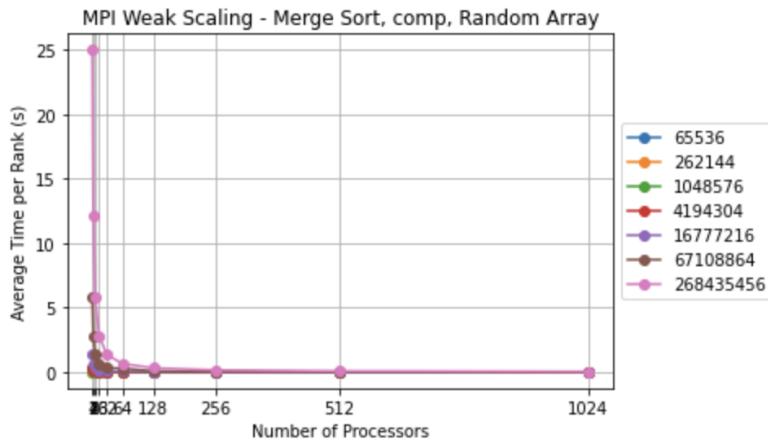
how these sorts work, total time is always going to grow as you increase processors because its an aggregate of all times over all processors. Additionally, I wanted to begin by focusing on the main function times for everything because I thought a good introduction to the analysis is how the program as a whole ran on average.

RANDOM INPUT ARRAY

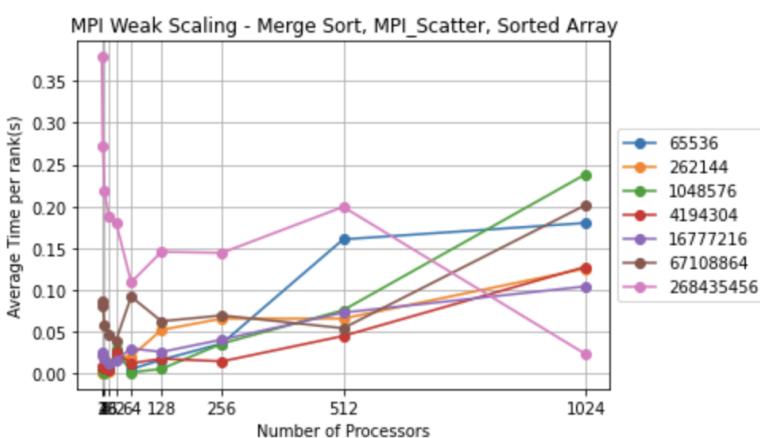
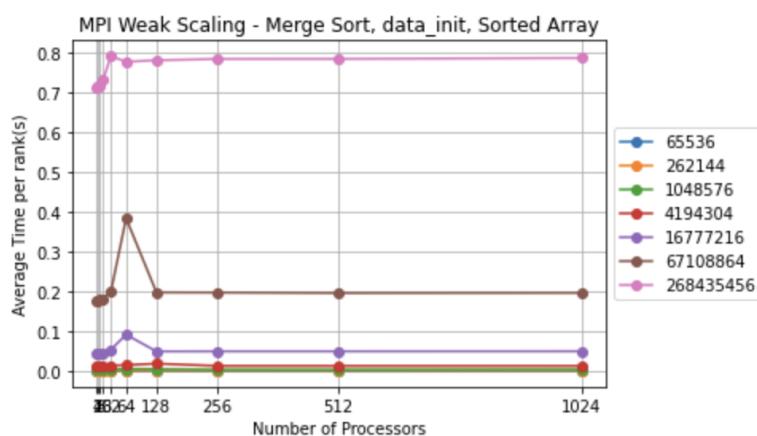
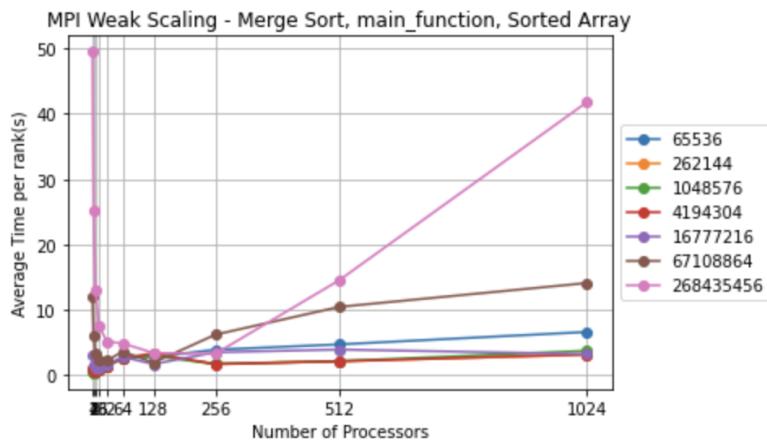
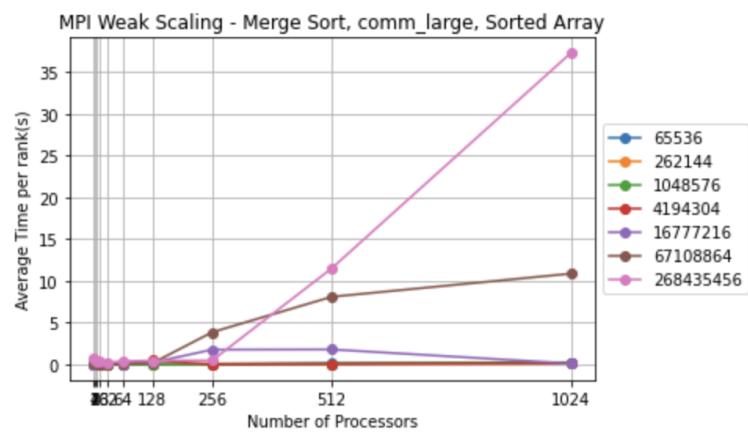
The analysis begins with weak scaling, specifically with MPI. For weak scaling, I began with the input type of a random array, where each element is a randomly generated number from 0 to n, where n is the size of the array. The most obvious trend between all of the graphs in this section is that as you increase the number of elements, the algorithm performs more slowly. This is shown best with the largest input size of 2^{28} having the largest graph, as signified by the pink line. It is interesting to note that the computation time for all the times looks relatively the same as we parallelize, and that most of the variation in times comes from the communication. I measured both the MPI_Scatter and MPI_Gather functions, and the MPI_Gather function takes much longer than the scatter. In the sorted array input type, we see very similar behavior, where larger input sizes seemed to have a much longer time to complete. The same behavior is present in the reverse sorted array as well. Finally, in the data initialized with 1% of the data perturbed, we see the similar data. We can assume this due to merge sort typically always breaking the arrays down into a single element subarray and merging them back together, so we should expect relatively consistent behavior.

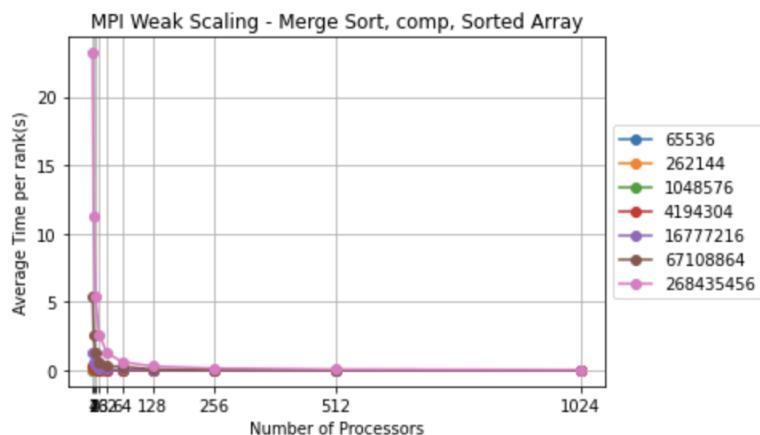
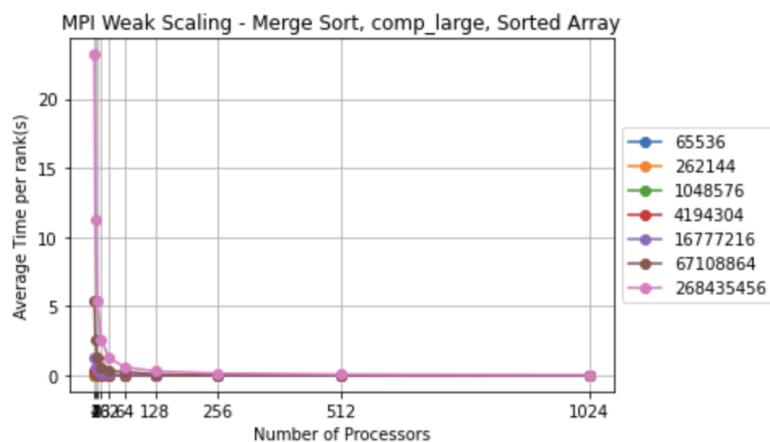
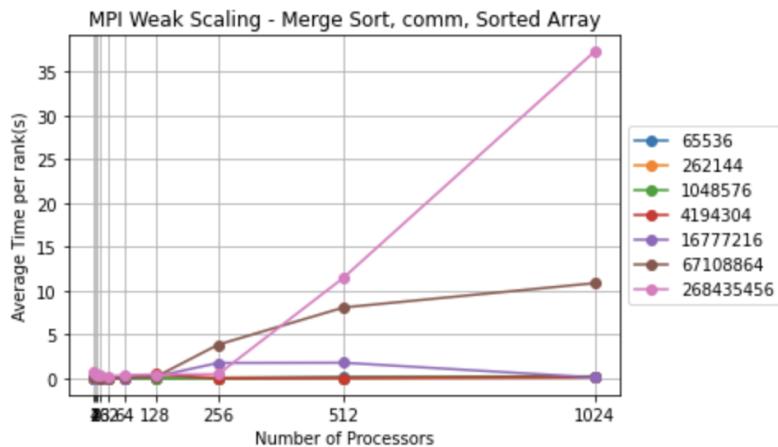




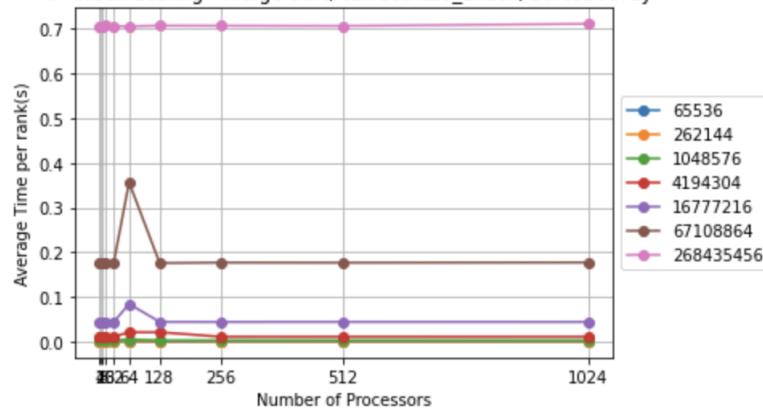


SORTED INPUT ARRAY

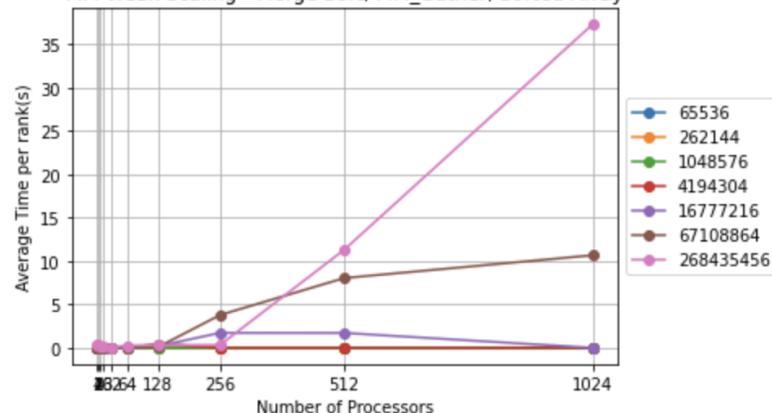




MPI Weak Scaling - Merge Sort, correctness_check, Sorted Array

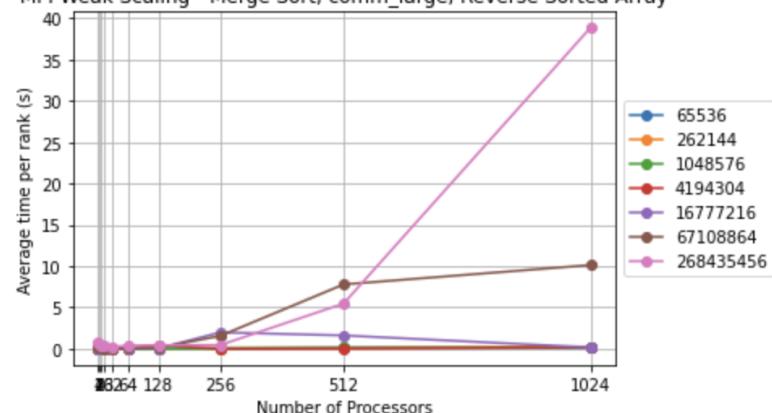


MPI Weak Scaling - Merge Sort, MPI_Gather, Sorted Array

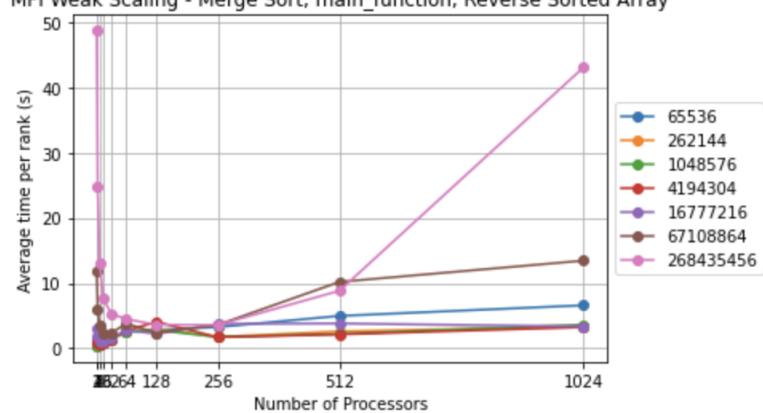


REVERSE SORTED INPUT ARRAY

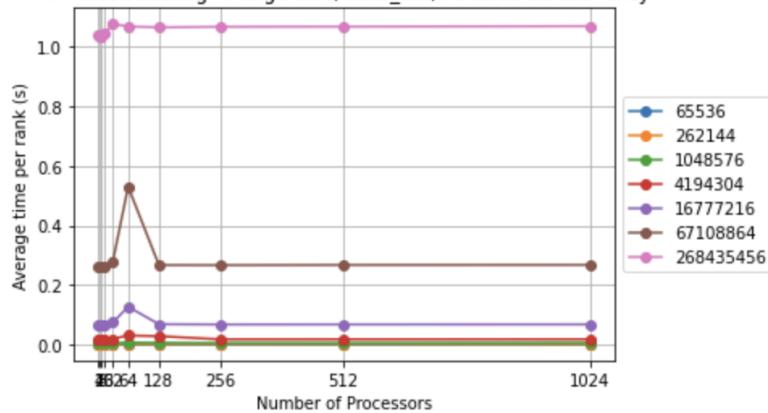
MPI Weak Scaling - Merge Sort, comm_large, Reverse Sorted Array



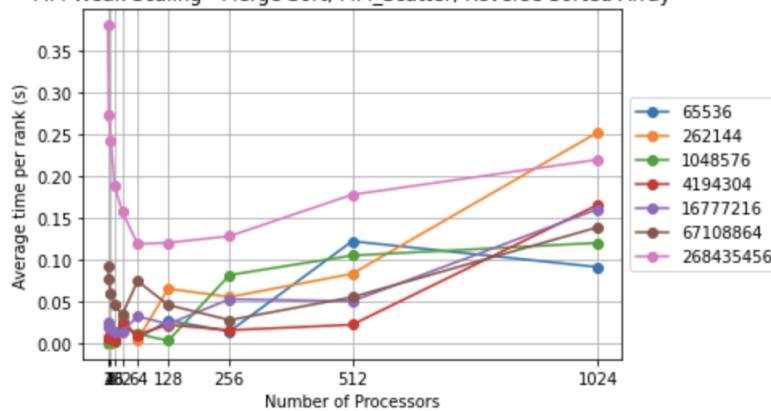
MPI Weak Scaling - Merge Sort, main_function, Reverse Sorted Array



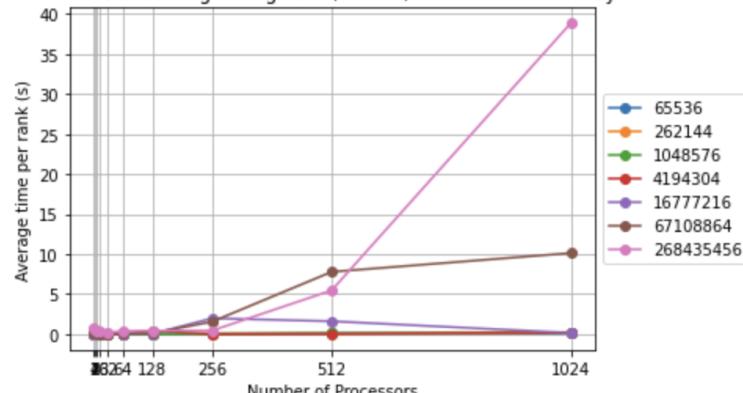
MPI Weak Scaling - Merge Sort, data_init, Reverse Sorted Array



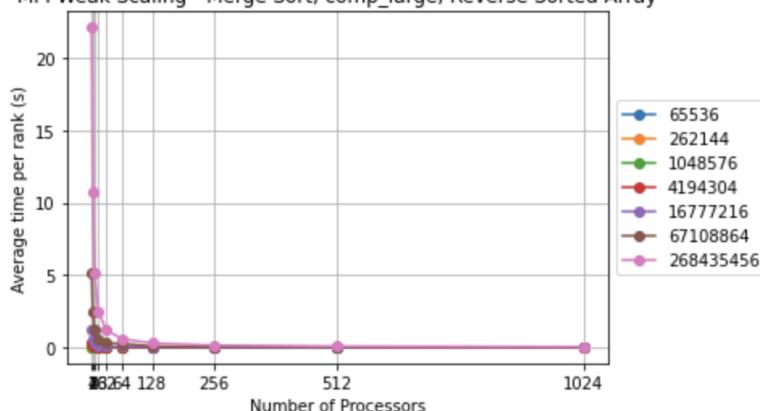
MPI Weak Scaling - Merge Sort, MPI_Scatter, Reverse Sorted Array



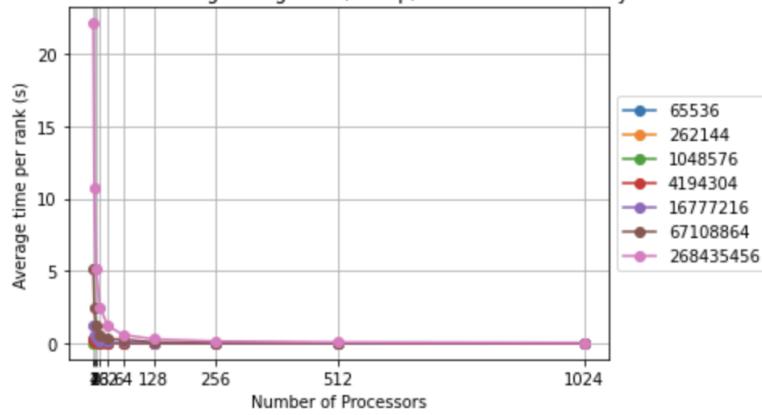
MPI Weak Scaling - Merge Sort, comm, Reverse Sorted Array



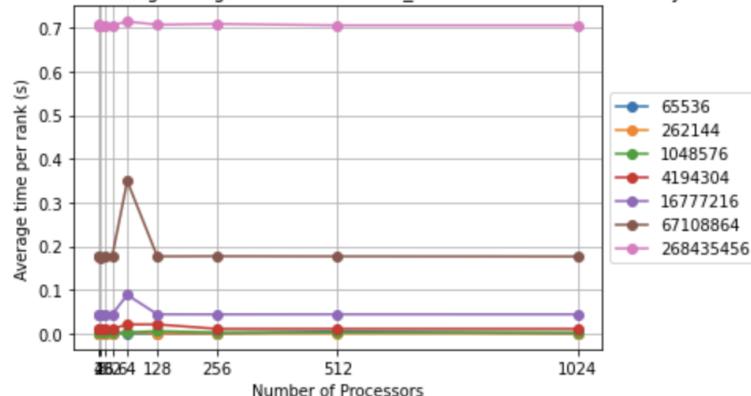
MPI Weak Scaling - Merge Sort, comp_large, Reverse Sorted Array



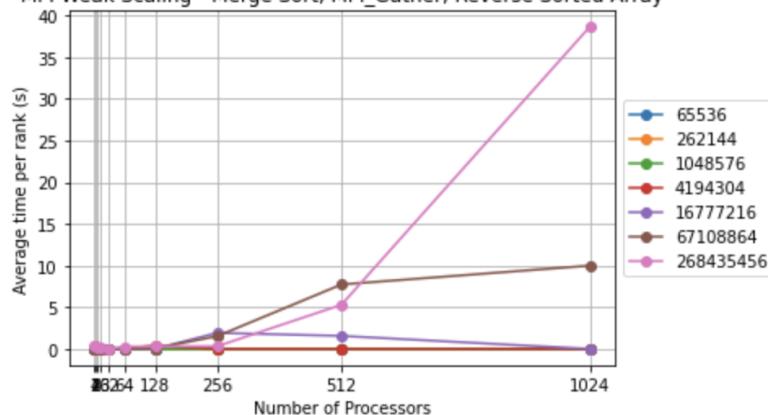
MPI Weak Scaling - Merge Sort, comp, Reverse Sorted Array



MPI Weak Scaling - Merge Sort, correctness_check, Reverse Sorted Array

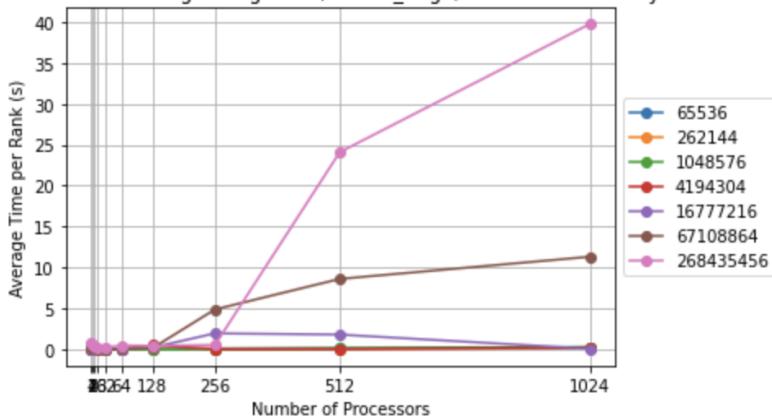


MPI Weak Scaling - Merge Sort, MPI_Gather, Reverse Sorted Array

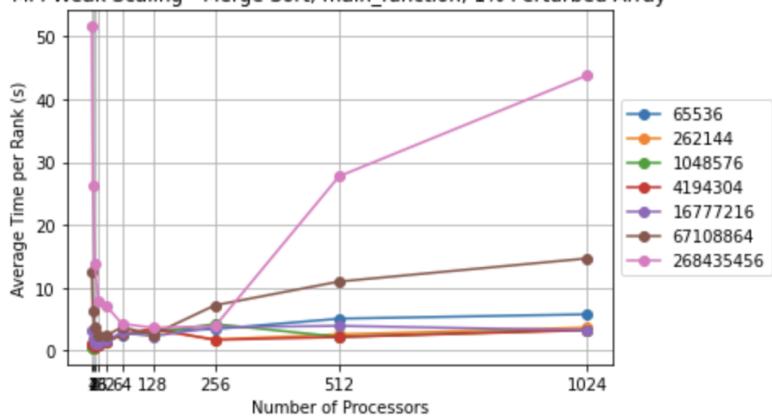


1% PERTURBED INPUT ARRAY

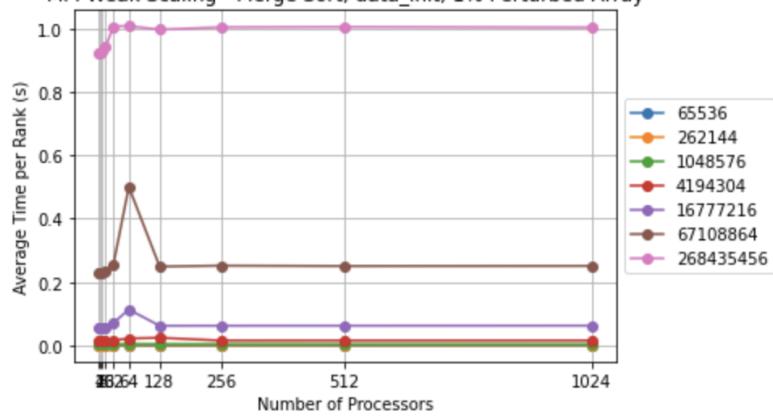
MPI Weak Scaling - Merge Sort, comm_large, 1% Perturbed Array



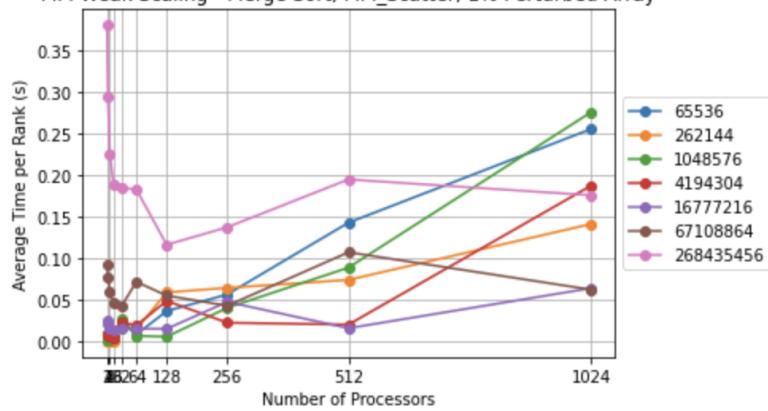
MPI Weak Scaling - Merge Sort, main_function, 1% Perturbed Array

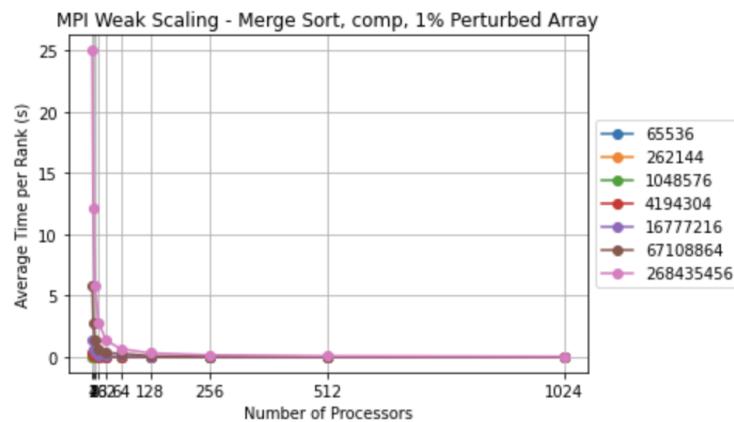
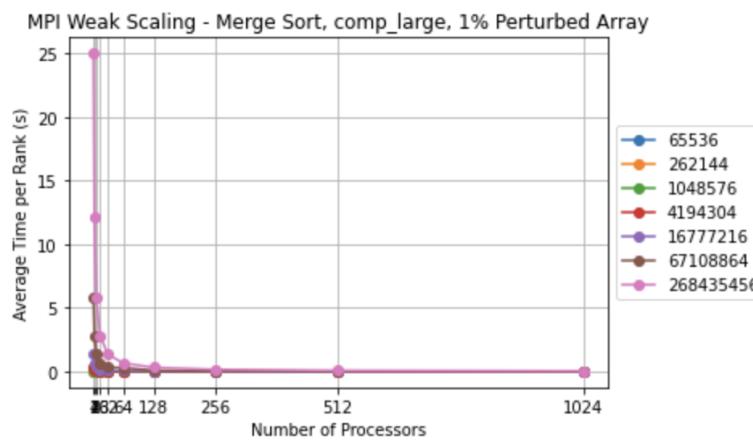
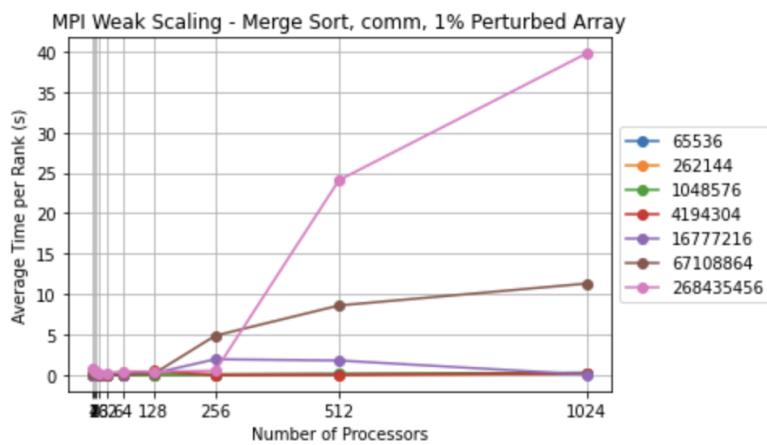


MPI Weak Scaling - Merge Sort, data_init, 1% Perturbed Array

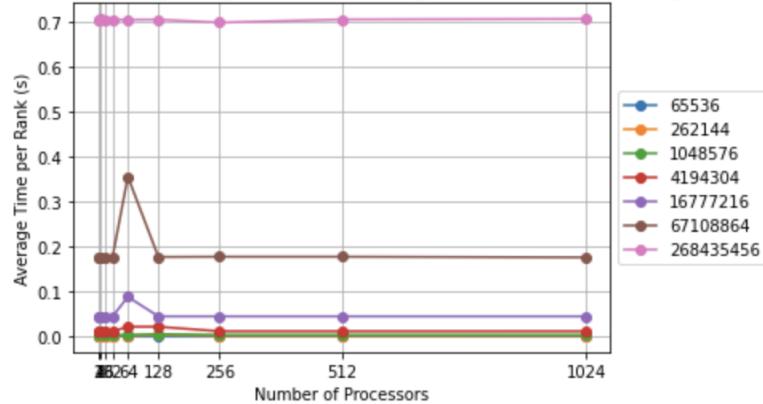


MPI Weak Scaling - Merge Sort, MPI_Scatter, 1% Perturbed Array

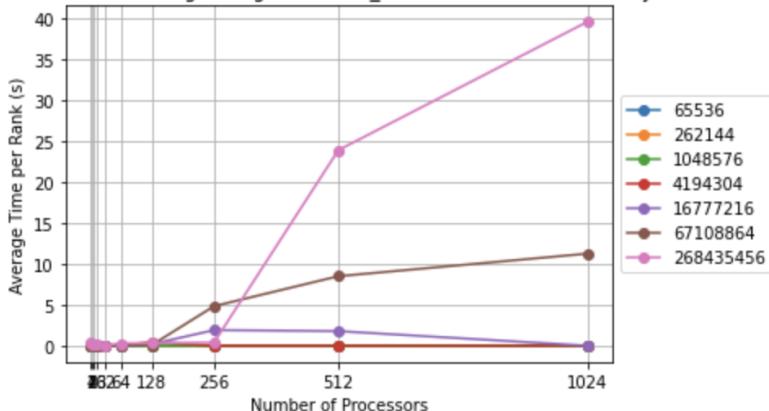




MPI Weak Scaling - Merge Sort, correctness_check, 1% Perturbed Array



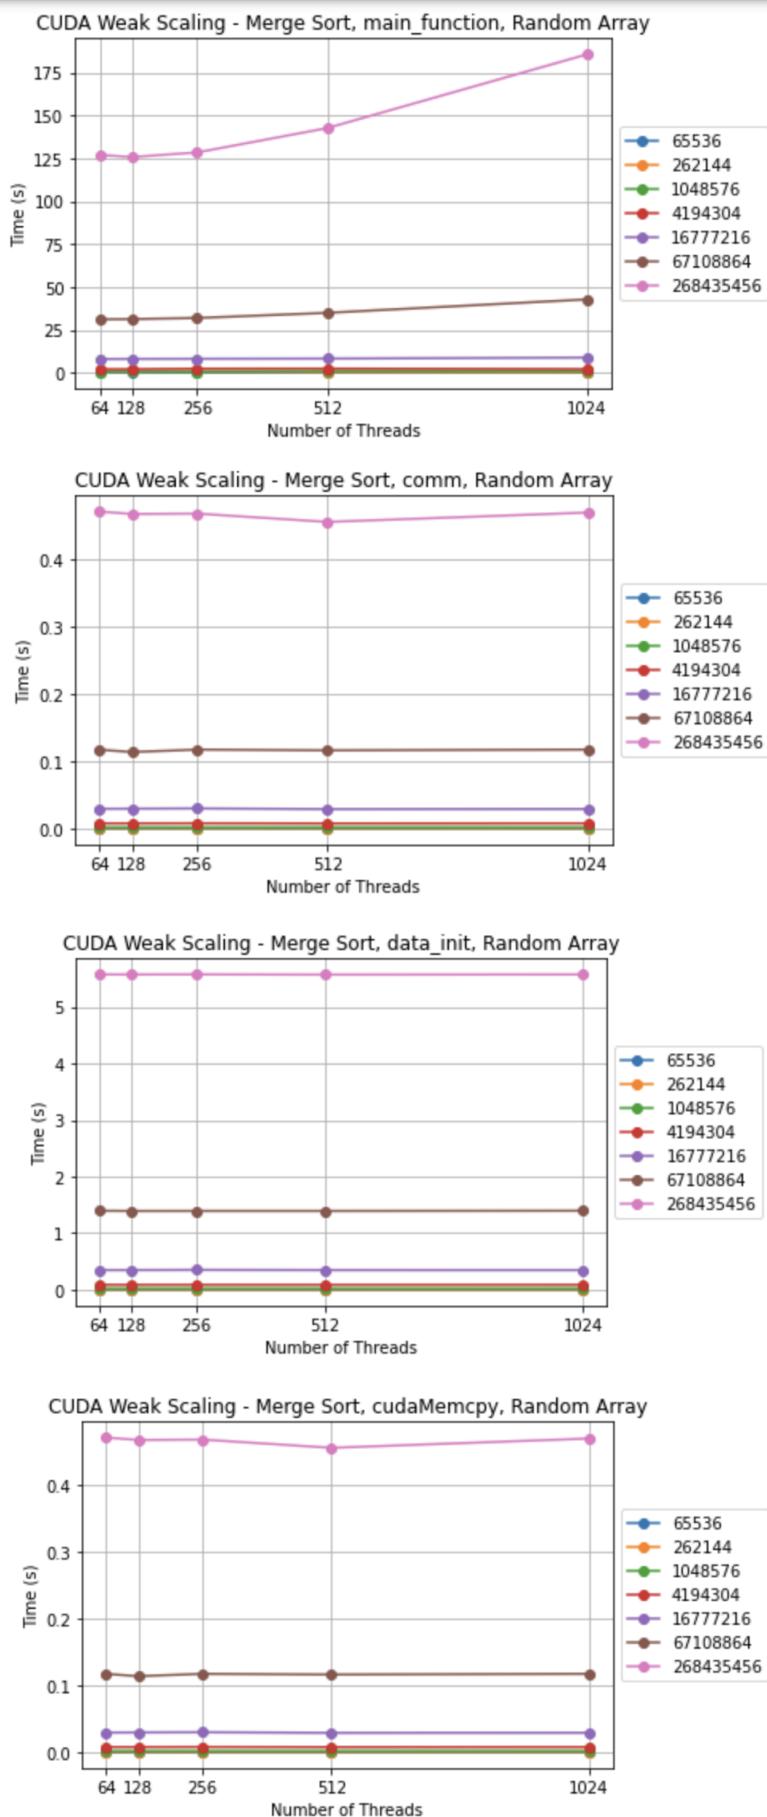
MPI Weak Scaling - Merge Sort, MPI_Gather, 1% Perturbed Array

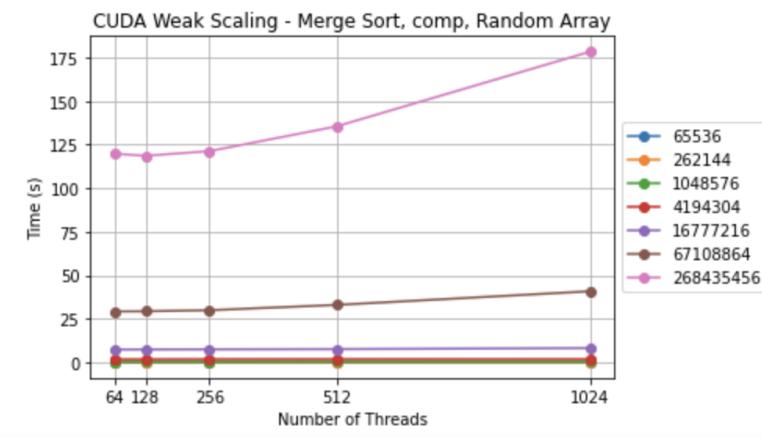
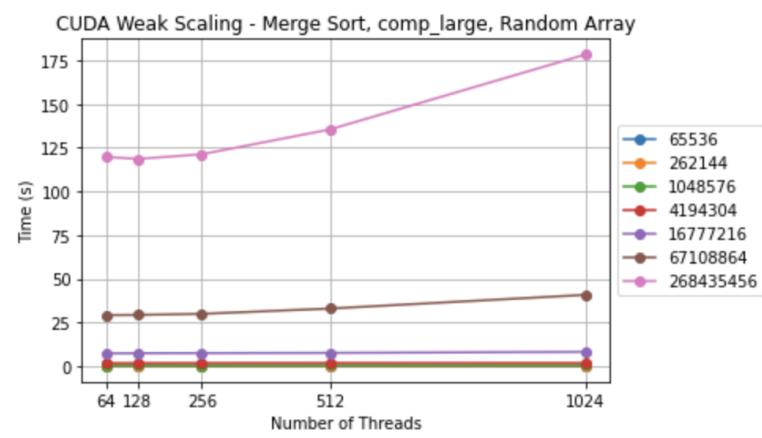
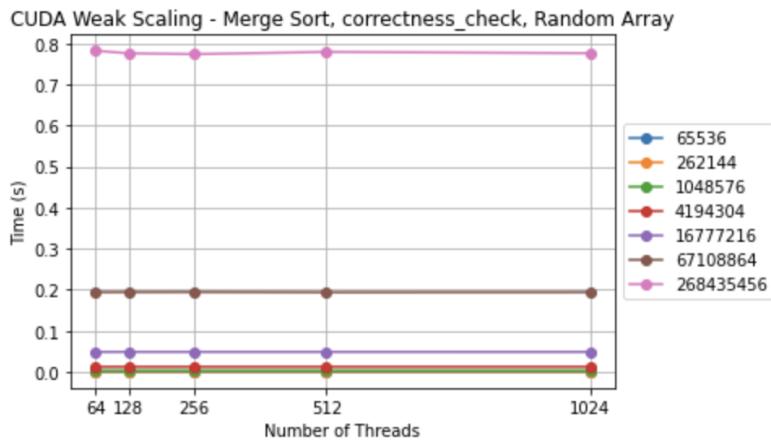
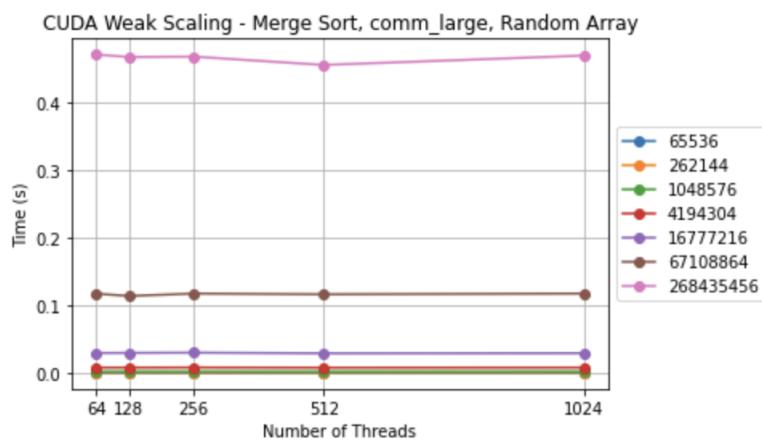


CUDA

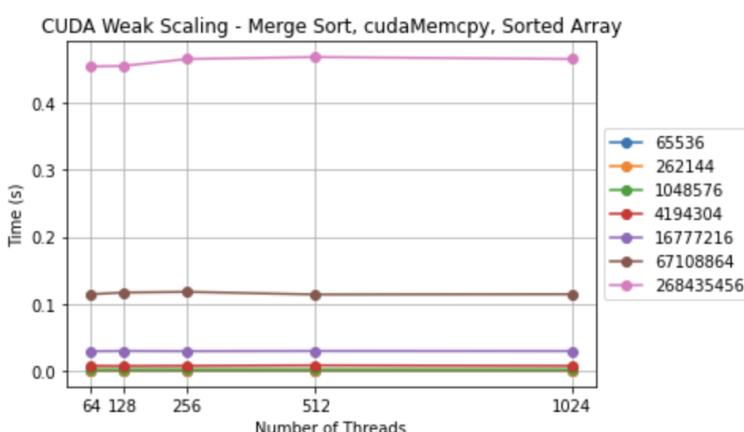
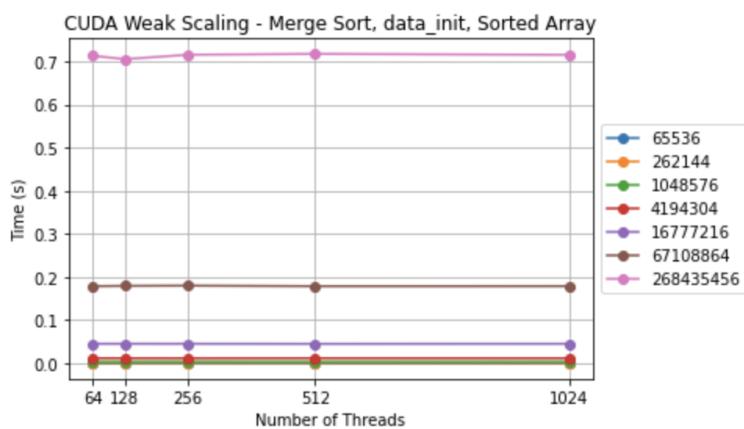
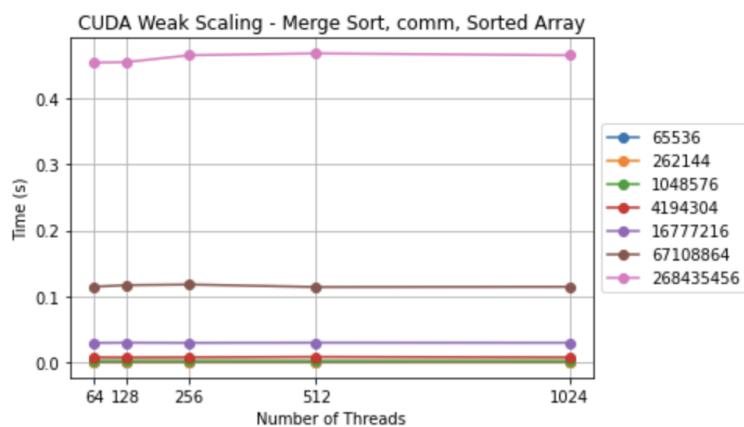
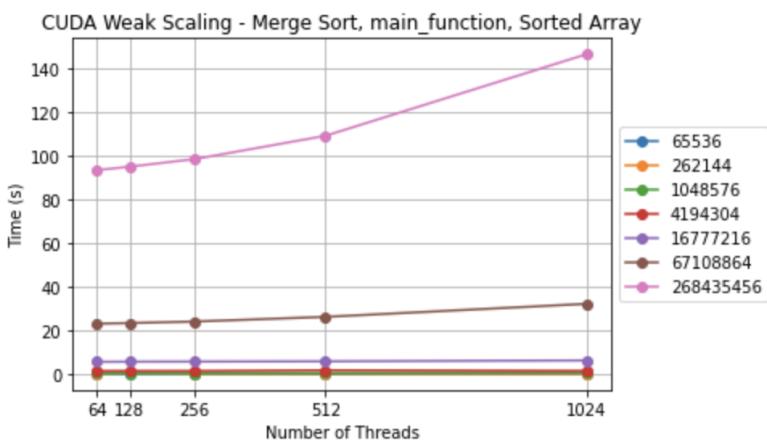
For the CUDA weak scaling of the different input types, we had a lot of the same activity as the MPI implementation. As we increased the number of elements in the input array, there was more time needed for the merge sort algorithm to complete. It also seems that the CUDA implementation takes a bit longer to complete than the MPI implementation. We can see this most notably with the 2^{28} th line (the pink), where it takes over 120 seconds to complete the main_function, whereas the MPI implementation had a maximum of 40 seconds to complete. We also see a lot steadier increase in time for the CUDA implementation. It is interesting to see that all of the different performance regions increase quite steadily. Another big difference with CUDA is how it handles communication. Rather than speaking to multiple processors like in MPI, CUDA handles only cudaMemcpy, so we see rather constant communication times between all of the different input sizes.

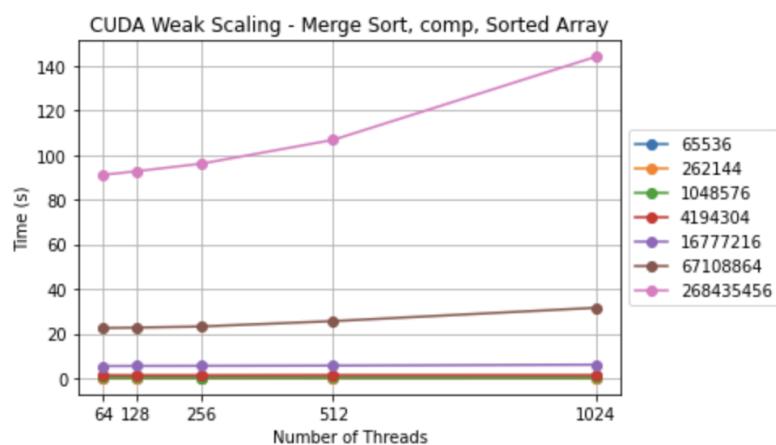
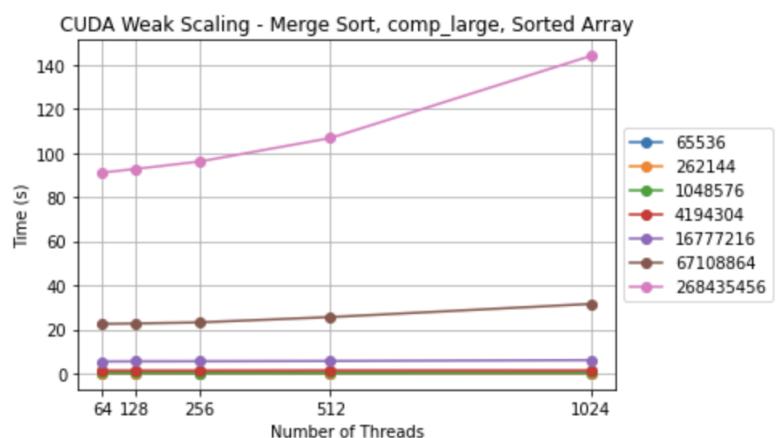
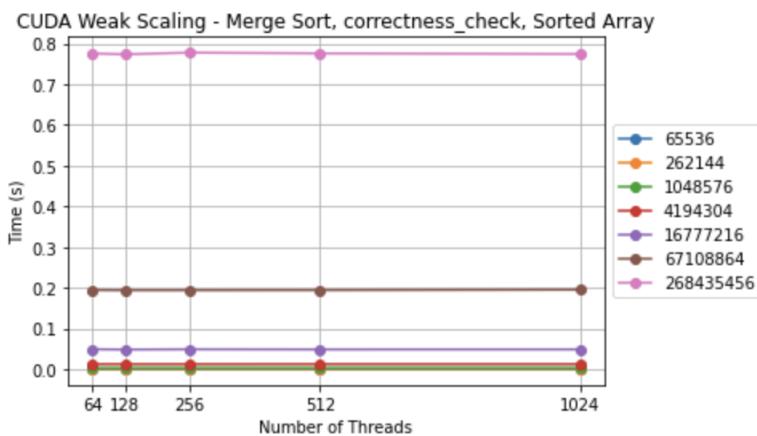
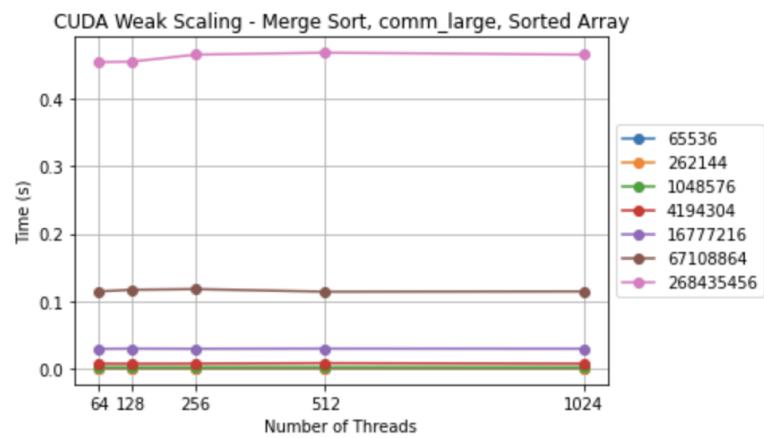
Random Array





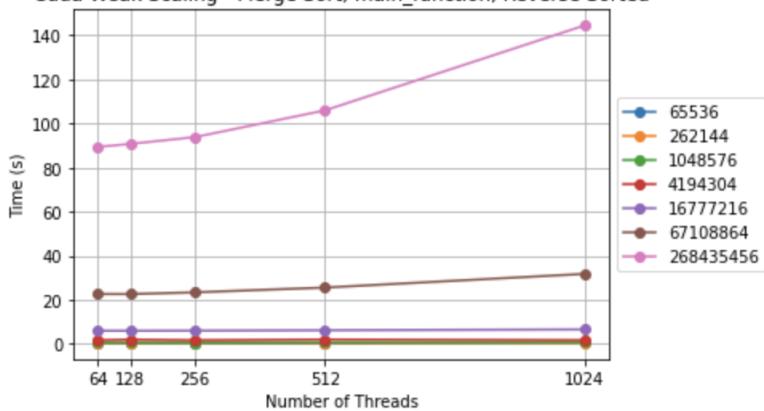
Sorted Array



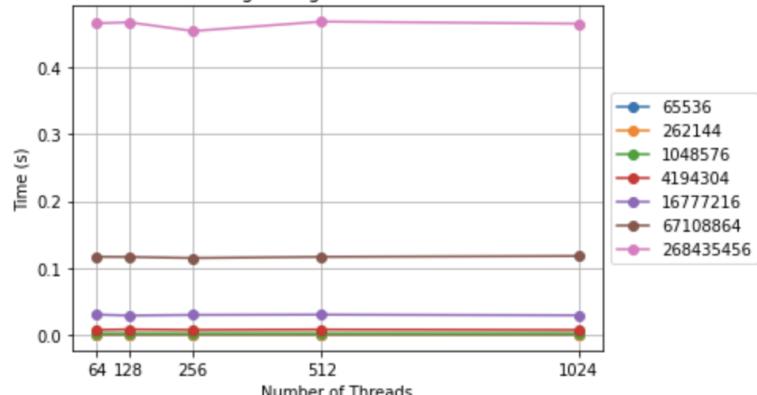


Reverse Sorted Array

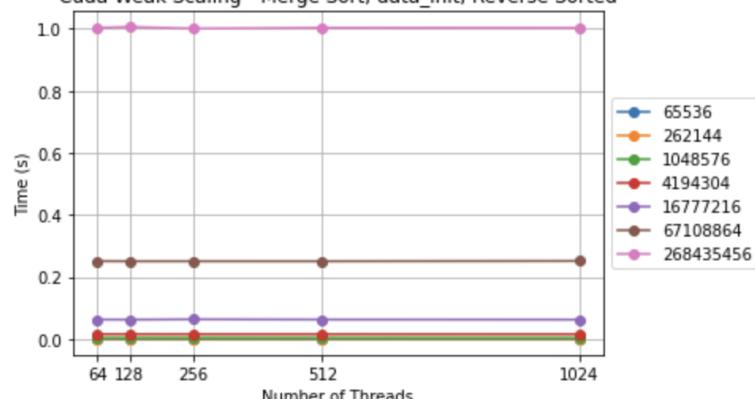
Cuda Weak Scaling - Merge Sort, main_function, Reverse Sorted



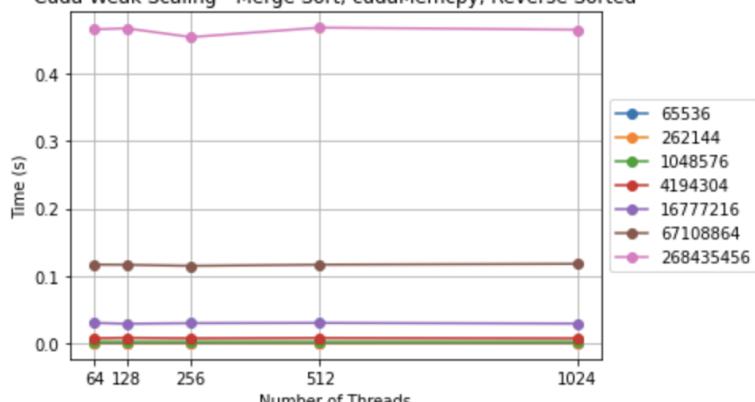
Cuda Weak Scaling - Merge Sort, comm, Reverse Sorted

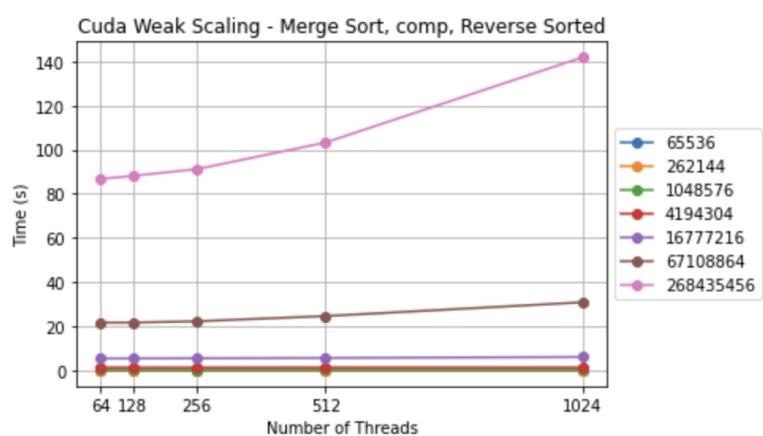
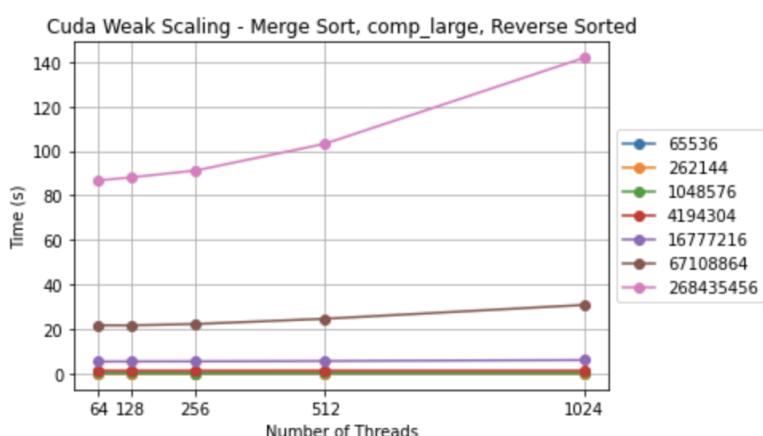
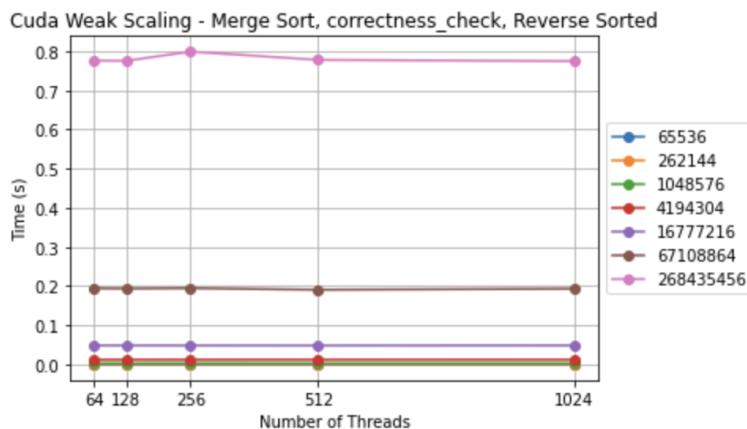
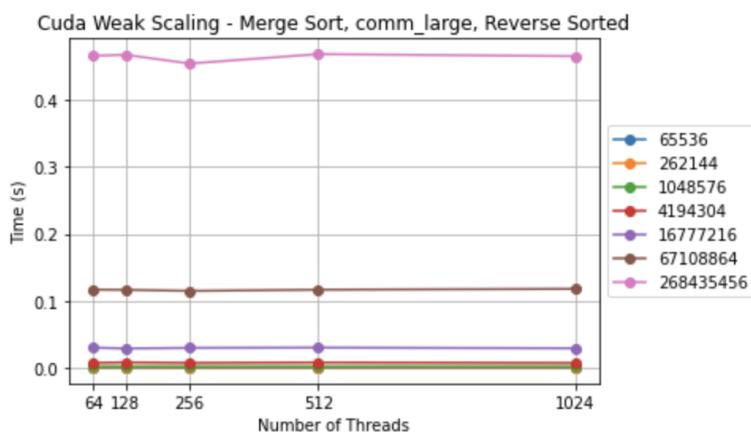


Cuda Weak Scaling - Merge Sort, data_init, Reverse Sorted



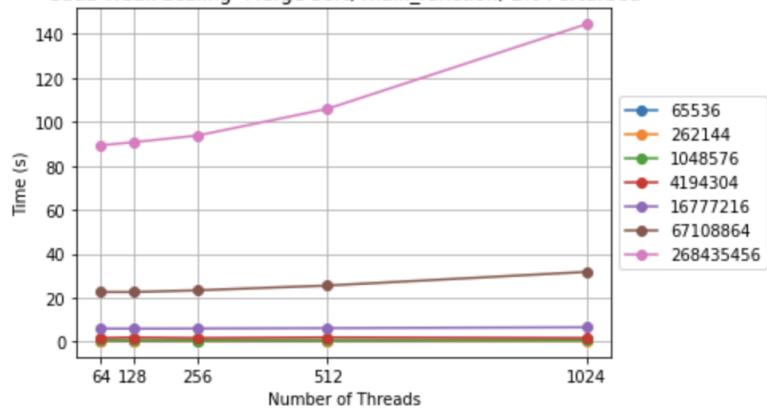
Cuda Weak Scaling - Merge Sort, cudaMemcpy, Reverse Sorted



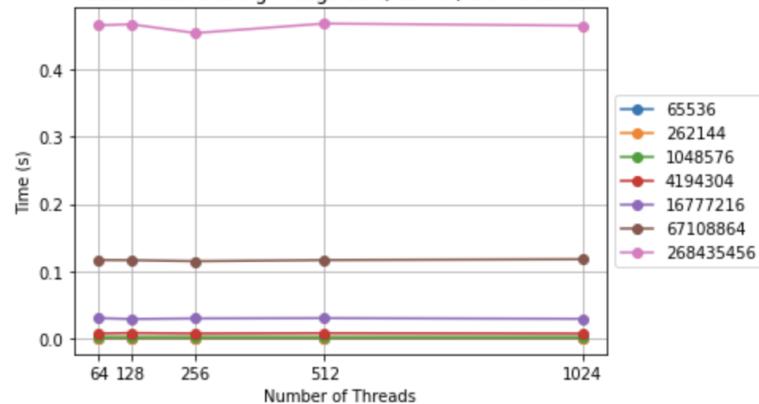


1% Perturbed Array

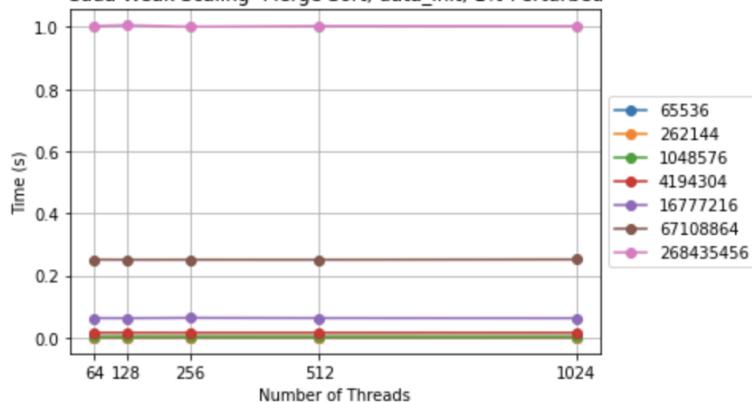
Cuda Weak Scaling- Merge Sort, main_function, 1% Perturbed



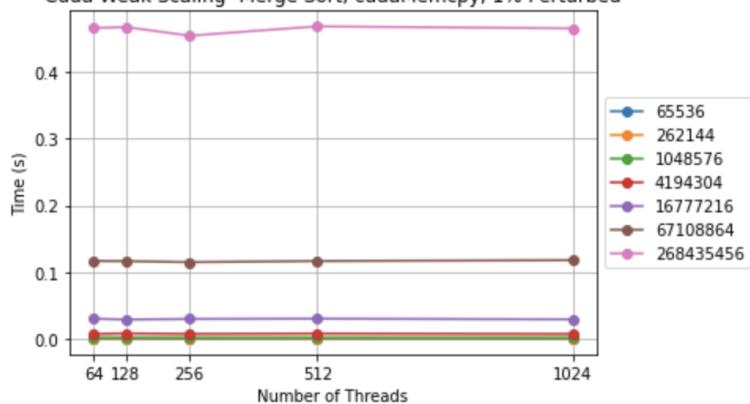
Cuda Weak Scaling- Merge Sort, comm, 1% Perturbed



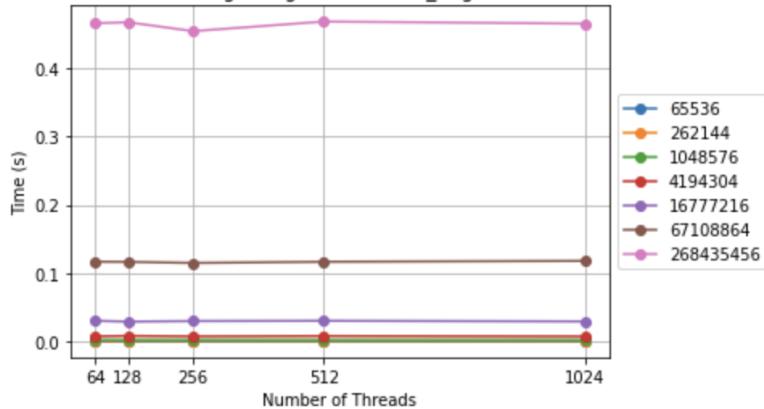
Cuda Weak Scaling- Merge Sort, data_init, 1% Perturbed



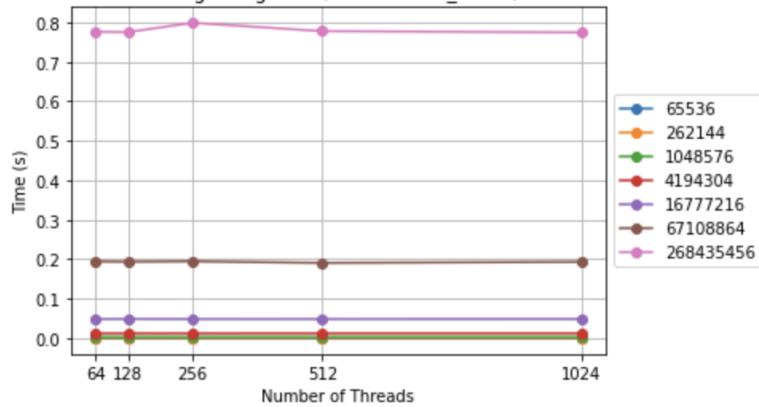
Cuda Weak Scaling- Merge Sort, cudaMemcpy, 1% Perturbed



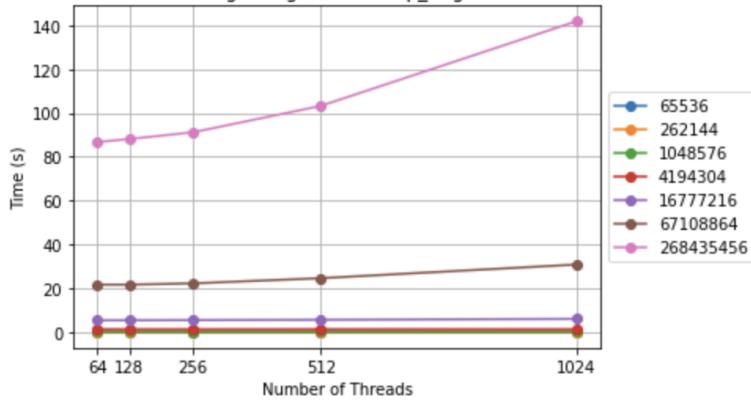
Cuda Weak Scaling- Merge Sort, comm_large, 1% Perturbed



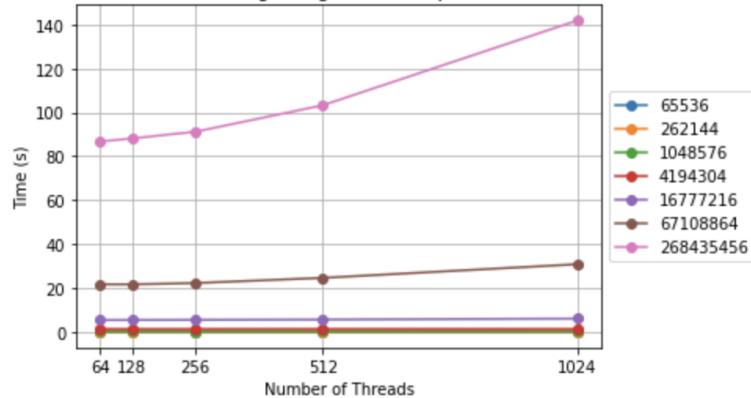
Cuda Weak Scaling- Merge Sort, correctness_check, 1% Perturbed



Cuda Weak Scaling- Merge Sort, comp_large, 1% Perturbed



Cuda Weak Scaling- Merge Sort, comp, 1% Perturbed

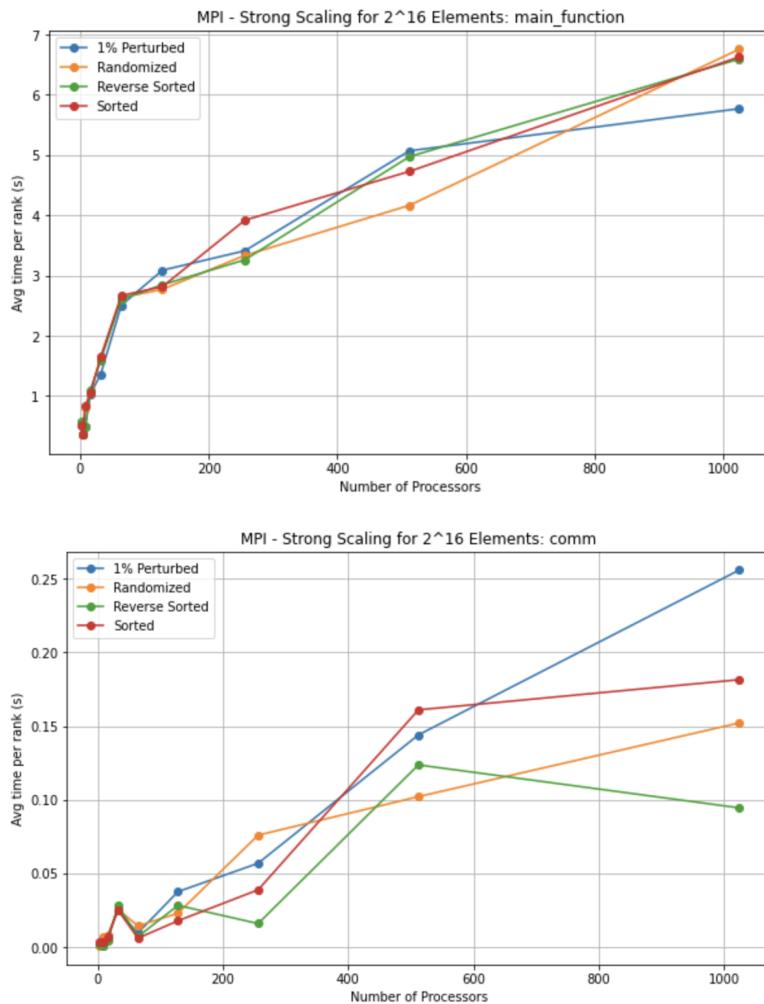


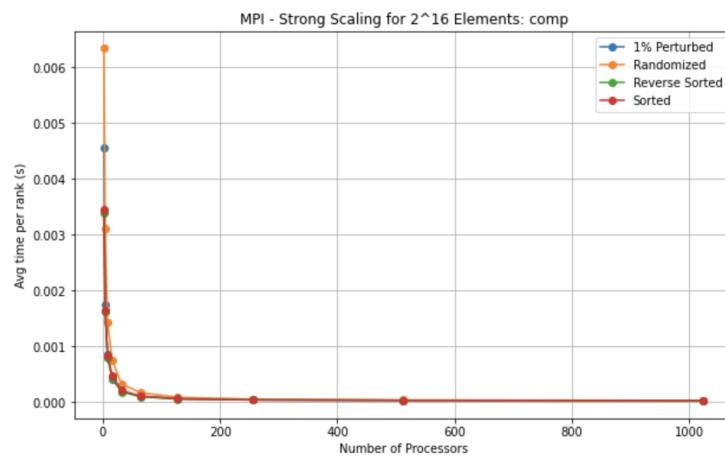
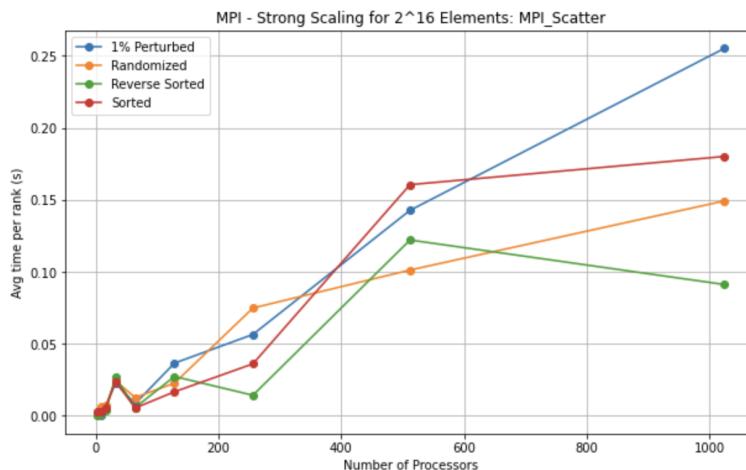
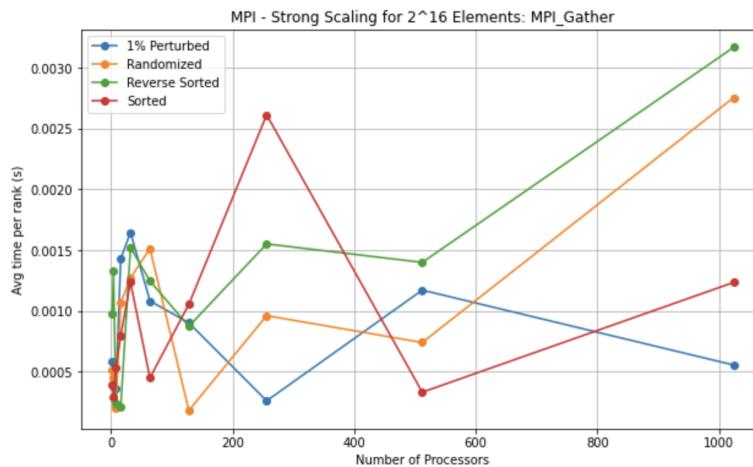
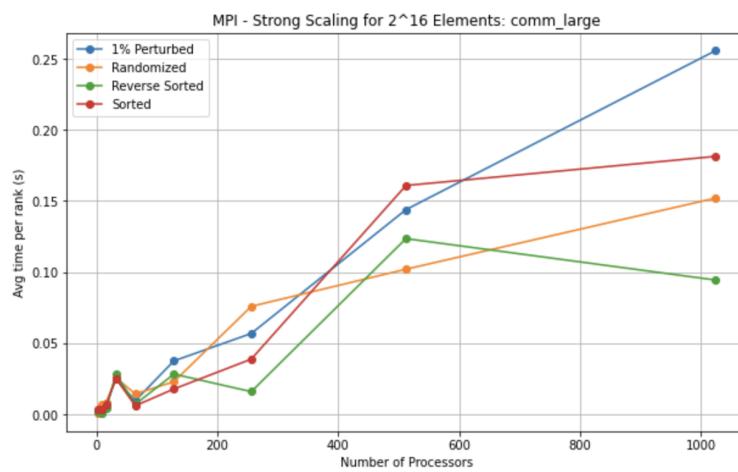
Strong Scaling

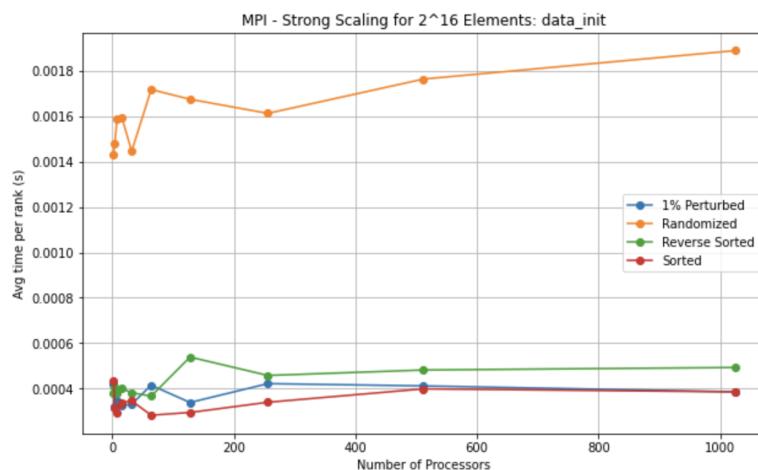
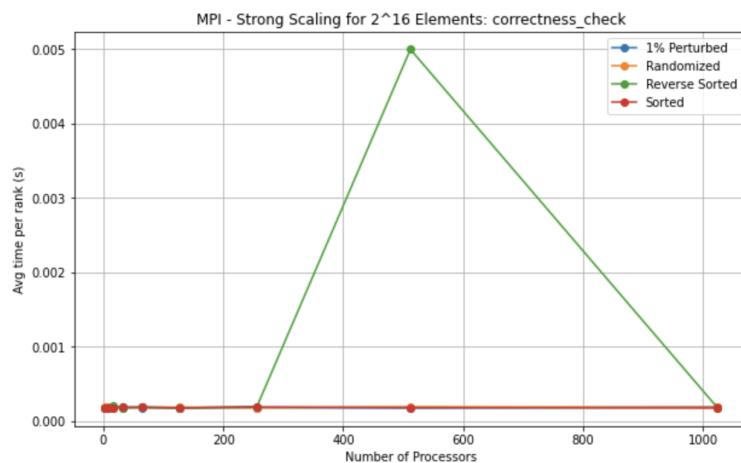
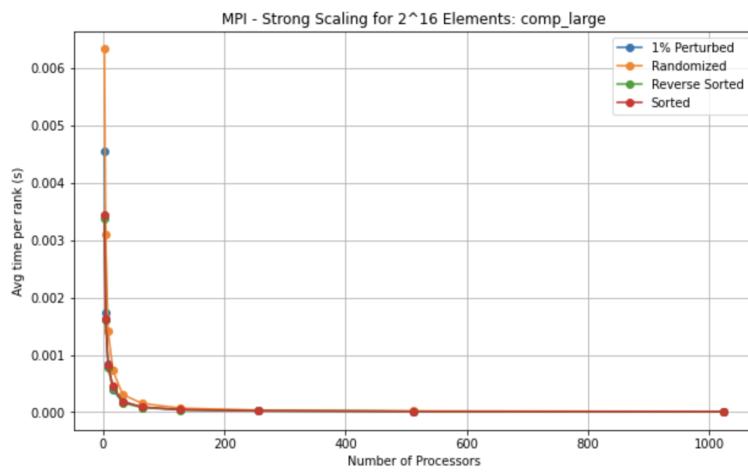
MPI

We see some similar behavior as the weak scaling here in terms of shape of the graph. The most notable one here is the comp and comp_large regions, which have the same exponentially decreasing graph, where each input type has relatively the same behavior. This is more or less expected given the nature of merge sort, where each input is eventually recursively broken down into single input subarrays. Especially considering that each graph has the same number of elements, we can understand why the computation is so much closer between the lines for this strong scaling. It is interesting because this strong scaling shows that MPI_Scatter is much slower than MPI_Gather, which is the opposite than from the weak scaling. We also see that as we increase the number of elements, our behavior becomes more erratic, which is most easily seen in the main_function graph of each input size. We're also seeing that the time peaks around 128 processors before decreasing and then further increasing again. For communication, we are seeing some steady increase in time as well, but overall pretty fast across the ranks.

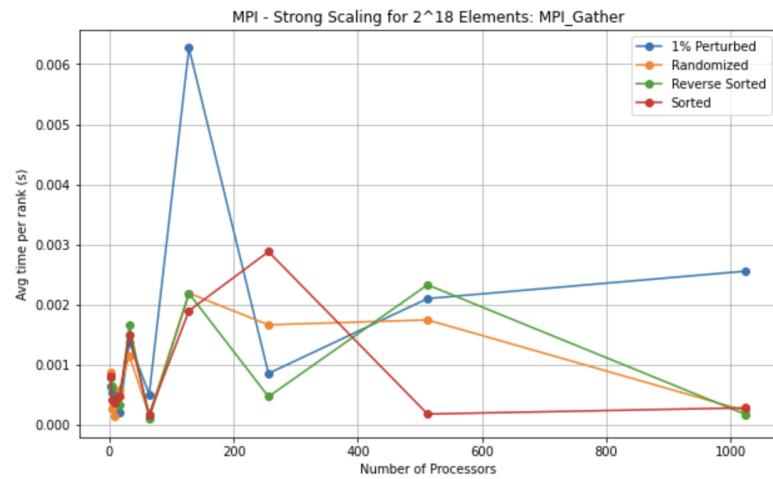
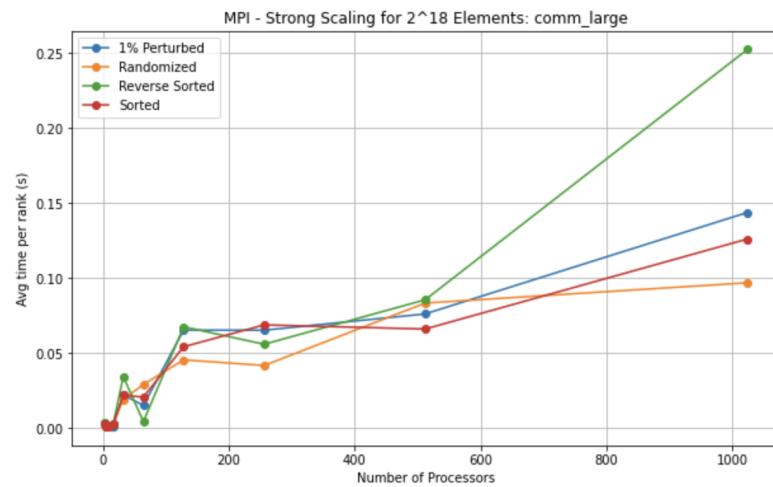
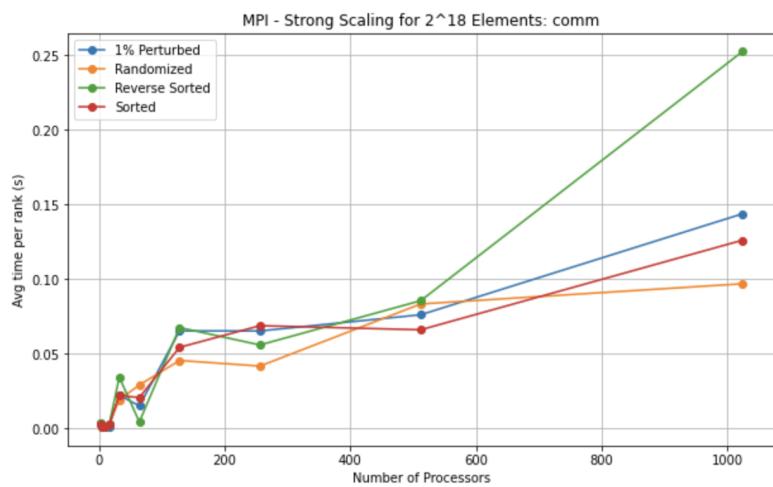
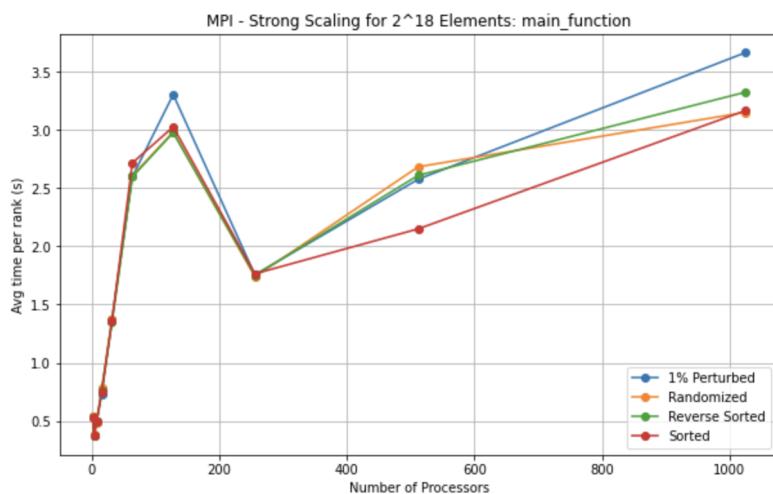
2^{16} Elements

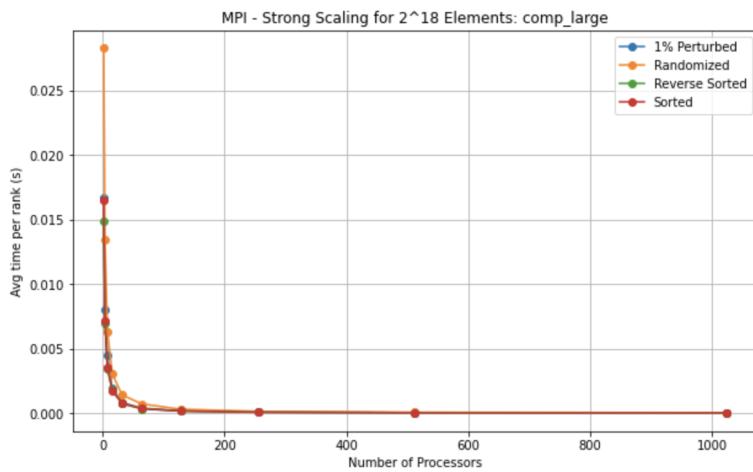
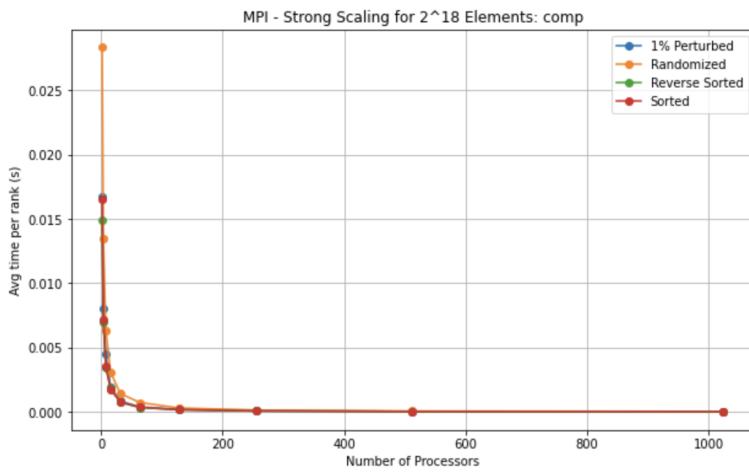
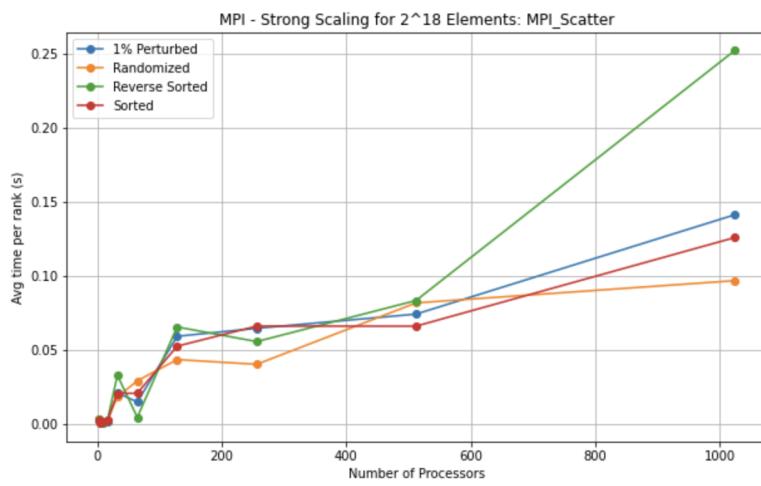


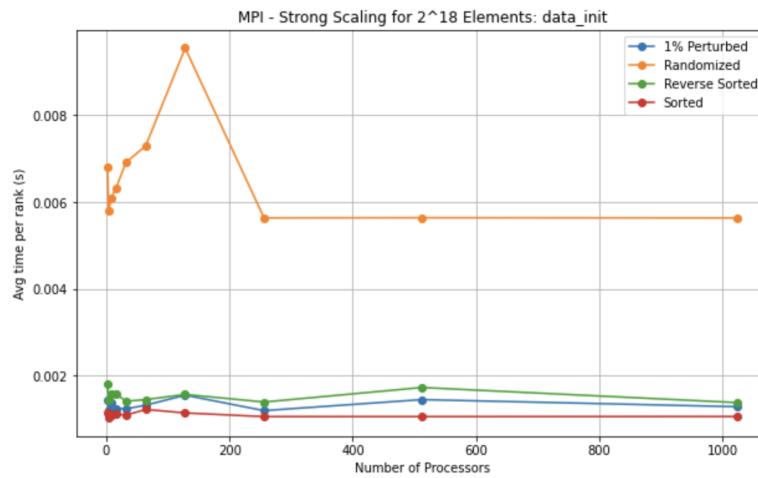
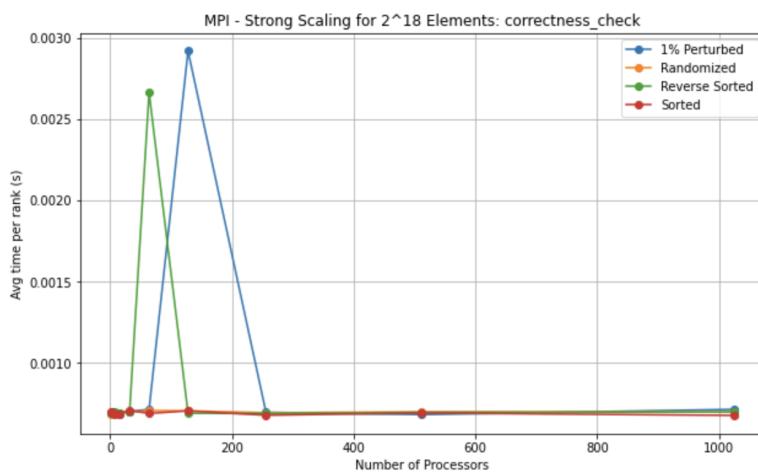




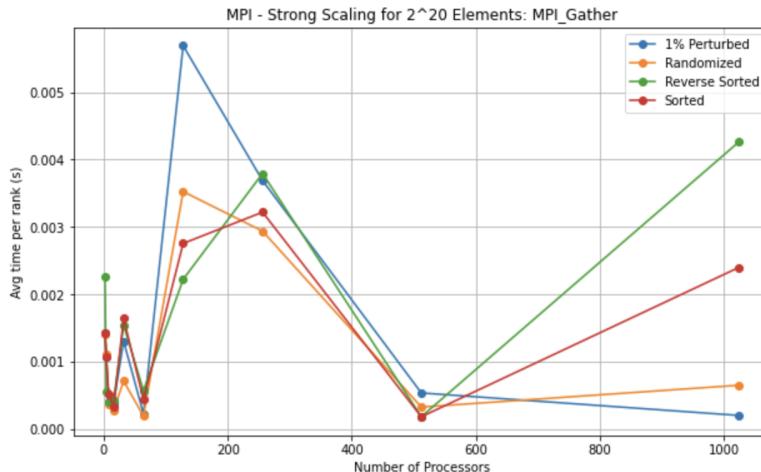
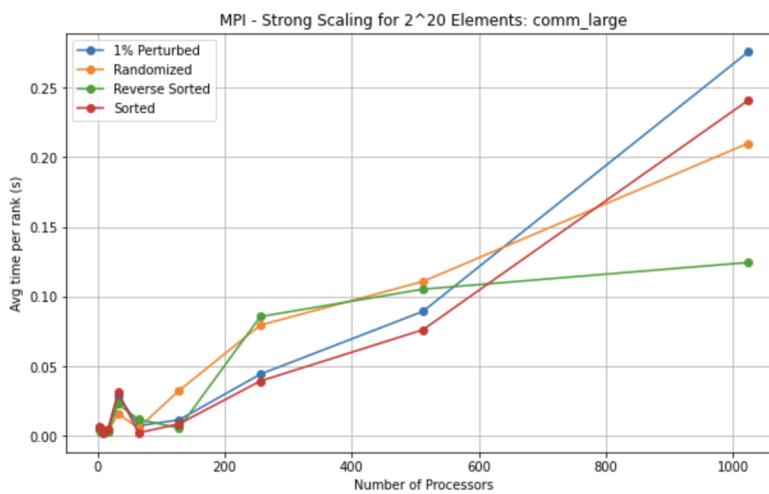
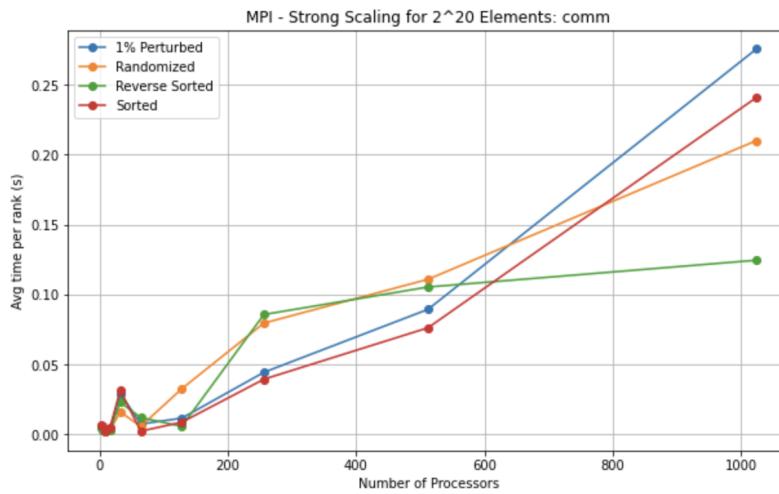
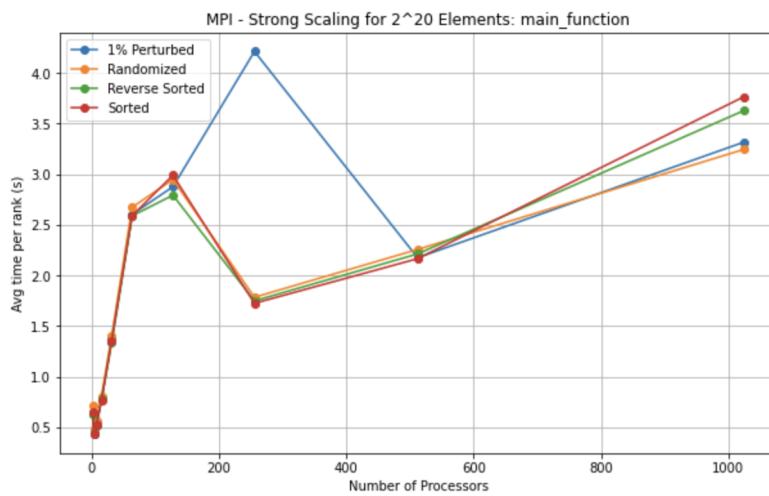
2 ^18 Elements

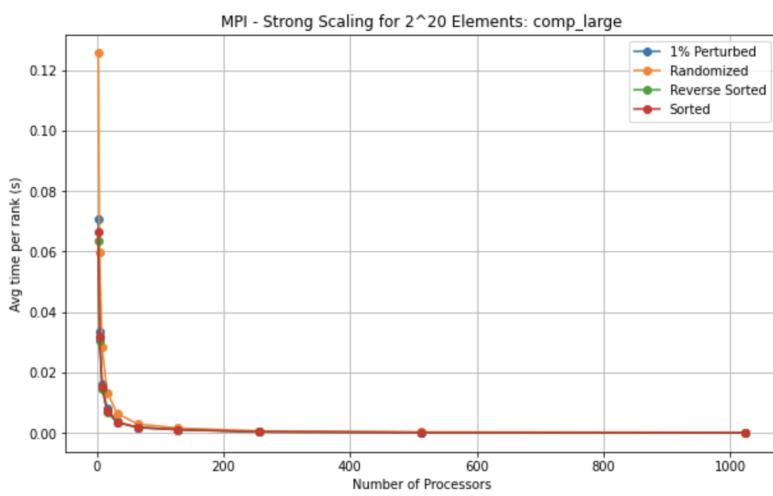
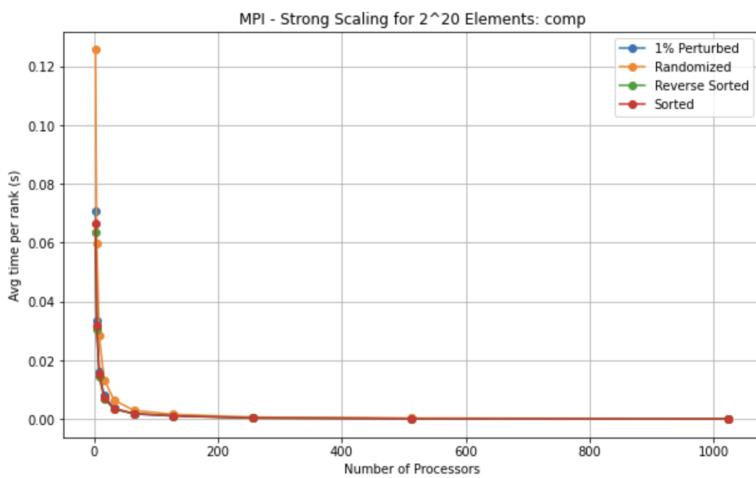
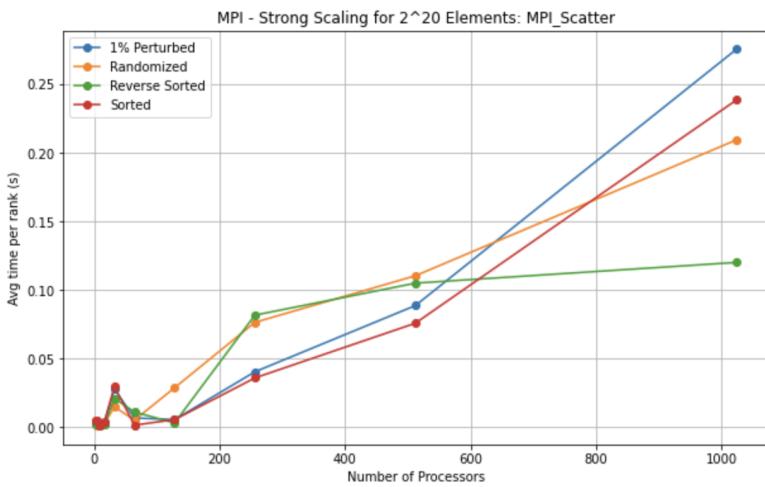


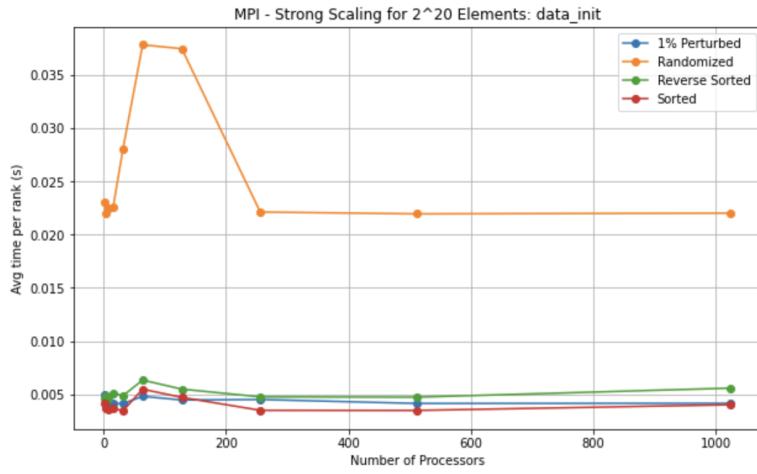
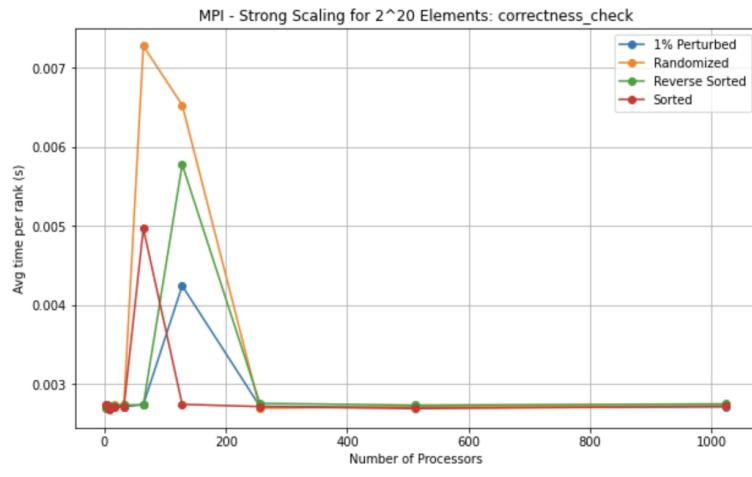




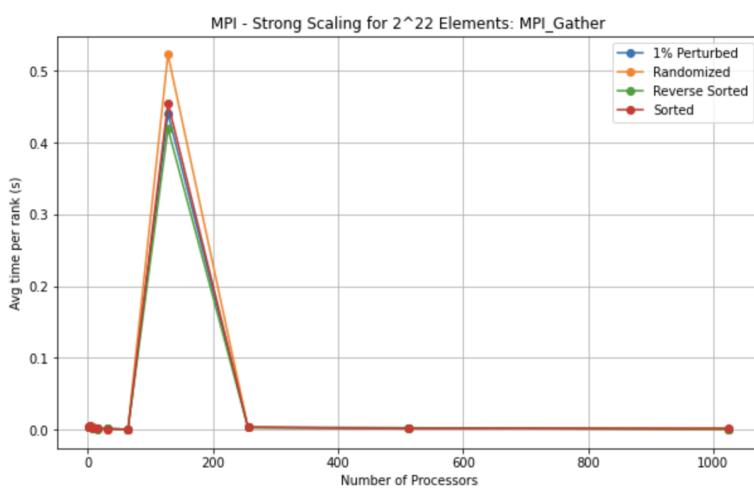
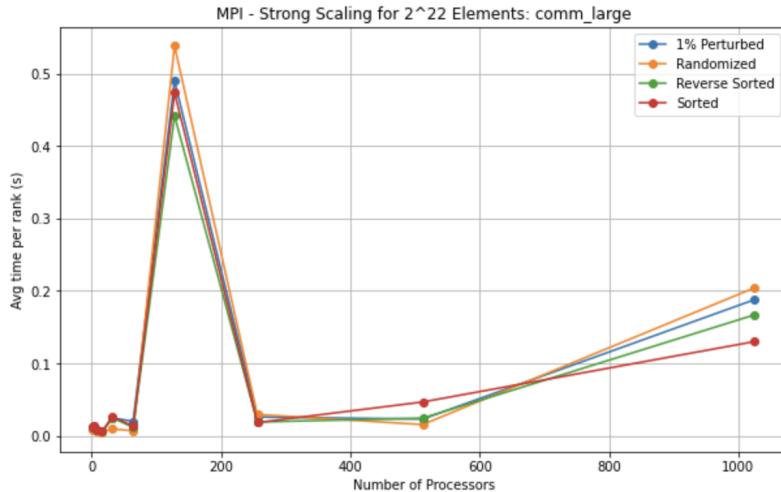
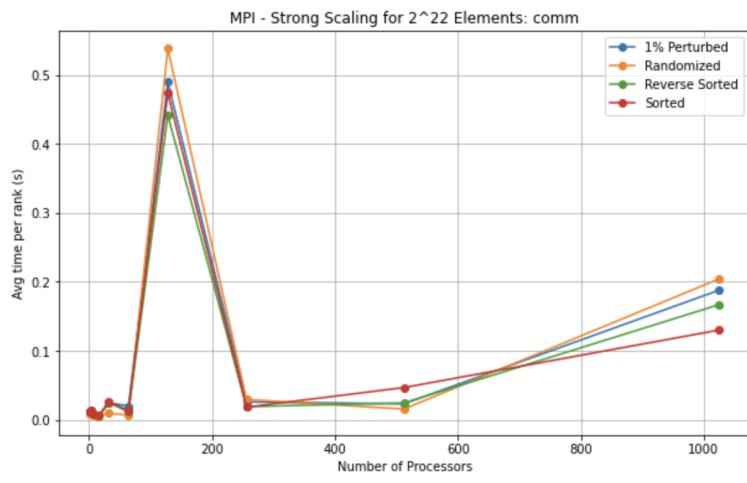
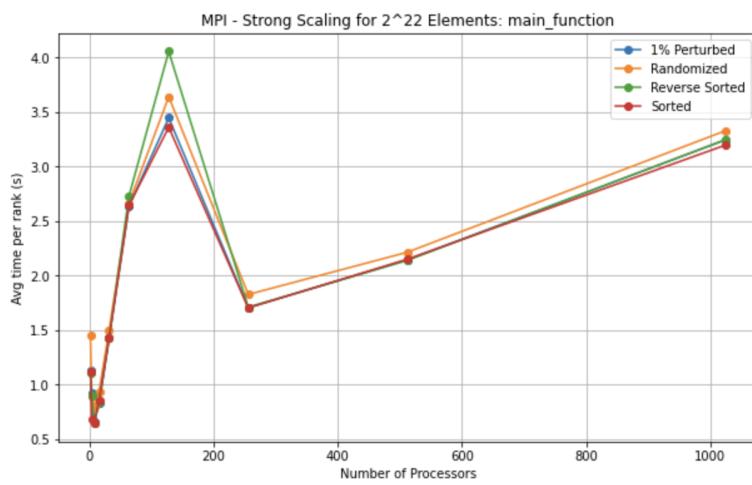
2²⁰ Elements

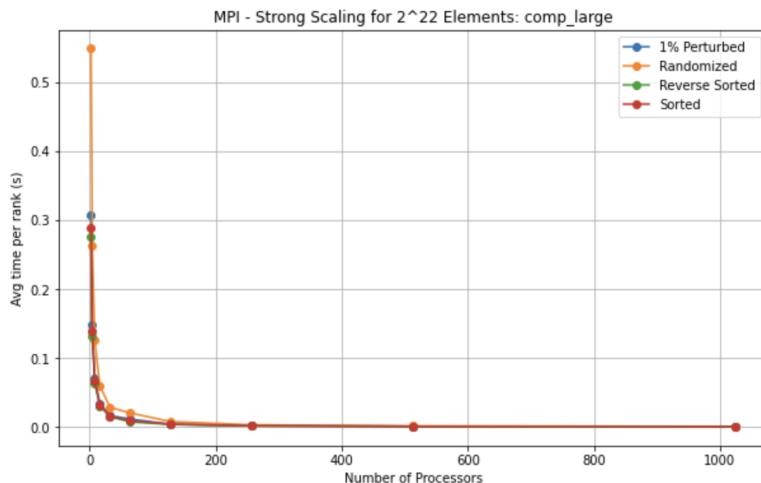
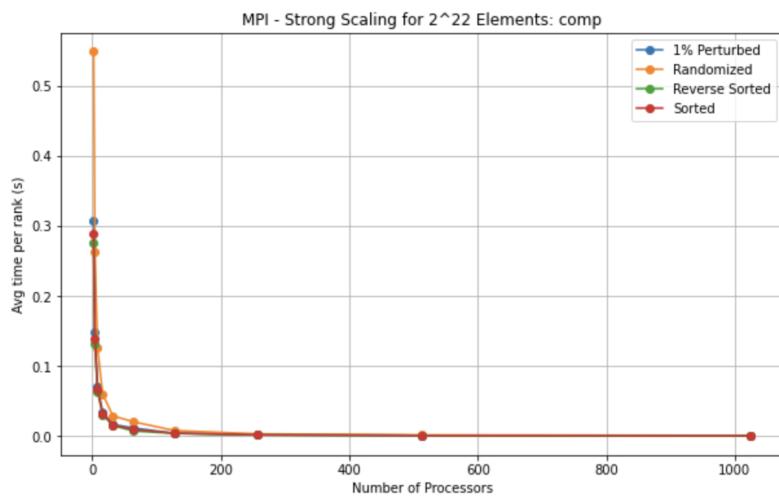
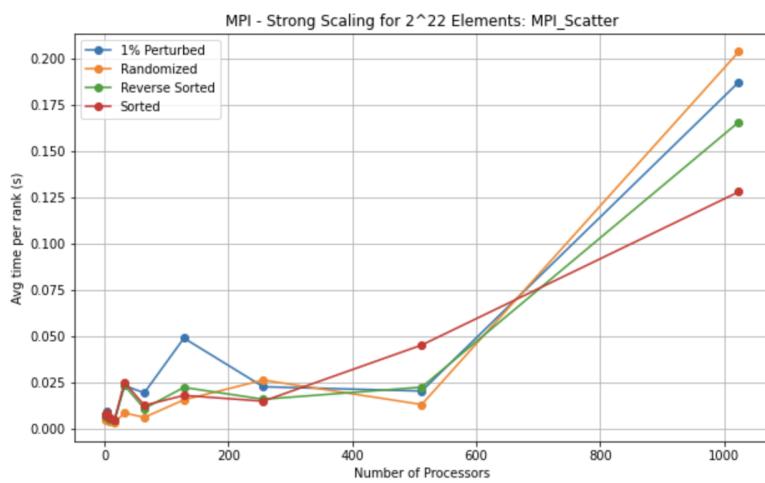


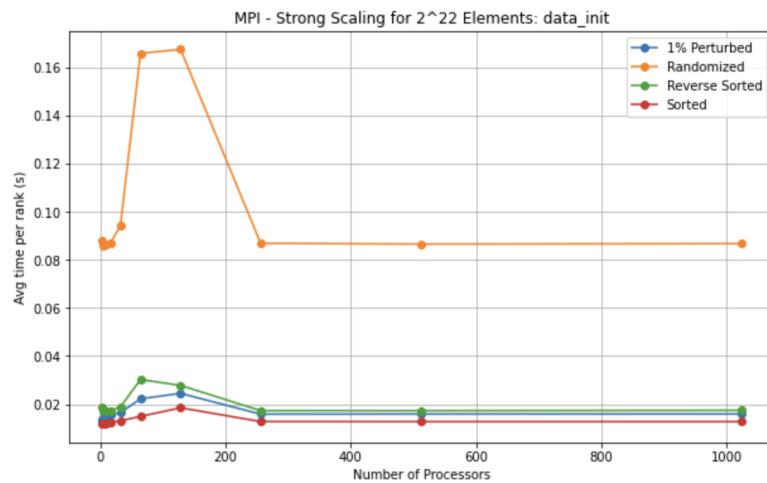
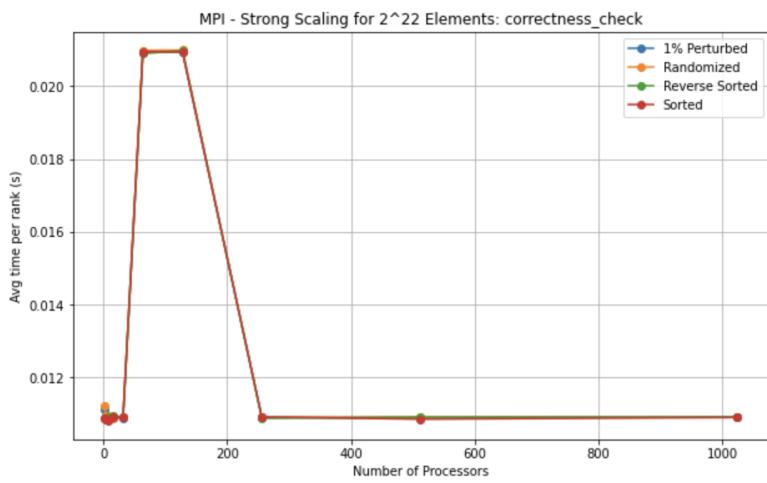




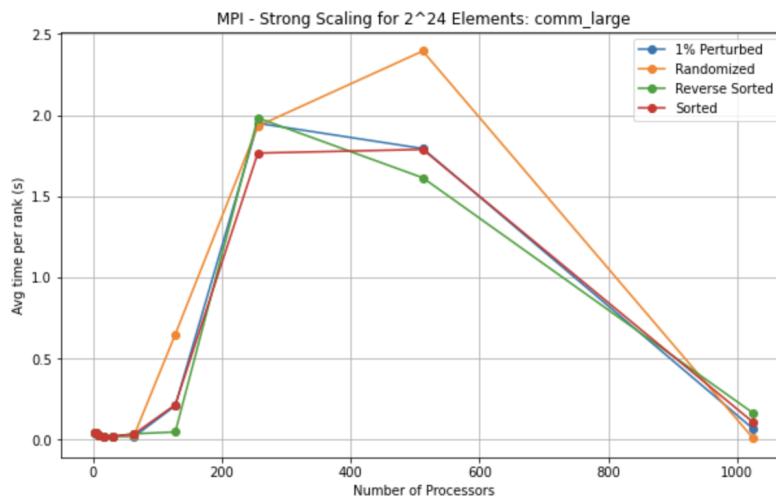
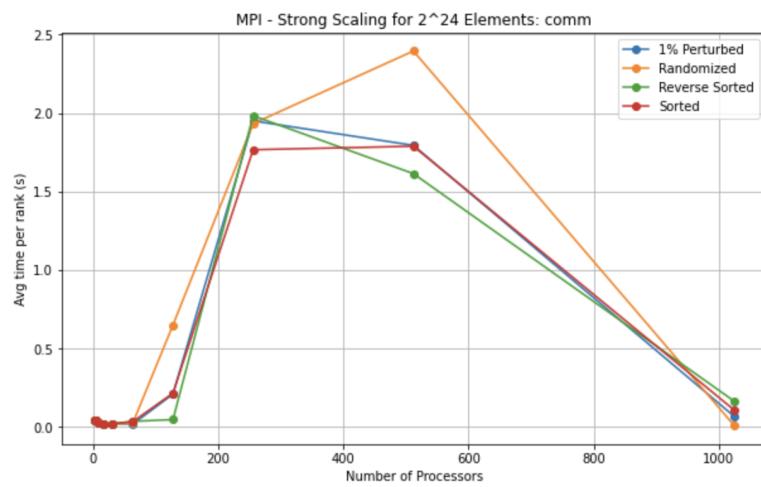
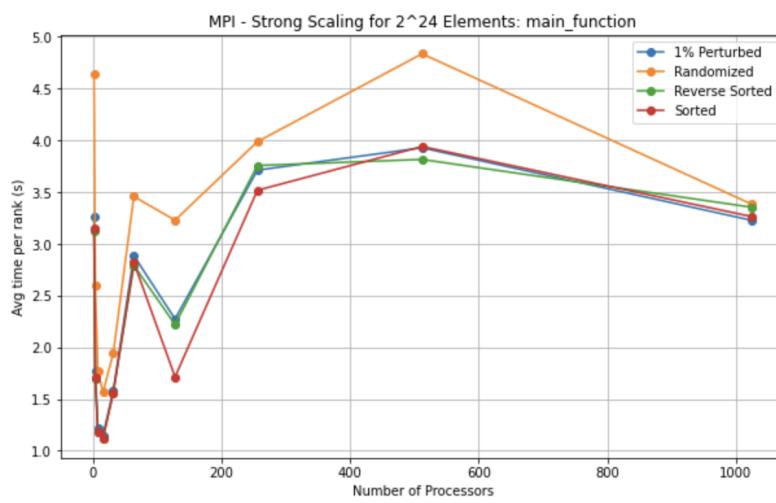
2²² Elements

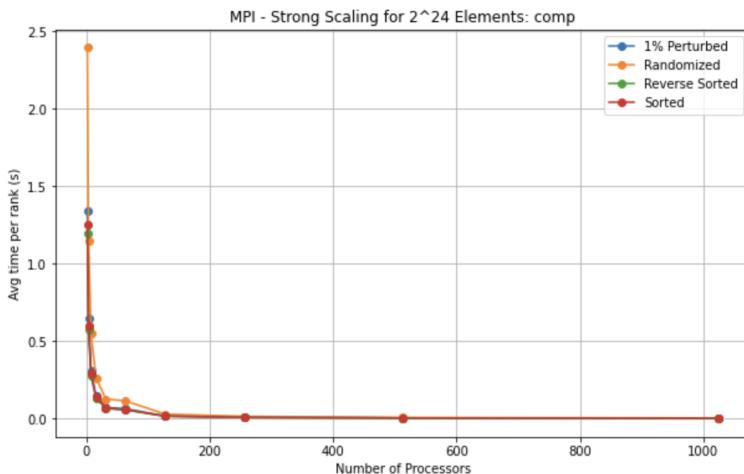
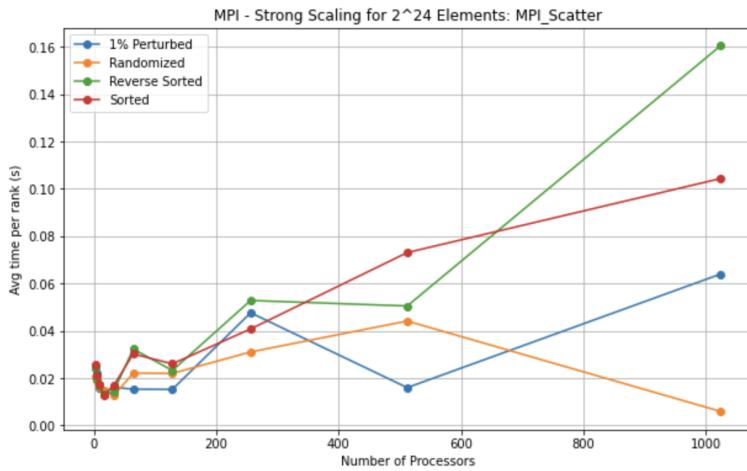
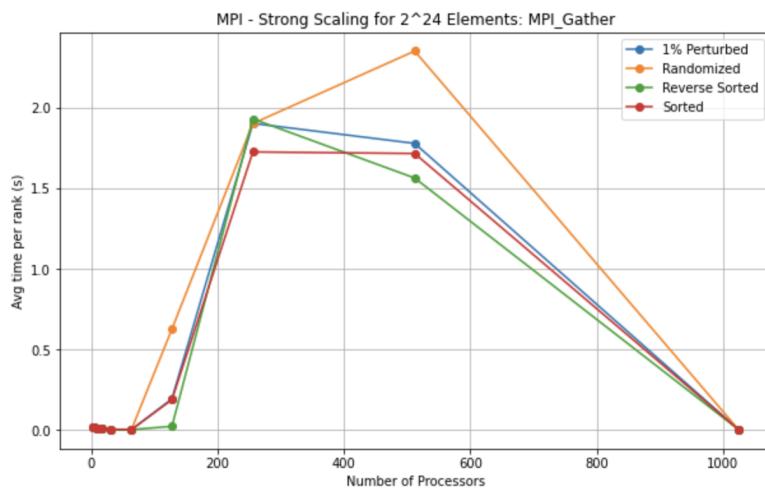


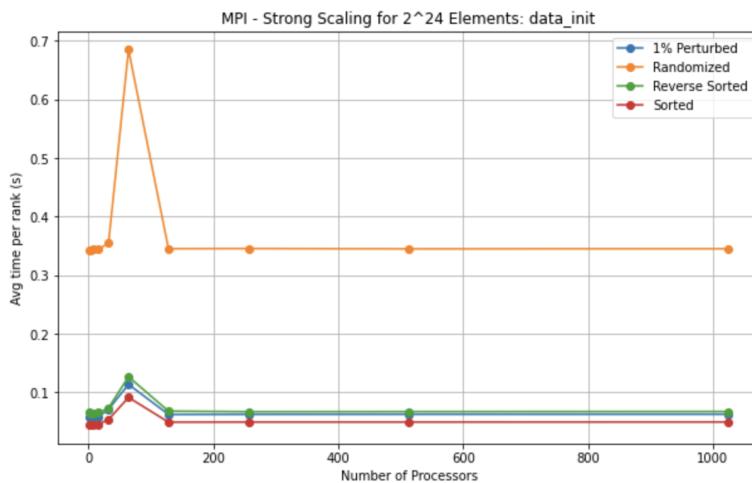
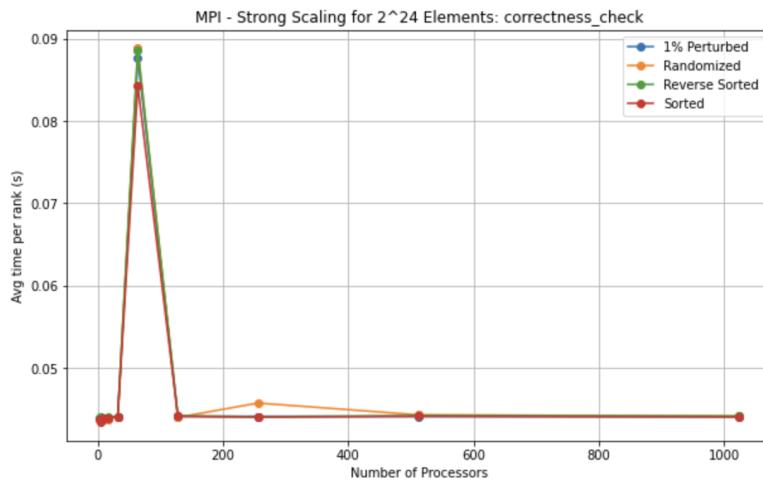
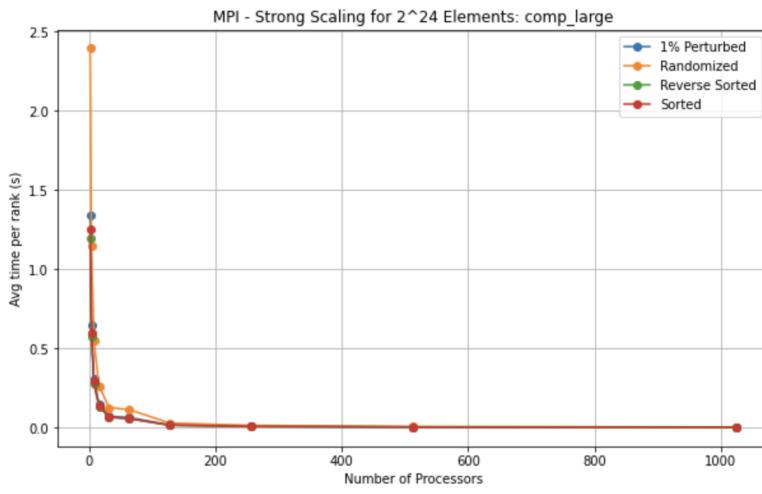




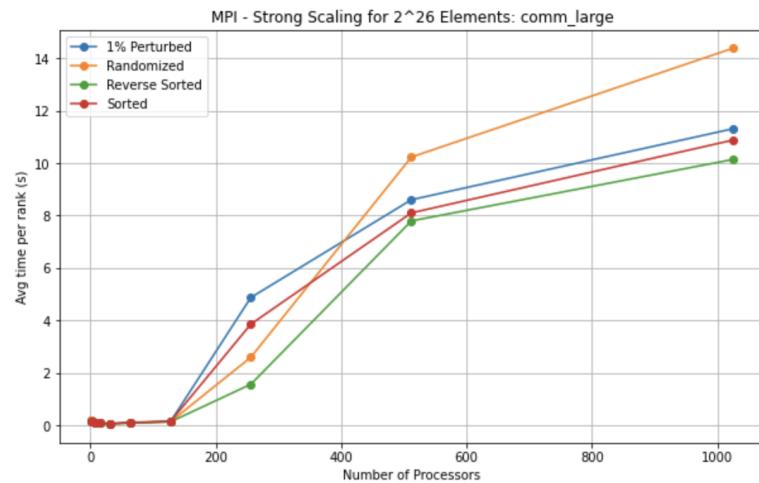
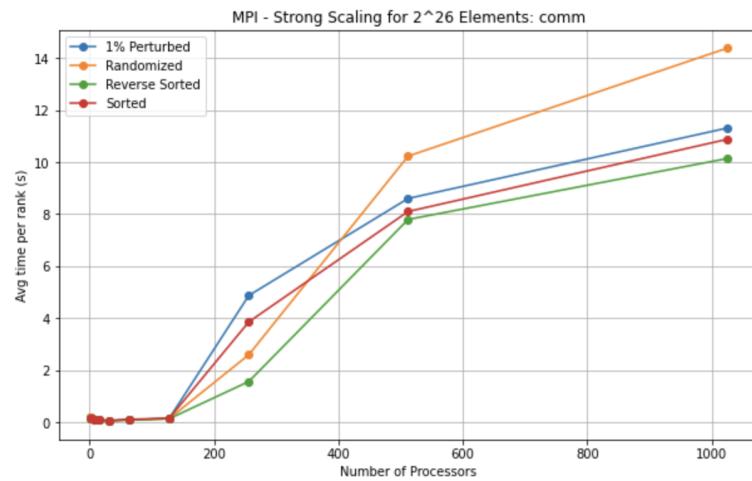
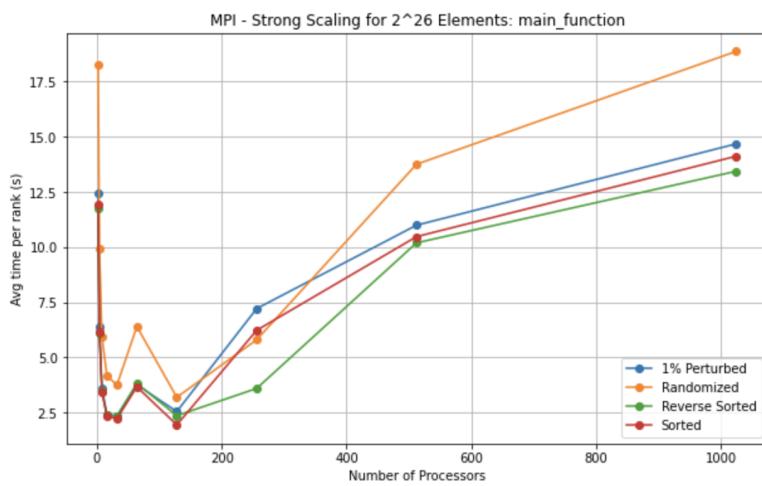
2²⁴ Elements

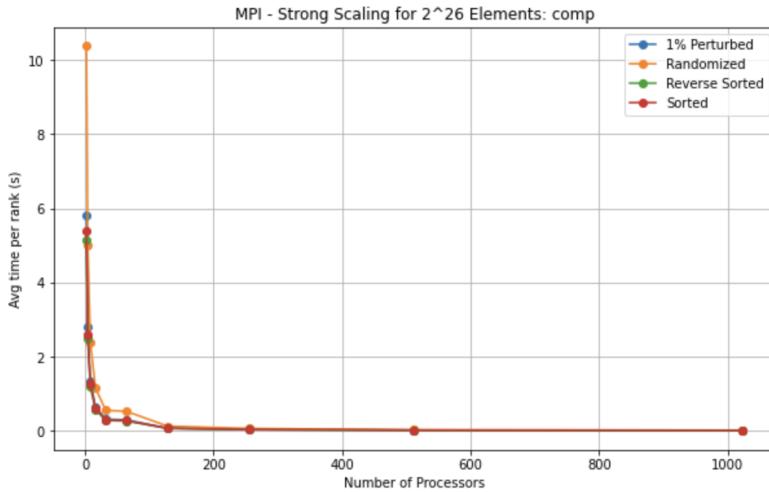
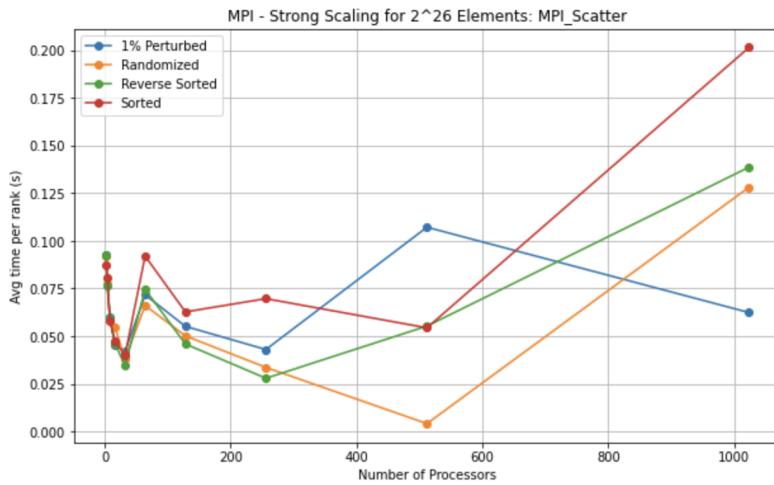
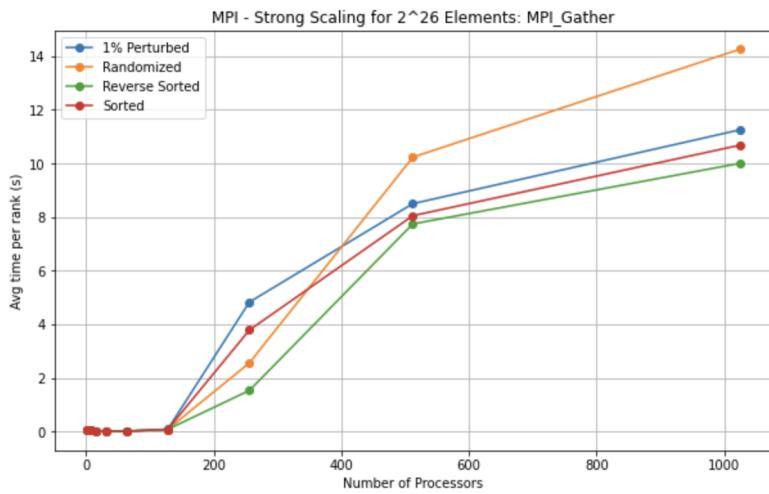


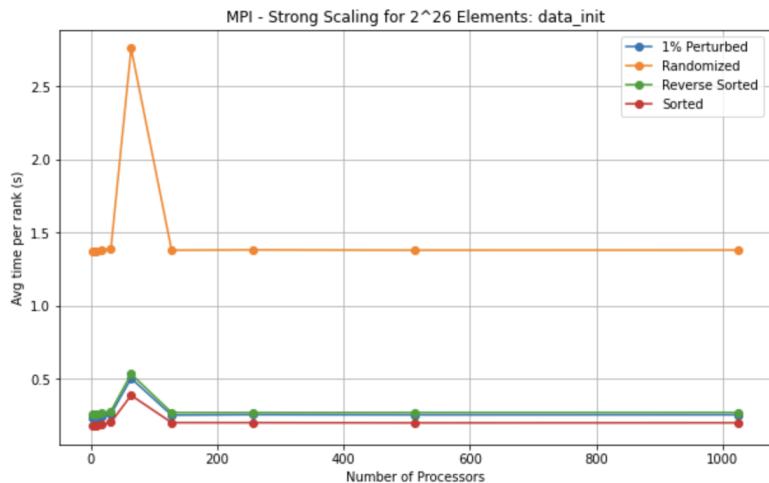
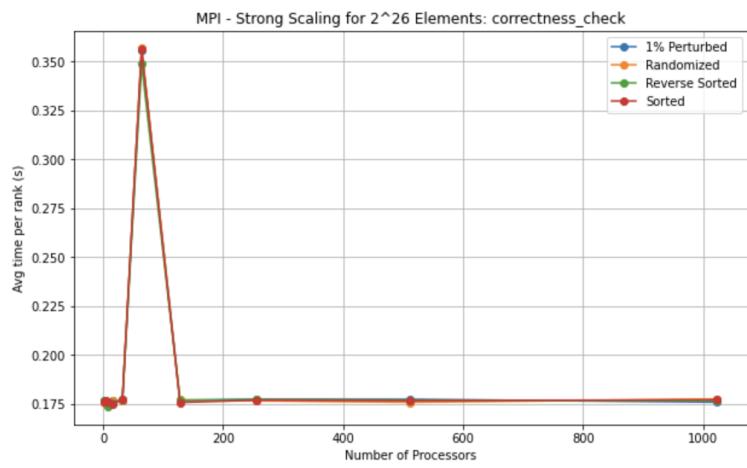
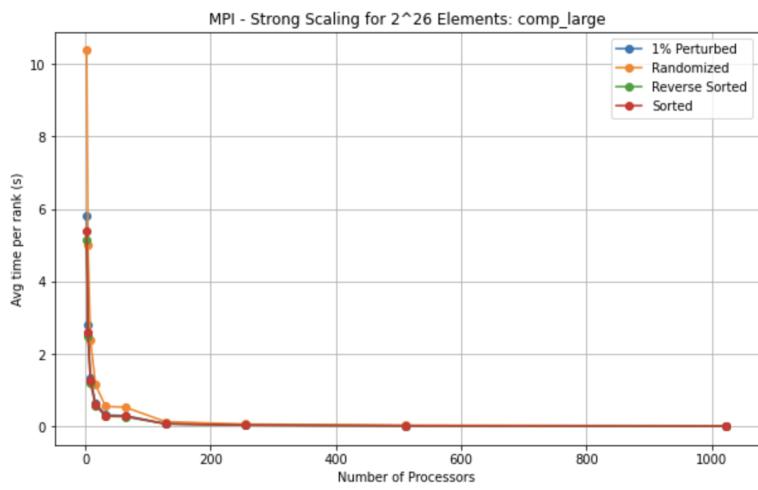




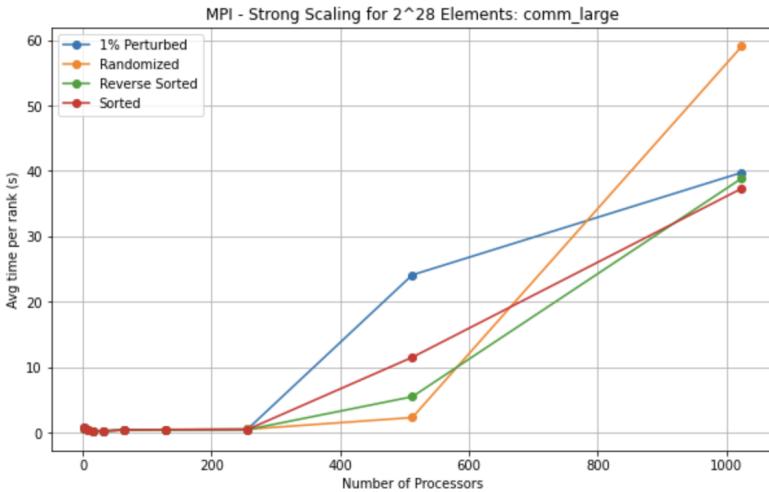
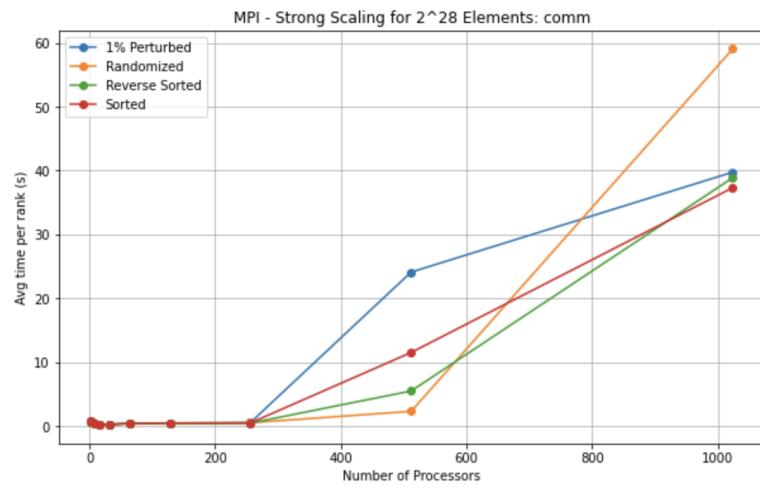
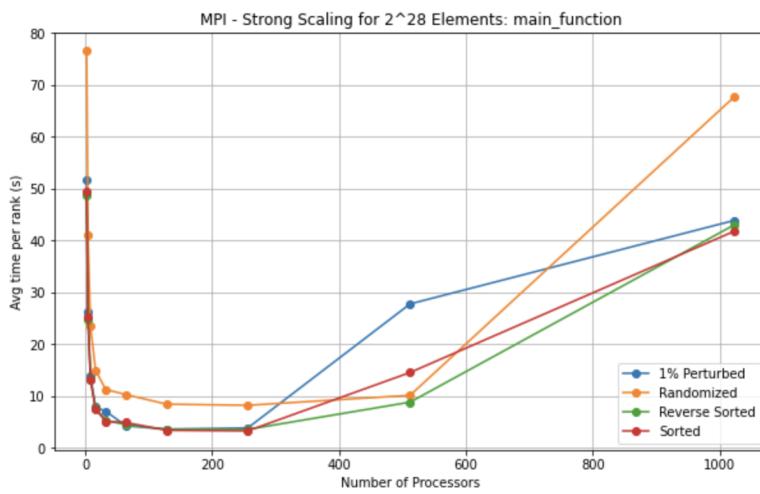
2²⁶ Elements

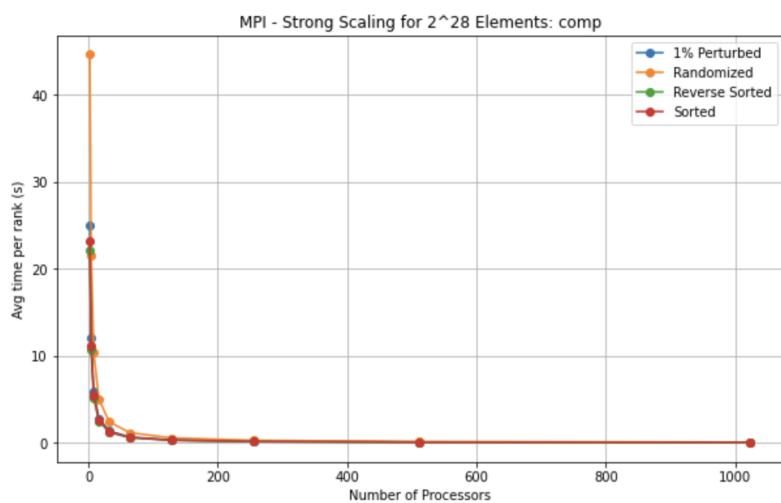
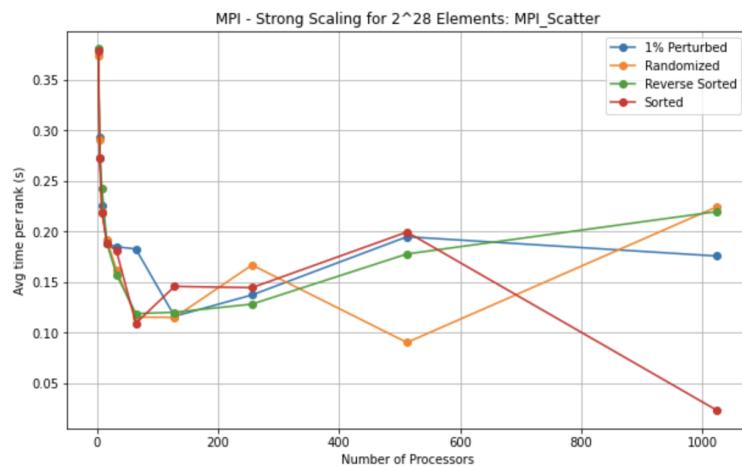
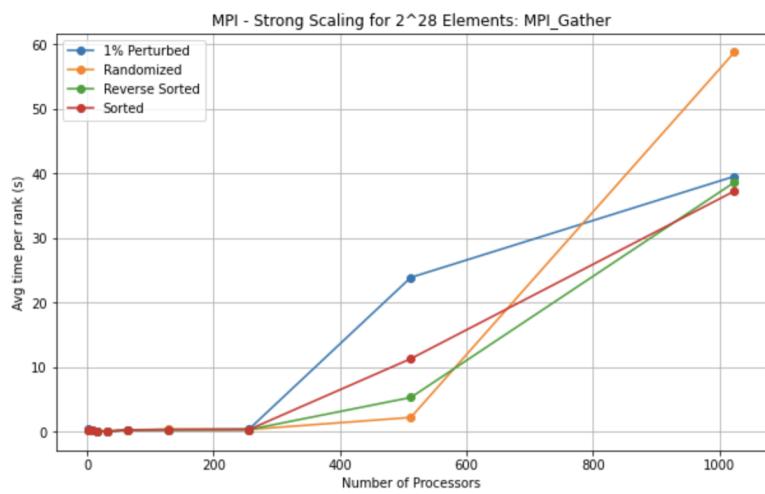


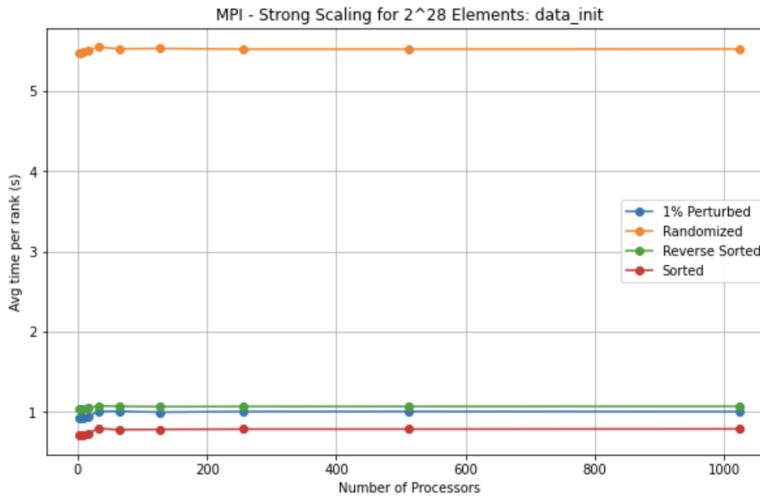
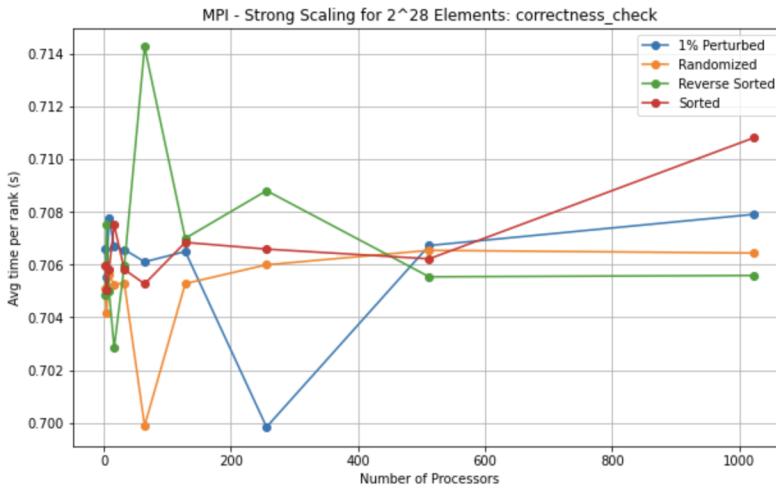
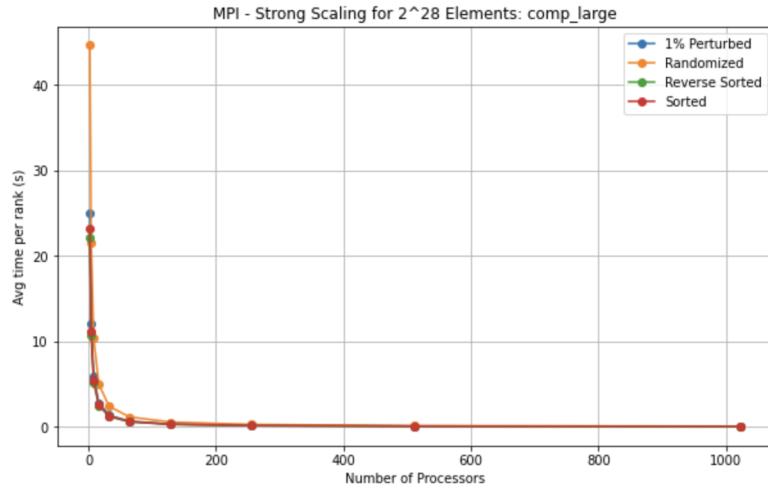




2²⁸ Elements







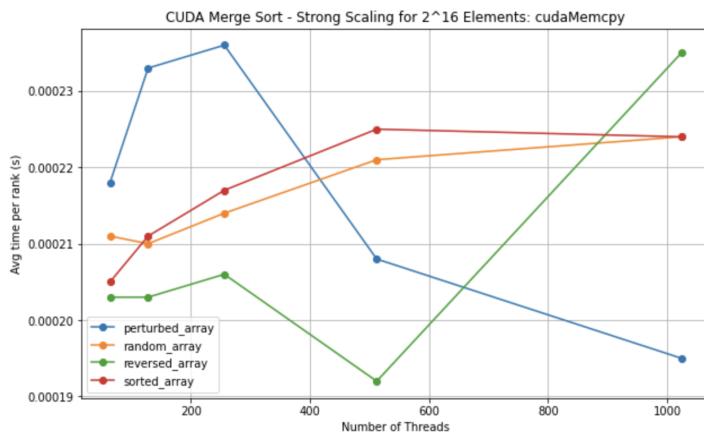
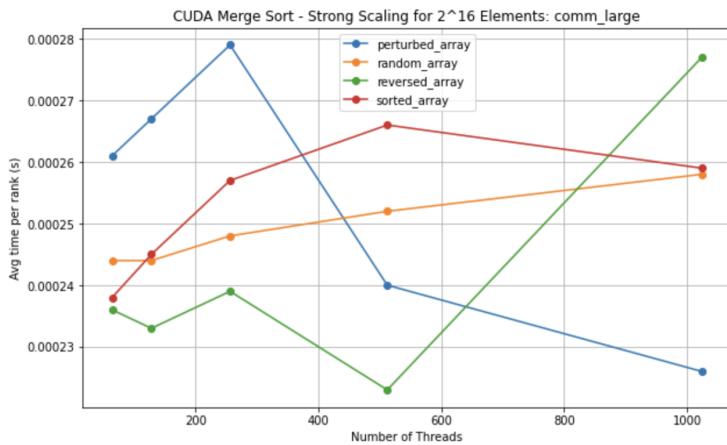
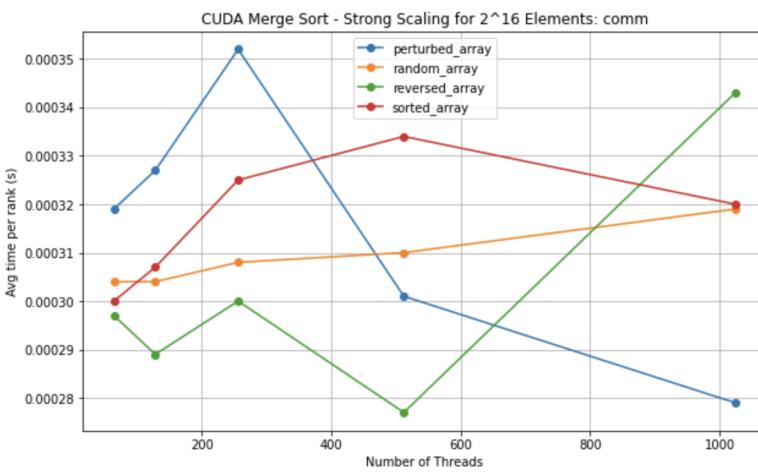
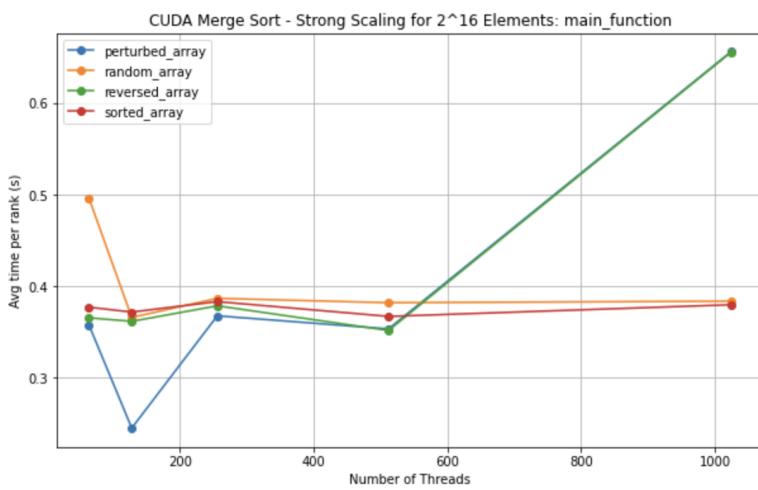
CUDA

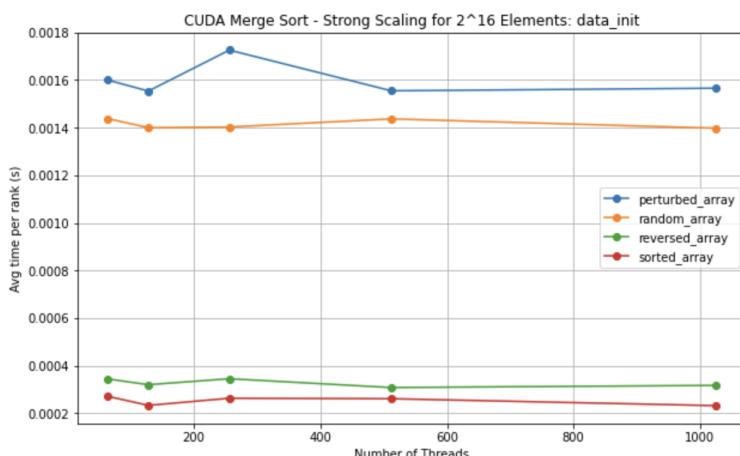
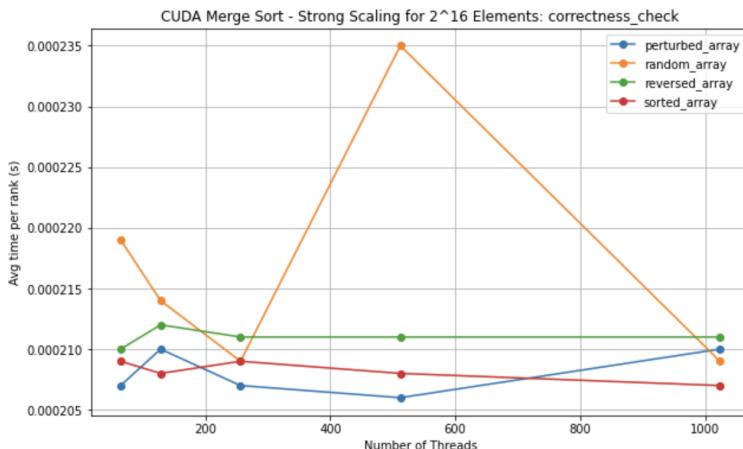
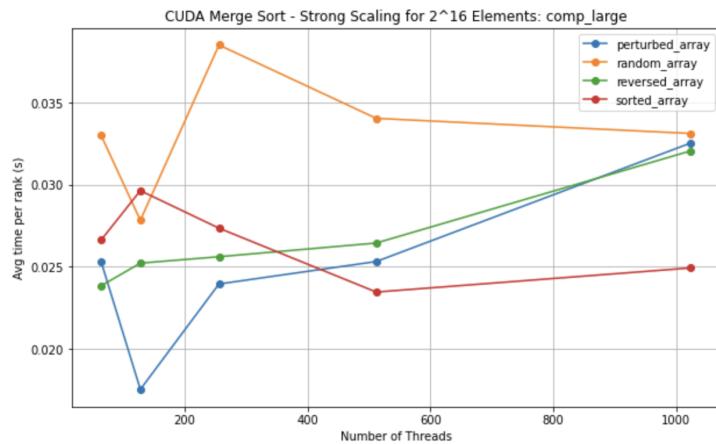
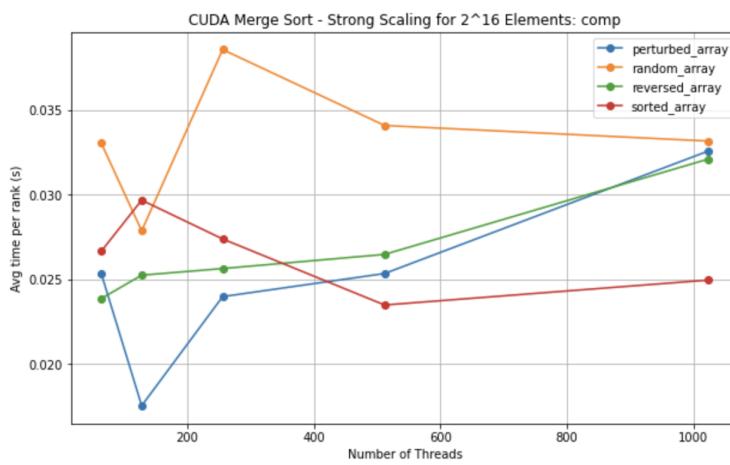
Looking at the graphs for CUDA strong scaling before, the data seems to be much more erratic. Rather than steady increases, we have many fluctuations as we increase the number of threads at a certain size, regardless of the input type. There doesn't seem to be an input type that outperforms the others. This is rather expected because merge sort will continue to break down the elements, regardless of the input type. It does look like the random array input does take a little more time than the other inputs, as it generally sits above the rest of the input types. This is followed by the 1% perturbed array. However, it is important to note that the scale of these graphs tend to be very close together, which can signify that there is no real tradeoff here based on input type. We definitely see that the comm region of the code,

which again is around the CUDAMemcpys, fluctuates quite a bit, regardless of input type. There is no definite winner here. We get the main differences in computation, but again due to scale, isn't very significant. It is interesting to see though what could potentially cause this. Random taking the longest makes sense because there are a lot of comparisons, but it may be concerning that the reverse sorted array is one of the quicker ones, since this should theoretically have the maximum number of comparisons to make. One thing to note is that these two slowest, the random and perturbed array, have the slowest data initialization timees as well, so this could be a cause.

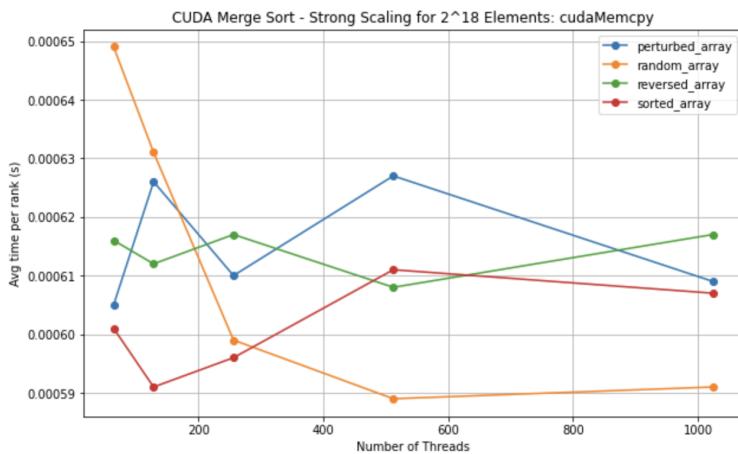
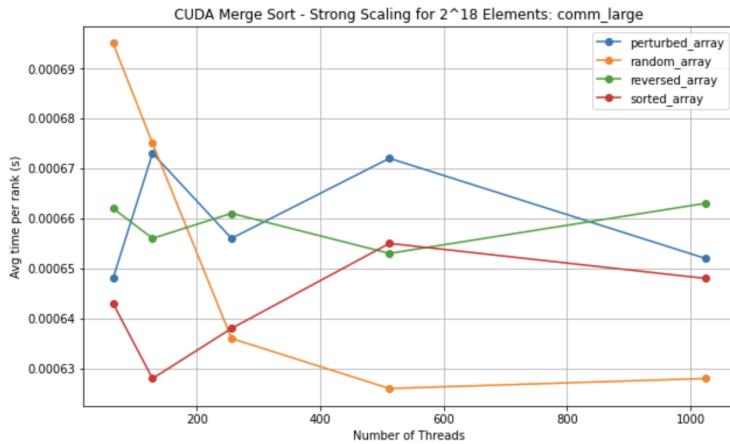
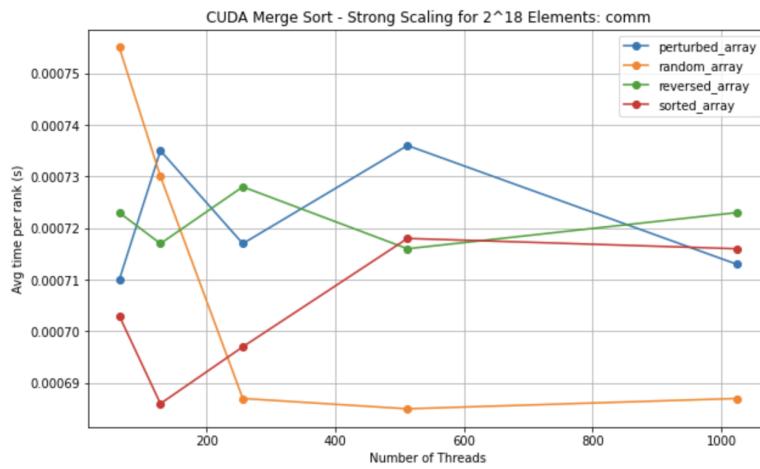
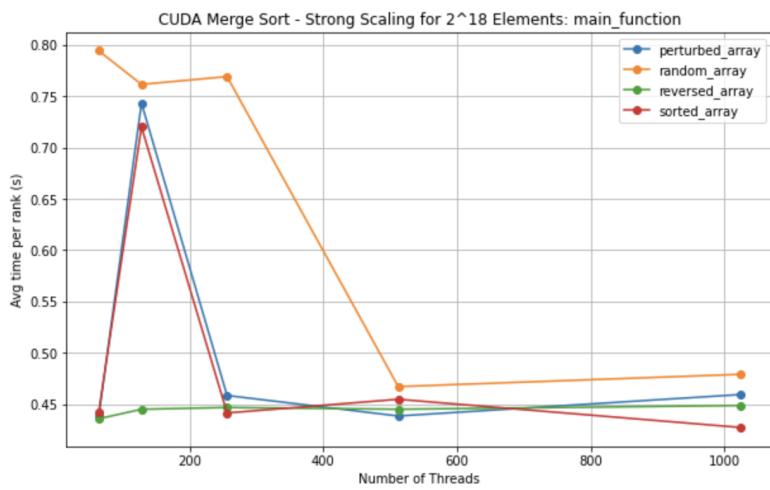
Looking at the graphs for CUDA strong scaling before, the data seems to be much more erratic. Rather than steady increases, we have many fluctuations as we increase the number of threads at a certain size, regardless of the input type. There doesn't seem to be an input type that outperforms the others. This is rather expected because merge sort will continue to break down the elements, regardless of the input type. It does look like the random array input does take a little more time than the other inputs, as it generally sits above the rest of the input types. This is followed by the 1% perturbed array. However, it is important to note that the scale of these graphs tend to be very close together, which can signify that there is no real tradeoff here based on input type. We definitely see that the comm region of the code, which again is around the CUDAMemcpys, fluctuates quite a bit, regardless of input type. There is no definite winner here. We get the main differences in computation, but again due to scale, isn't very significant. It is interesting to see though what could potentially cause this. Random taking the longest makes sense because there are a lot of comparisons, but it may be concerning that the reverse sorted array is one of the quicker ones, since this should theoretically have the maximum number of comparisons to make. One thing to note is that these two slowest, the random and perturbed array, have the slowest data initialization timees as well, so this could be a cause. We also see that the merge step was done on the CPU, which introduces a bit of a bottleneck as we increase the problem size and number of threads.

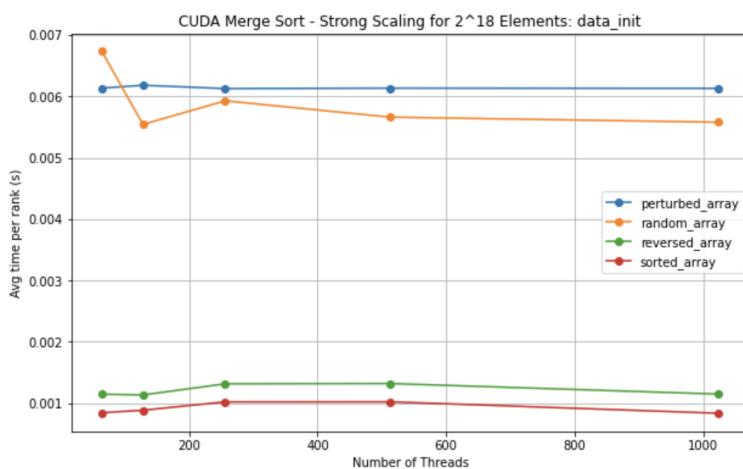
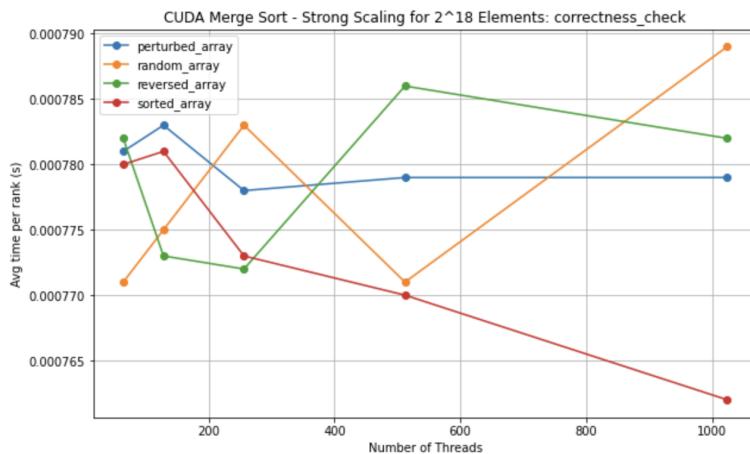
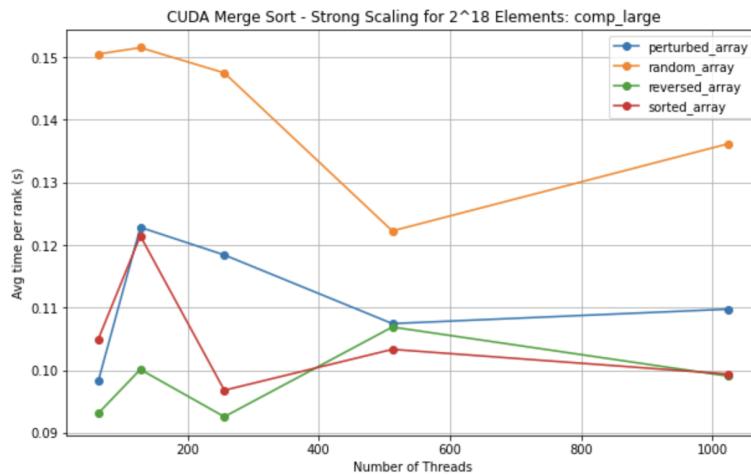
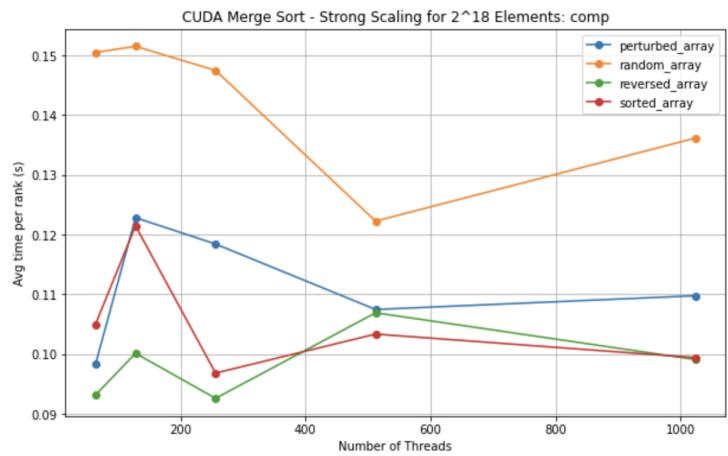
2¹⁶ Elements



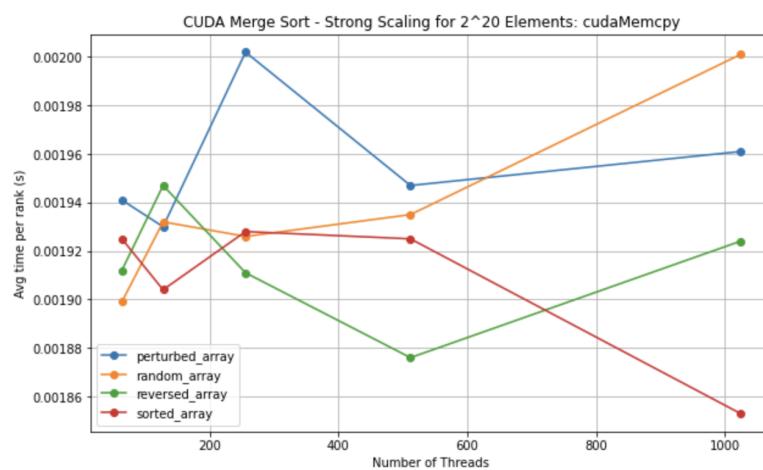
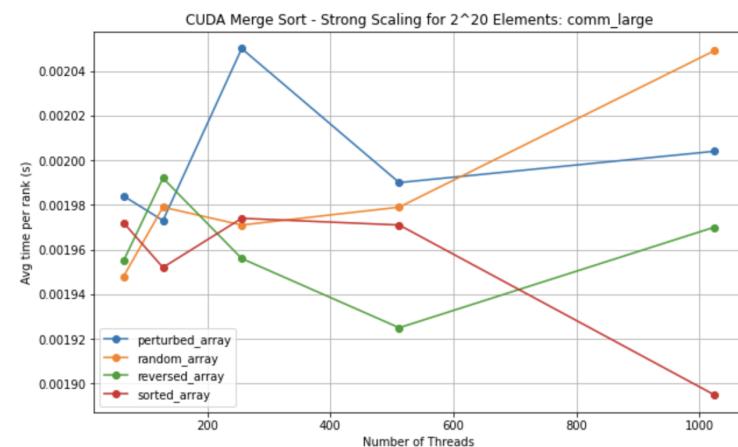
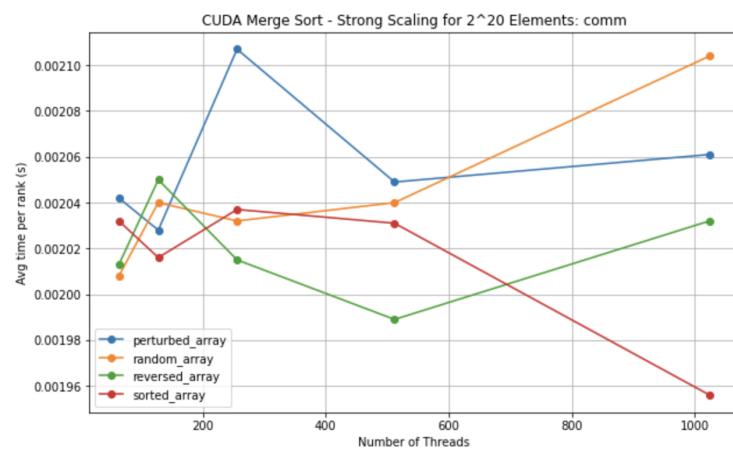
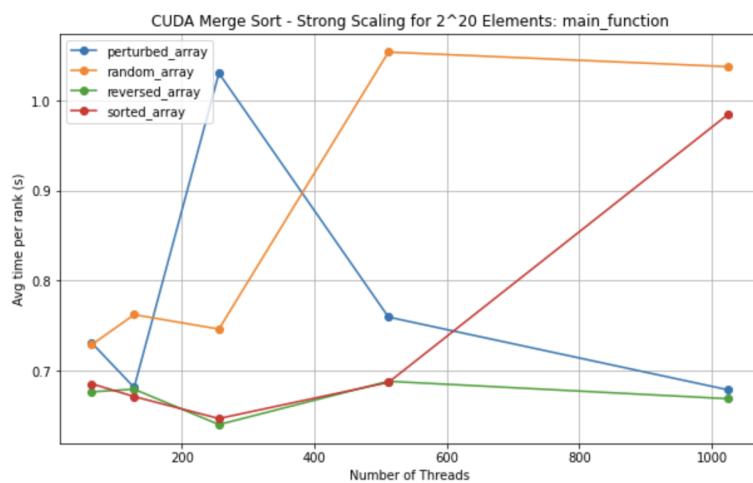


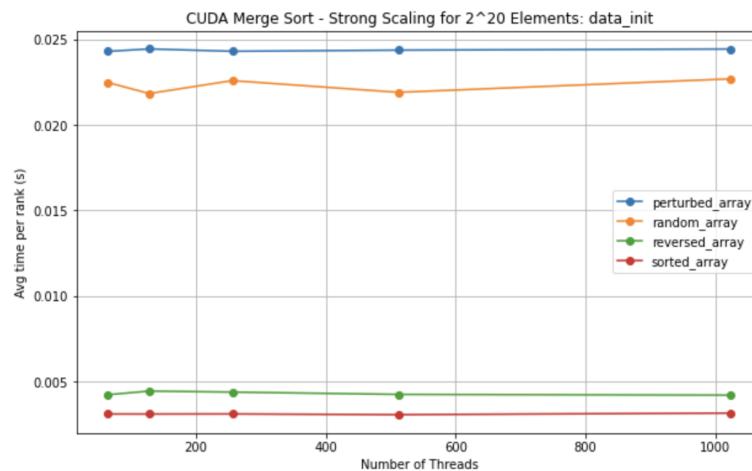
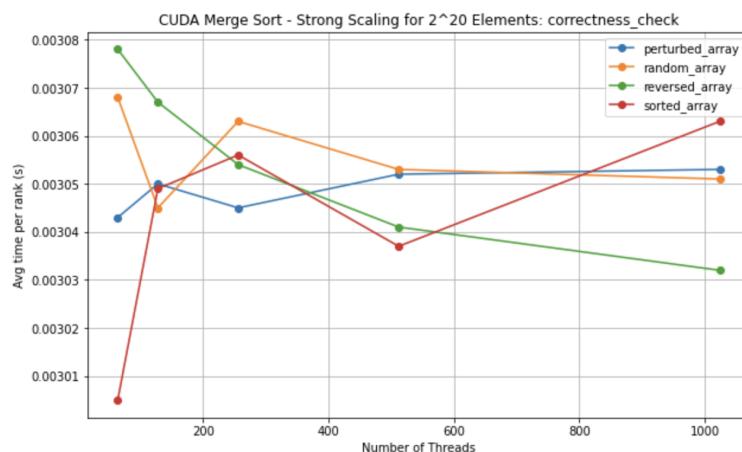
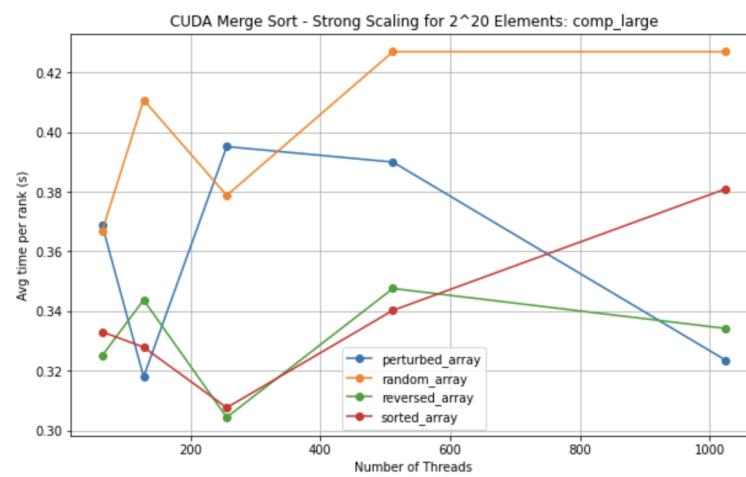
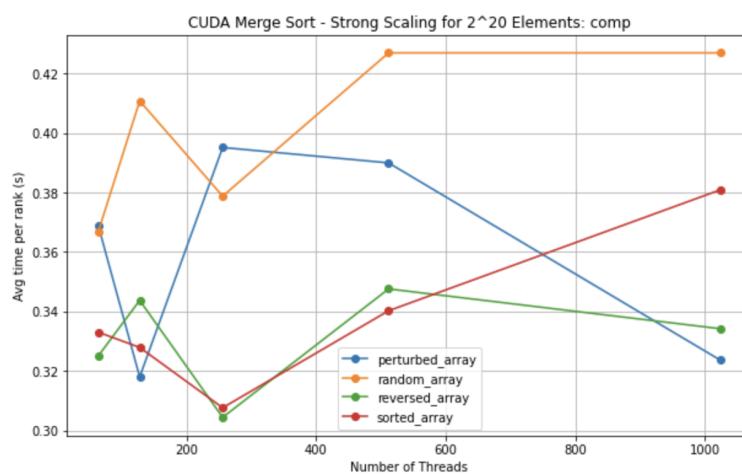
2 \wedge 18 Elements



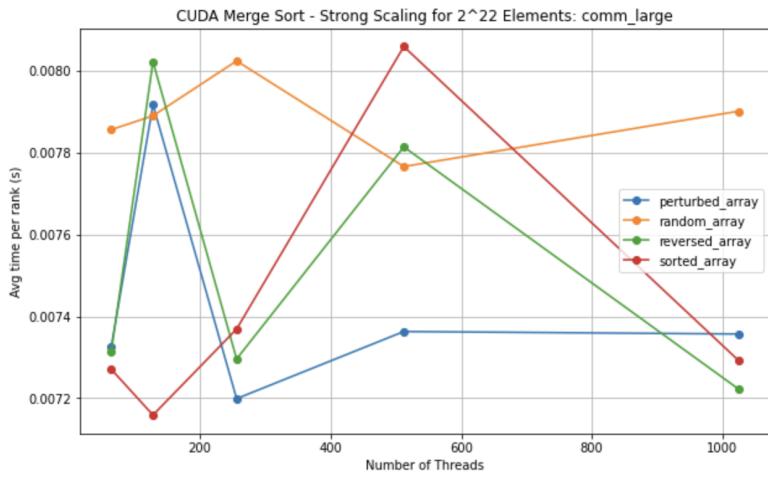
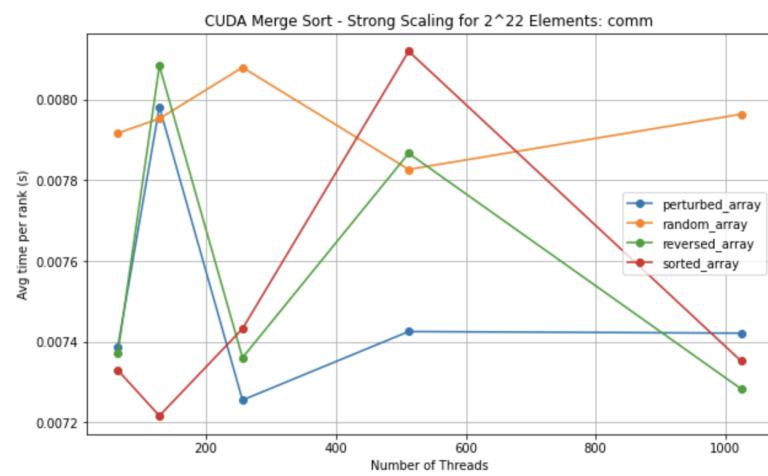
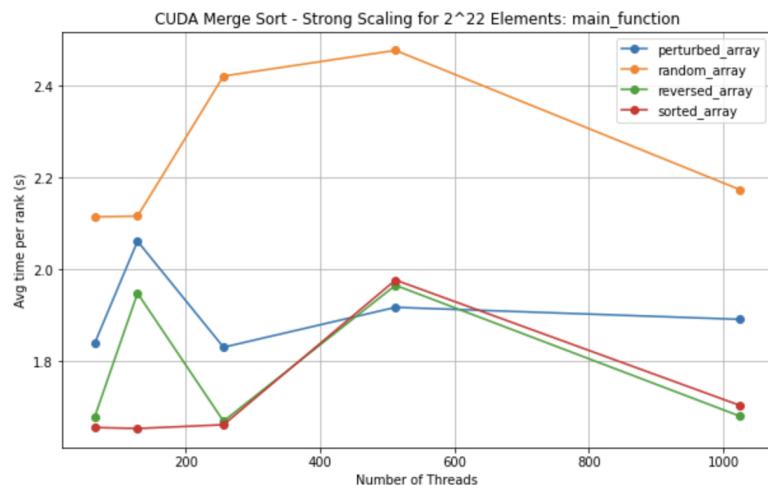


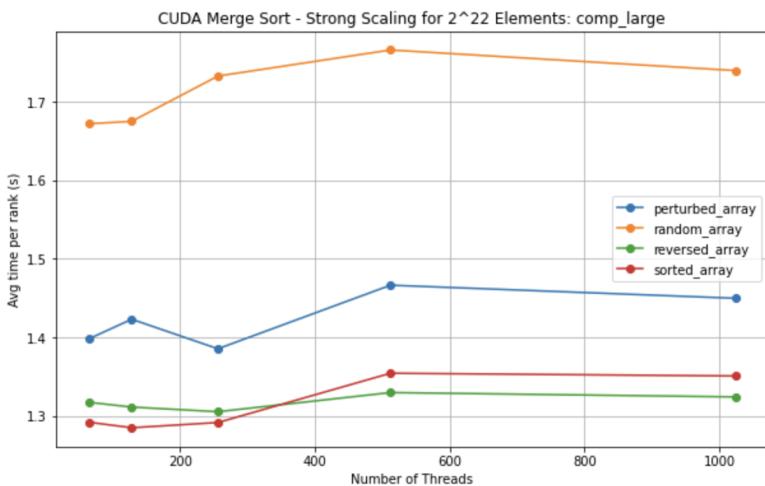
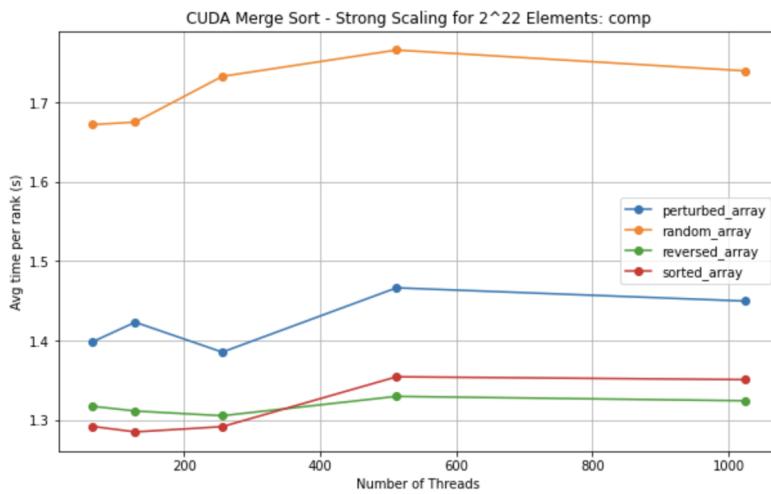
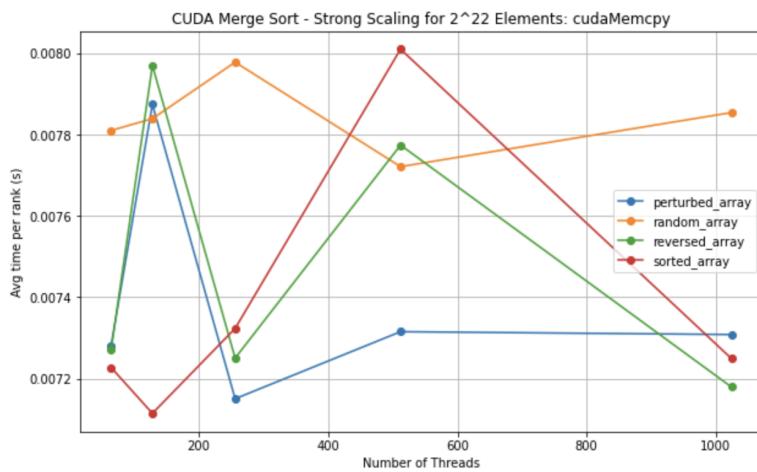
2^{20} Elements

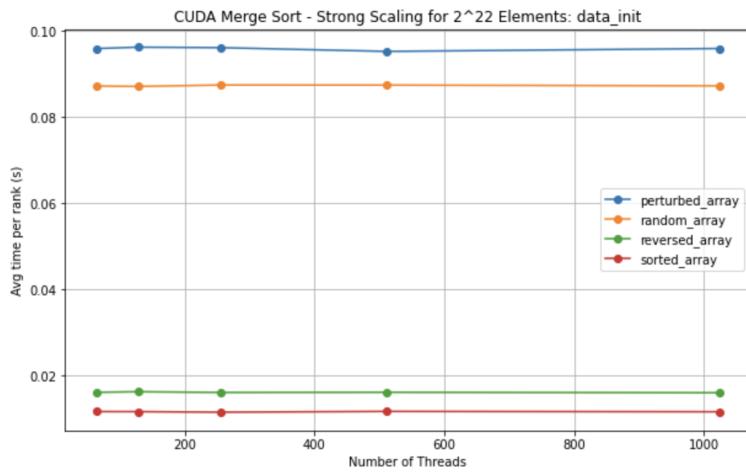
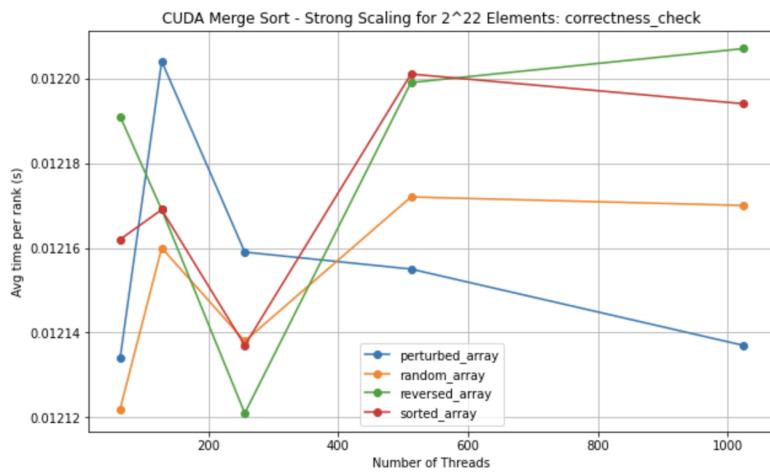




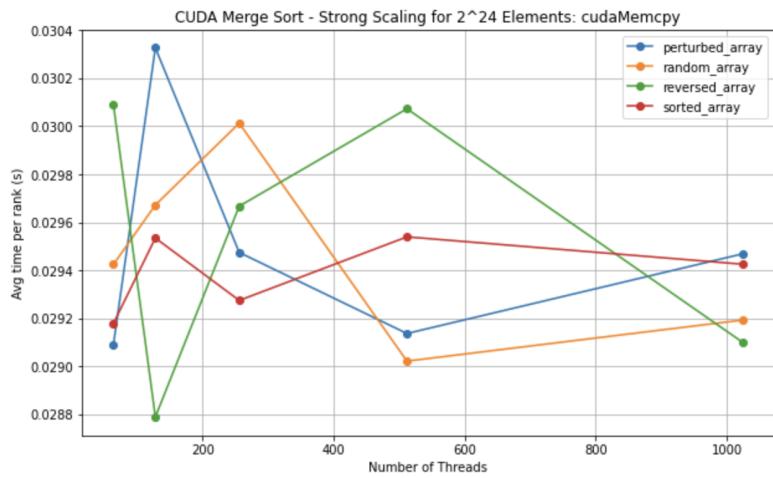
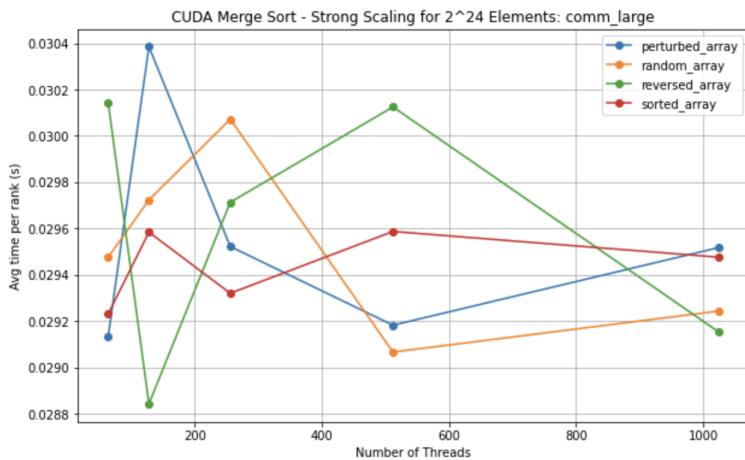
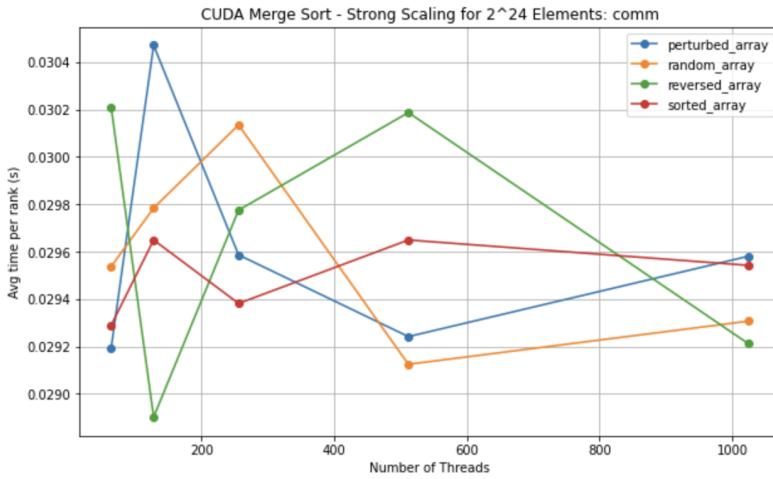
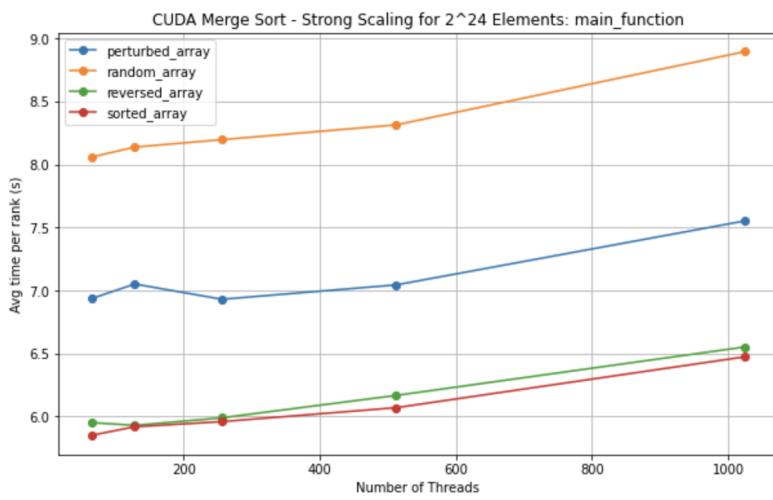
2^22 Elements

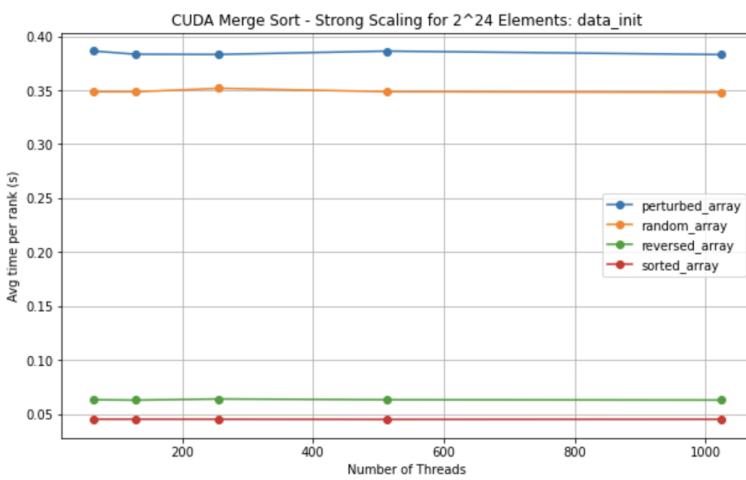
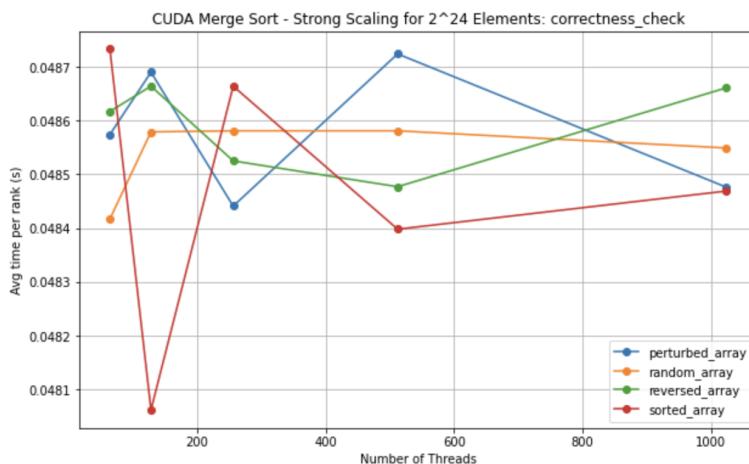
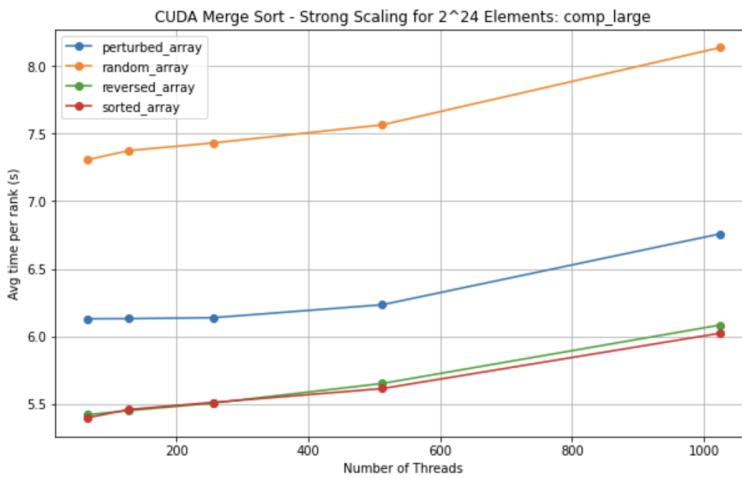
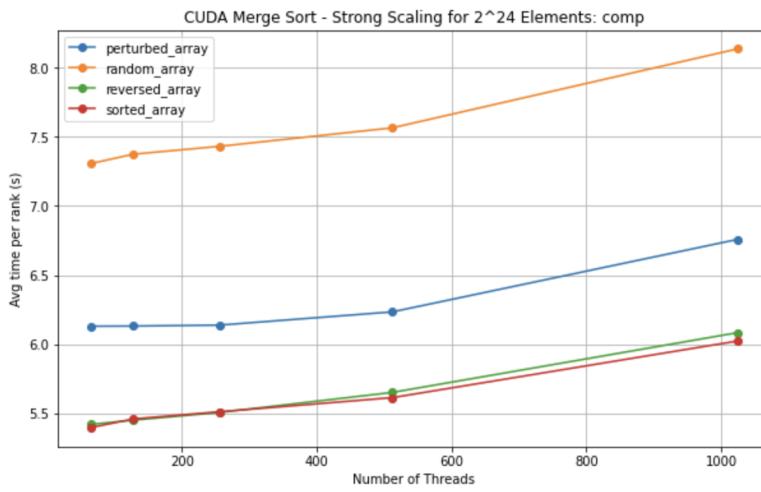




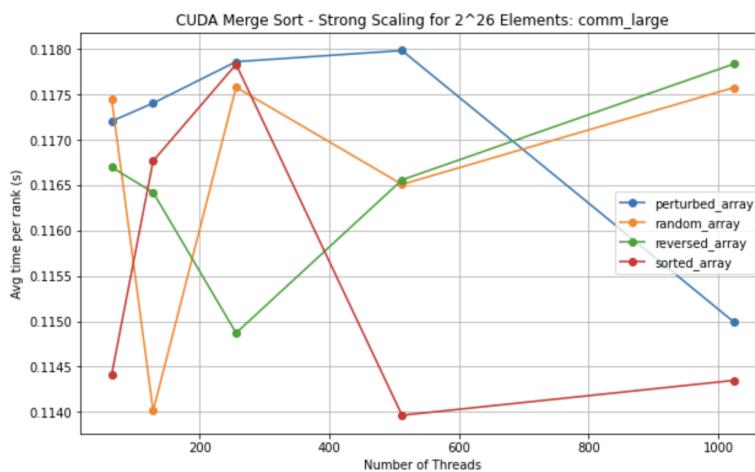
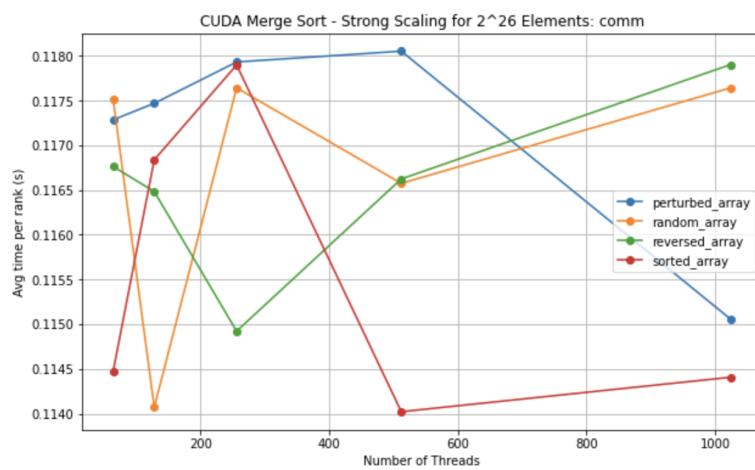
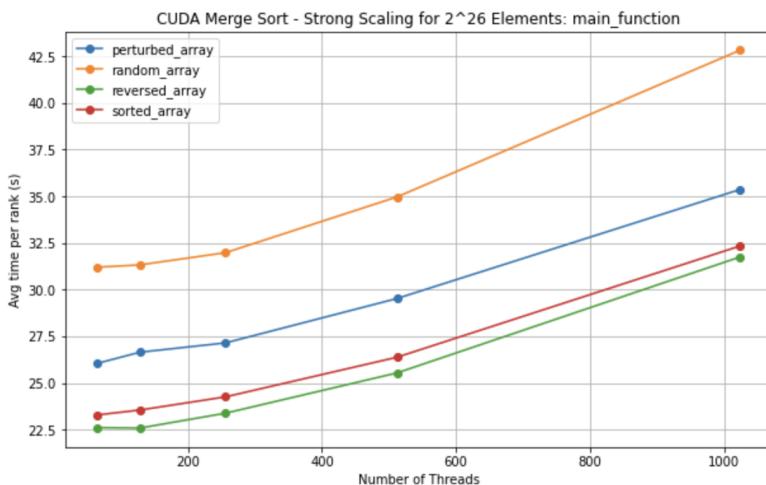


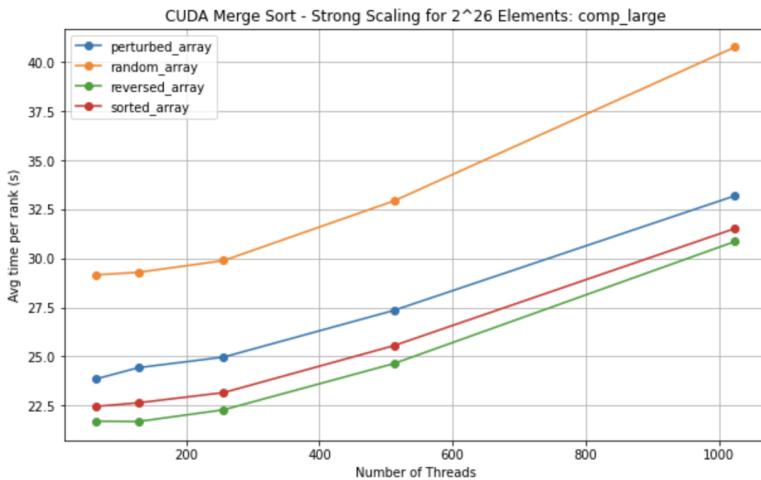
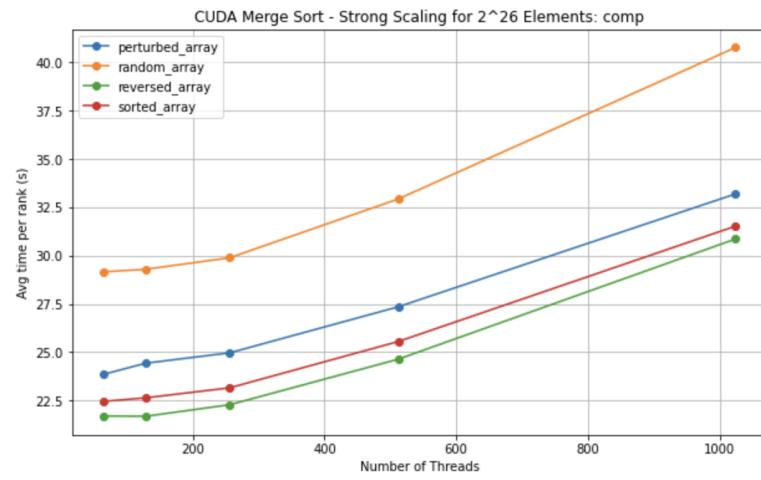
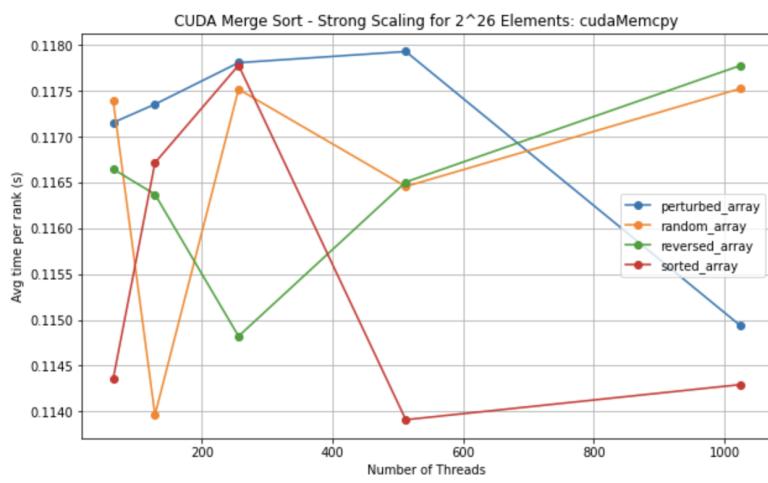
2^{24} Elements

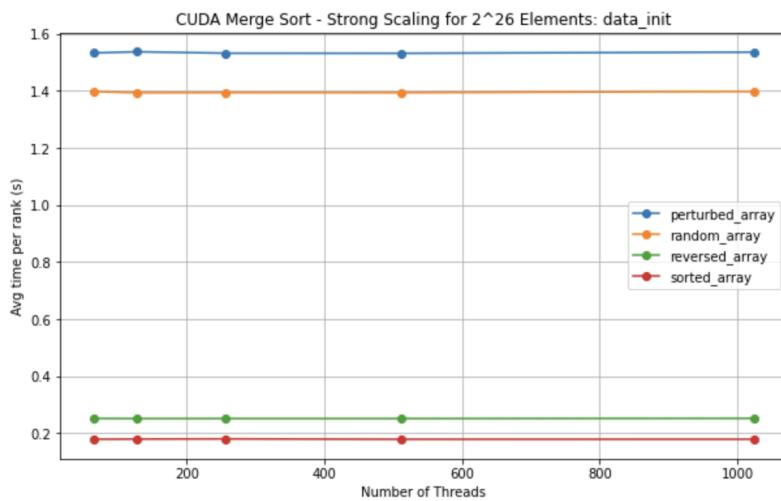
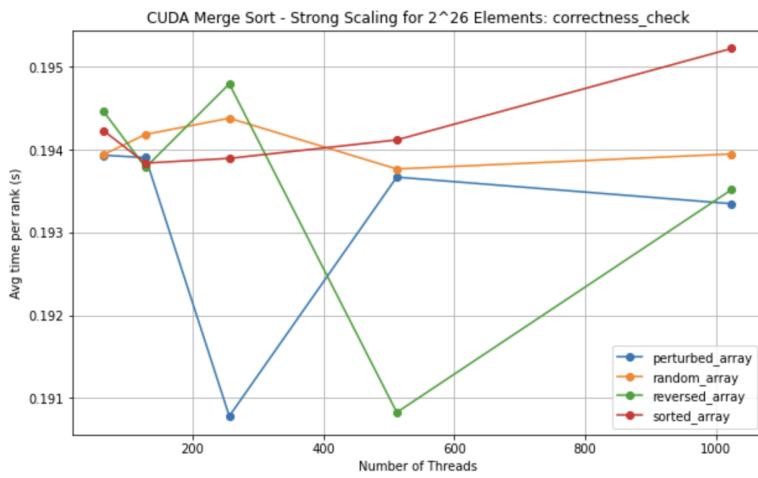




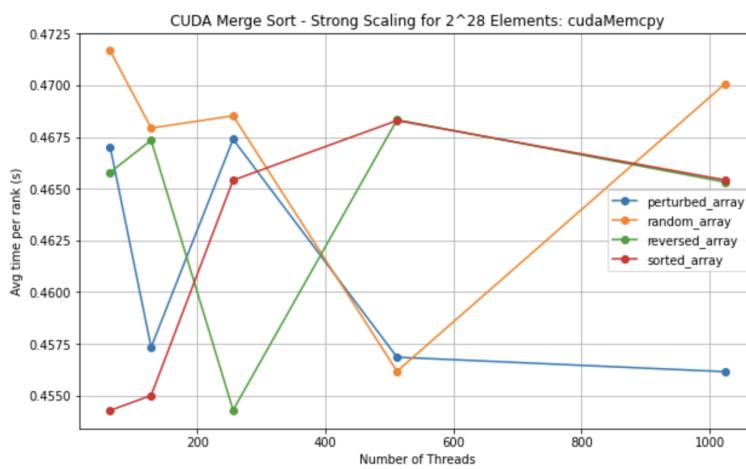
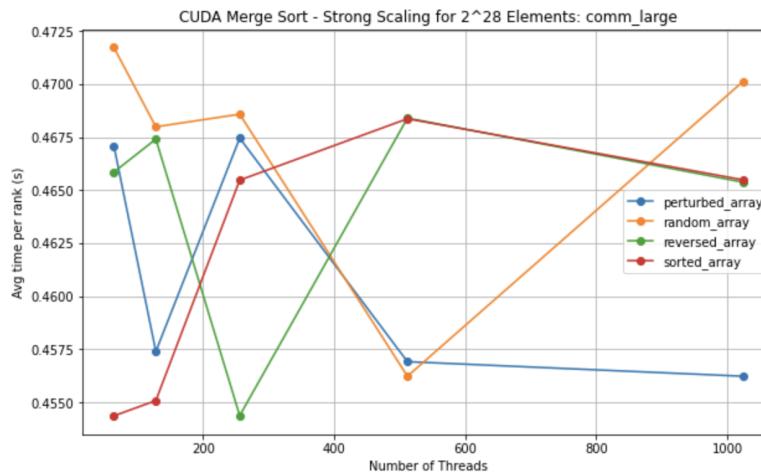
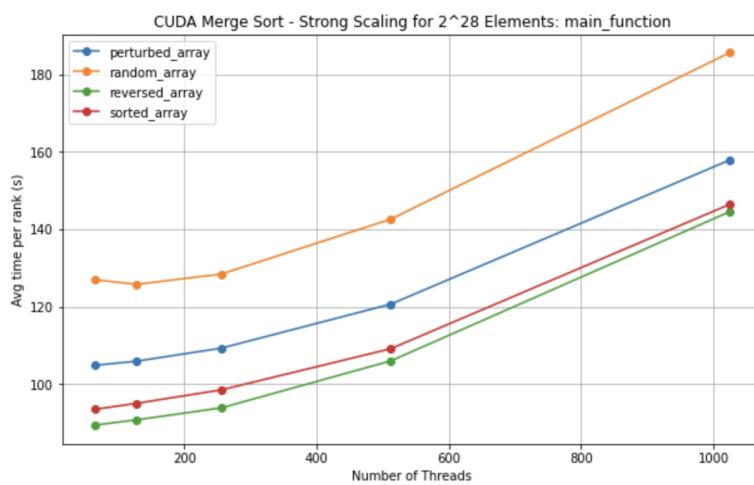
2^26 Elements

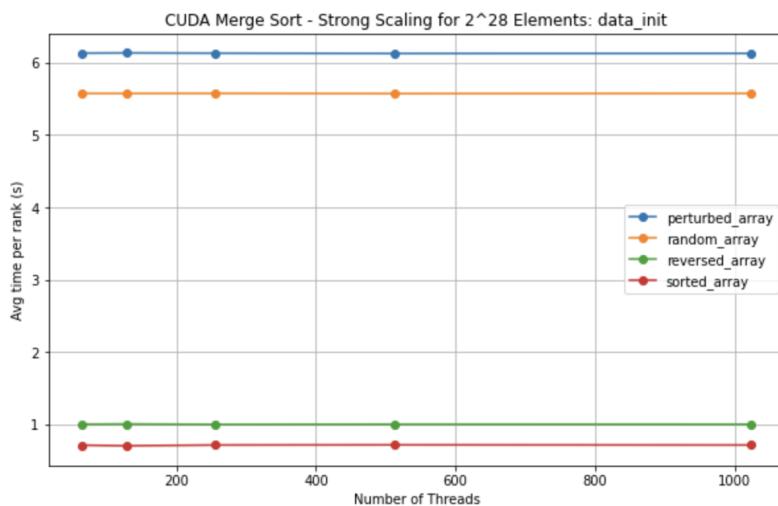
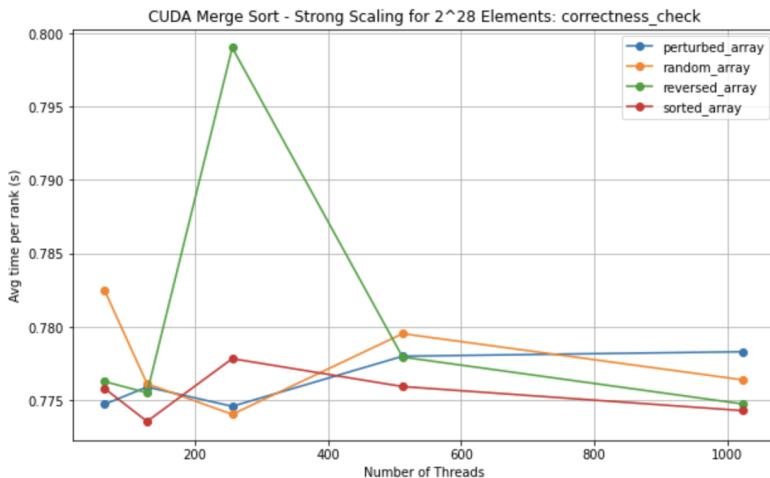
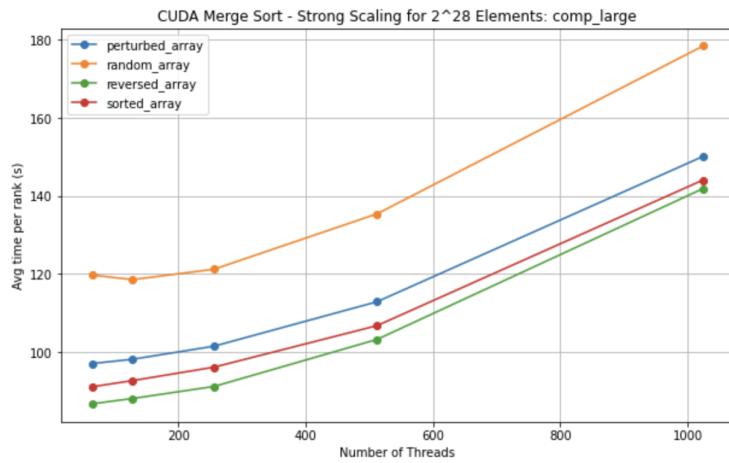
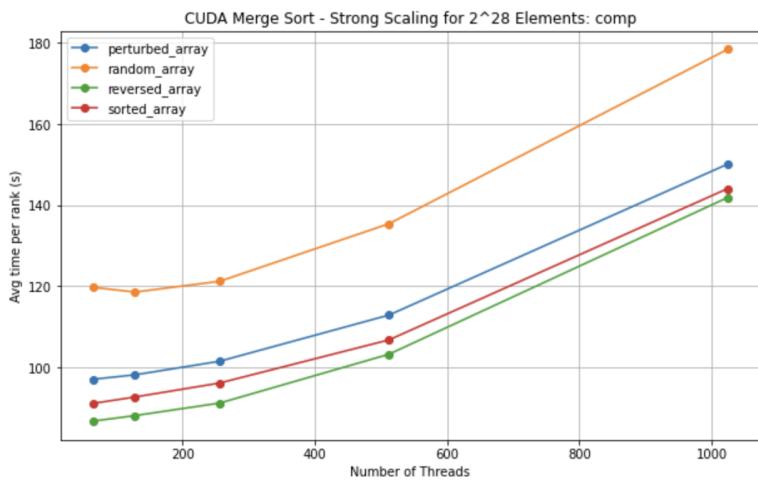






2^{28} Elements



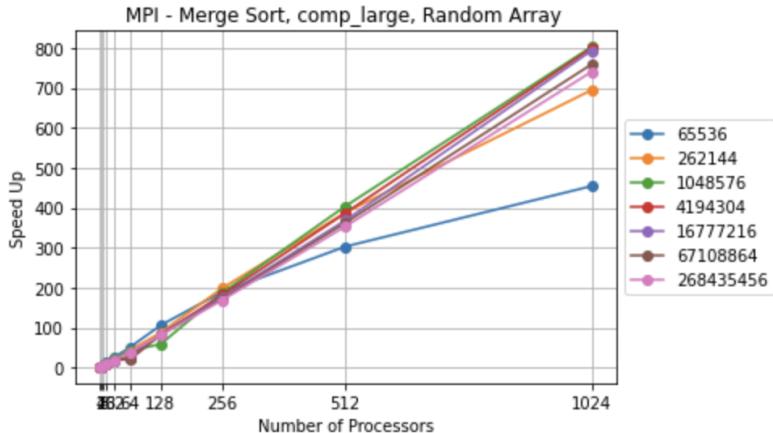
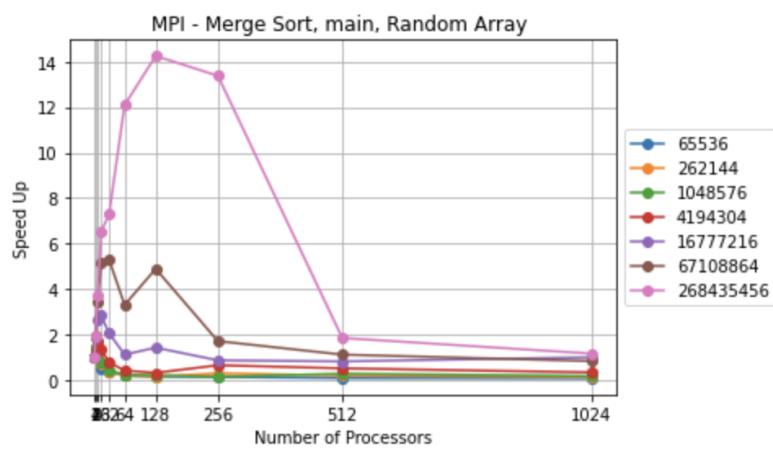


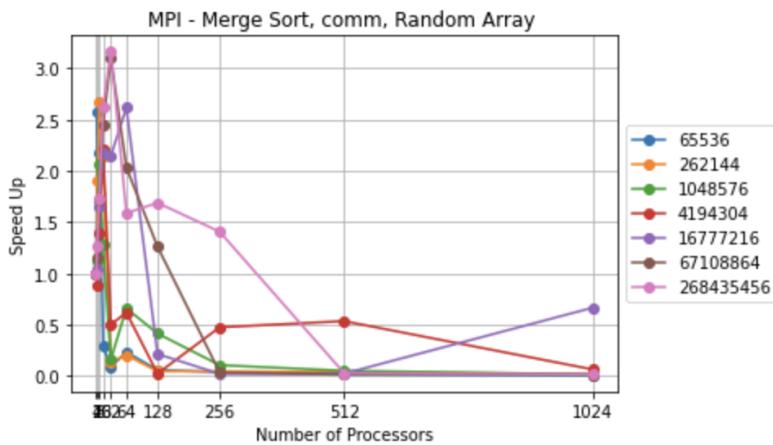
Speedup

MPI

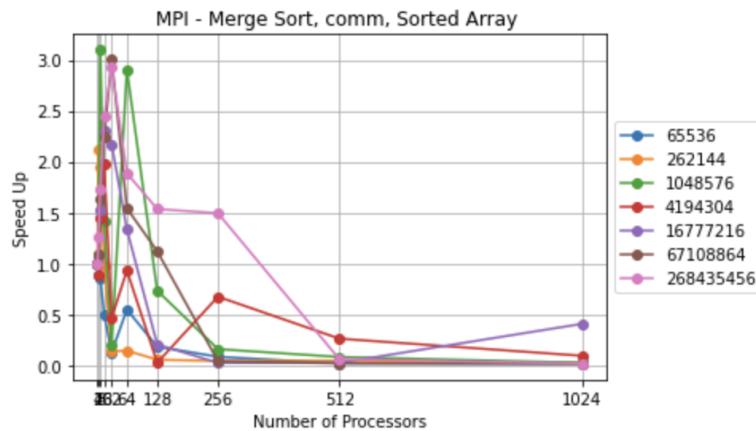
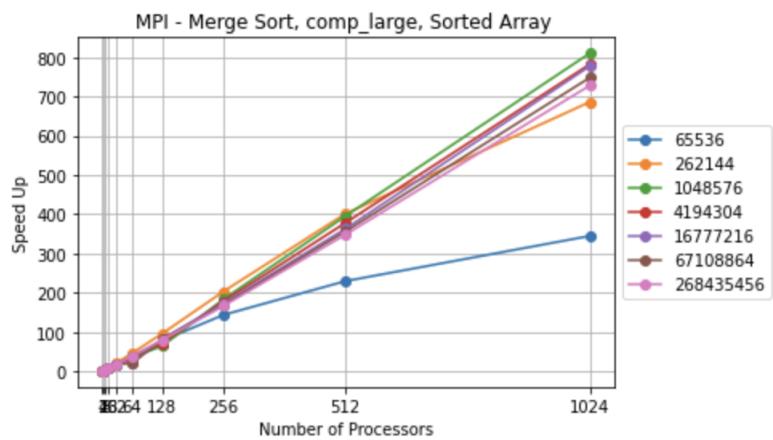
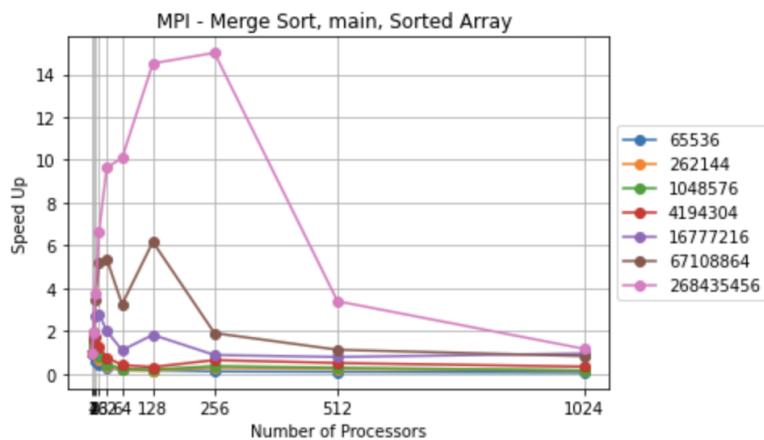
When looking at the computation speedups, we see a linear trend between the input sizes. This shows the importance of parallelism. For all sizes, as we increase the number of processors, the pure computation time gets better and better, this showing the higher speedup. We do see that speedup for the 2^{16} elements begins to taper off for the comp regions, indicating that the parallelism is good for very large sizes, but smaller sizes aren't as impacted by increasing processors. This makes sense because you would spend more percentage of your time increasing your processors and have more communication and face more overhead for the smaller values. If we look at the entire main region as a whole, we can see that the larger the input the size, the more speedup there is. This further reinforces that there is a need for parallelism, and a definite benefit. For all of the input sizes, there seems to be a peak speedup at 256 processors before decreasing again. After 256 processors, there may be so much overhead that we lose overall performance. It is a little difficult to look at it easily for the smaller sizes, but the speedup graphs tend to all show that a bigger input size benefits from this parallelism. We also see that for communication, speedup decreases for more processors, again reiterating that we are likely facing some significant overhead.

Random Array

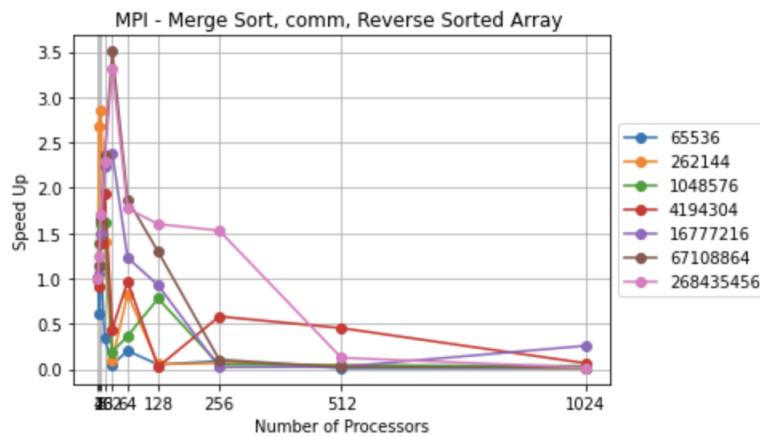
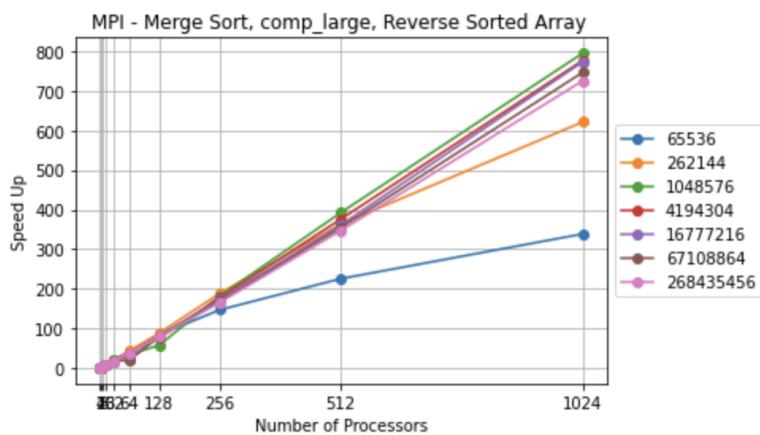
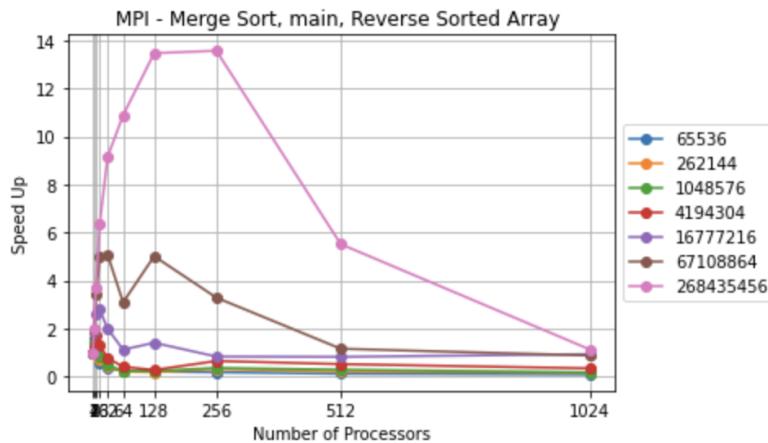




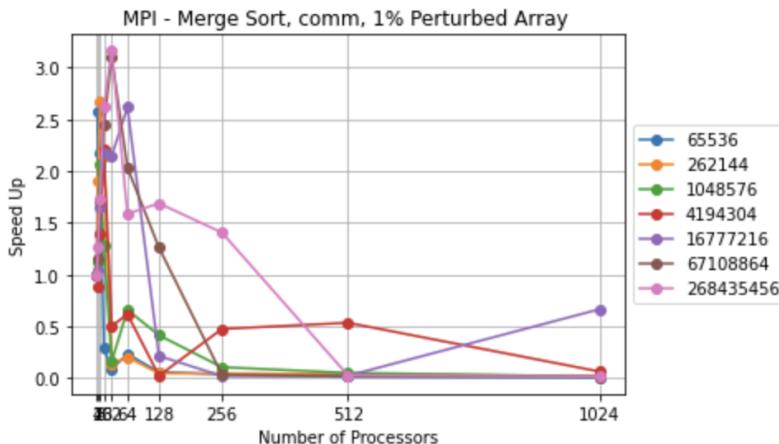
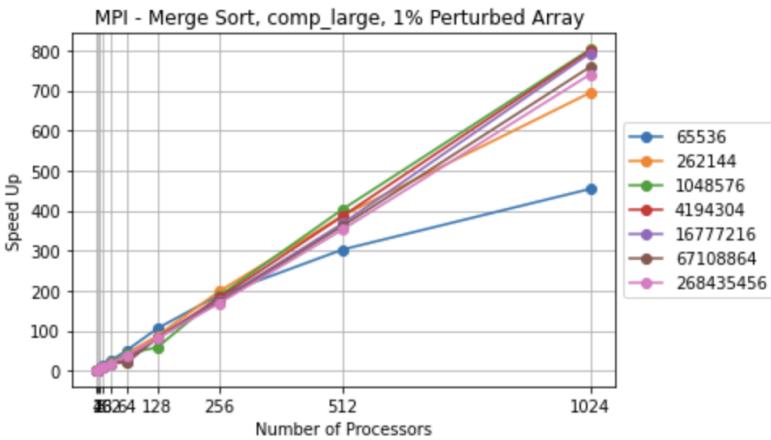
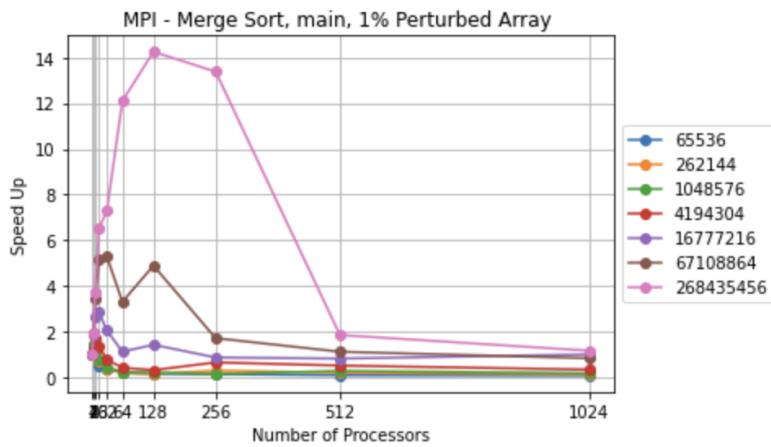
Sorted Array



Reverse Sorted Array



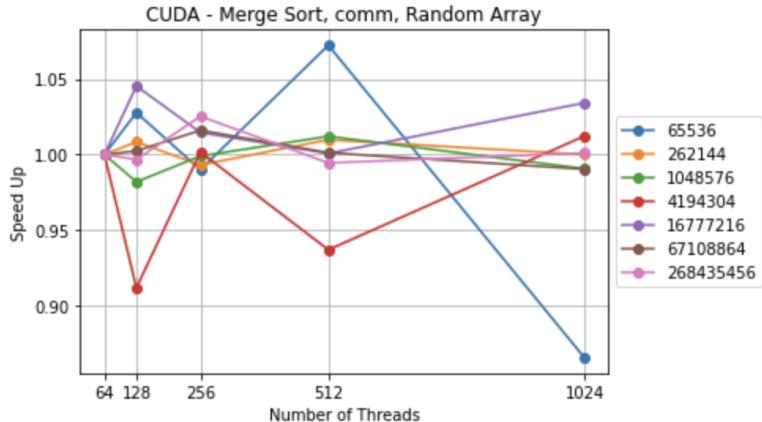
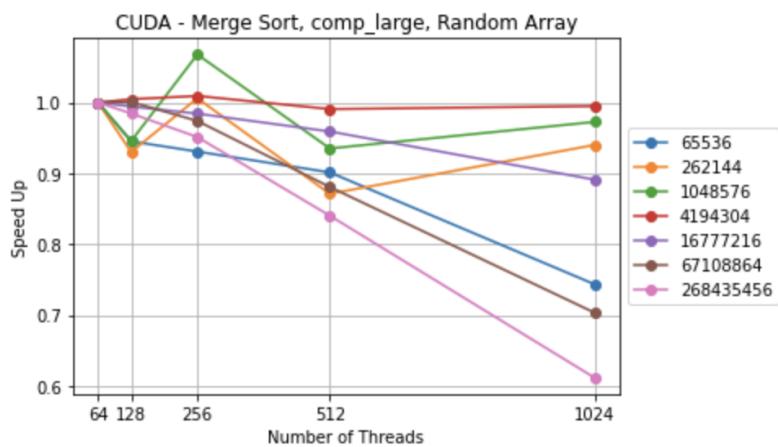
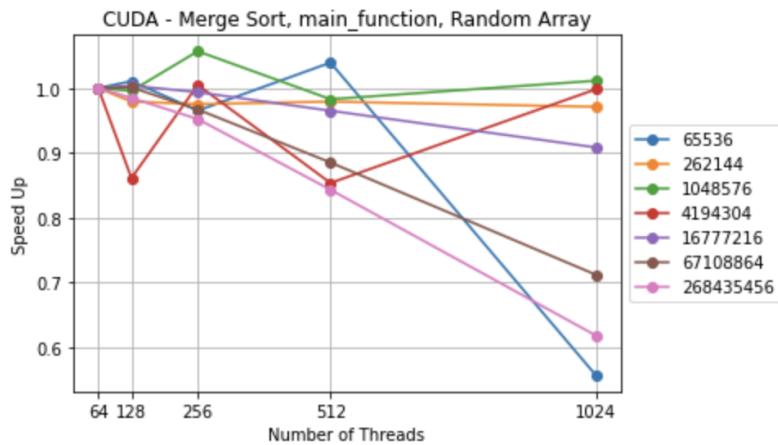
1% Perturbed Array



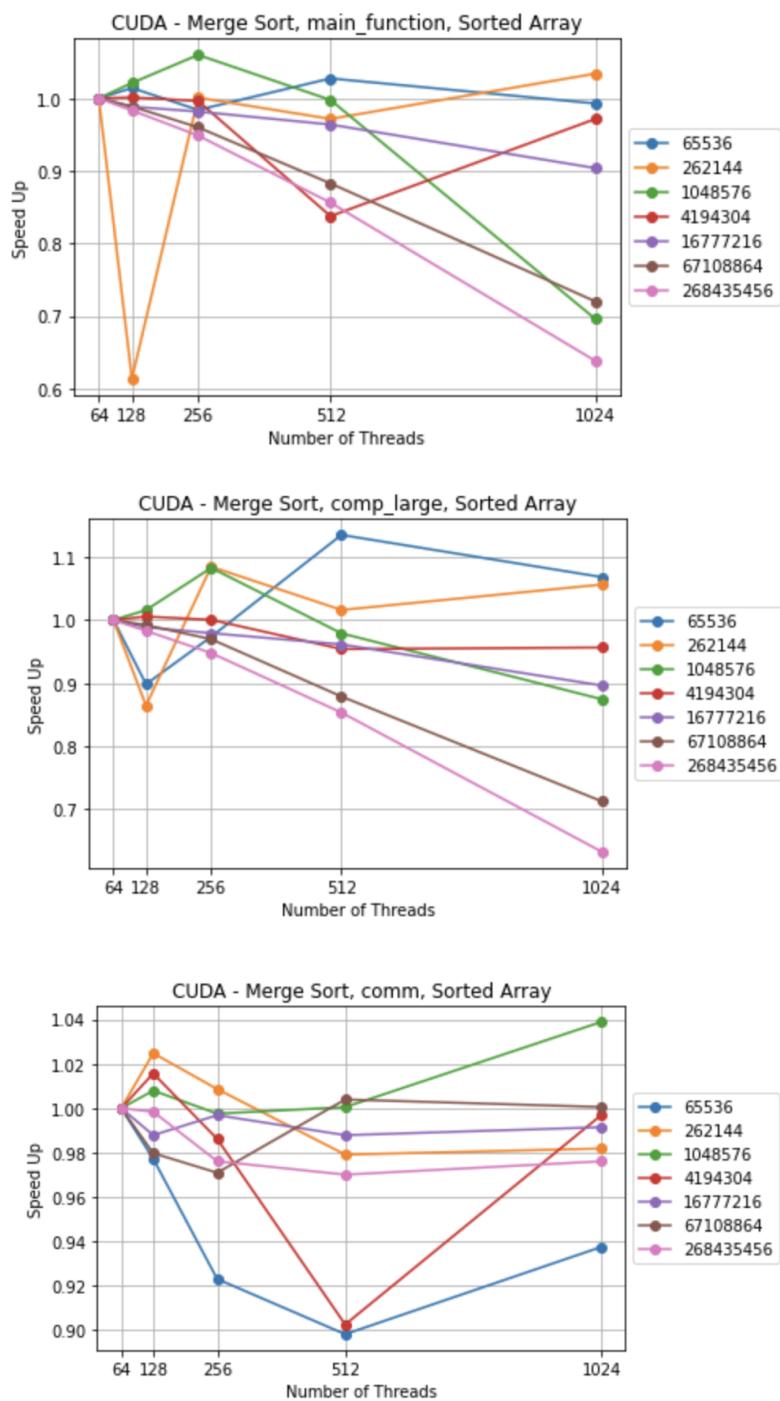
CUDA

In the CUDA speedup graphs, there doesn't seem to be any inherent benefit to parallelism. When we look at the main_function as a whole, we see that the speedup values tend to hover around 1.0. In fact, most of the lines tend to be below a value of 1, indicating that it is performing worse by being parallel. There also seems to be no pattern with the varying input size. We see that at the highest number of processors, the worst speedups are 2^{16} elements and 2^{28} elements, which are both the smallest and largest values respectively. It seems like as we increase processors, speedup for the communication is worst for the 2^{16} element, but the computation is worst for 2^{28} elements. Communication is hard to gauge completely because there is quite a bit of fluctuation here, where computation speedup hovers around 1.0 and steadily decreases for most of the input sizes. This could also be because theres no real inter-process communication for CUDA like there is for MPI.

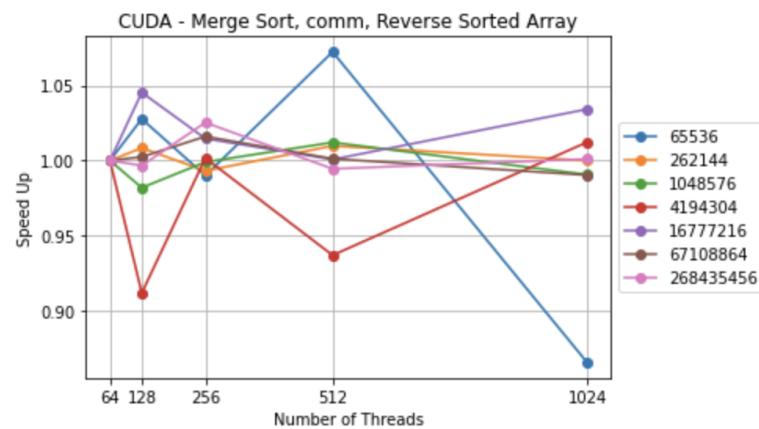
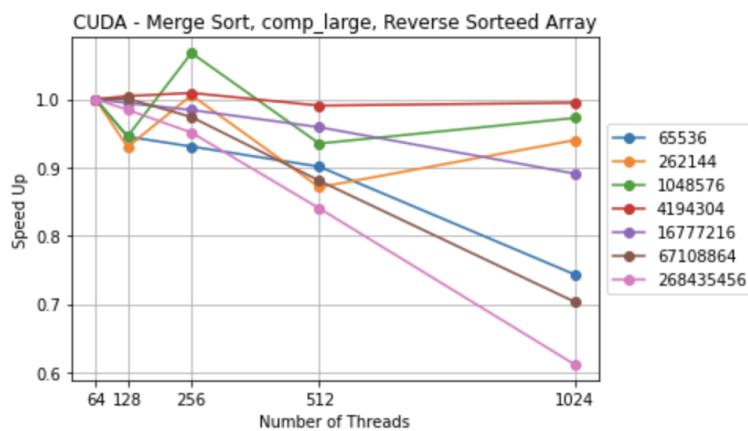
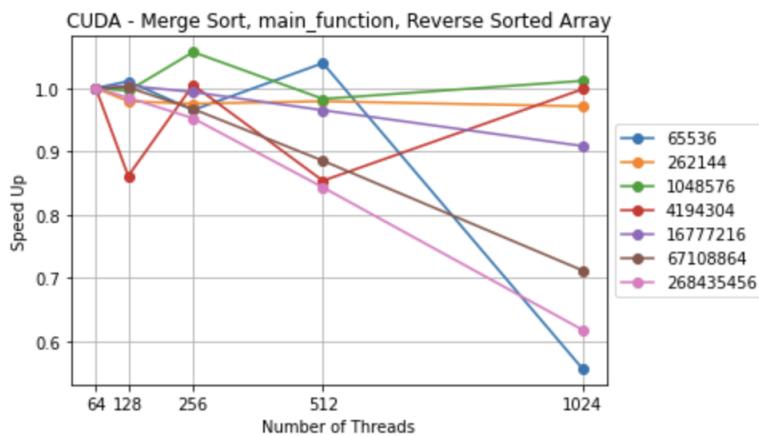
Random Array



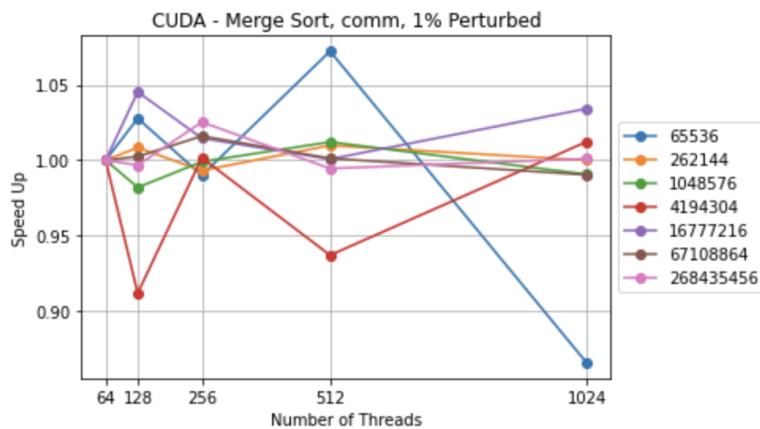
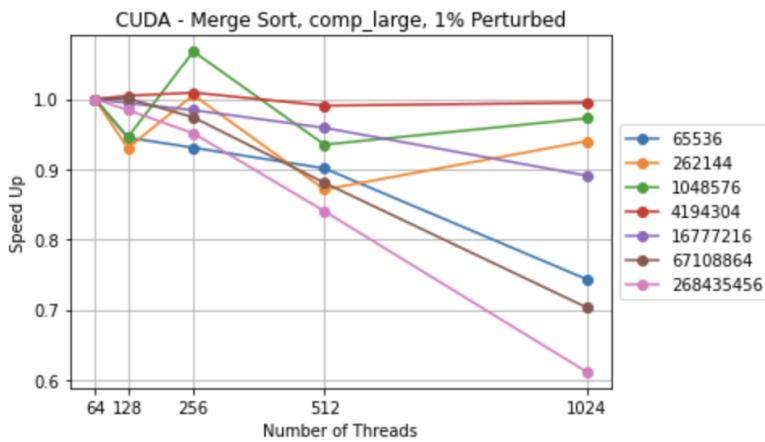
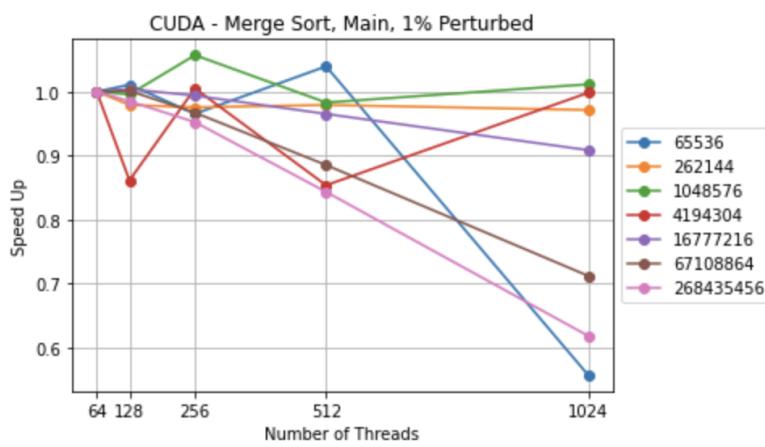
Sorted Array



Reverse Sorted Array



1% Perturbed Array



Bitonic Sort

Weak Scaling

MPI

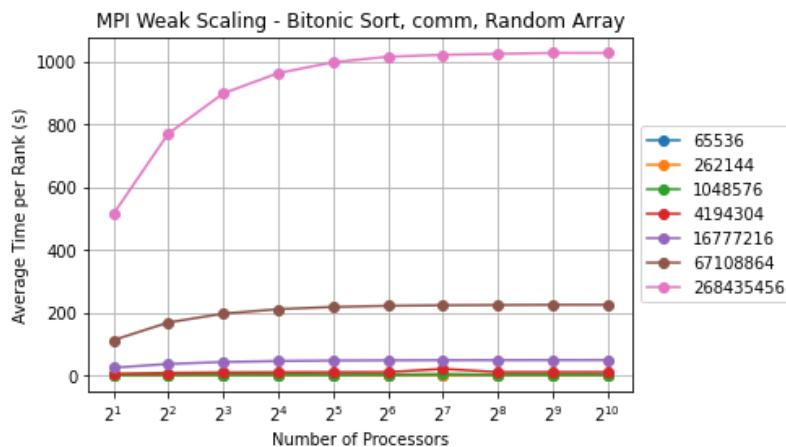
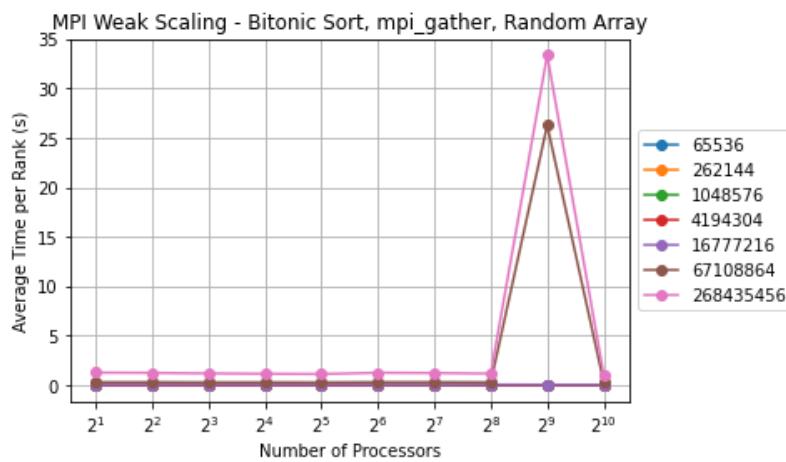
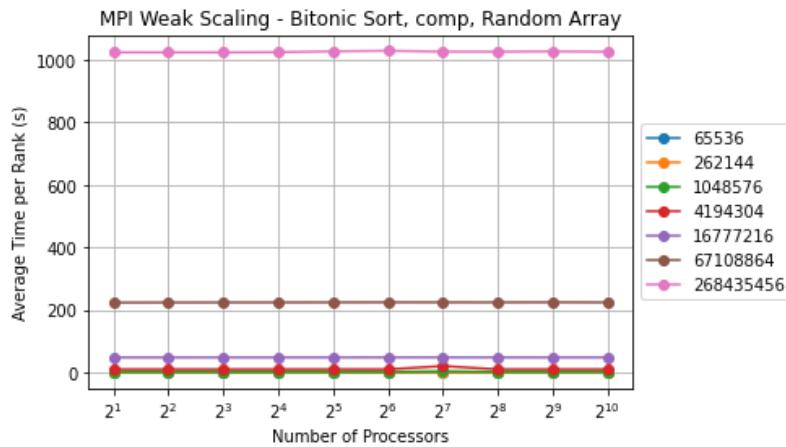
Like Shawn for merge sort, I also opted to assess the average time per rank in relation to the number of processors instead of total time since total time invariably increases with the addition of processors, as it represents an aggregate of all times across all processors.

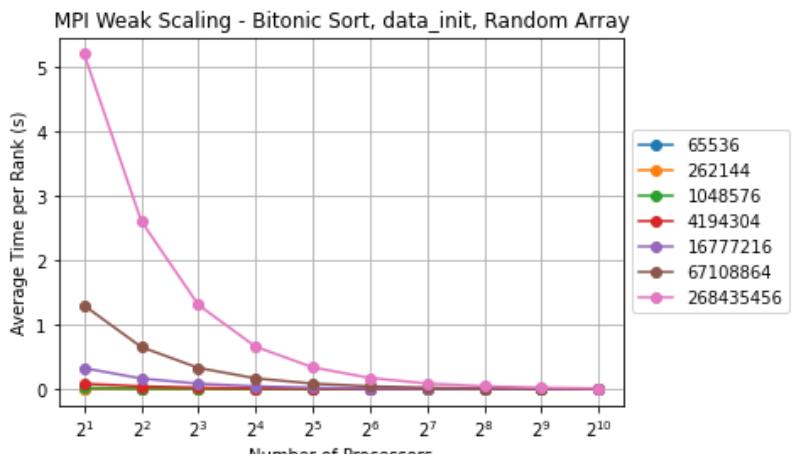
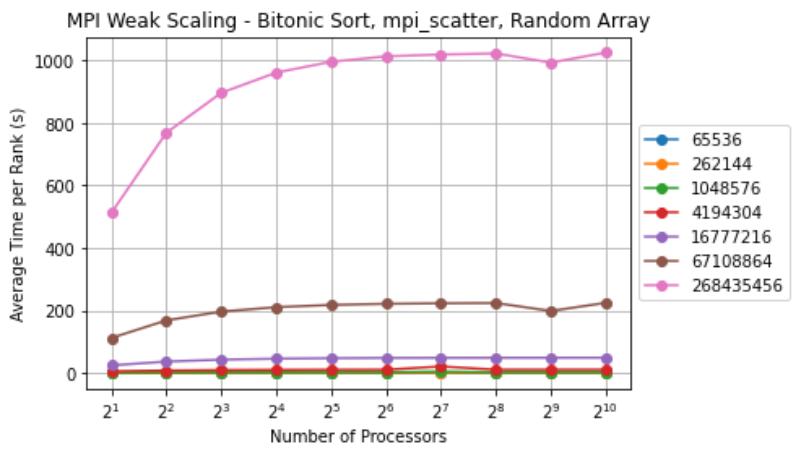
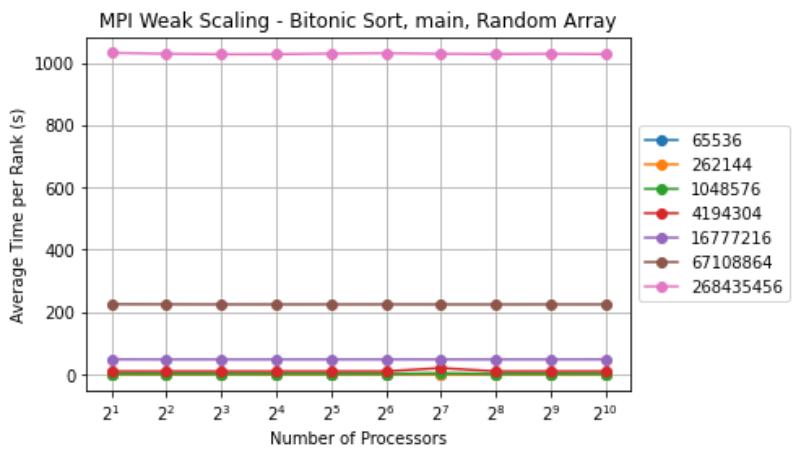
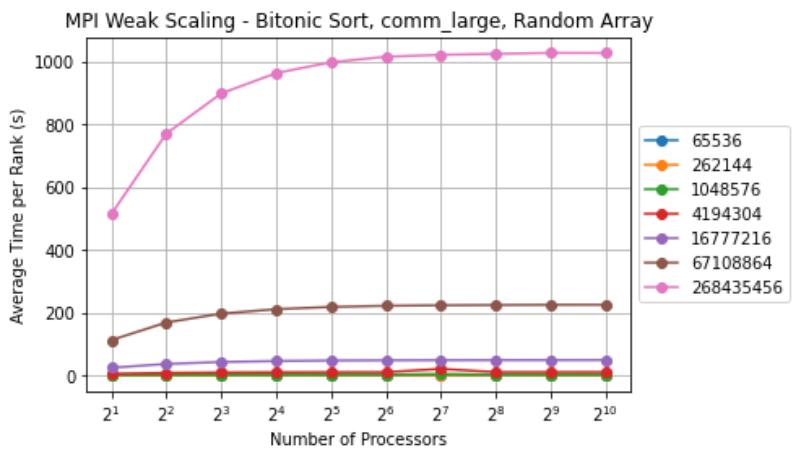
For this analysis, I examined weak scaling using 2^{16} elements across all four input types (sorted, reverse sorted, randomized, and 1% perturbed). It is evident that as the number of processors increases, the sorted array consistently exhibits the fastest runtime. This scenario represents the optimal condition for

bubble sort, as it involves zero swaps for elements. The efficiency of this input type remains largely unaffected by the increase in processors, given its inherently efficient nature. The next fastest is the 1% perturbed input, which demonstrates a noticeable increase in runtime as the number of processors grows. The subsequent rankings alternate between reverse sorted and randomized inputs, as expected, considering they require a greater number of swaps. Reverse sorted input is relatively insensitive to the number of processors, as it consistently involves the maximum number of swaps.

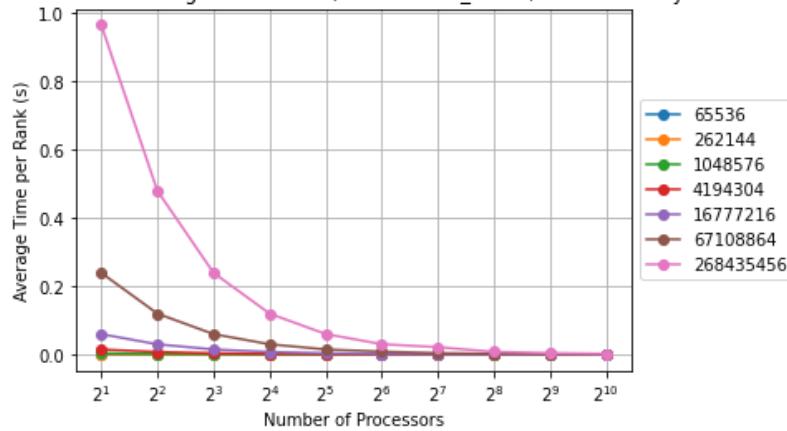
FIX ME!

RANDOM INPUT ARRAY

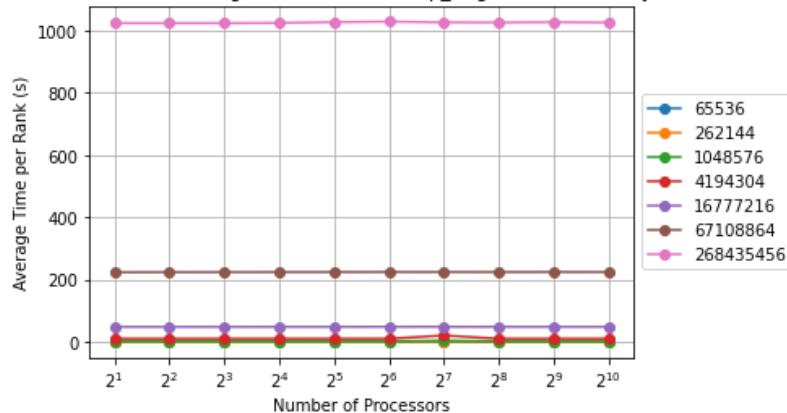




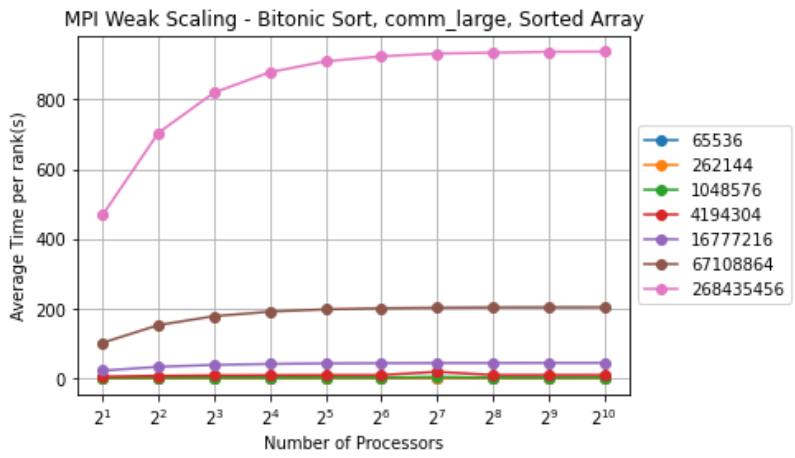
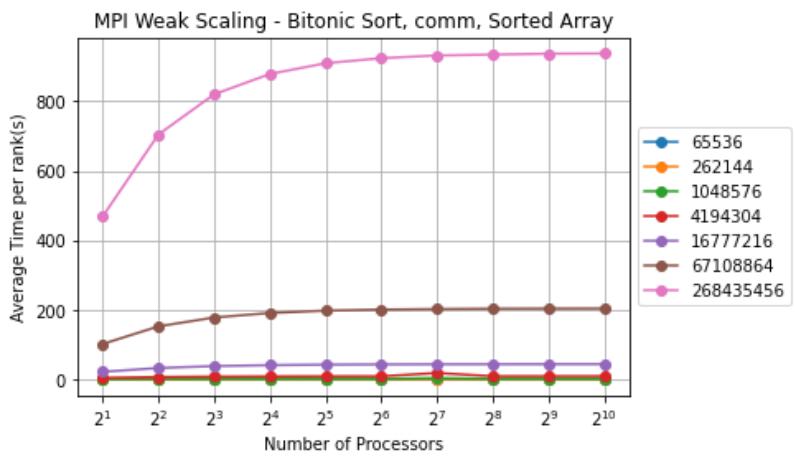
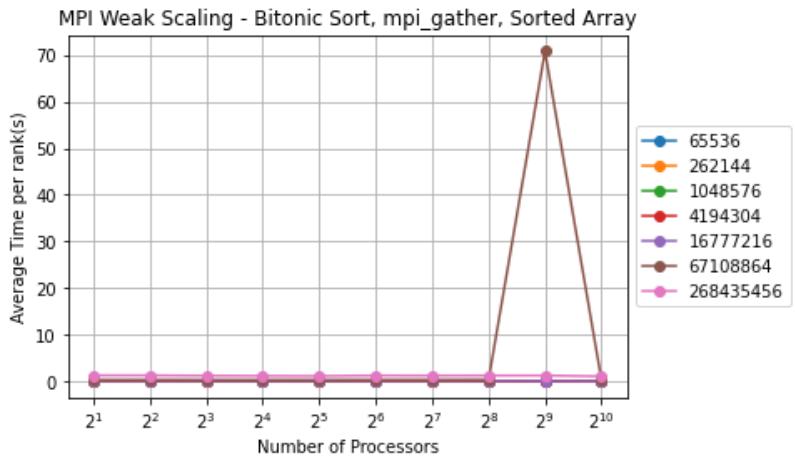
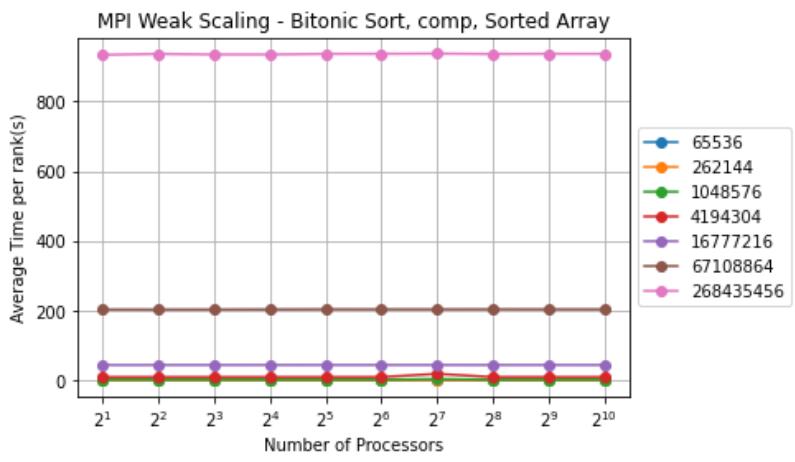
MPI Weak Scaling - Bitonic Sort, correctness_check, Random Array

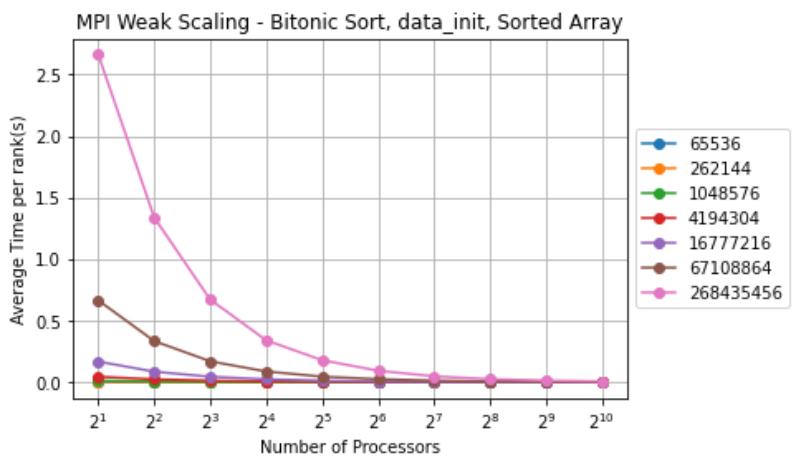
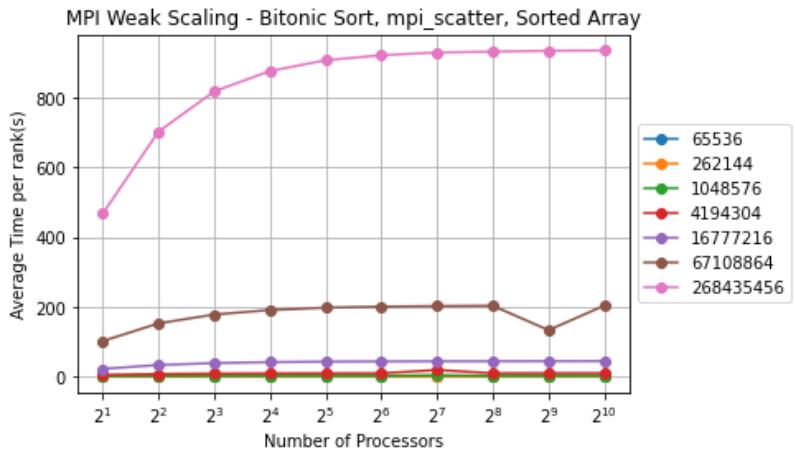
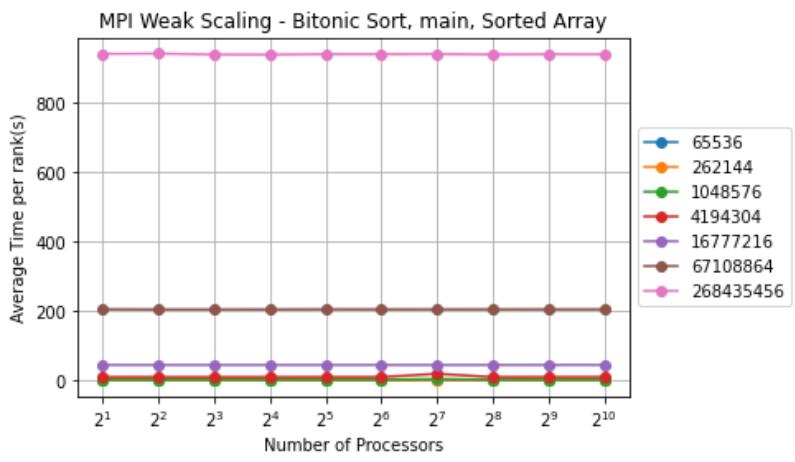


MPI Weak Scaling - Bitonic Sort, comp_large, Random Array

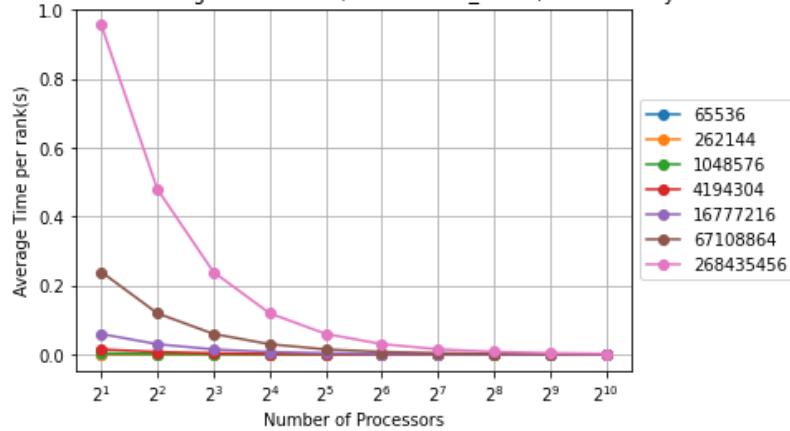


SORTED INPUT ARRAY

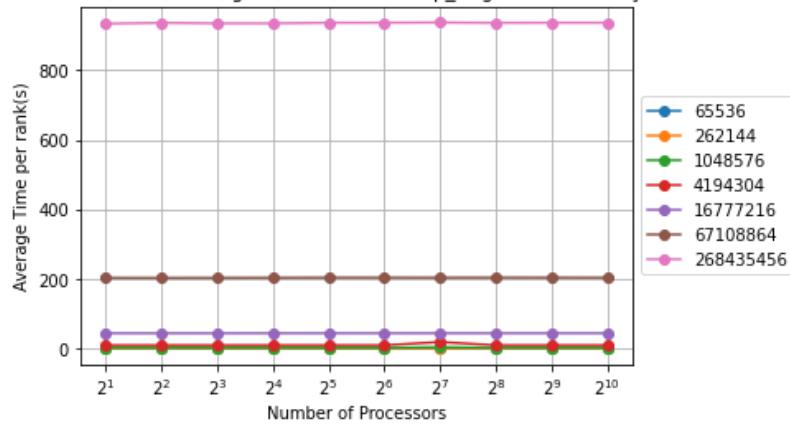




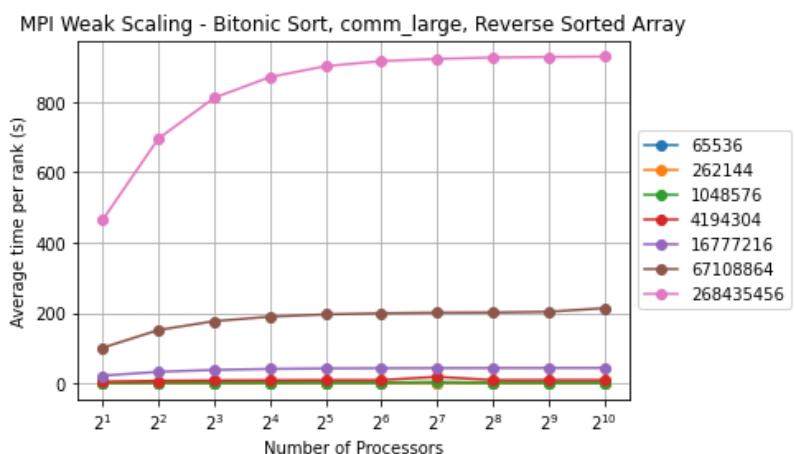
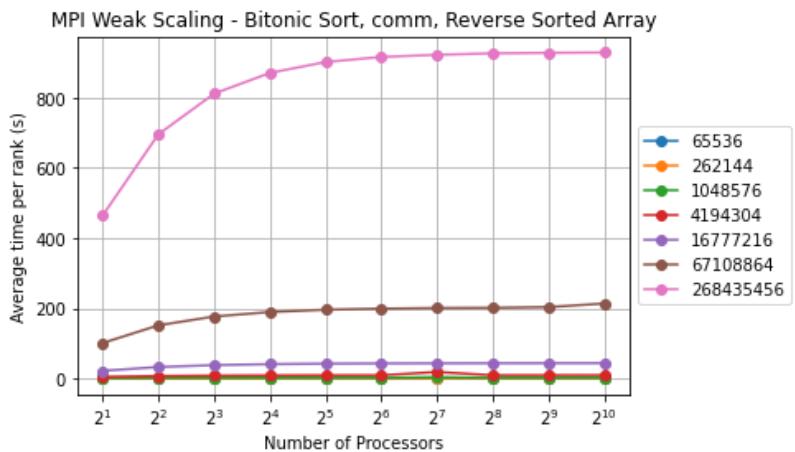
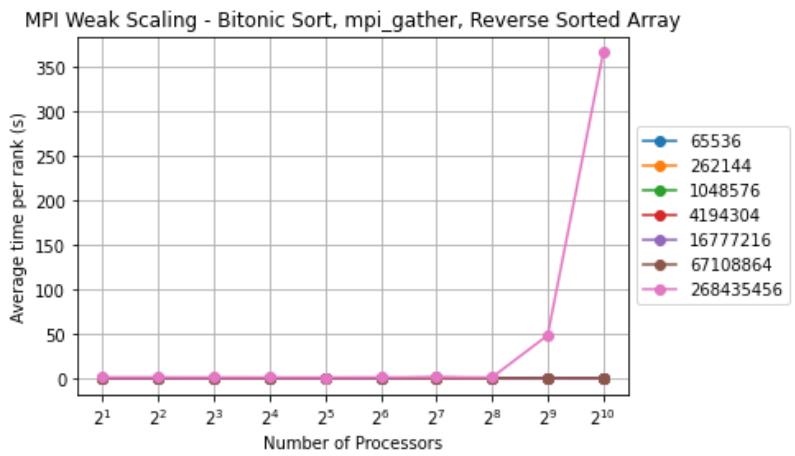
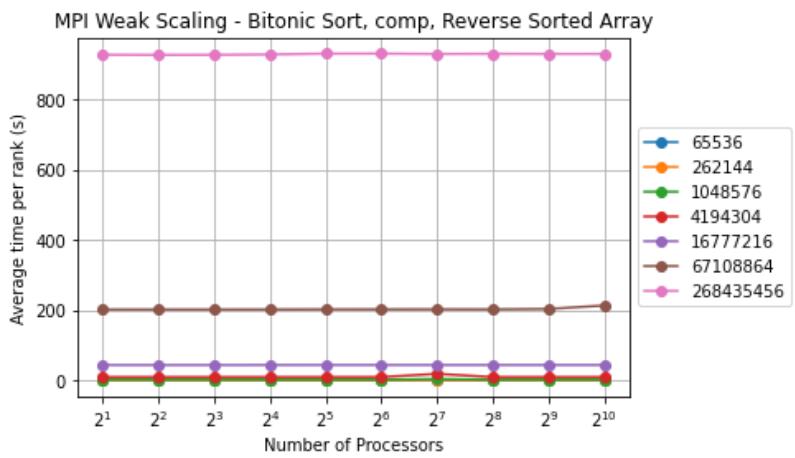
MPI Weak Scaling - Bitonic Sort, correctness_check, Sorted Array

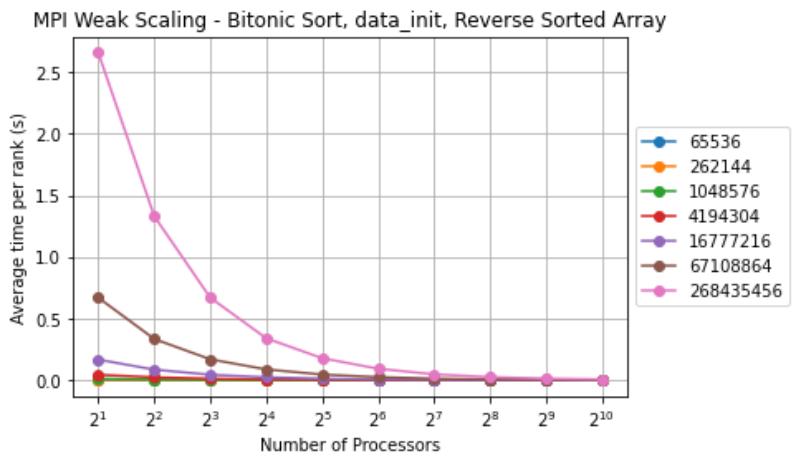
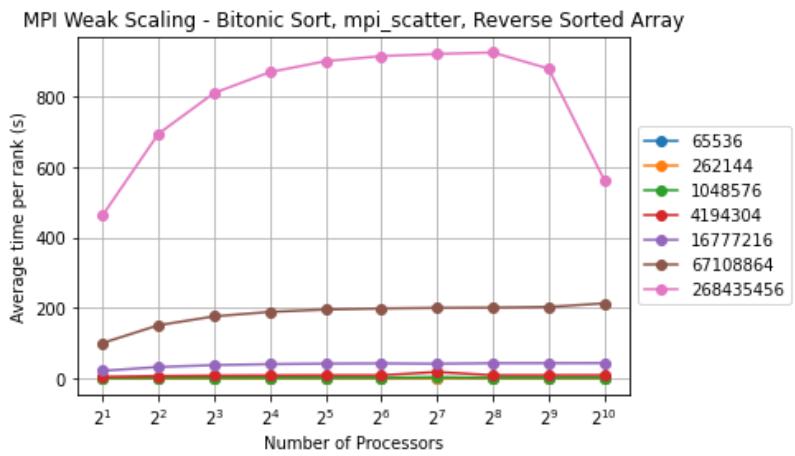
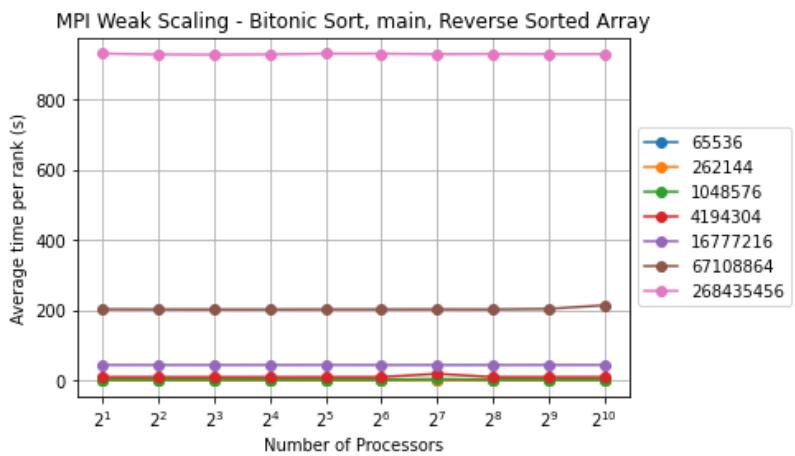


MPI Weak Scaling - Bitonic Sort, comp_large, Sorted Array

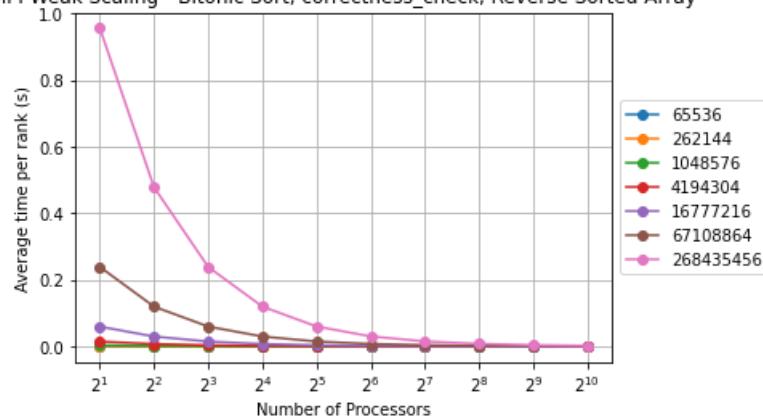


REVERSE SORTED INPUT ARRAY

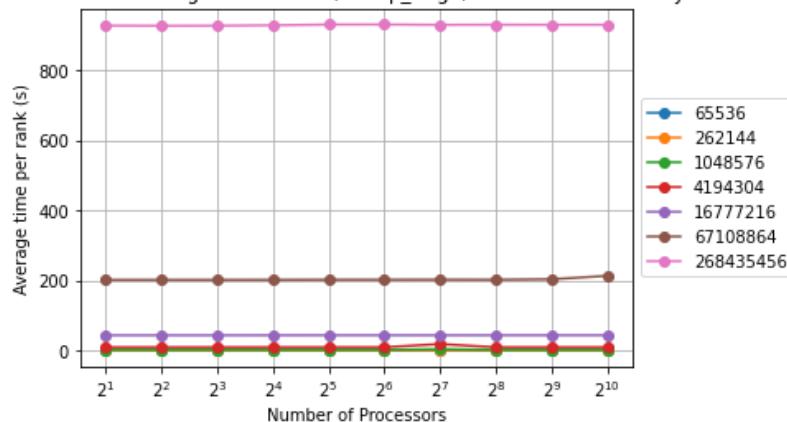




MPI Weak Scaling - Bitonic Sort, correctness_check, Reverse Sorted Array

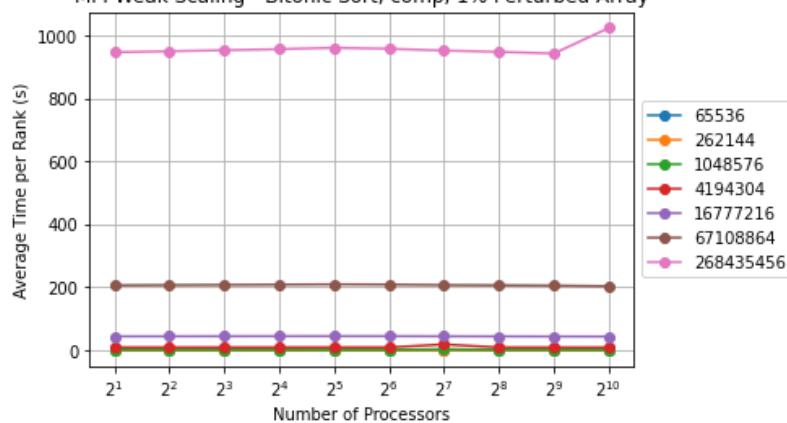


MPI Weak Scaling - Bitonic Sort, comp_large, Reverse Sorted Array

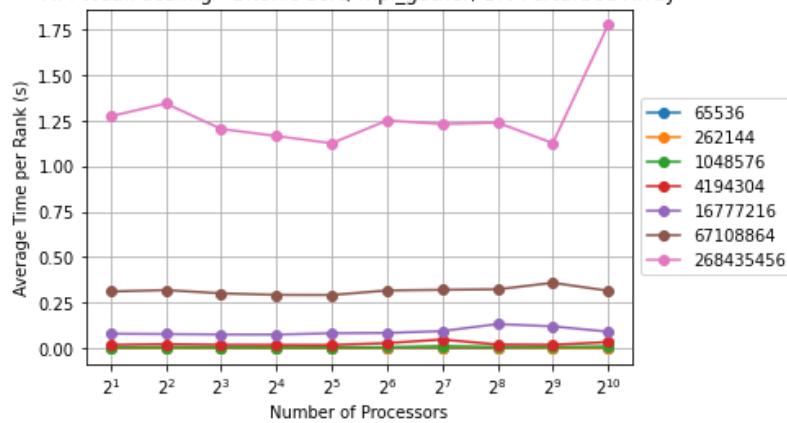


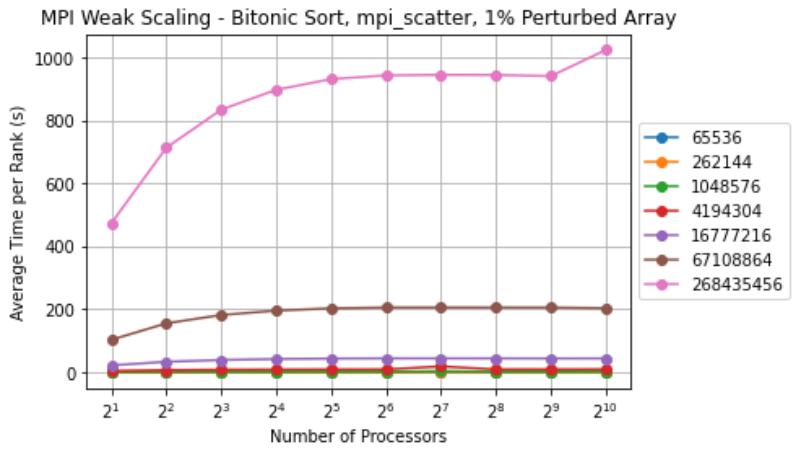
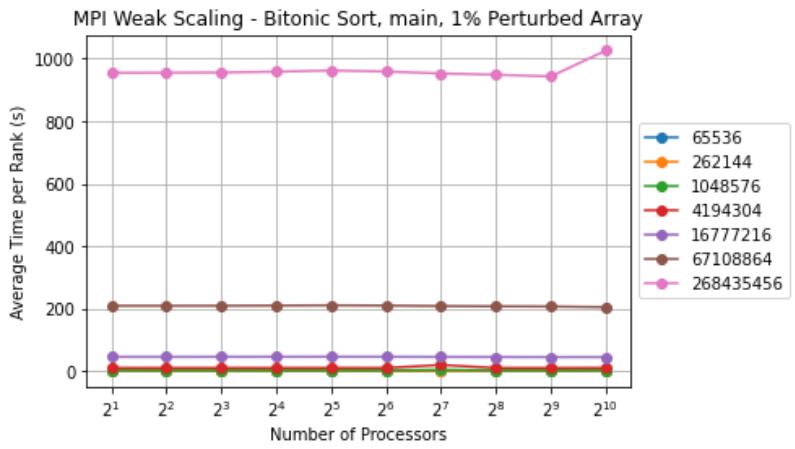
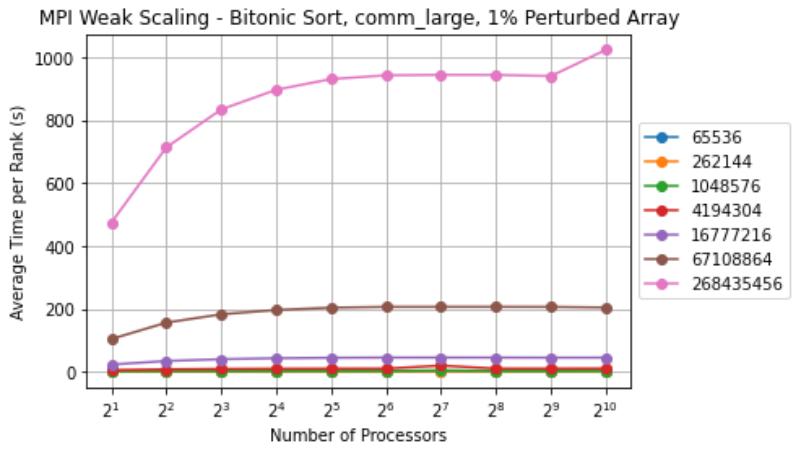
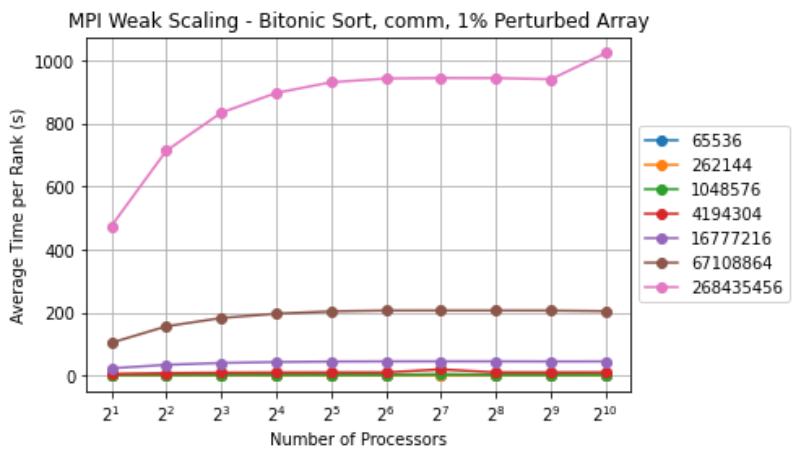
1% PERTURBED INPUT ARRAY

MPI Weak Scaling - Bitonic Sort, comp, 1% Perturbed Array

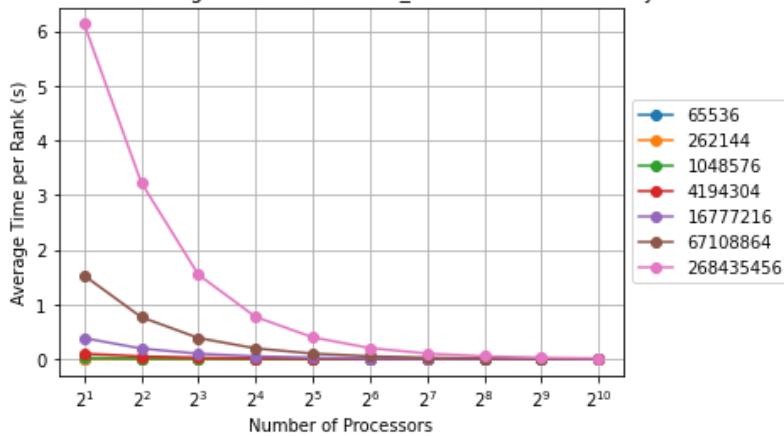


MPI Weak Scaling - Bitonic Sort, mpi_gather, 1% Perturbed Array

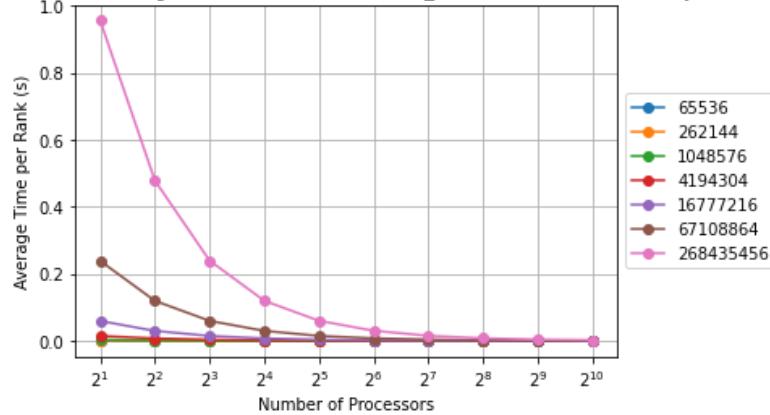




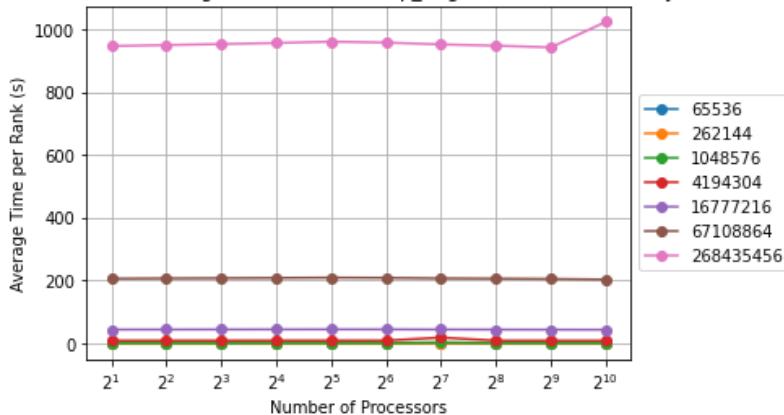
MPI Weak Scaling - Bitonic Sort, data_init, 1% Perturbed Array



MPI Weak Scaling - Bitonic Sort, correctness_check, 1% Perturbed Array



MPI Weak Scaling - Bitonic Sort, comp_large, 1% Perturbed Array



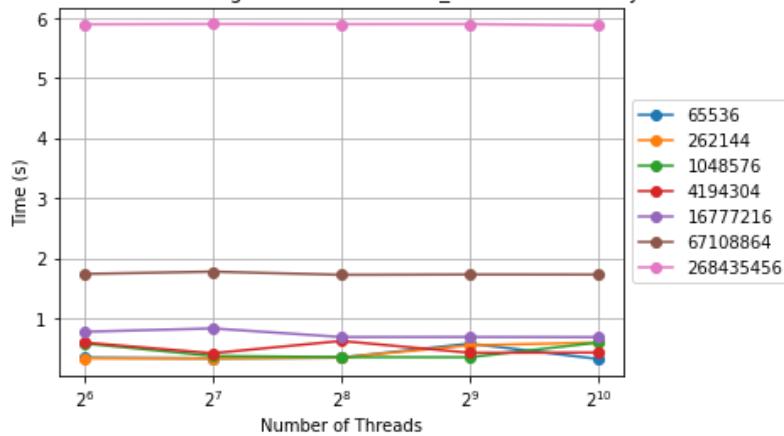
CUDA

For this analysis, I examined weak scaling using 2^{16} elements across three input types (sorted, reverse sorted, and randomized). The results are unexpected, as all average times increase with the number of threads used. One would anticipate a decrease for some time, given that parallelization should enhance efficiency. Furthermore, similar to MPI, we observe that "Sorted" consistently ranks as the fastest, as expected, since no swaps are required. "Reverse sorted" and "randomized" exhibit variations in their next fastest rankings, which aligns with expectations as they entail more swaps. The graph suggests that for this input size (65536), parallelization is not beneficial. With additional graphs featuring different input sizes, we can discern the general trend.

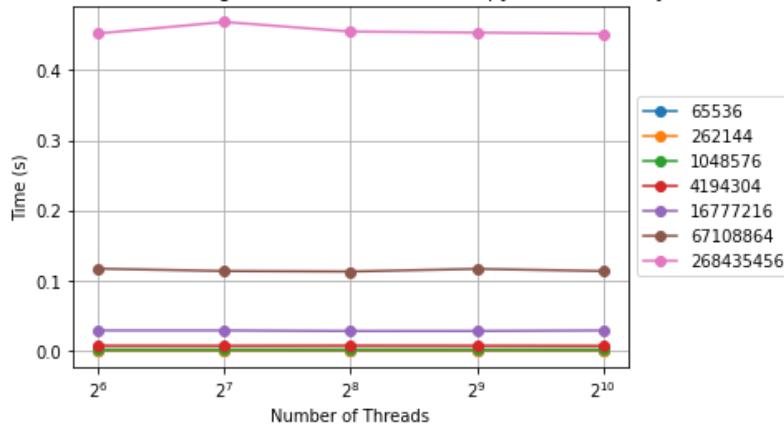
FIX ME!

RANDOM INPUT ARRAY

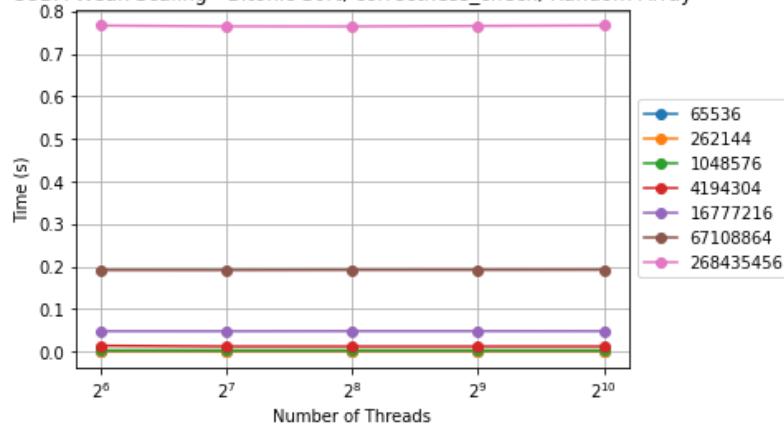
CUDA Weak Scaling - Bitonic Sort, data_init, Random Array



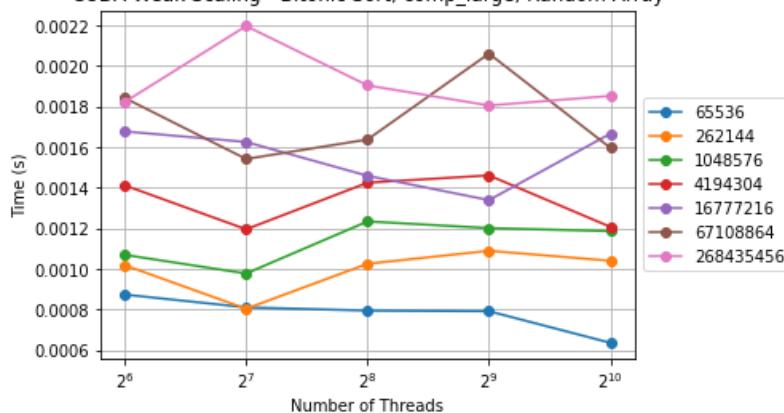
CUDA Weak Scaling - Bitonic Sort, cudaMemcpy, Random Array

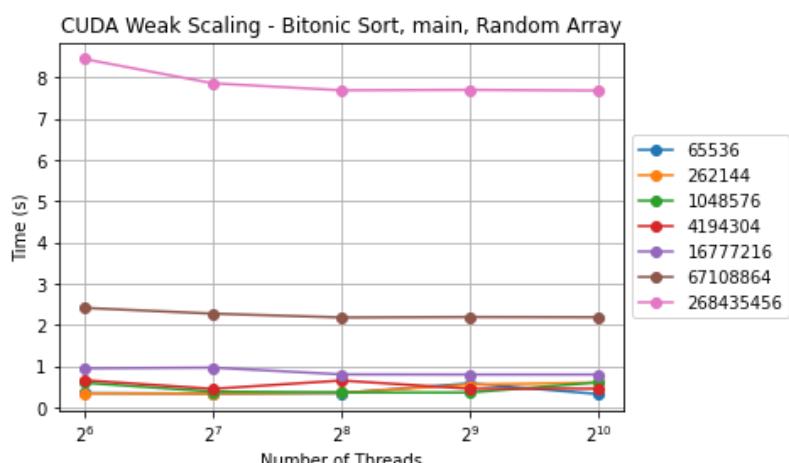
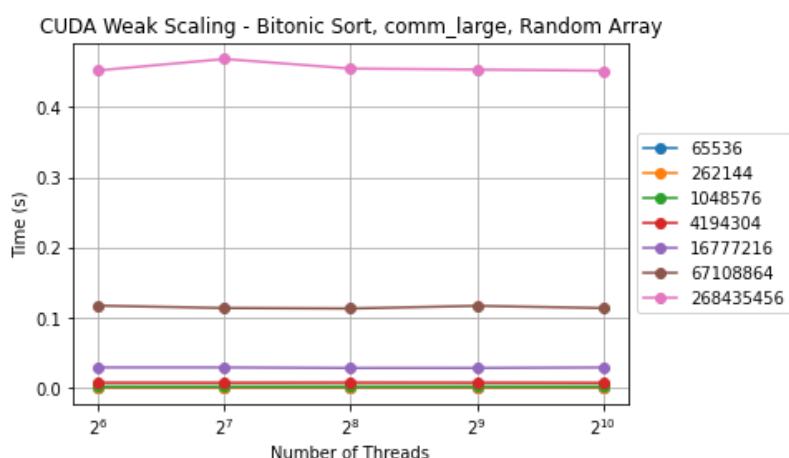
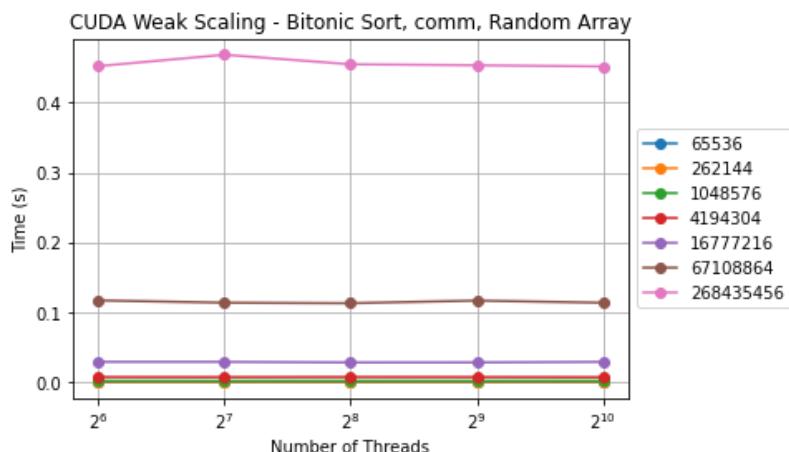
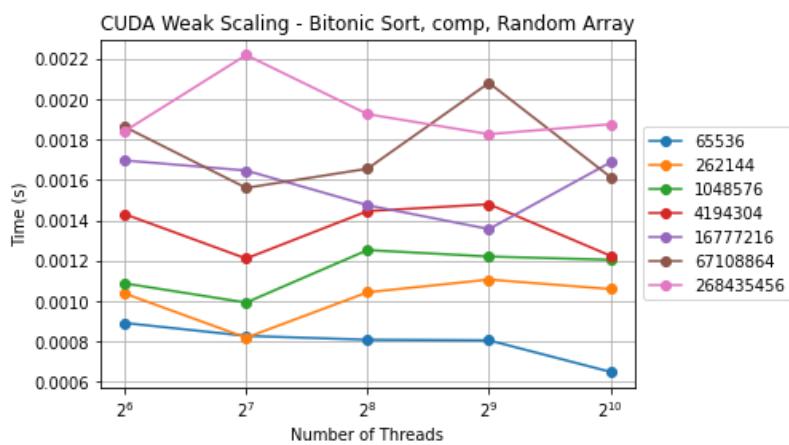


CUDA Weak Scaling - Bitonic Sort, correctness_check, Random Array

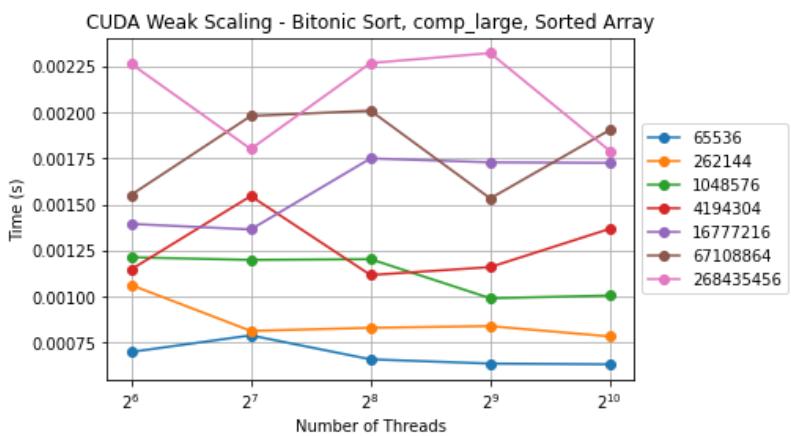
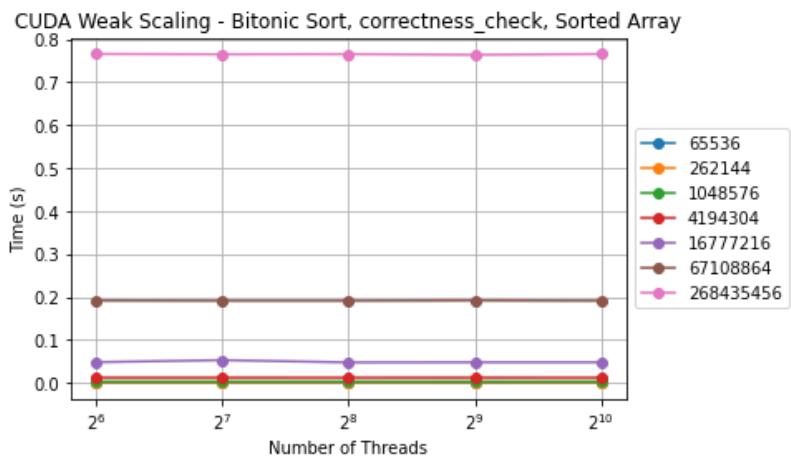
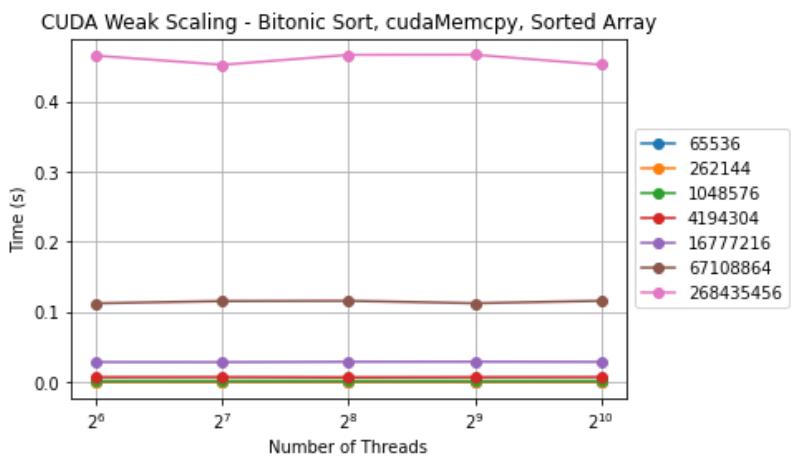
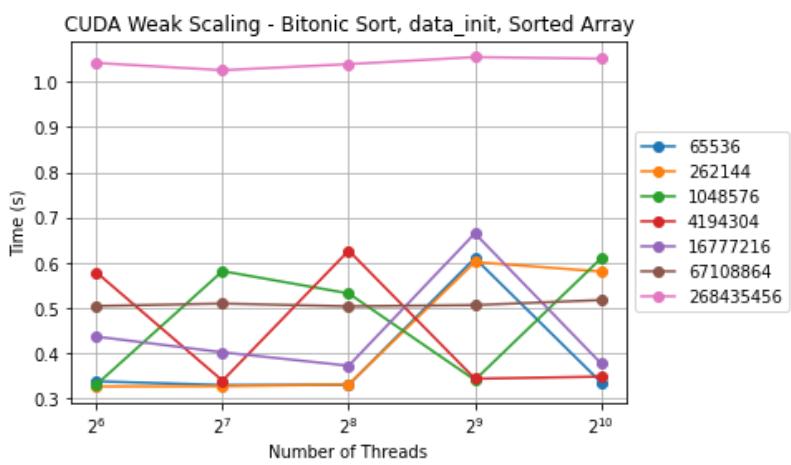


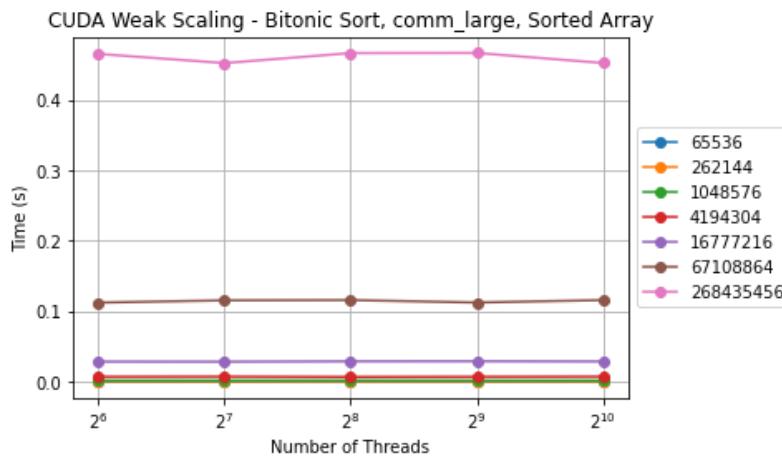
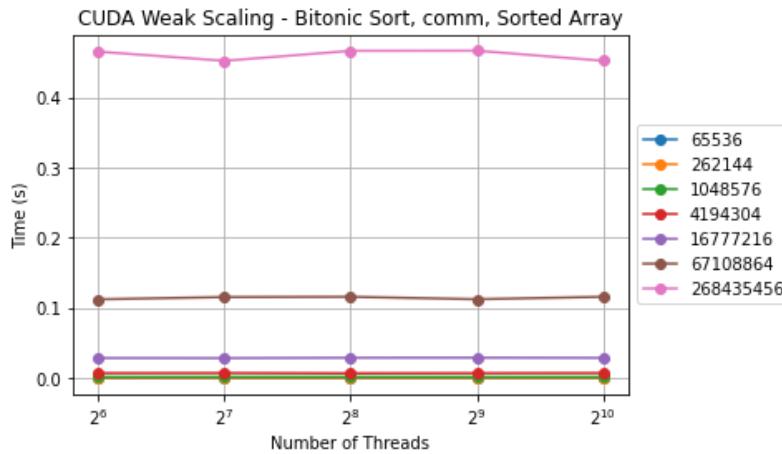
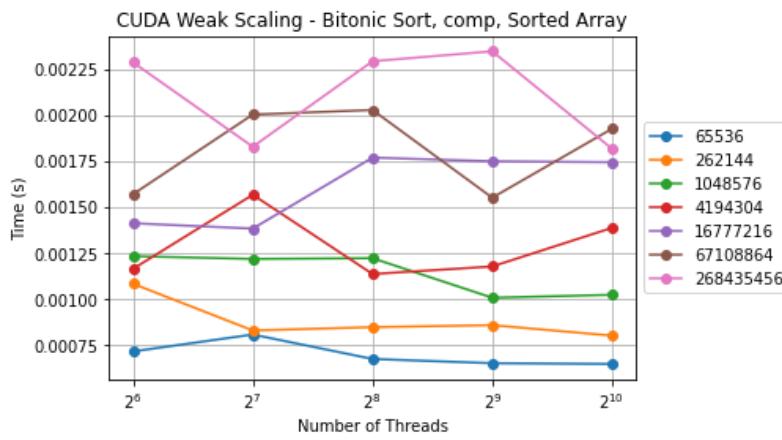
CUDA Weak Scaling - Bitonic Sort, comp_large, Random Array





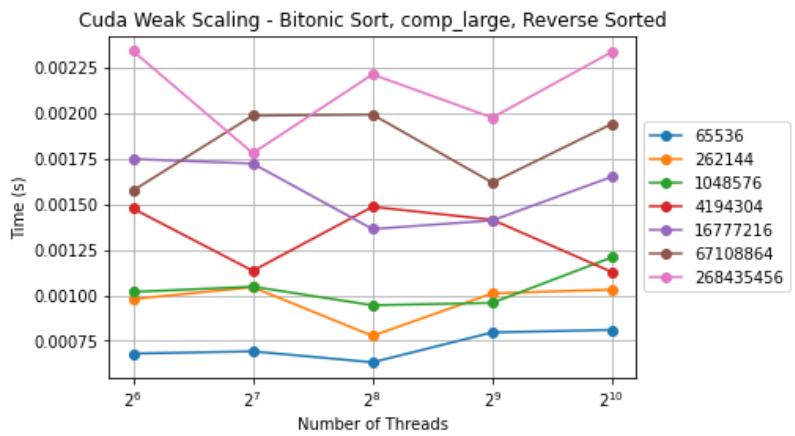
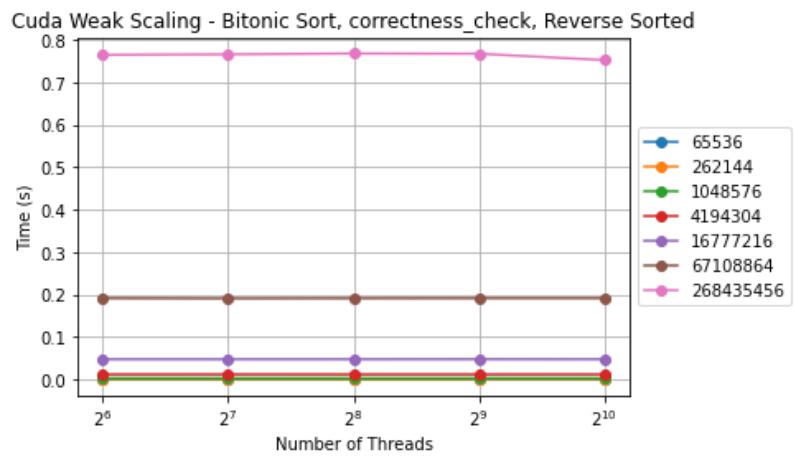
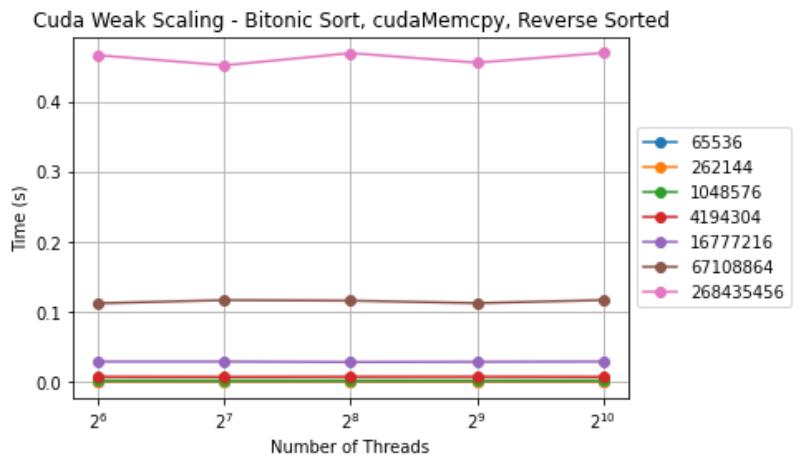
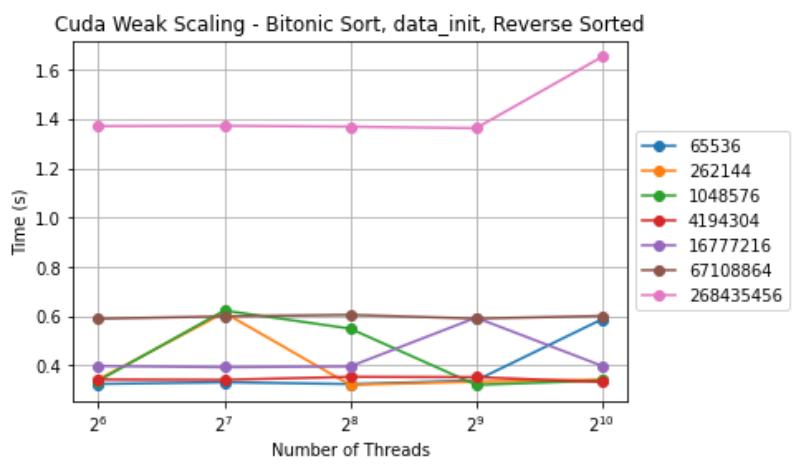
SORTED INPUT ARRAY

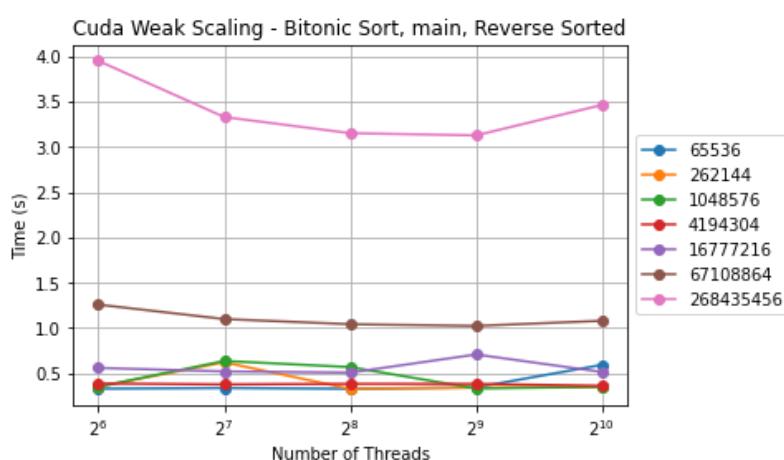
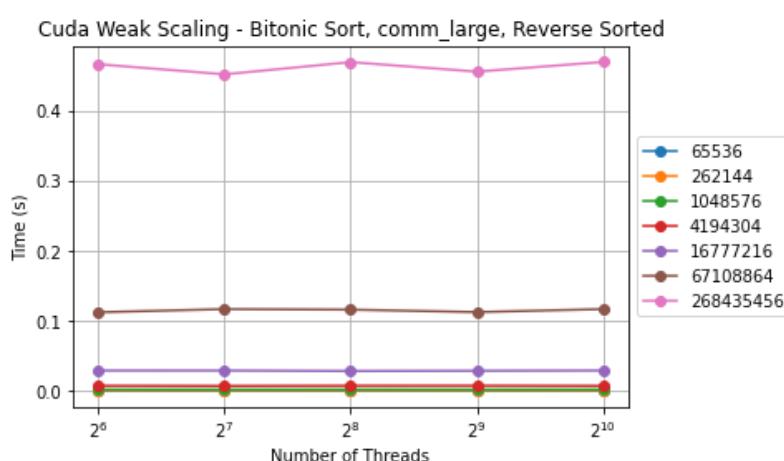
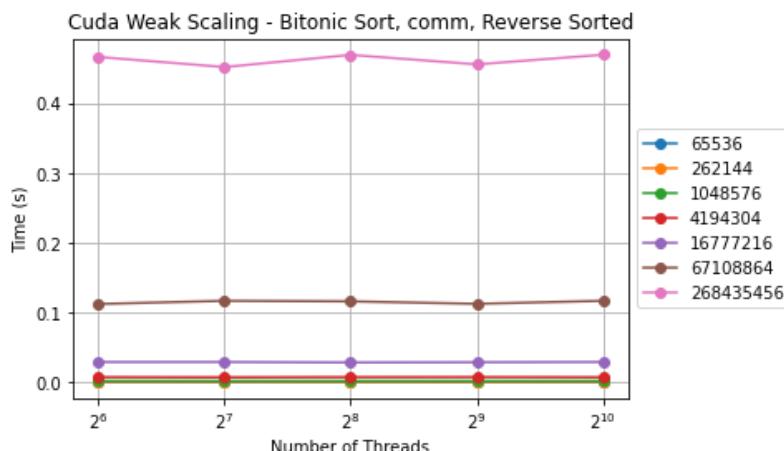
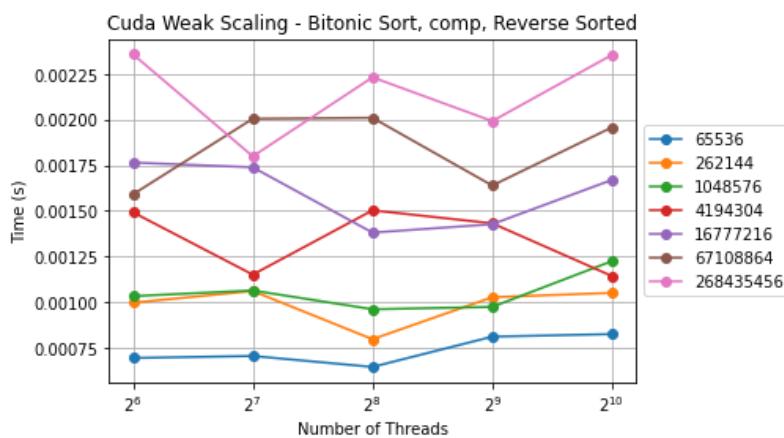




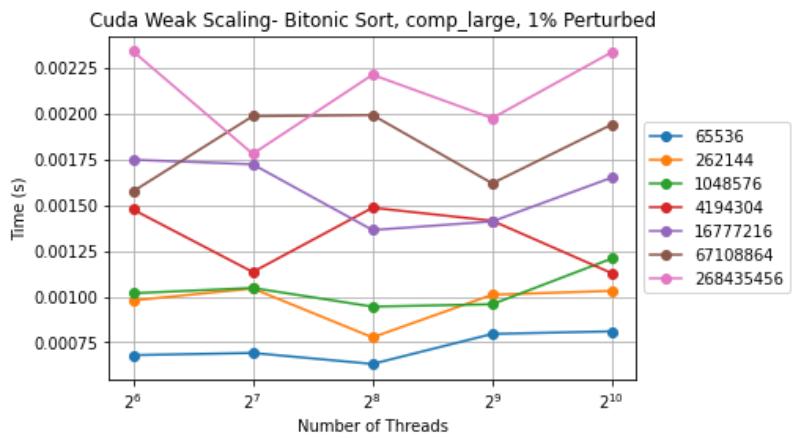
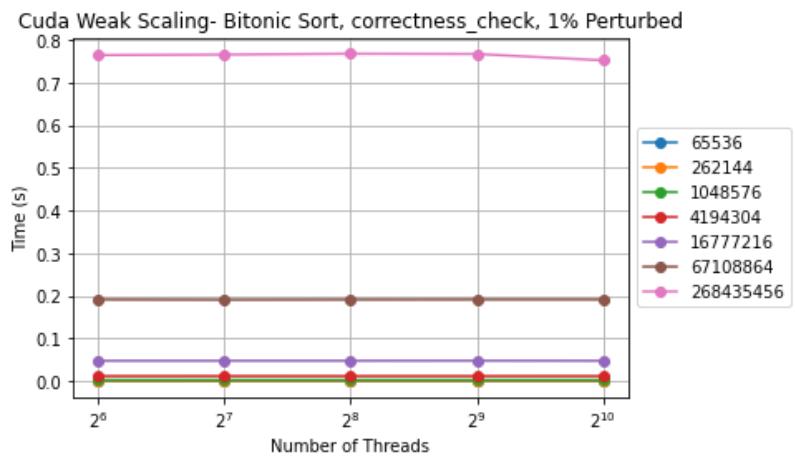
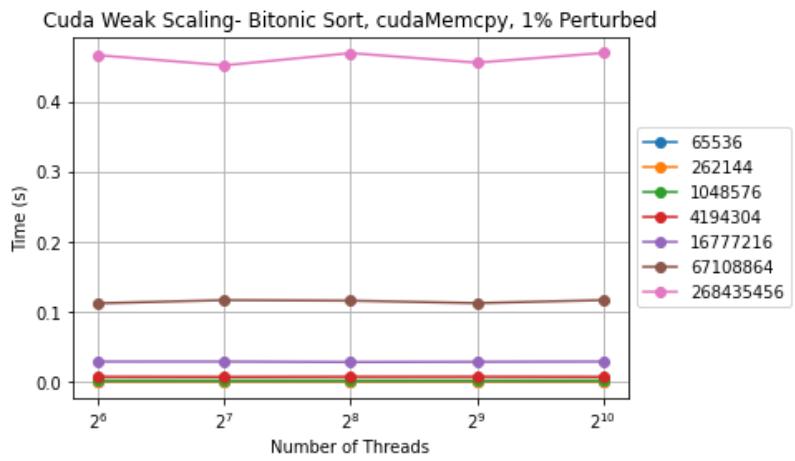
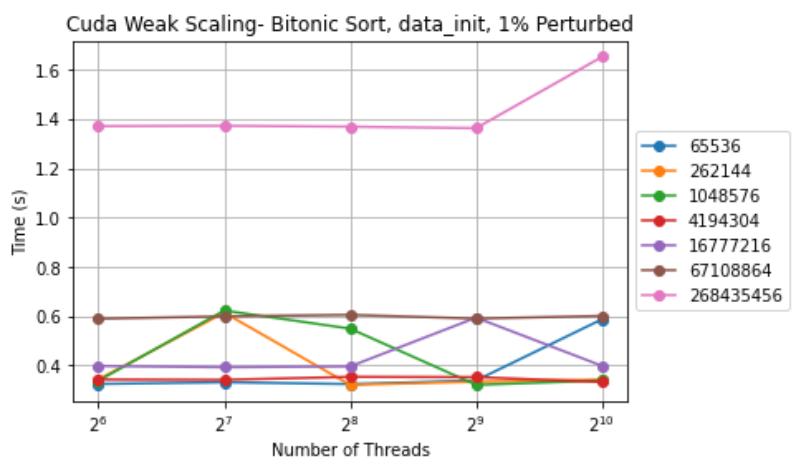
Average-Time-main-Weak Scaling

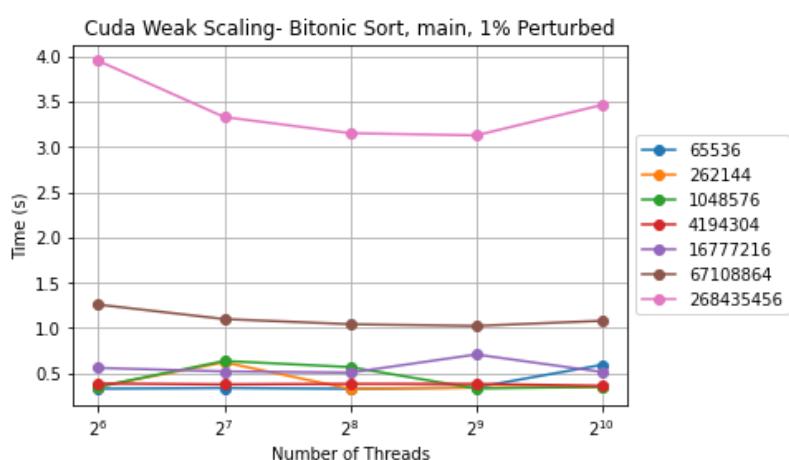
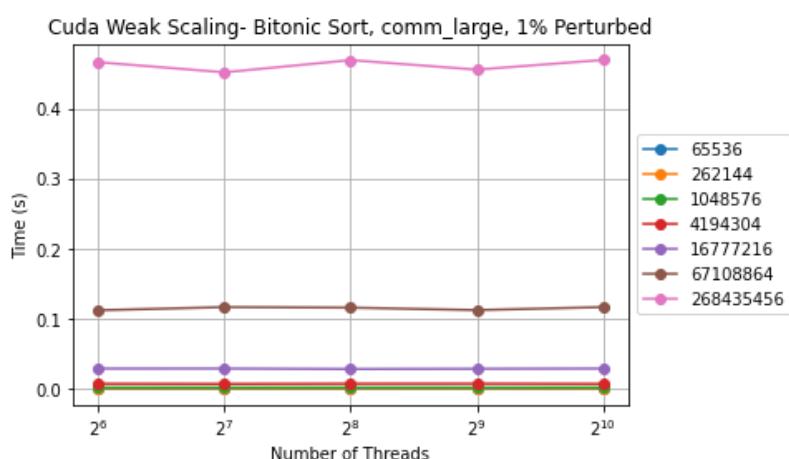
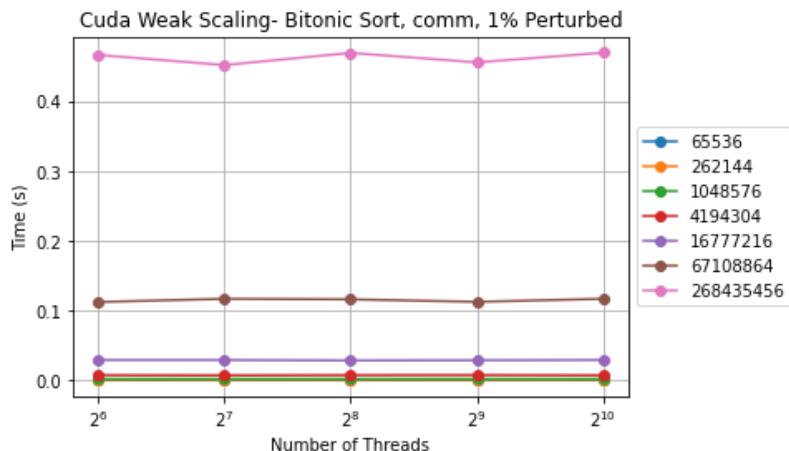
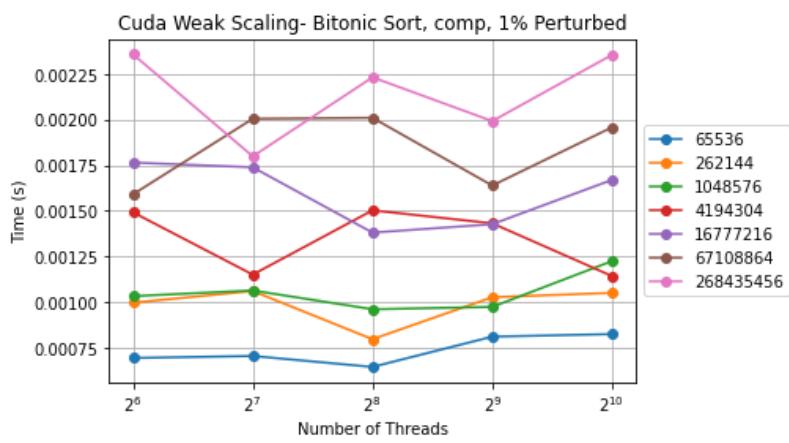
REVERSE SORTED INPUT ARRAY





1% PERTURBED INPUT ARRAY





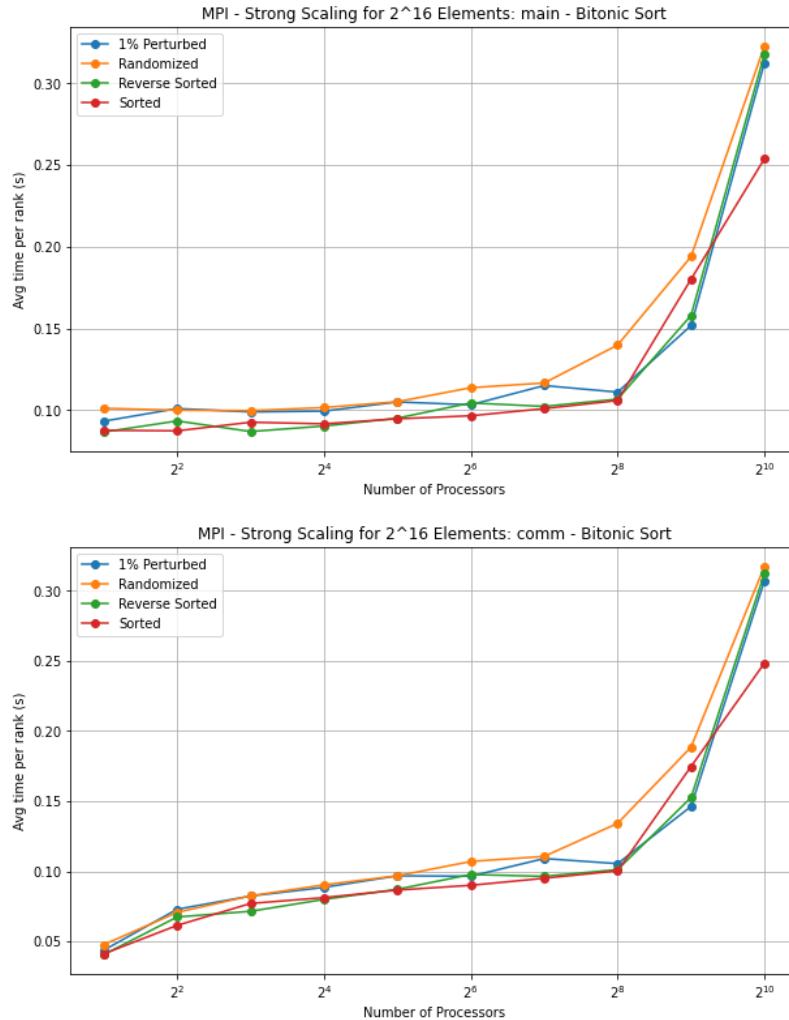
Weak Scaling

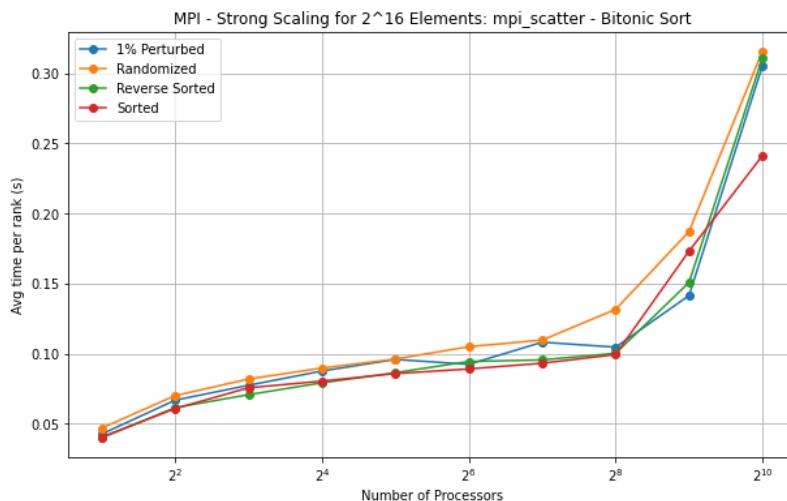
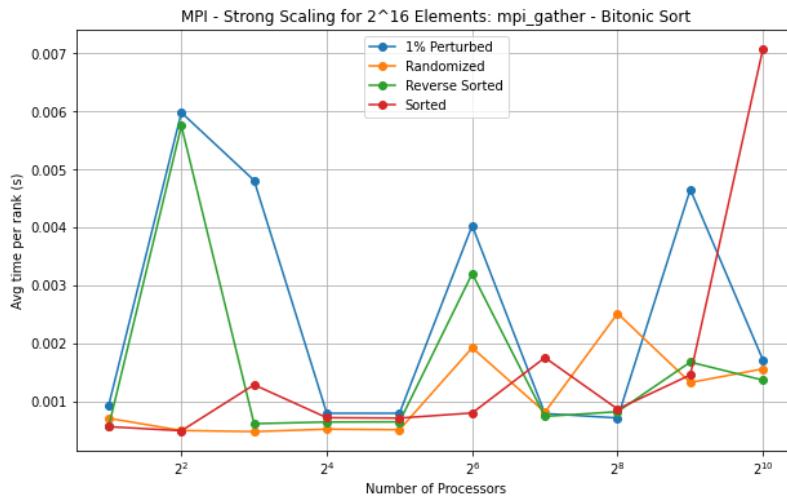
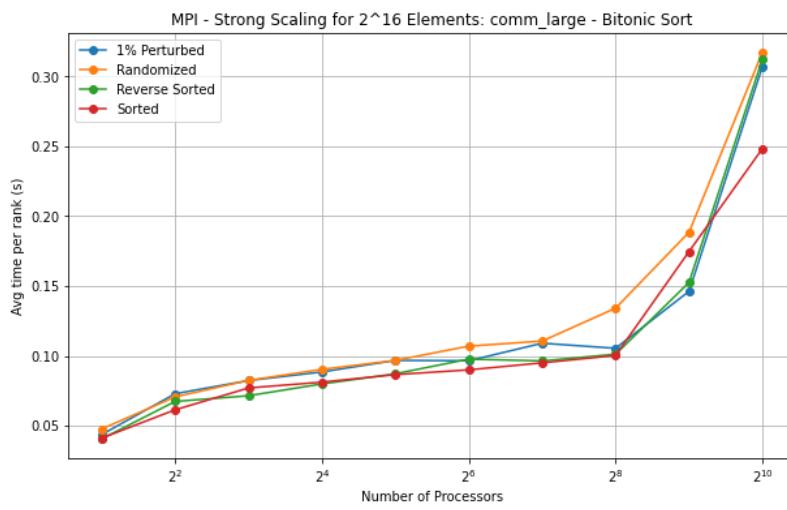
MPI

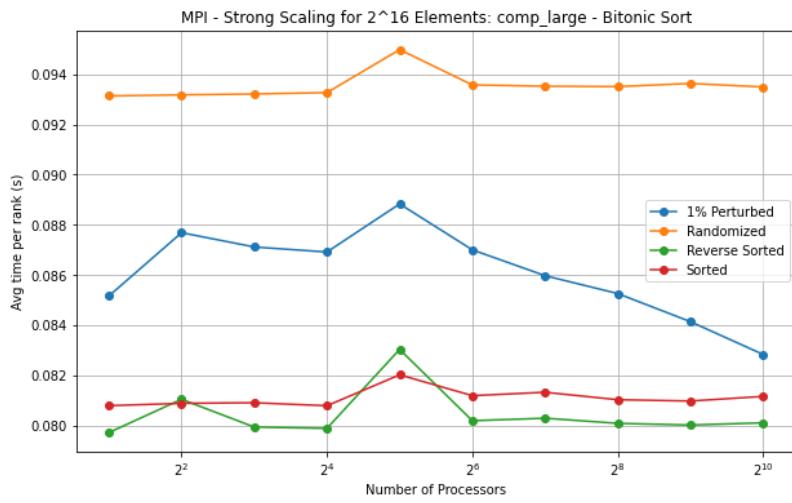
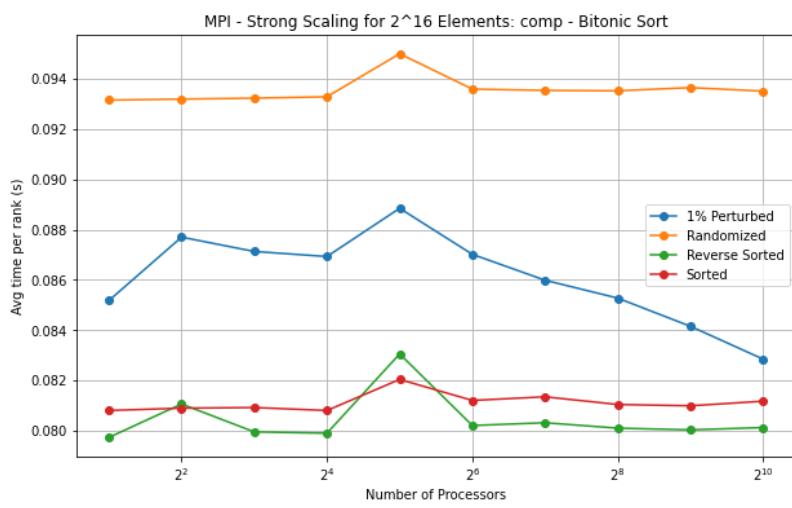
For the analysis, weak scaling was examined across the sorted input type. 2, 4, and 8 processors were run and it showed that there was an increase in average time. However, this was interesting as usually as the number of processors increases, the time should decrease due to parallelization and efficiency. However, since smaller numbers were run it might not be as easy to tell. Therefore, as more trials with more data is run, then the graph will show parallelization. The rest of the input types with the number of processors and threads will be run as the next time in order to gather more accurate data creating the various graphs.

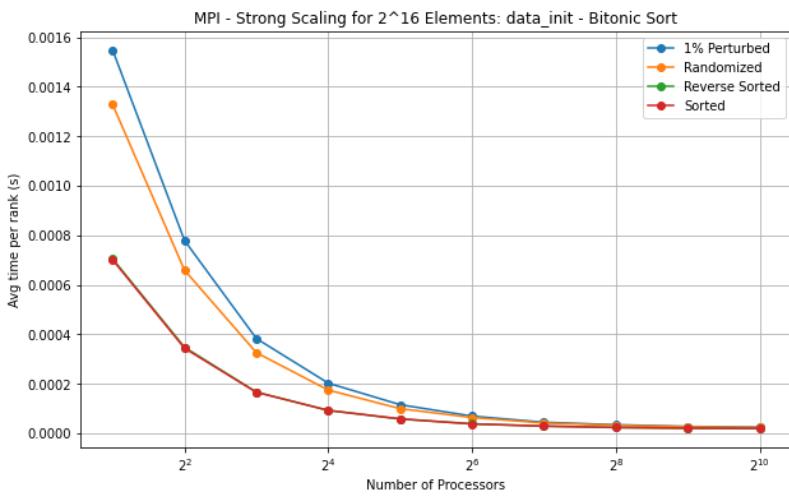
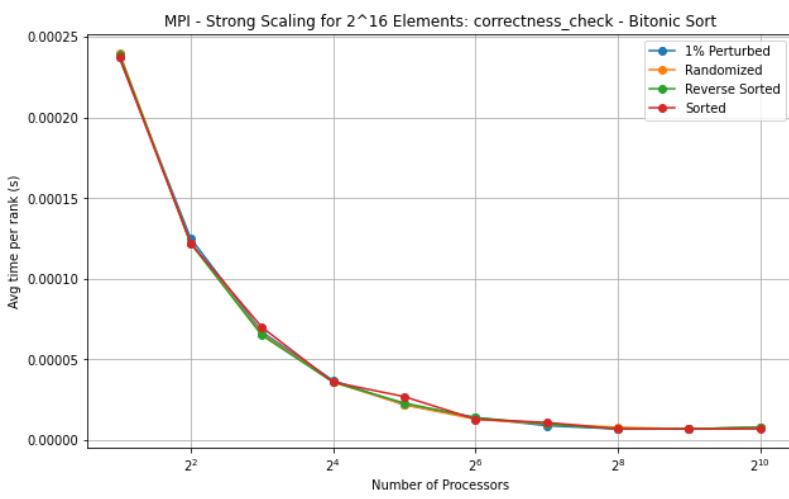
FIX ME!

2^{16} Elements

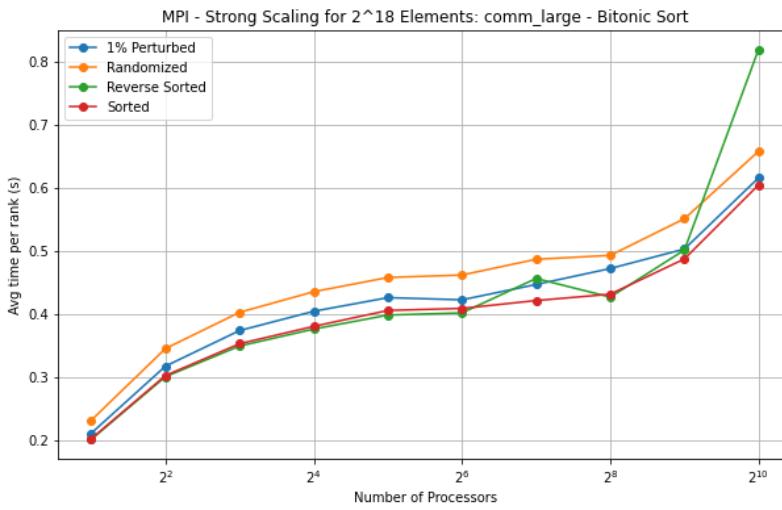
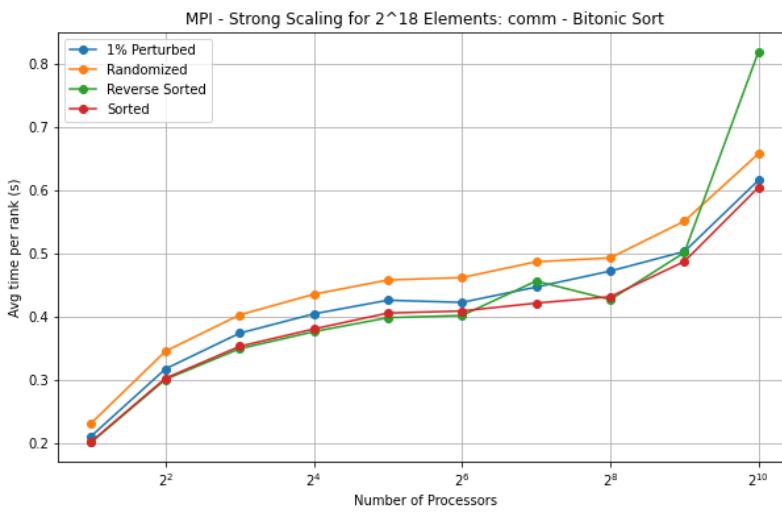
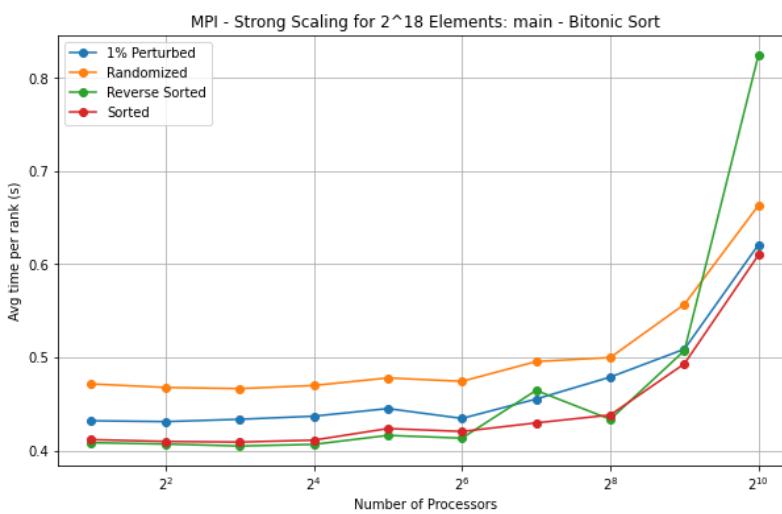


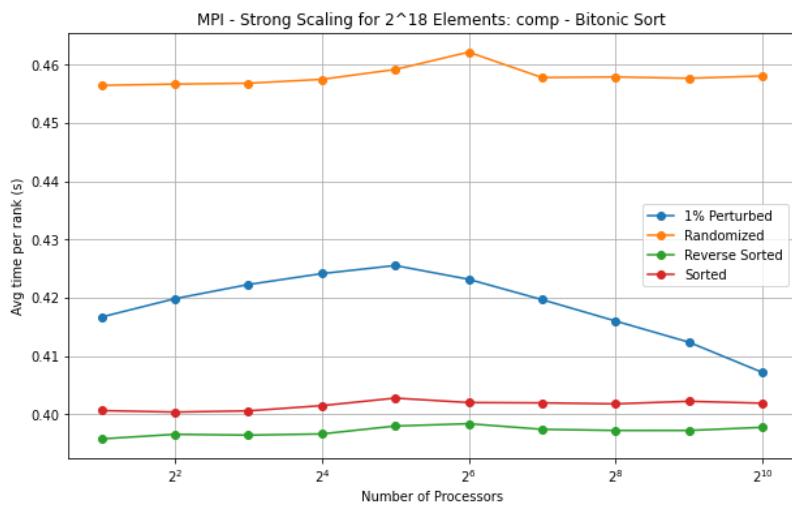
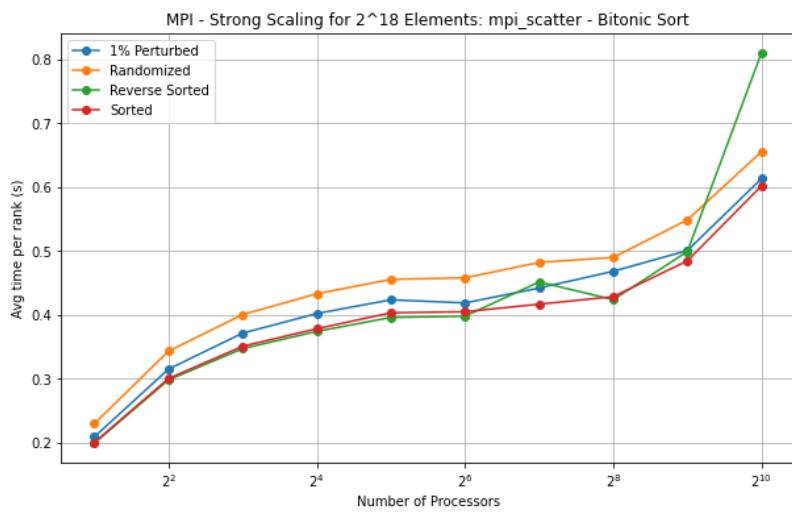
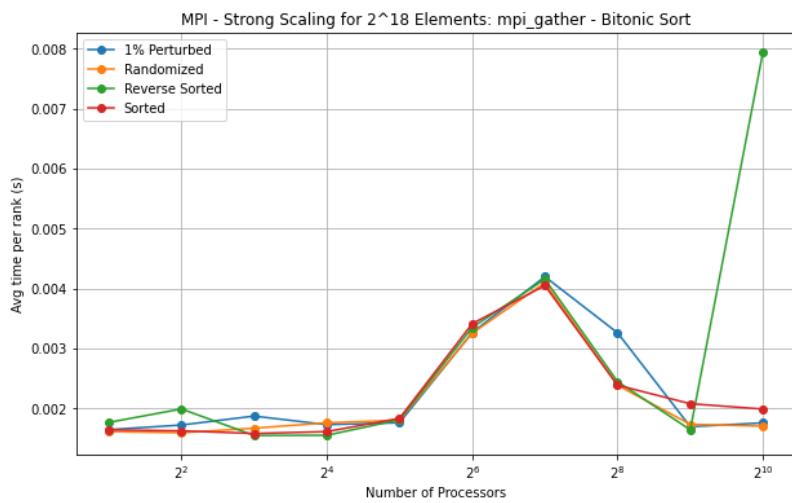


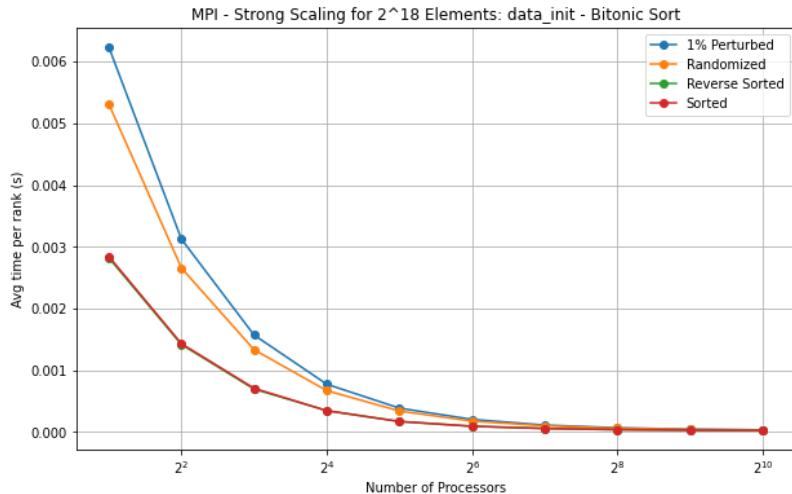
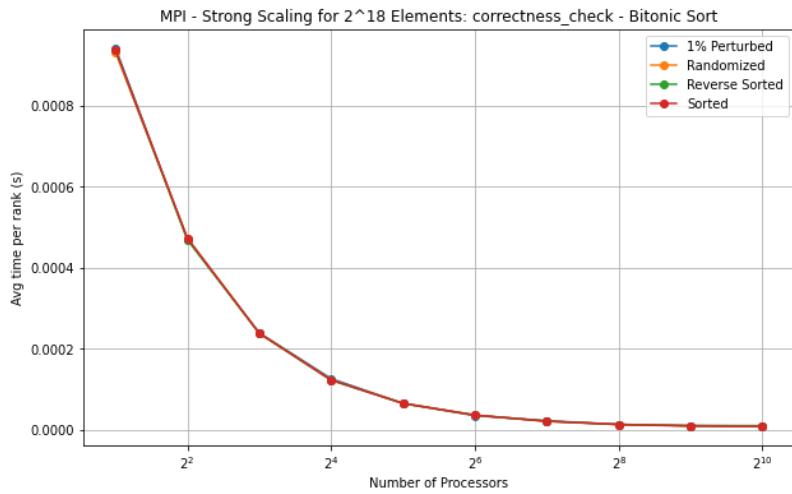
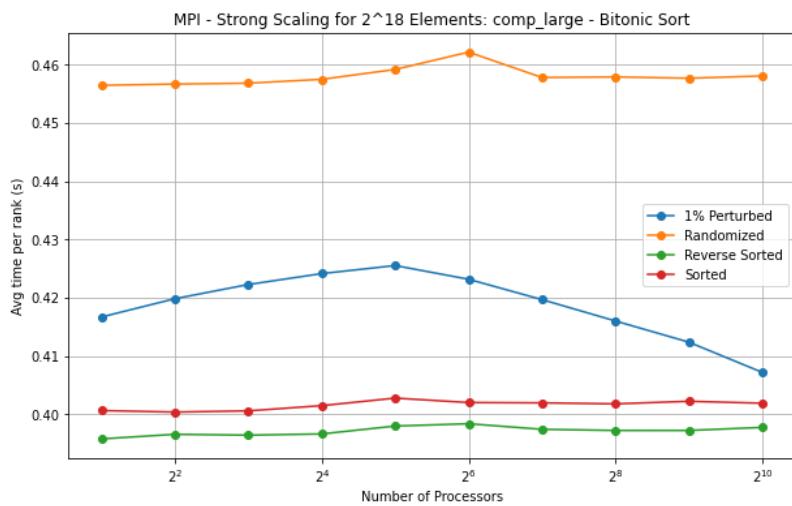




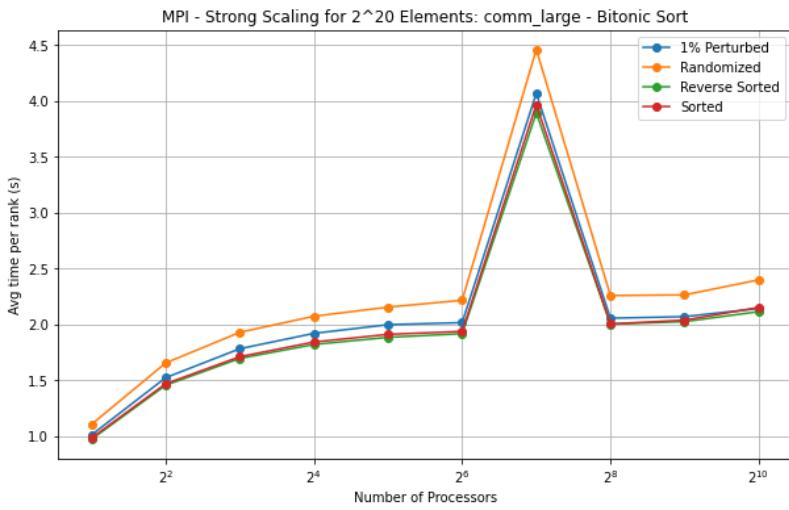
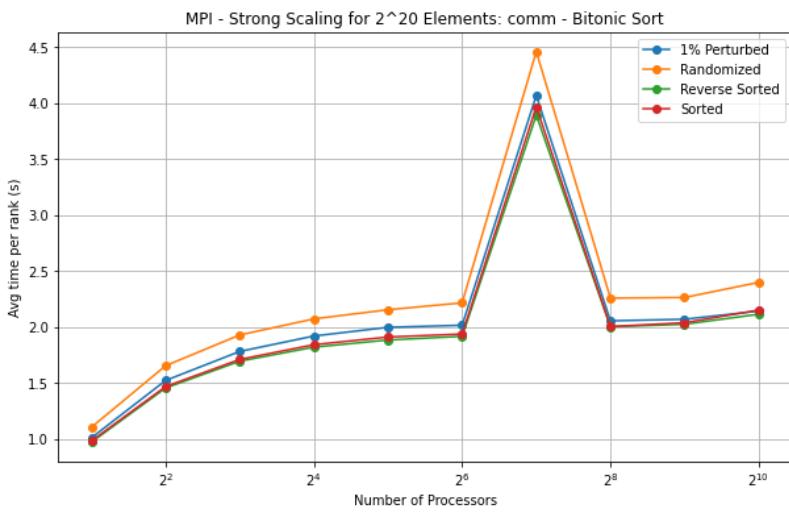
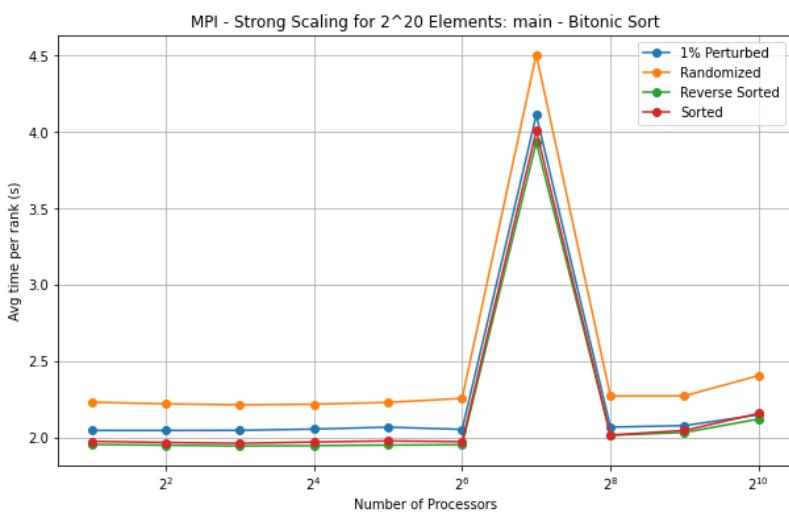
2^{18} Elements

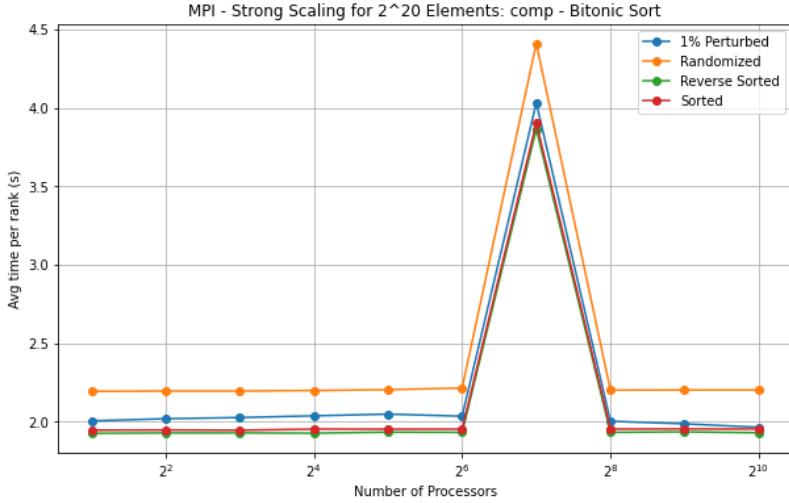
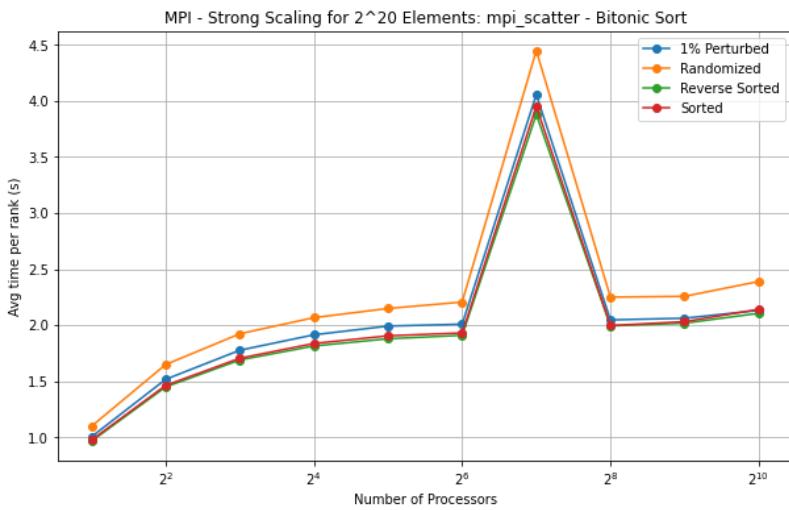
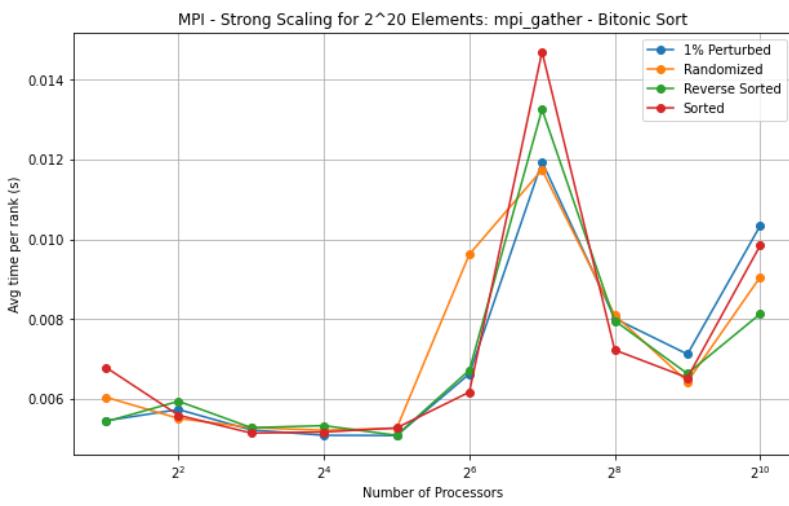


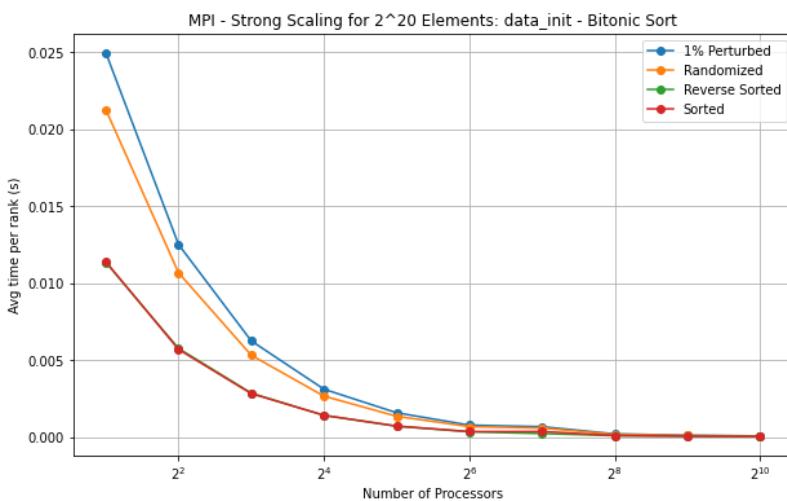
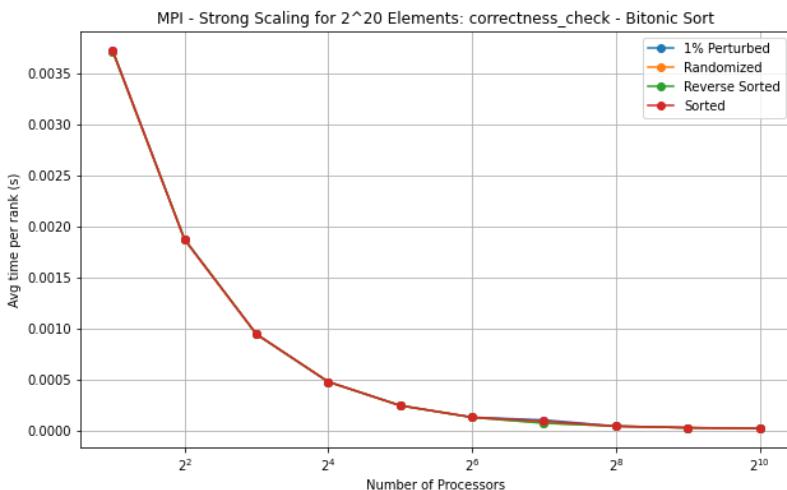
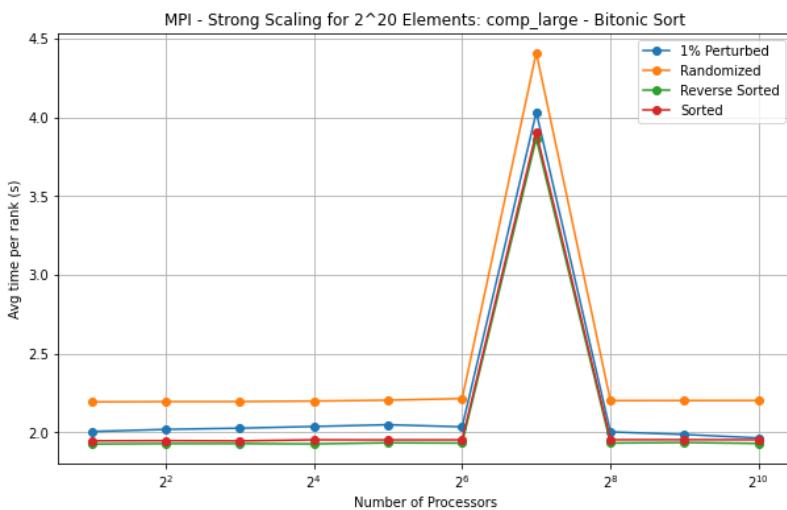




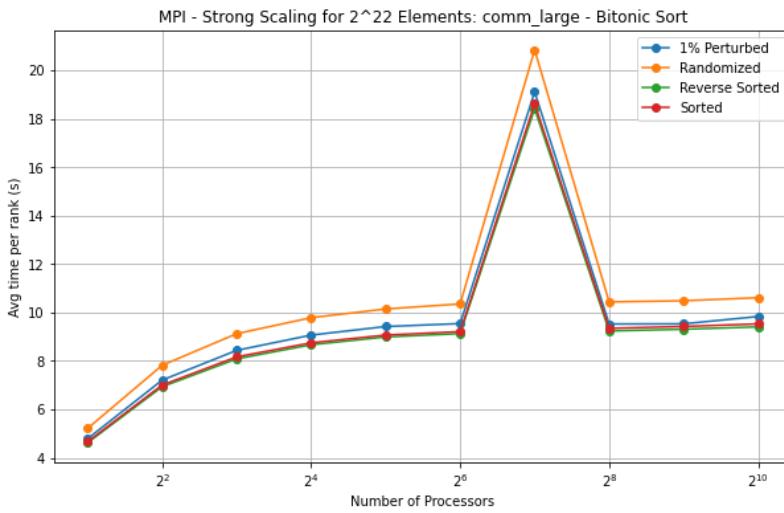
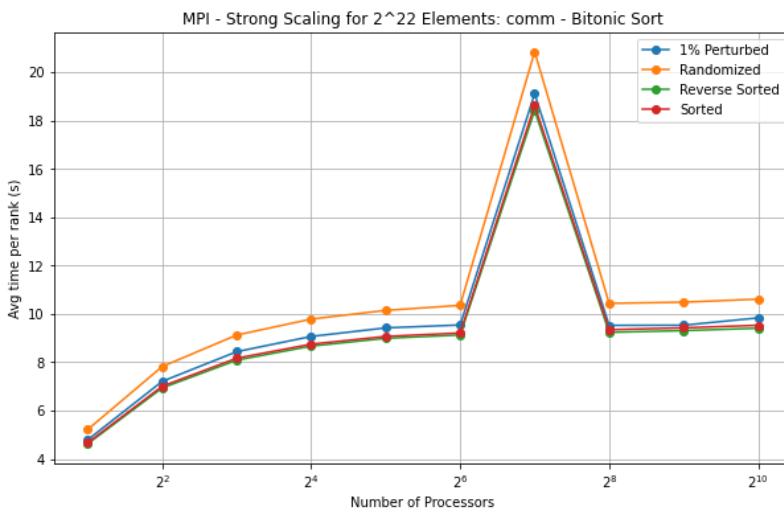
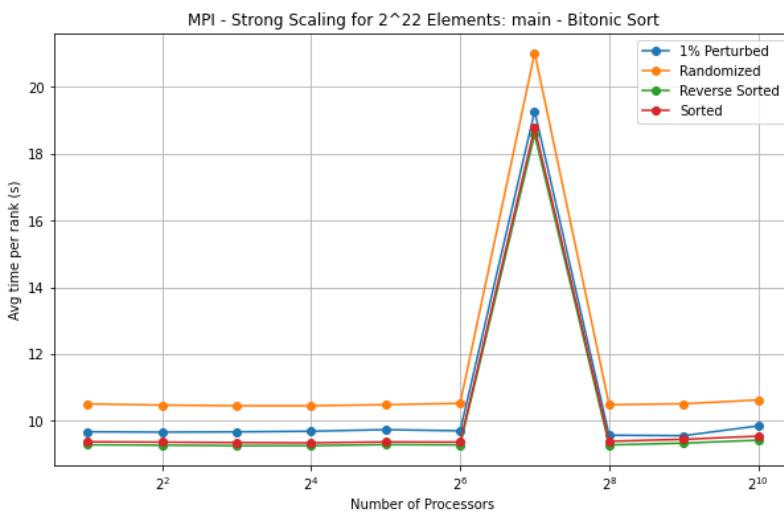
2^{20} Elements

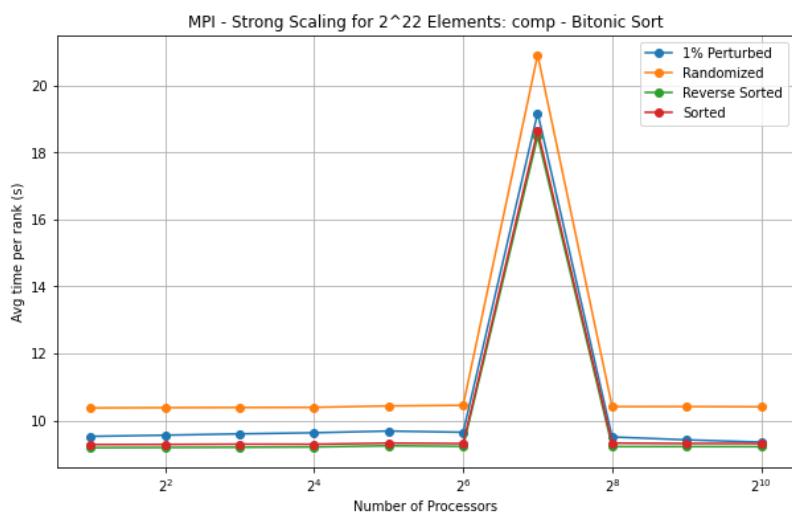
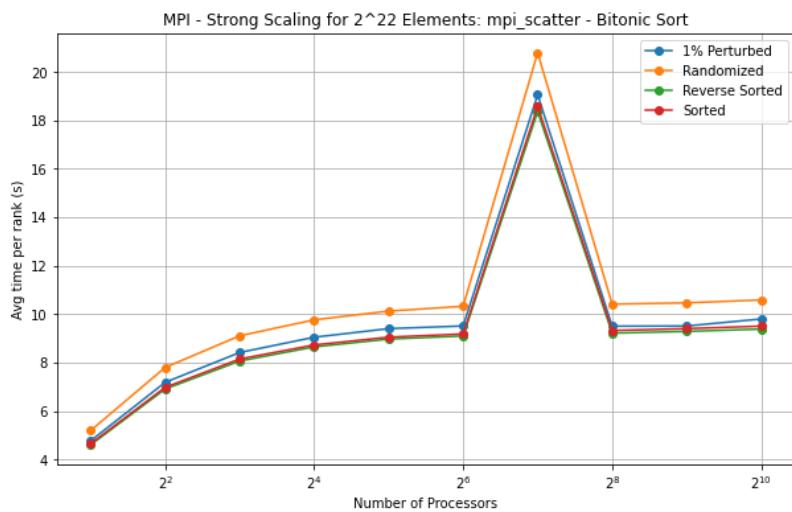
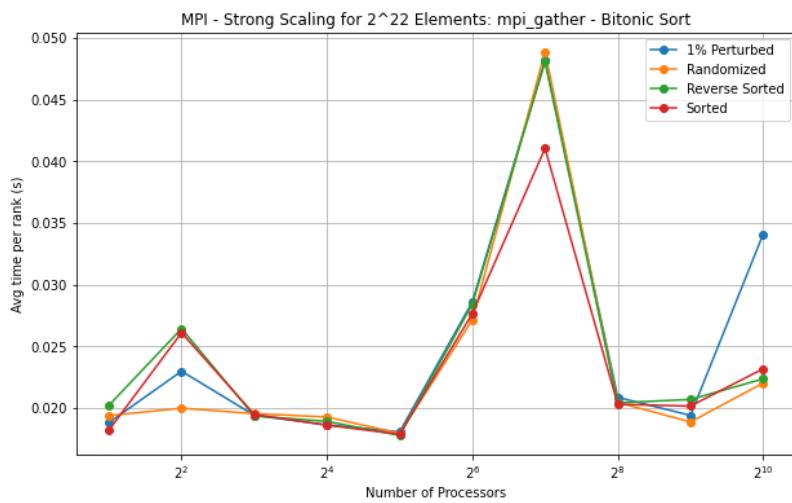


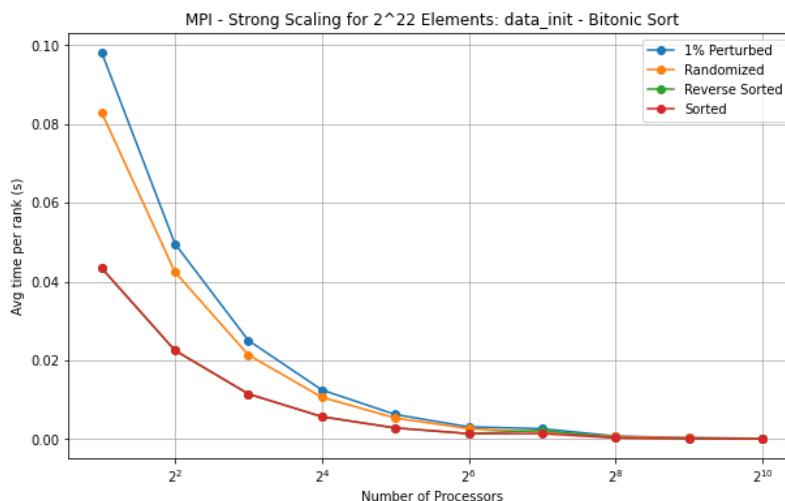
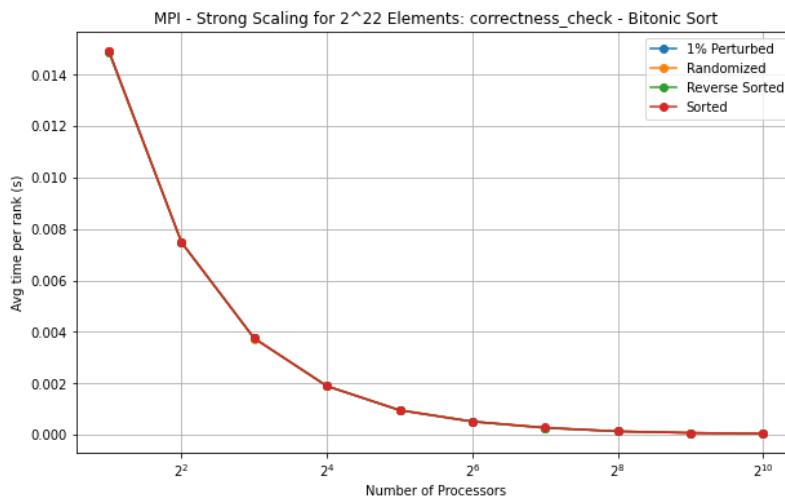
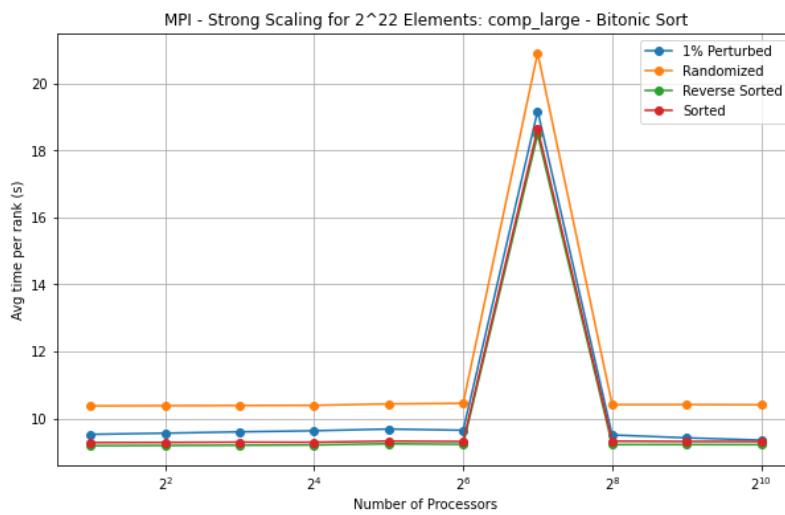




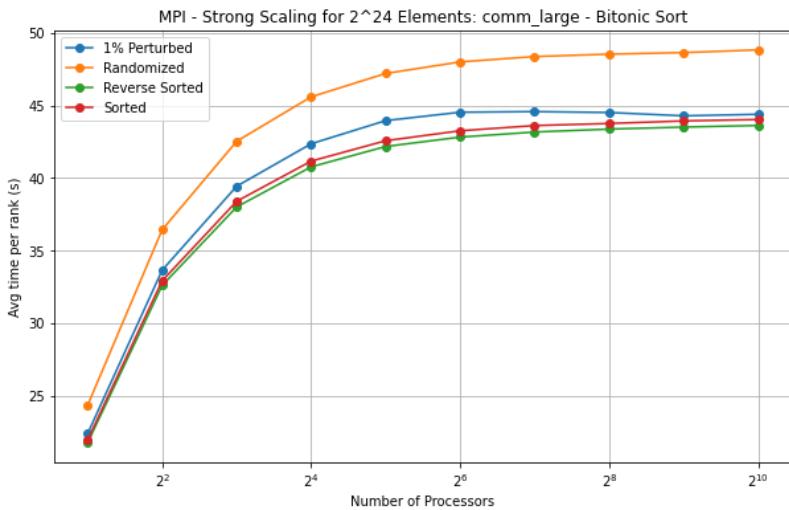
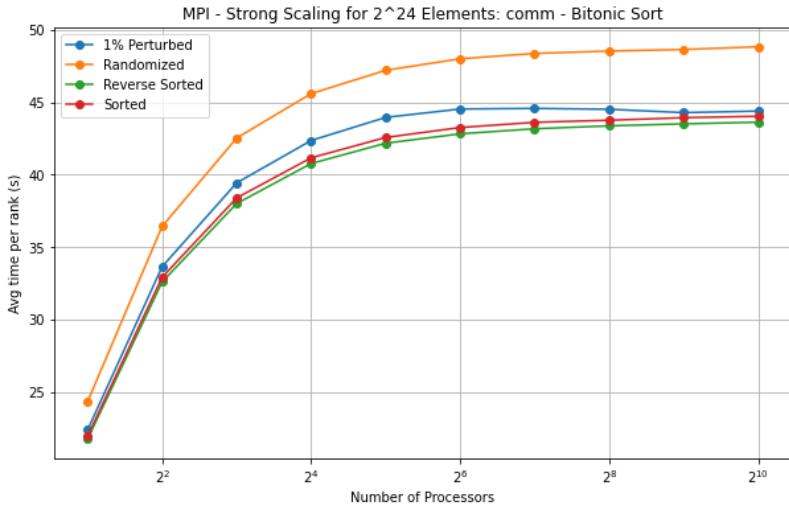
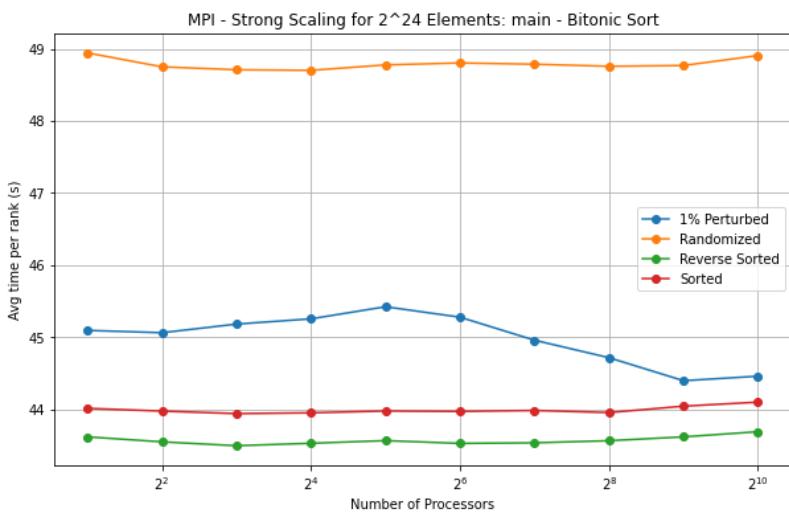
2²² Elements

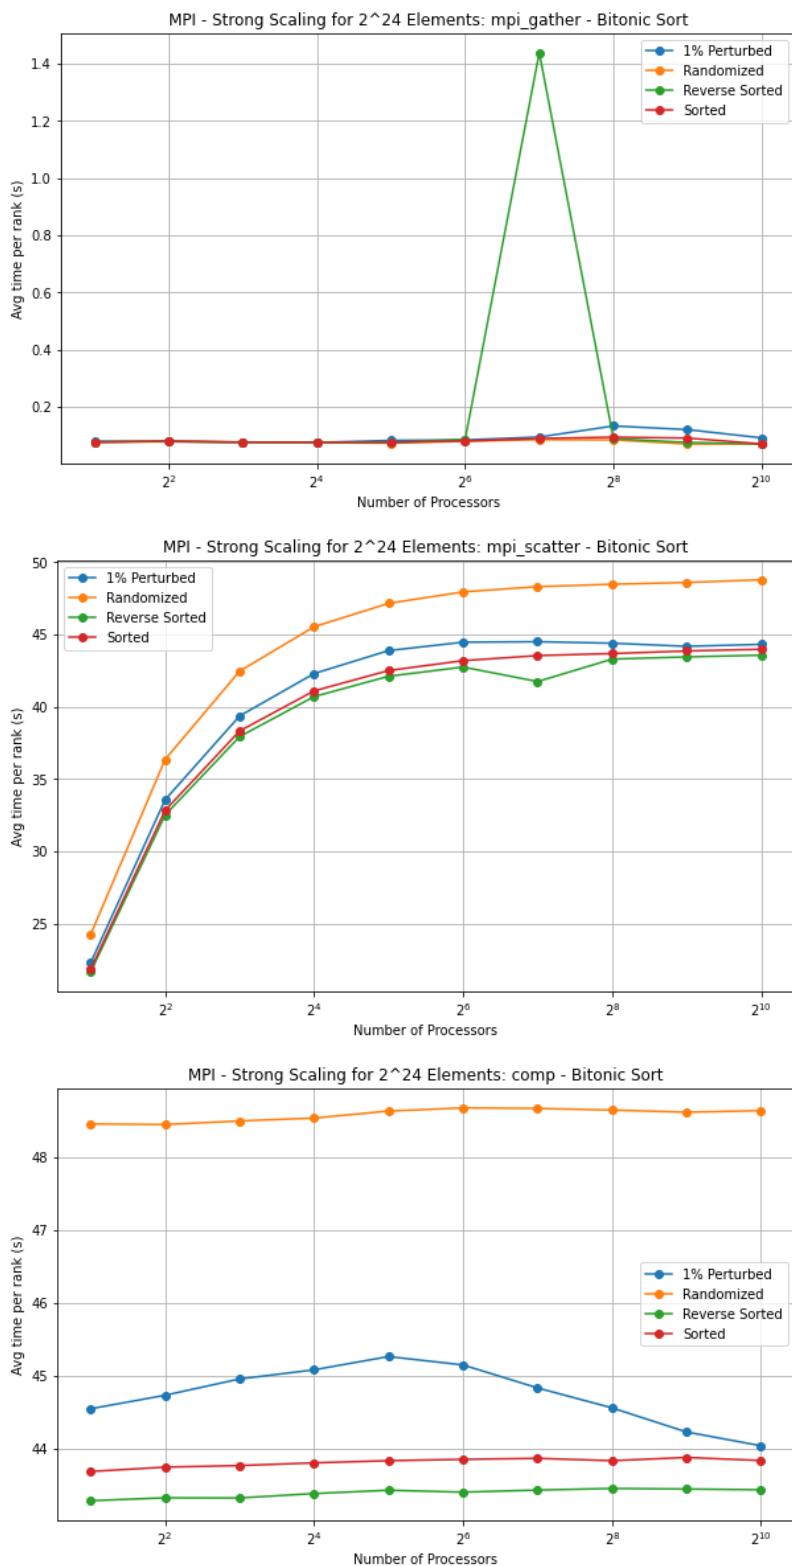


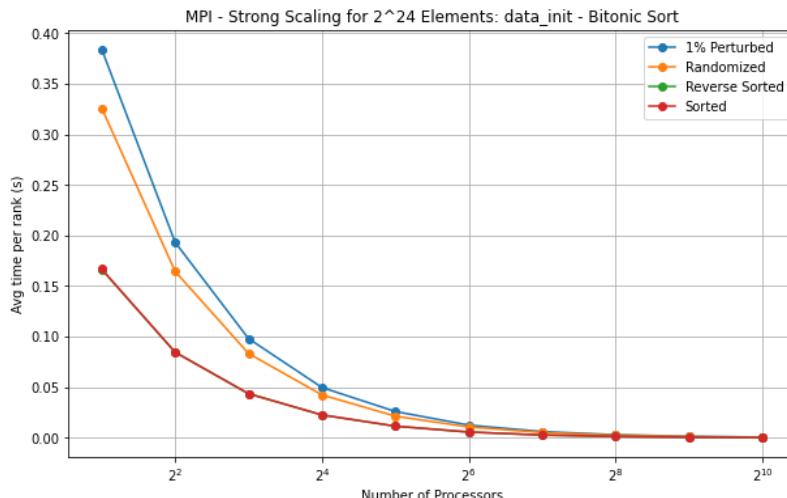
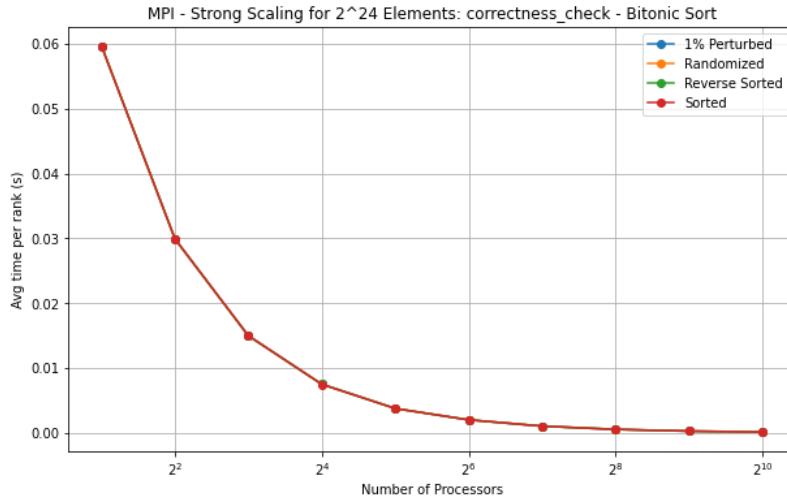
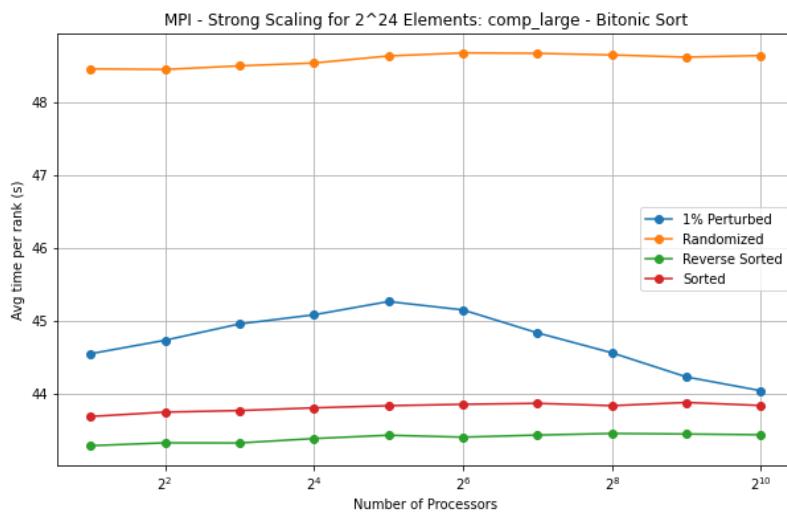




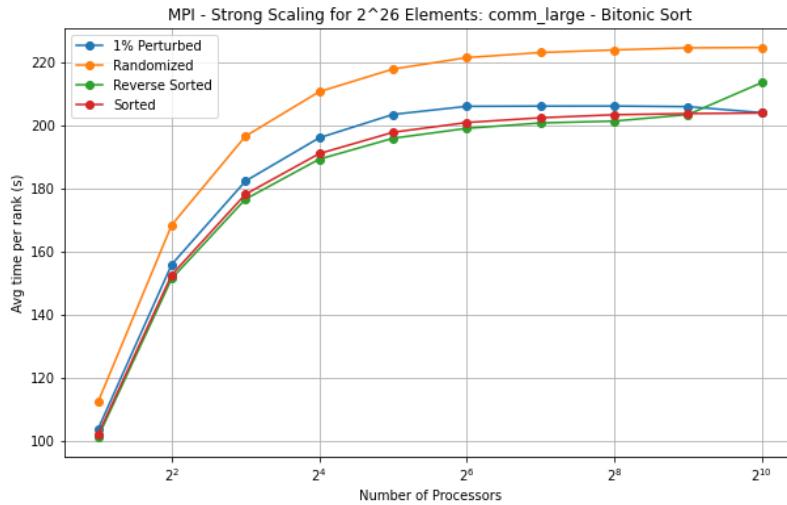
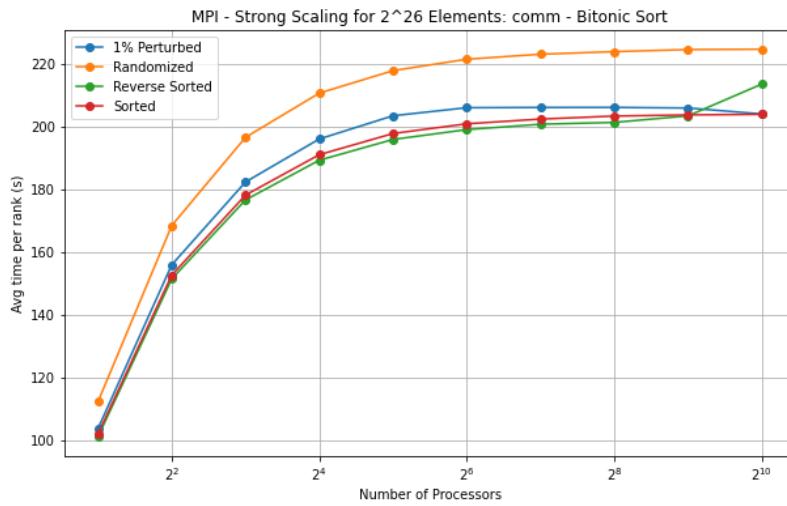
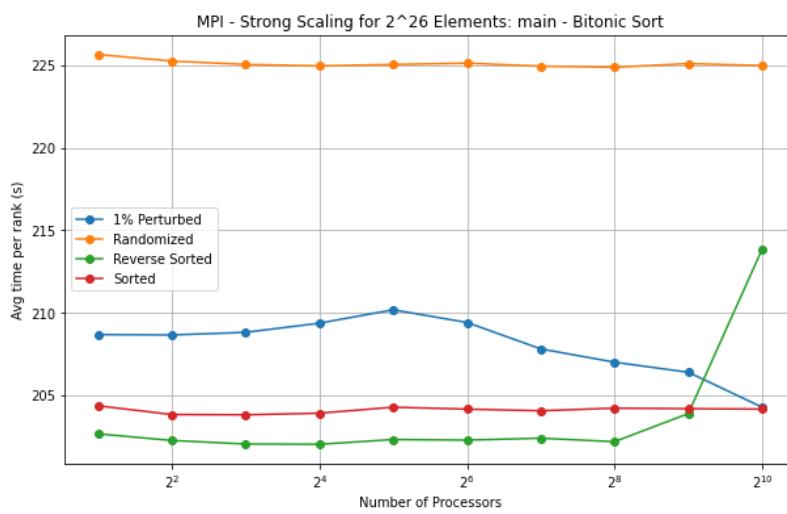
2^{24} Elements

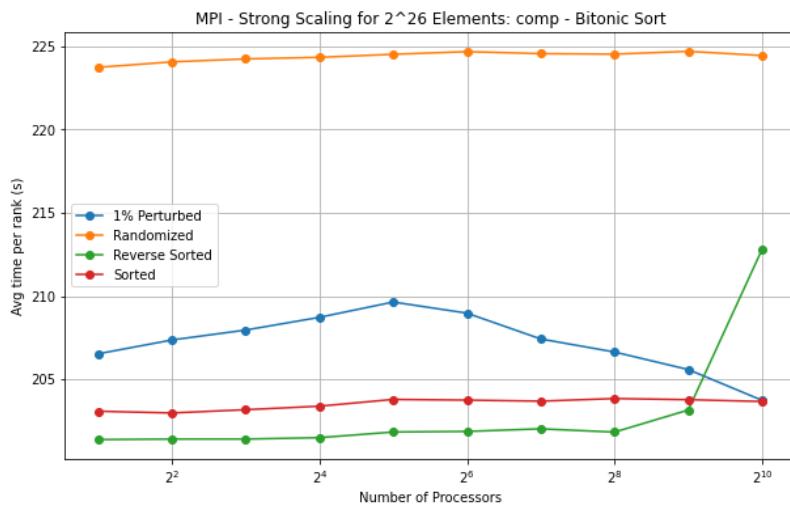
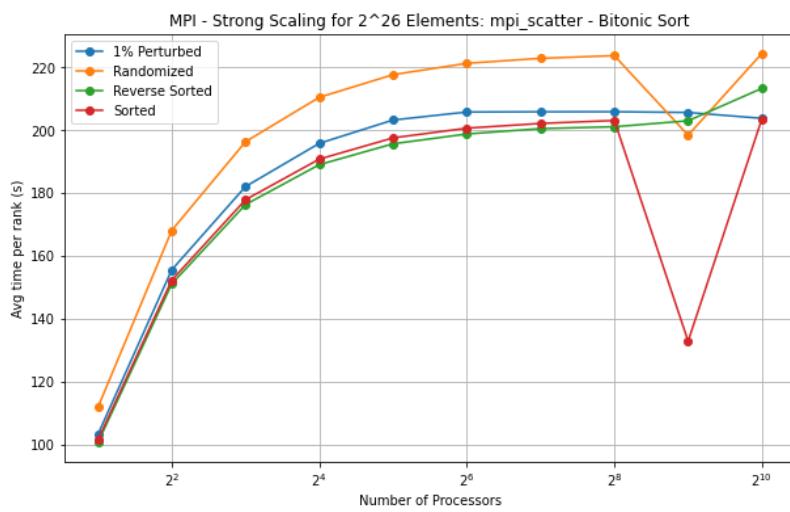
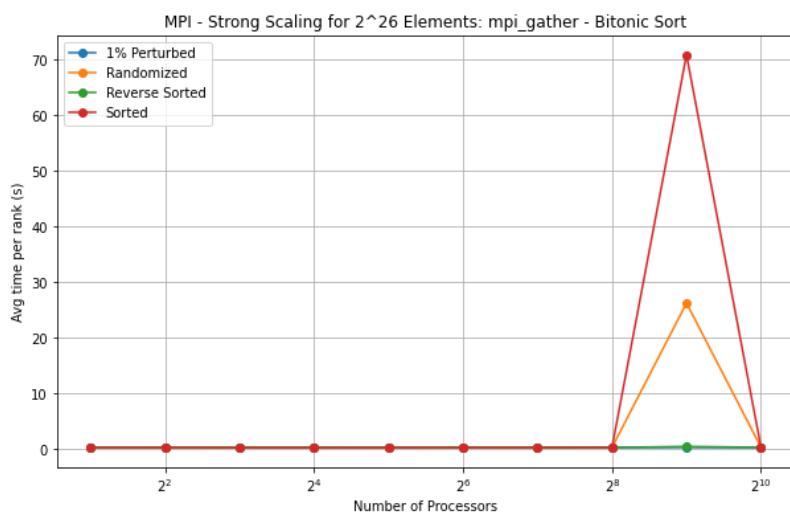


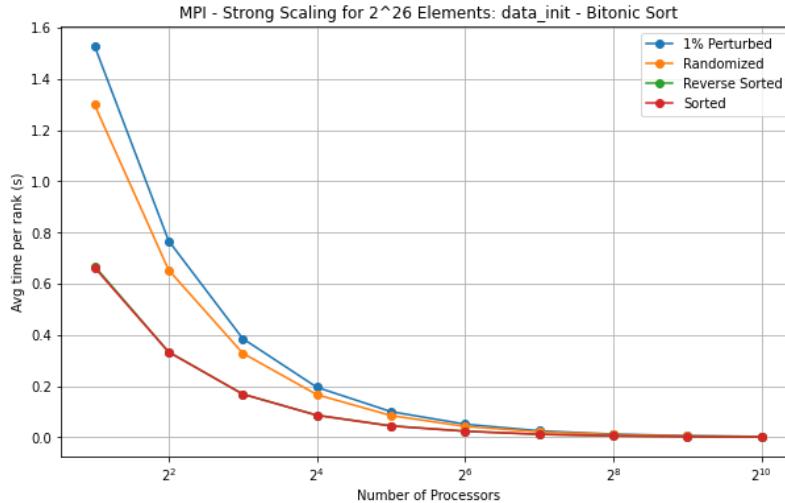
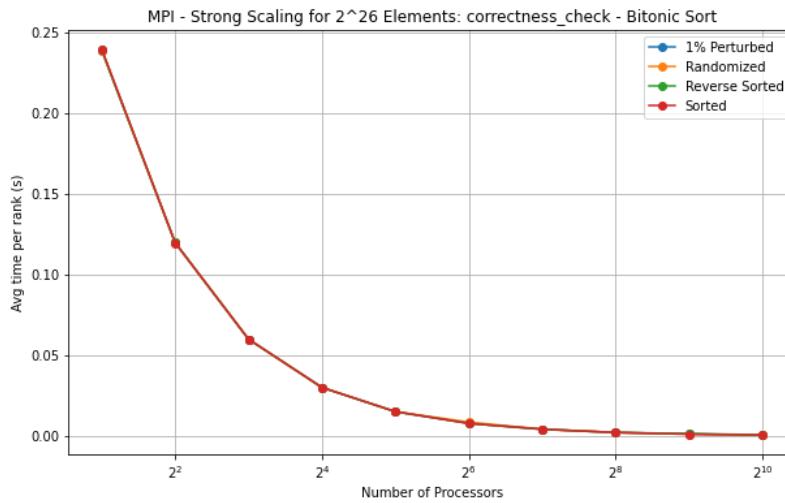
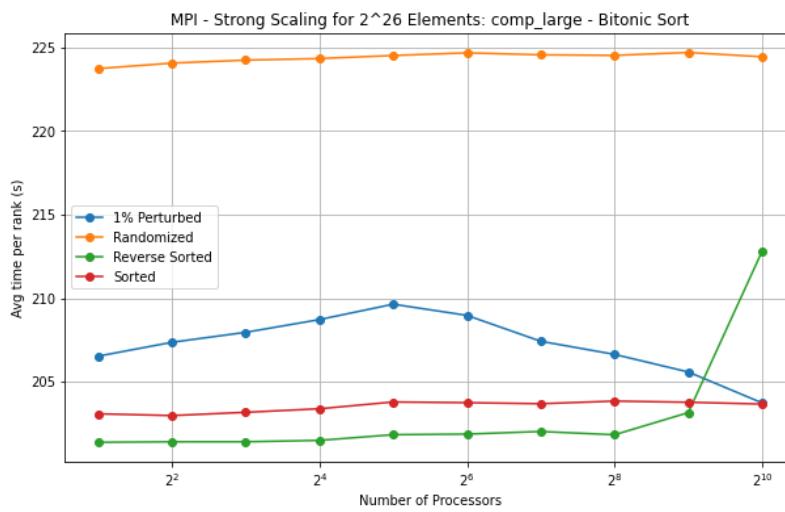




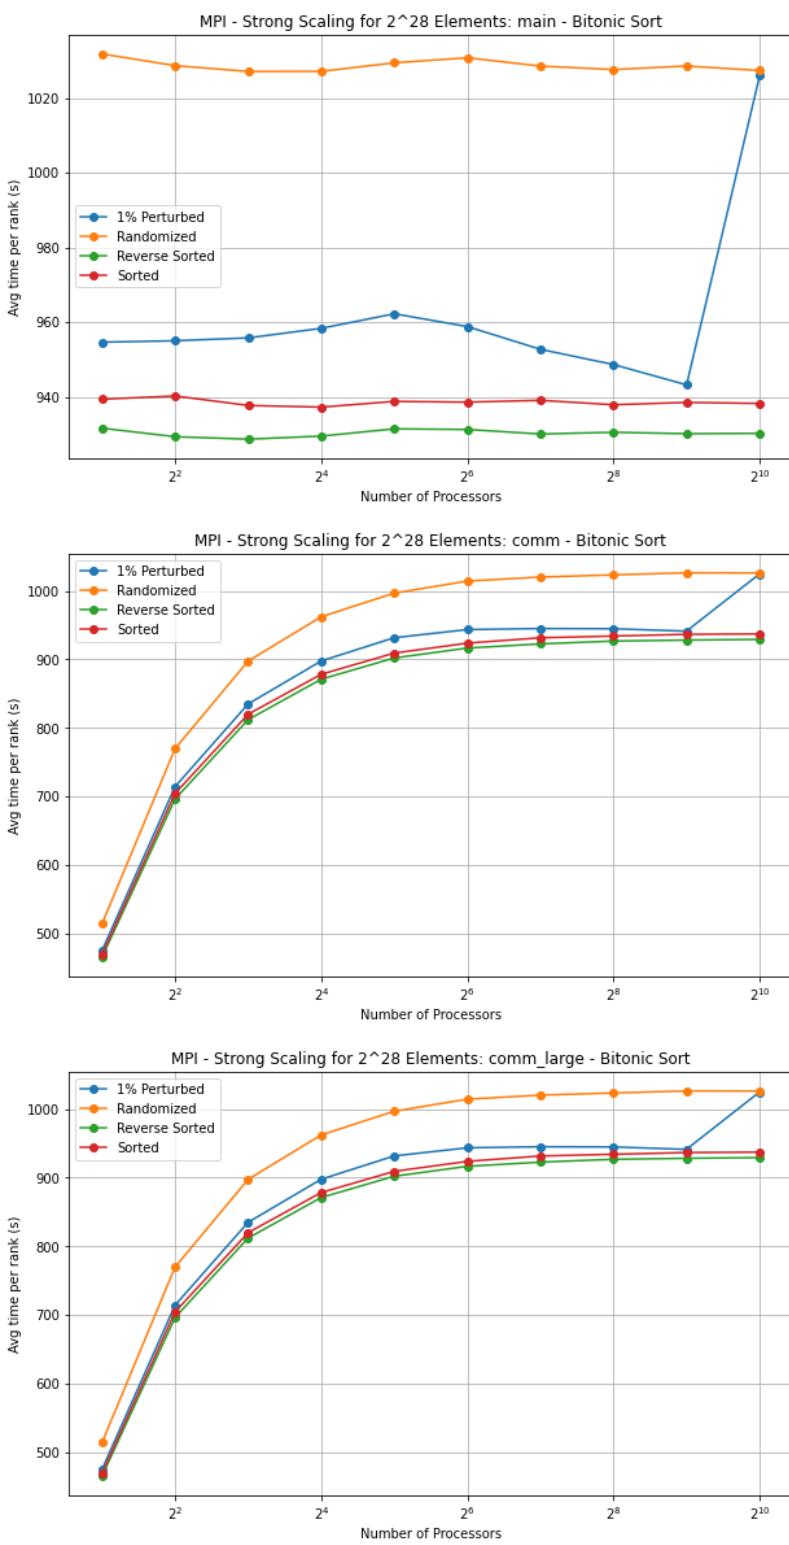
2^{26} Elements

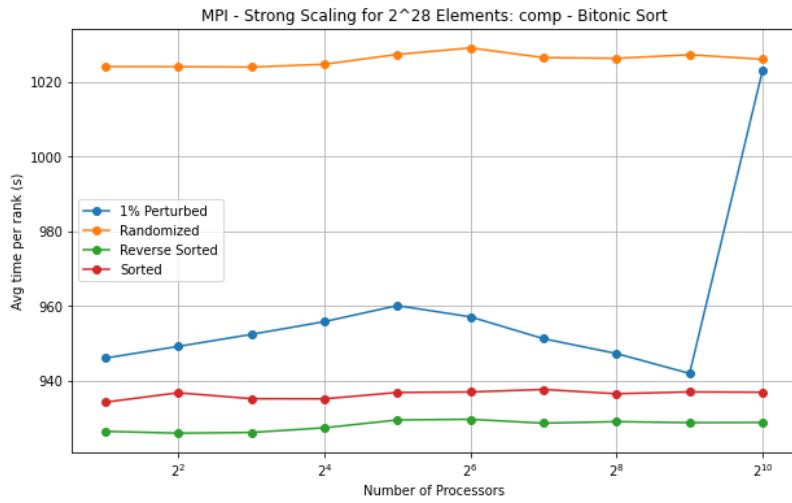
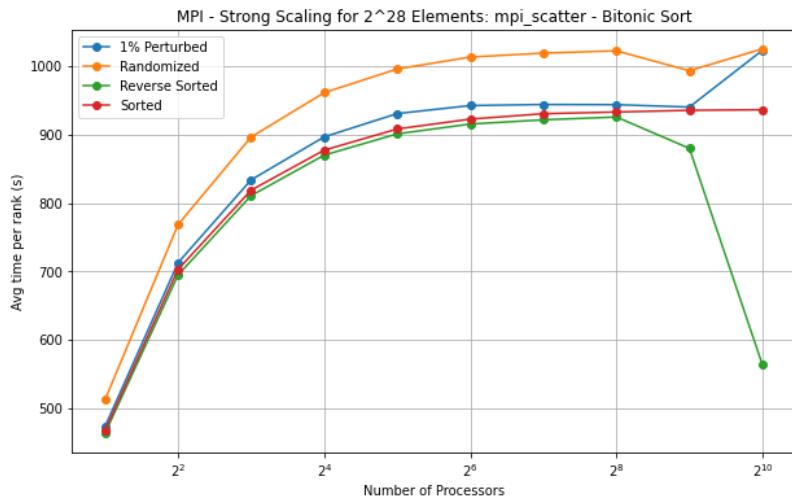
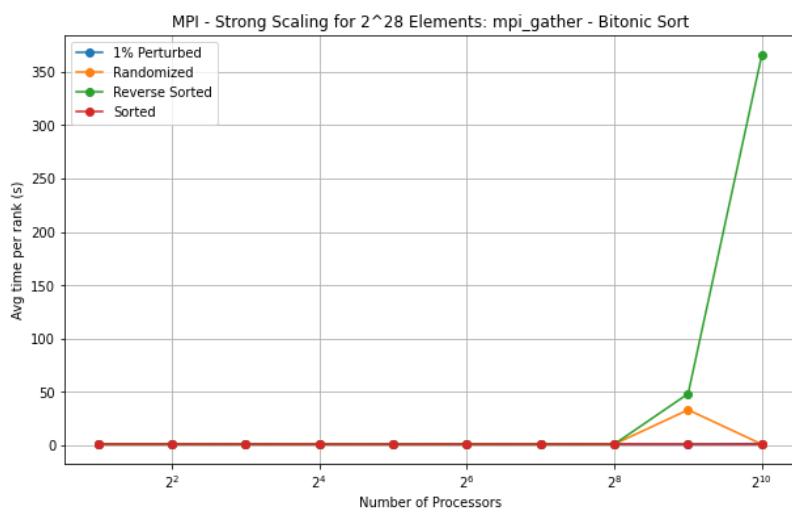


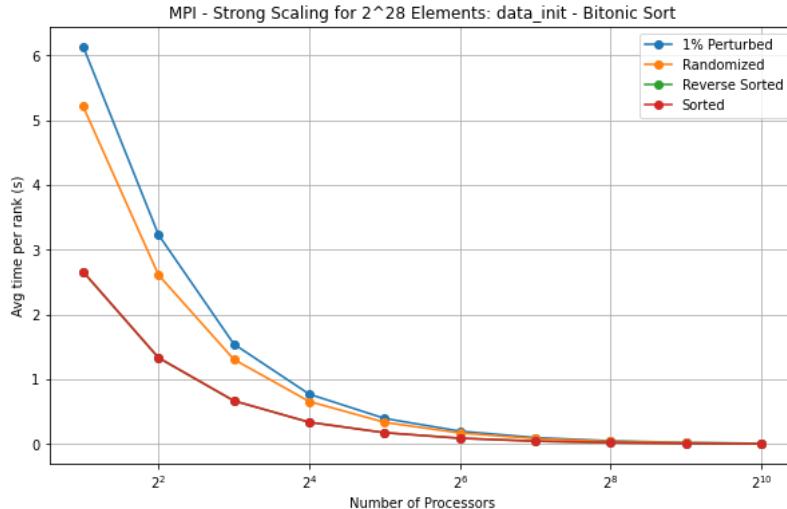
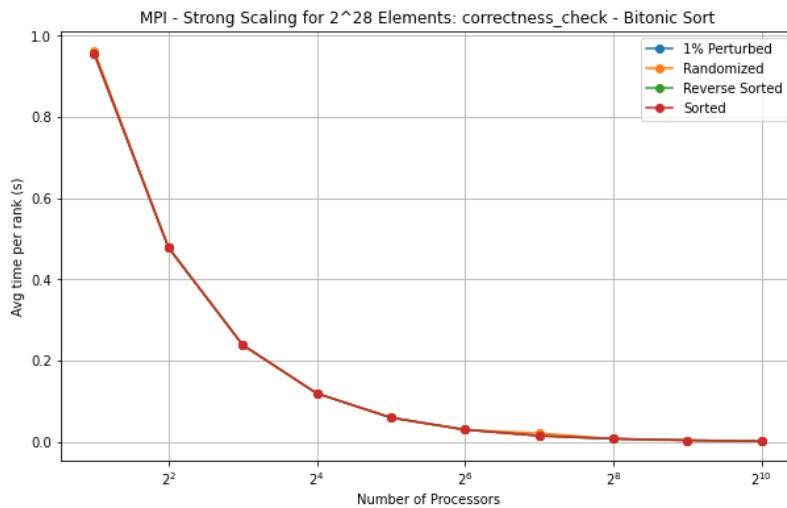
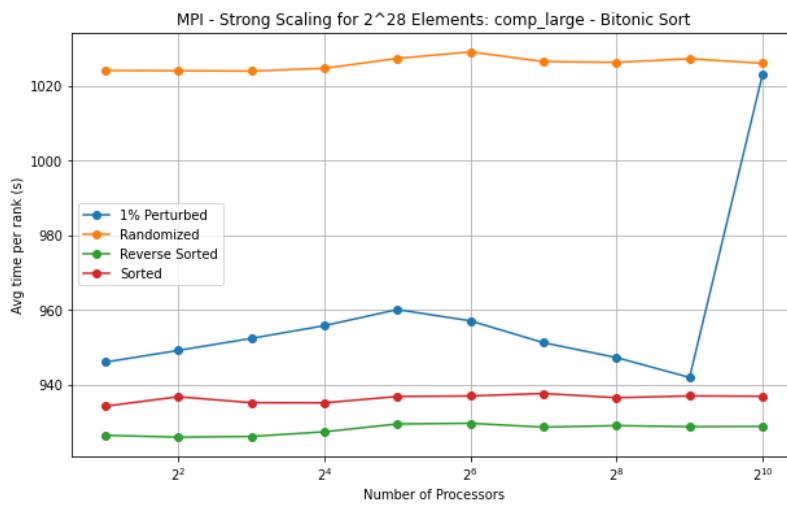




2^{28} Elements





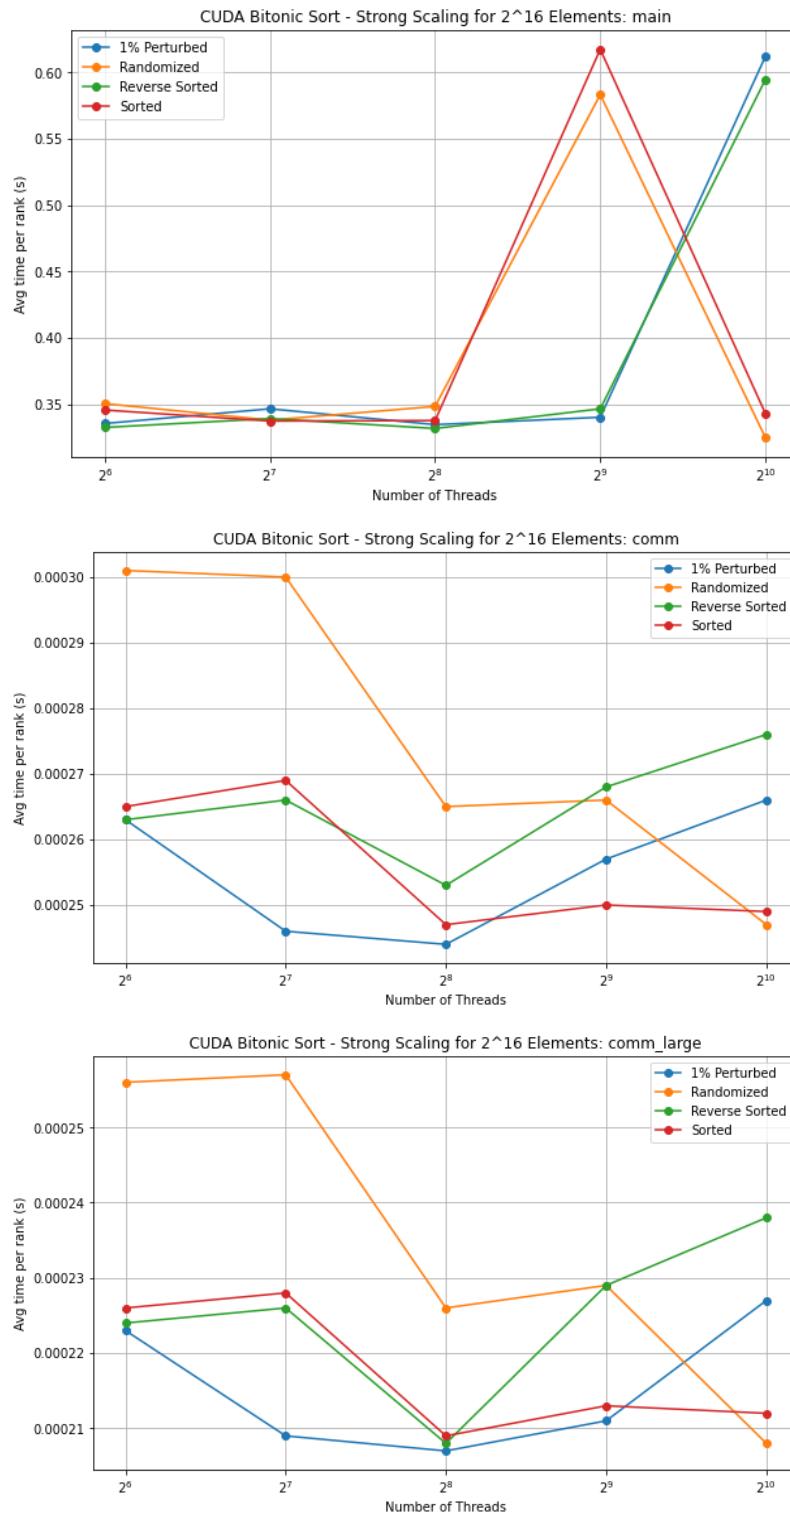


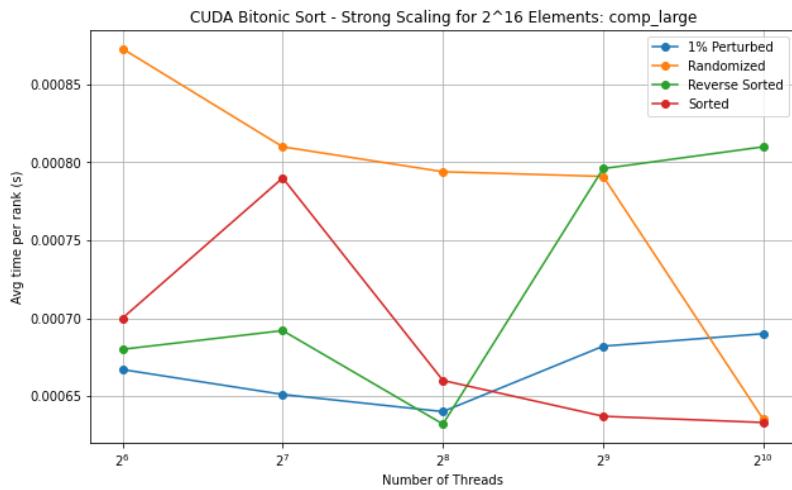
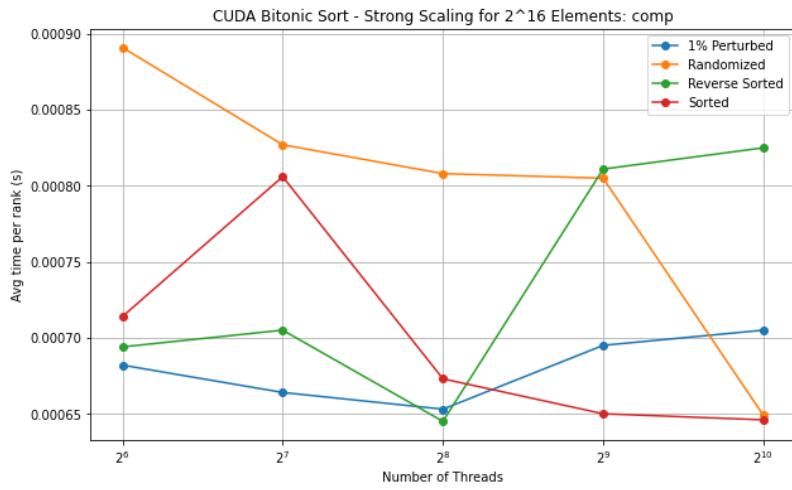
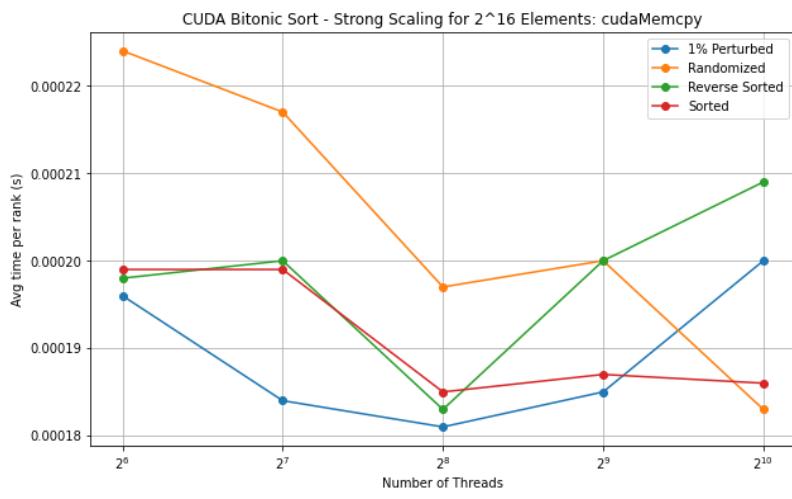
CUDA

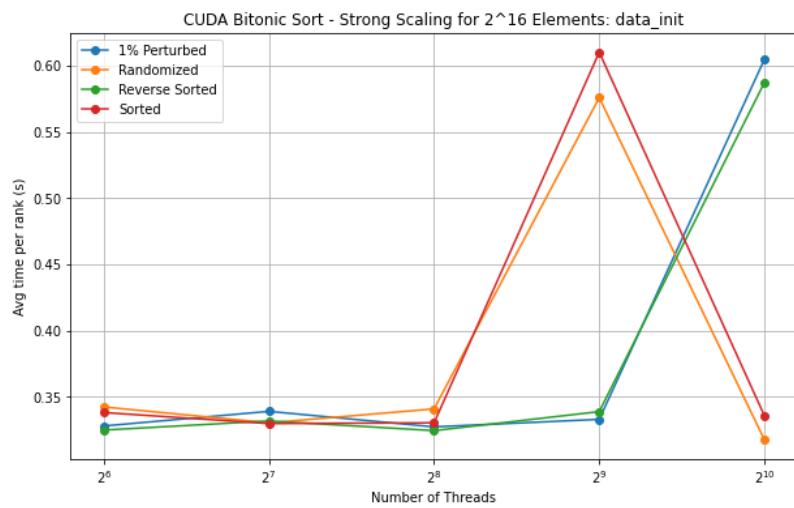
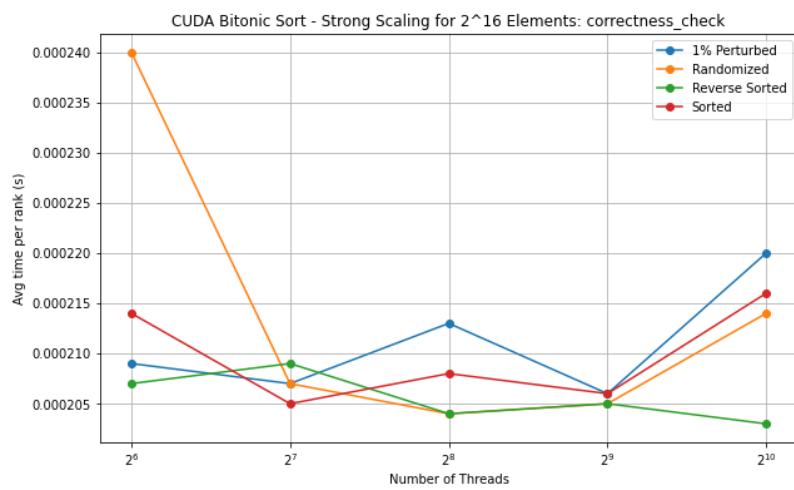
For the analysis, for CUDA only two cali files with 16 were run so only two points are showing. It does seem that one of the points has a higher average time. Due to the small number of trials, not much can be observed from the graph. The jobs were queued on grace portal so the data was not fully able to be collected for this implementation. However, as the number of processors increases, the average time should decrease as the the parallelization of the tasks has been implemented. Therefore, the rest of the input types will be run

FIX ME!

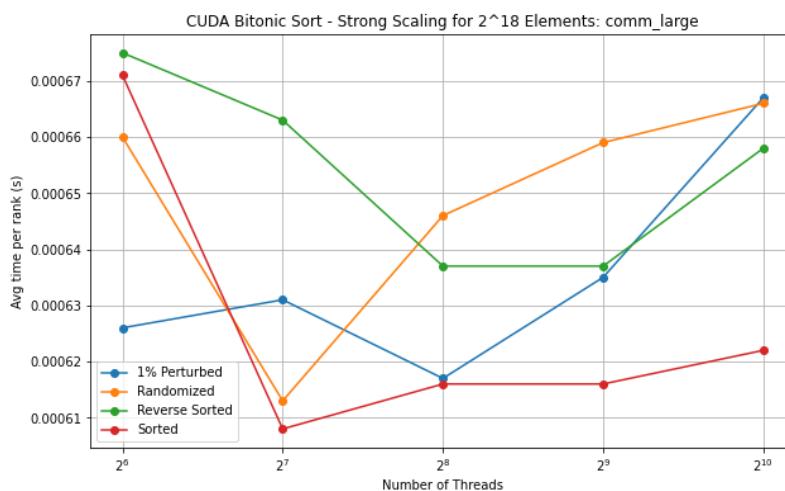
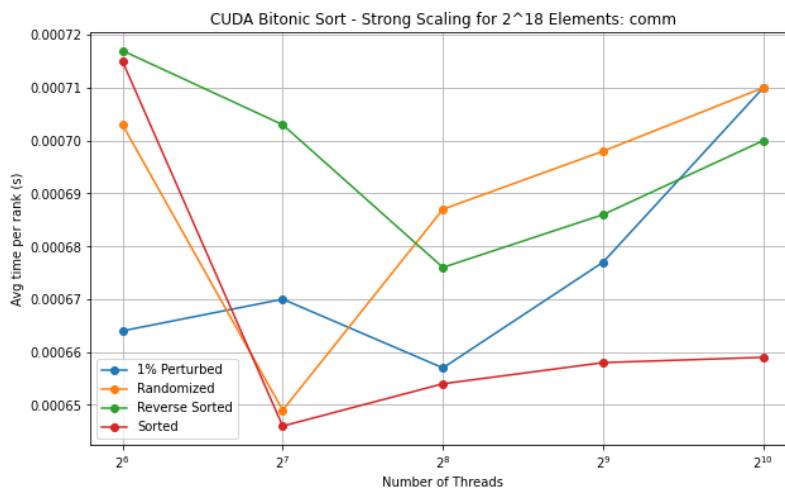
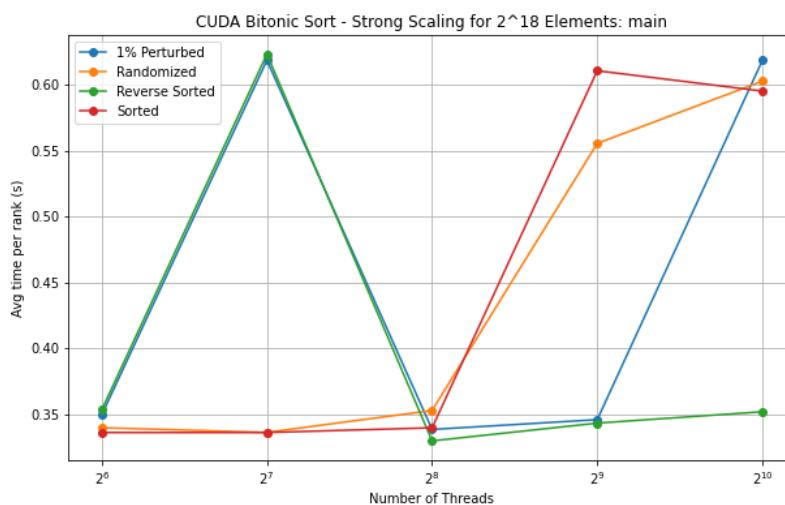
2¹⁶ Elements

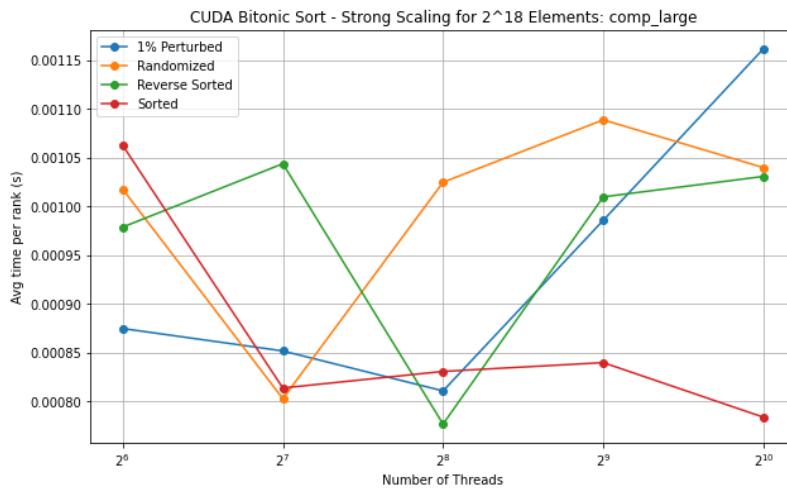
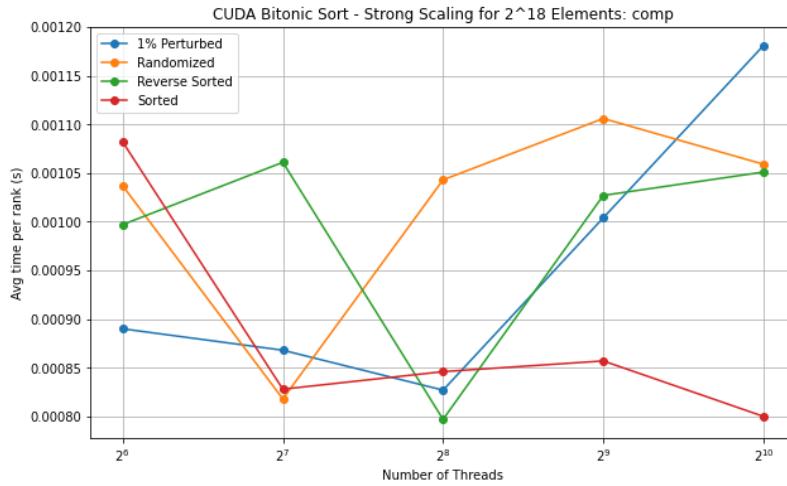
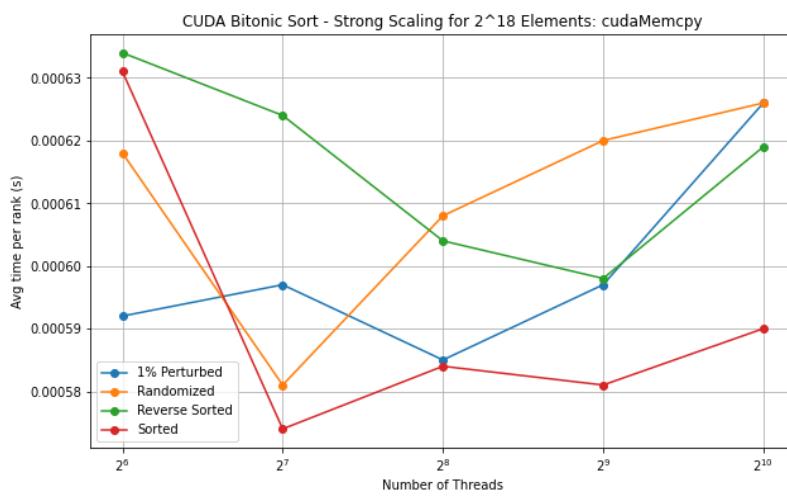


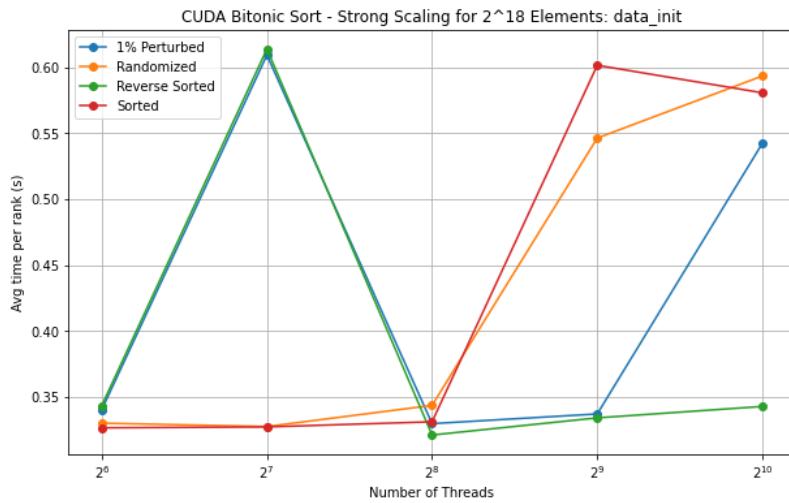
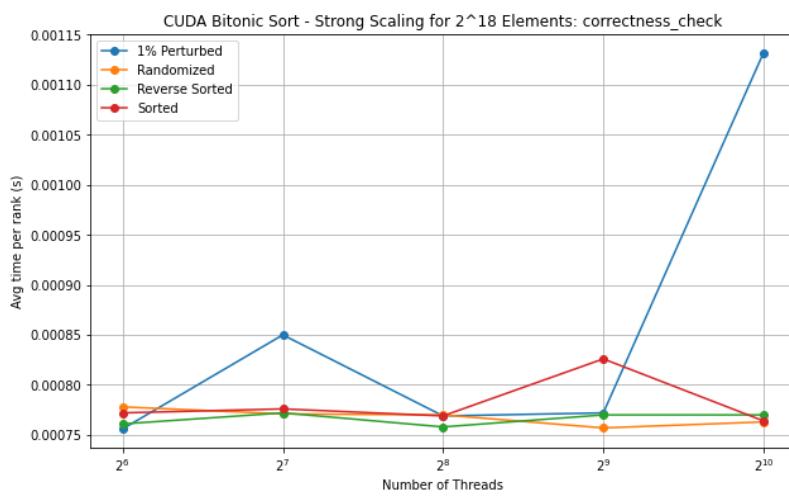




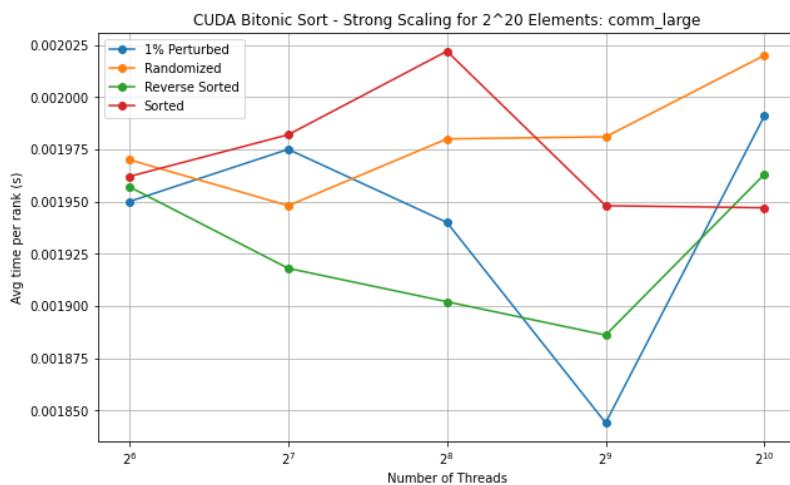
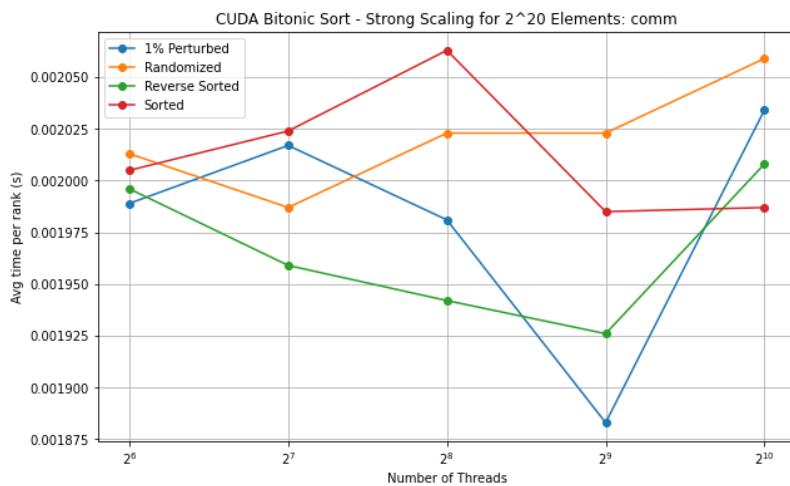
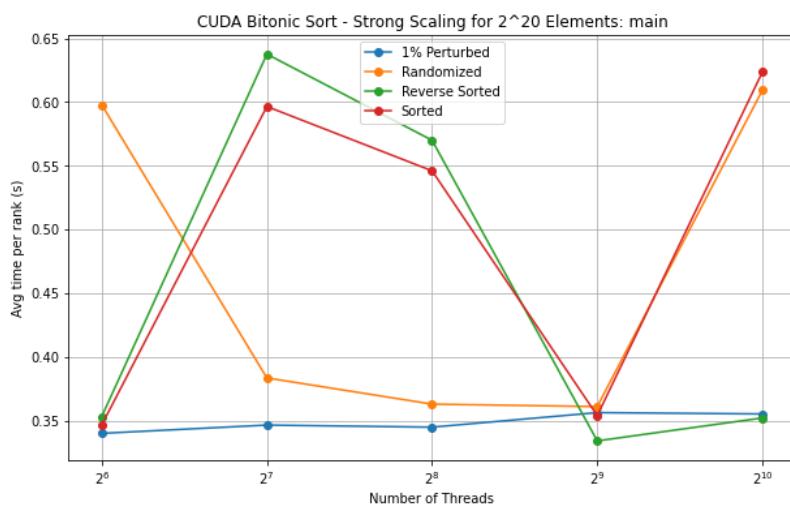
2^{18} Elements

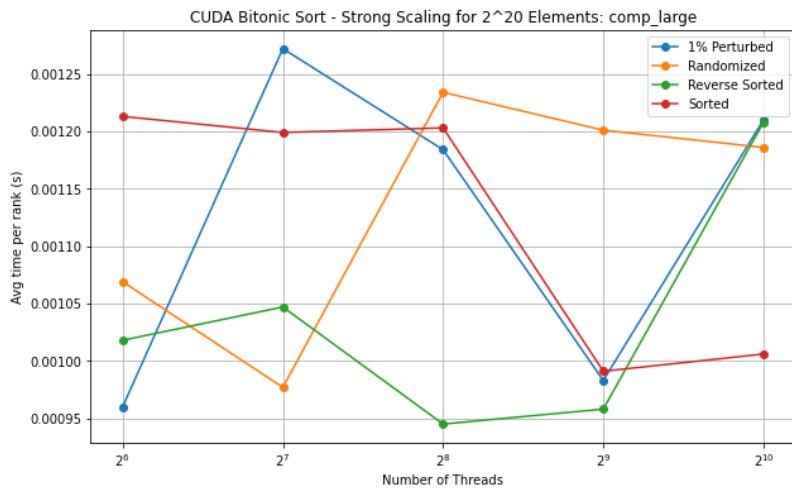
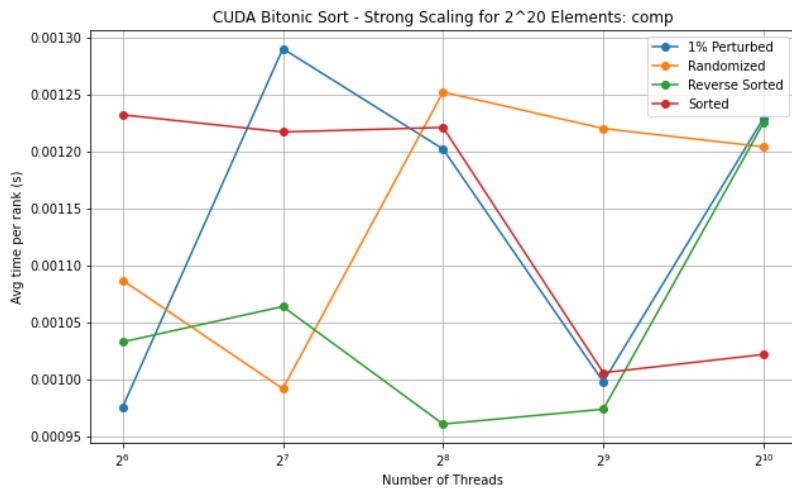
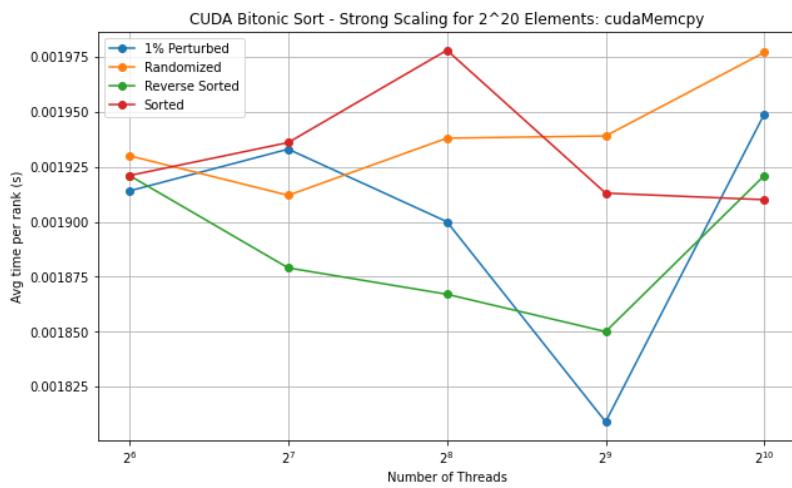


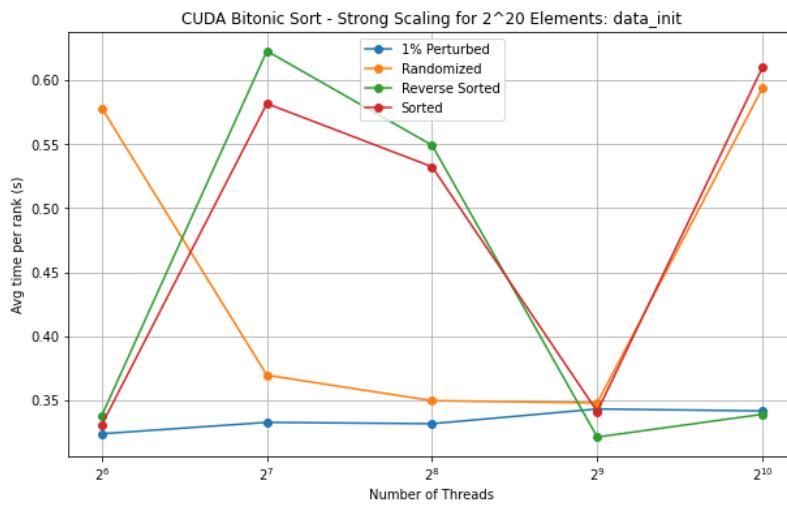
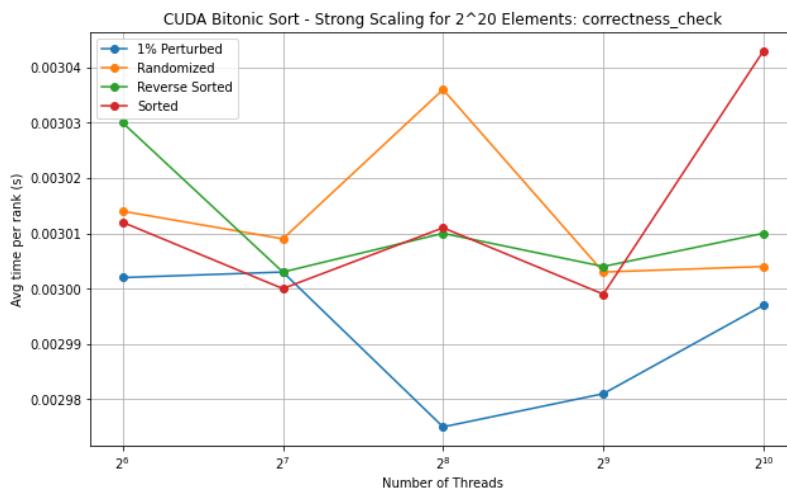




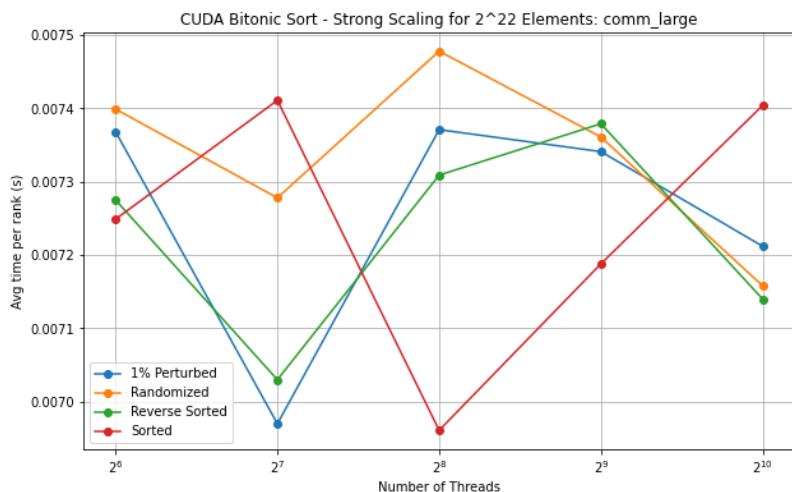
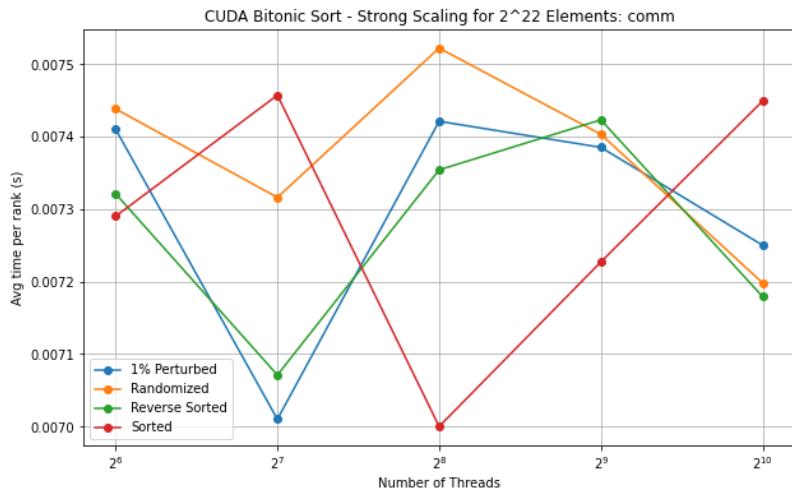
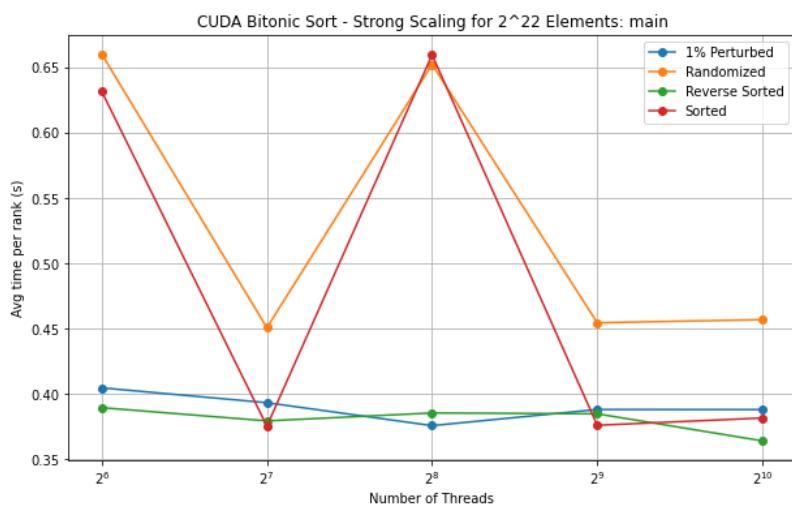
2^{20} Elements

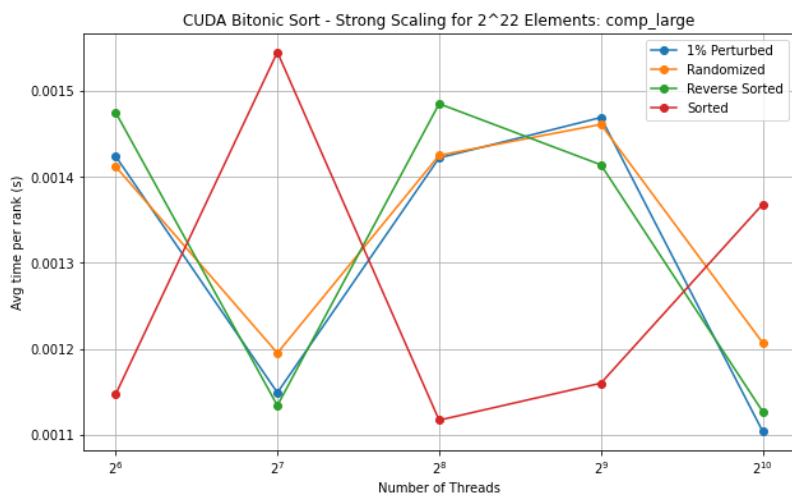
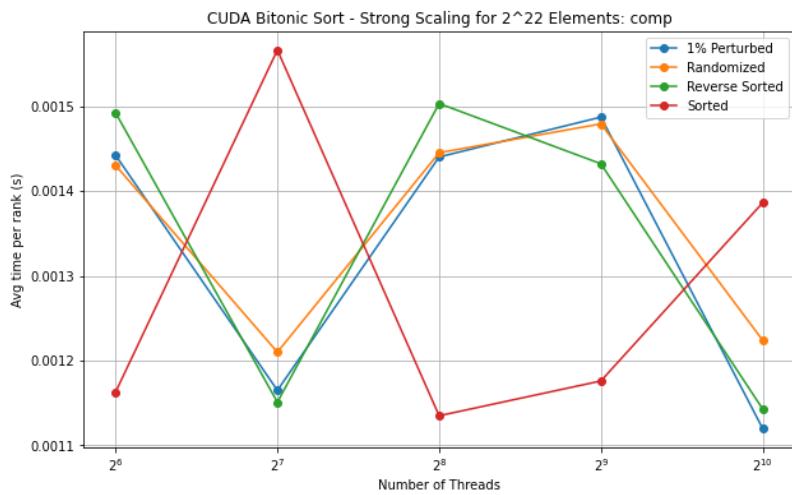
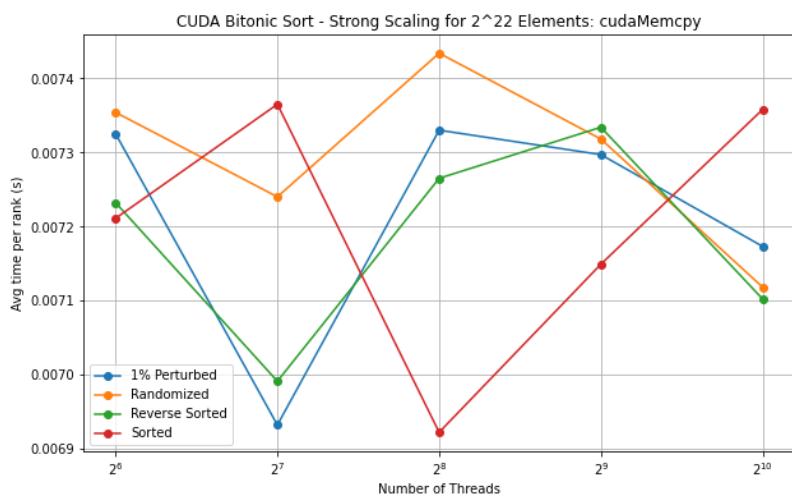


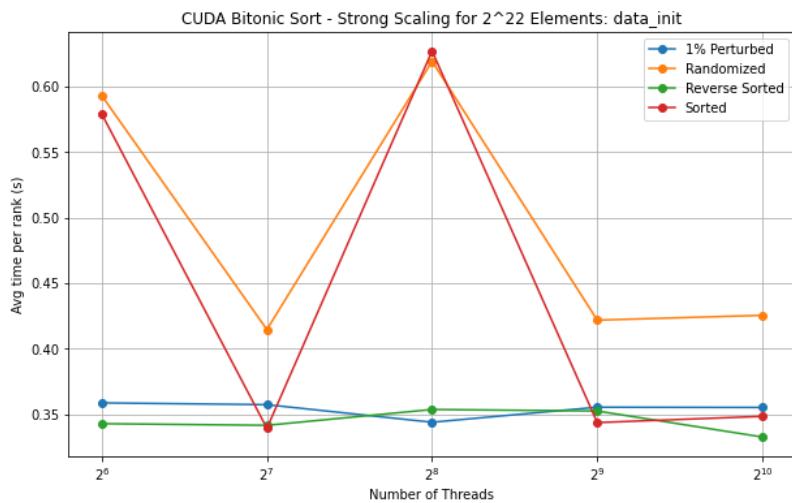
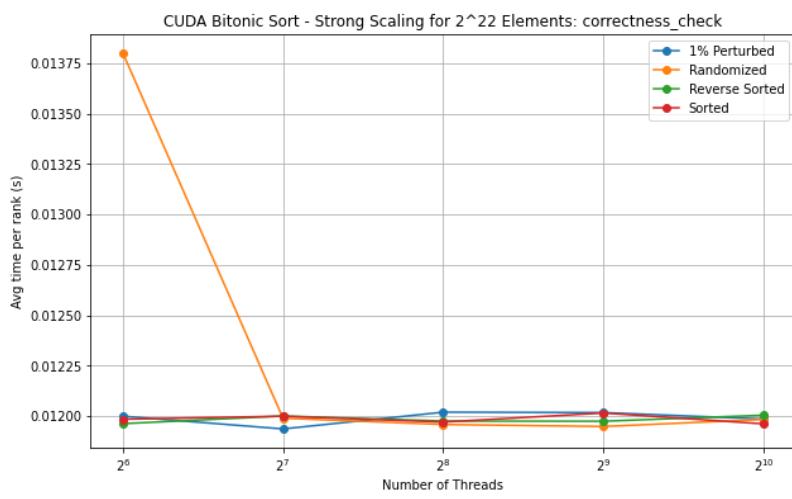




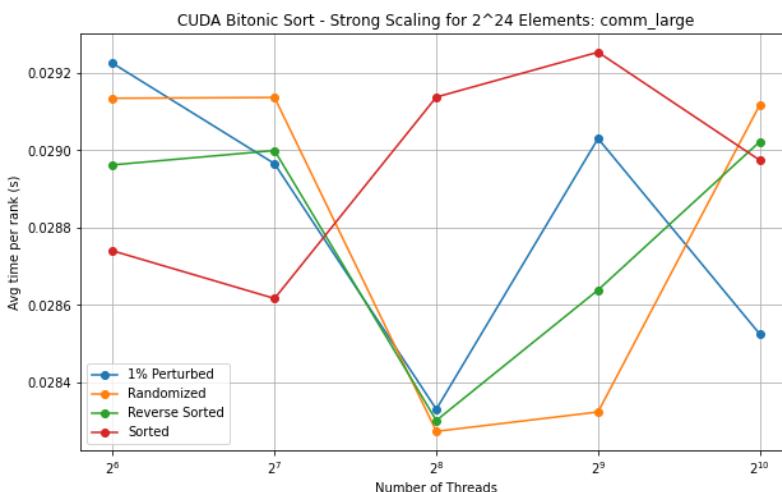
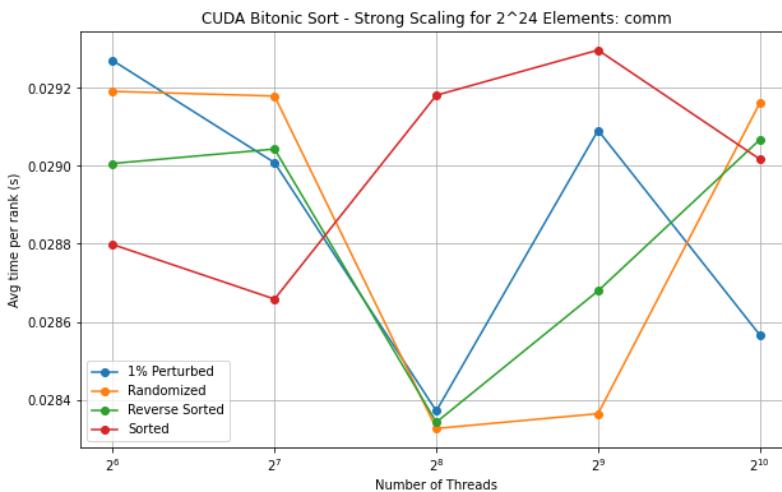
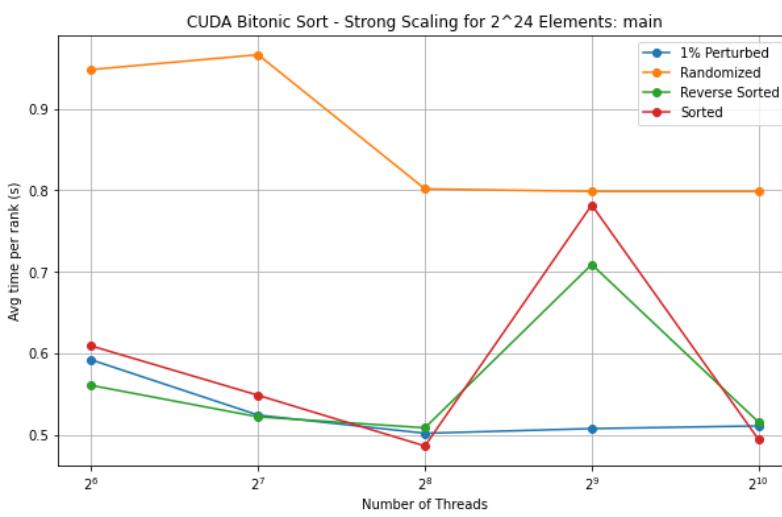
2^{22} Elements

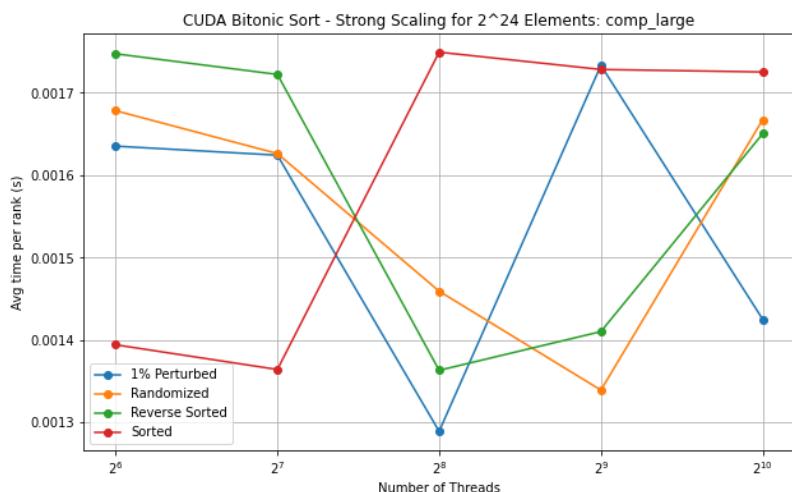
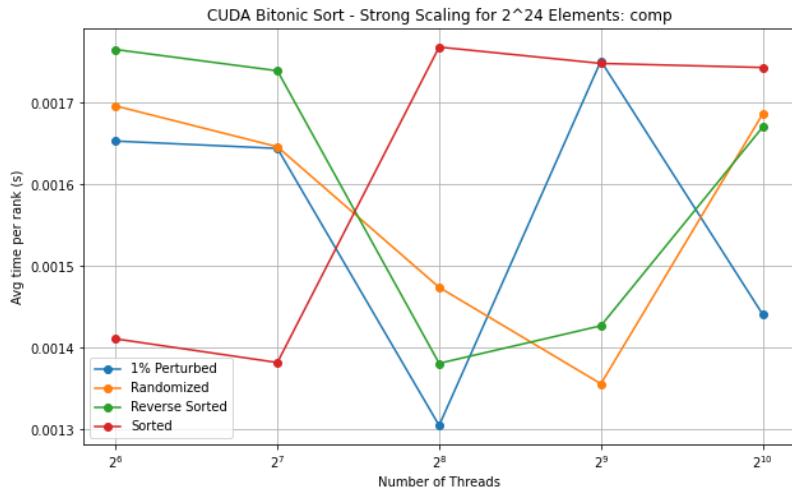
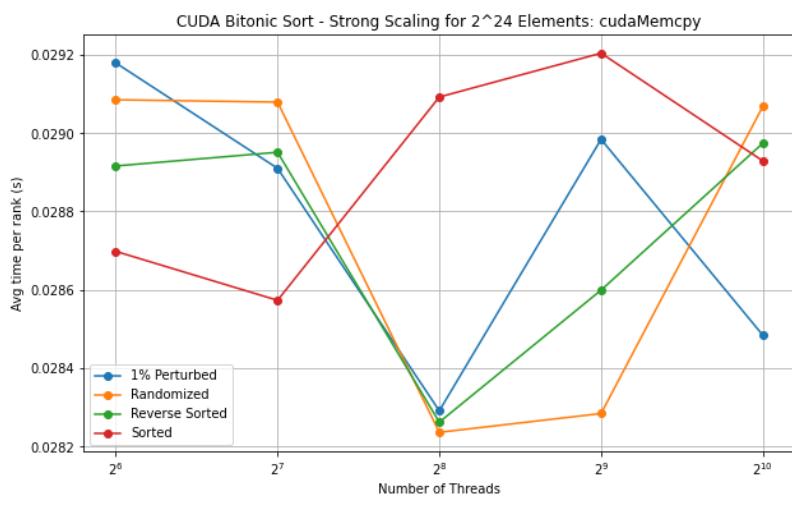


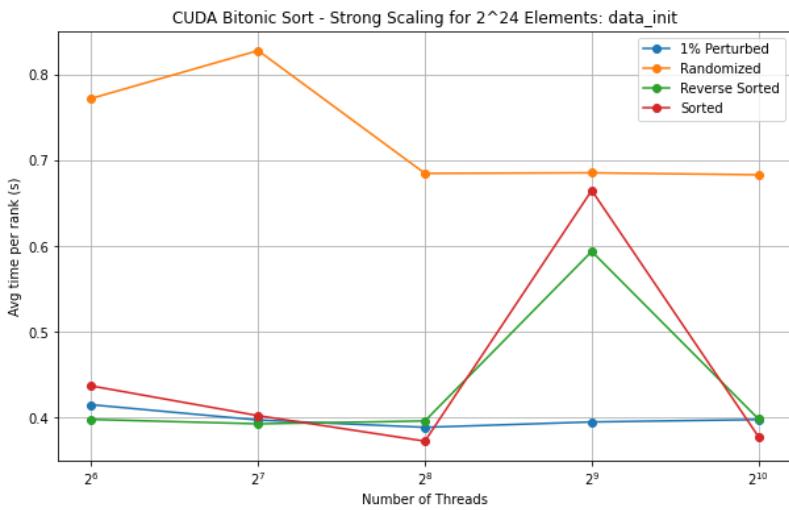
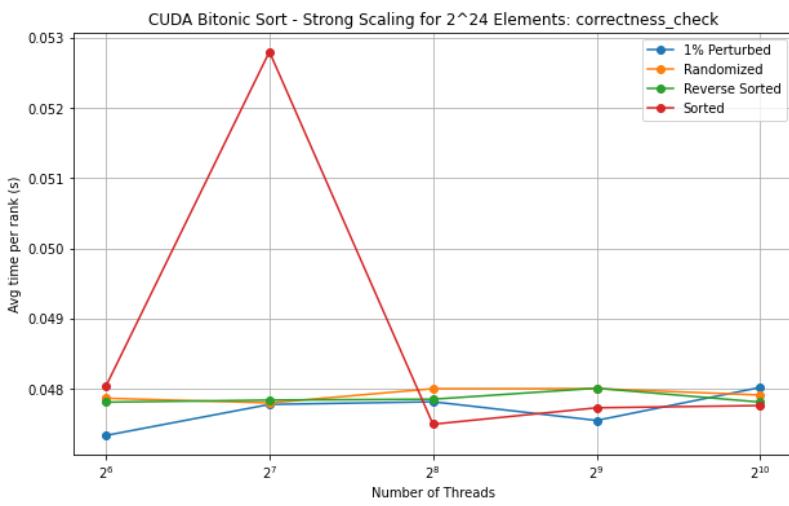




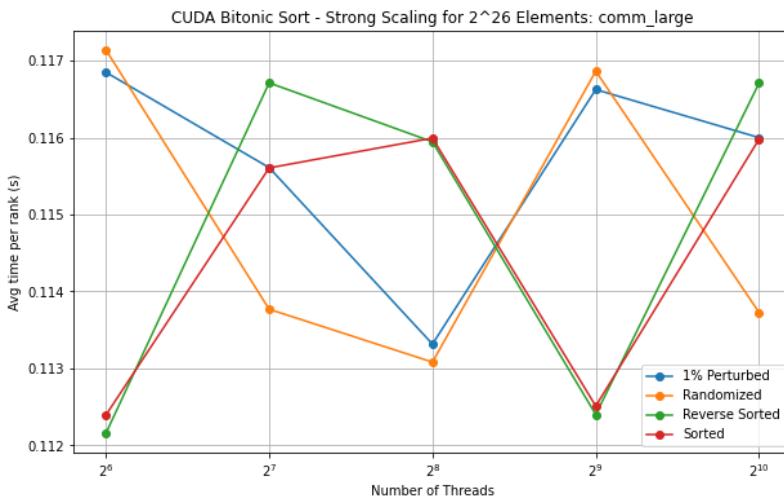
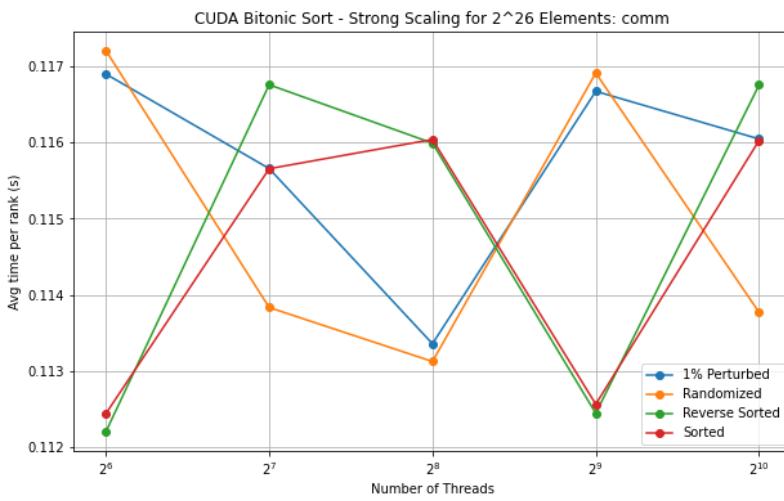
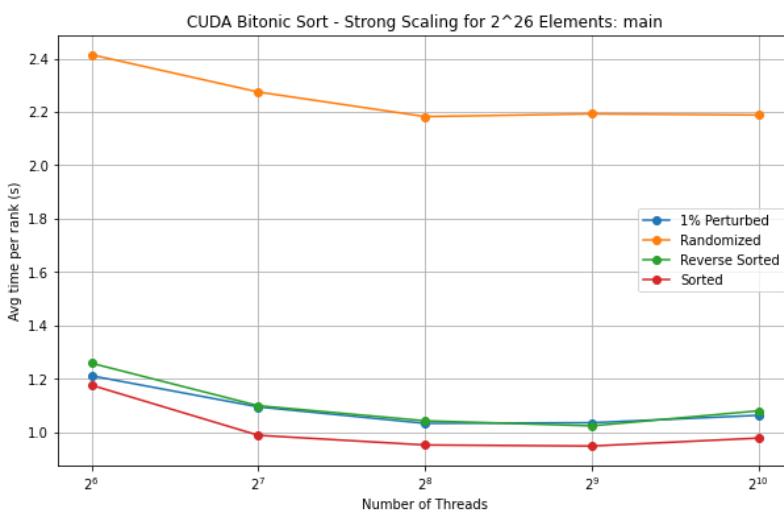
2^{24} Elements

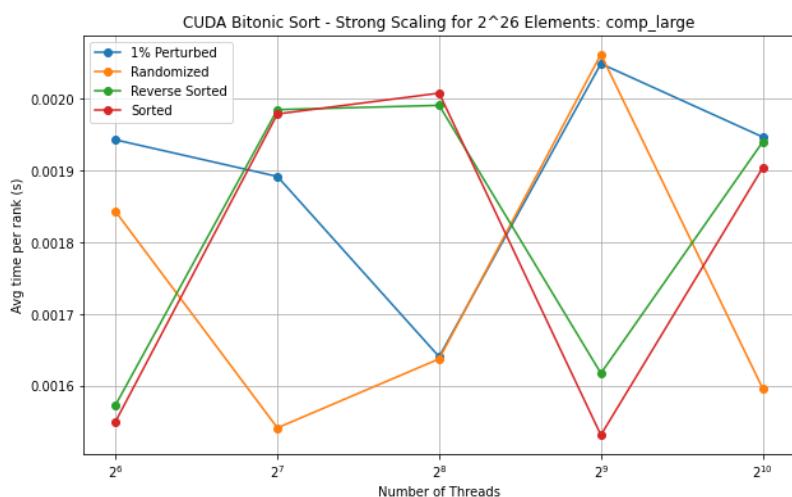
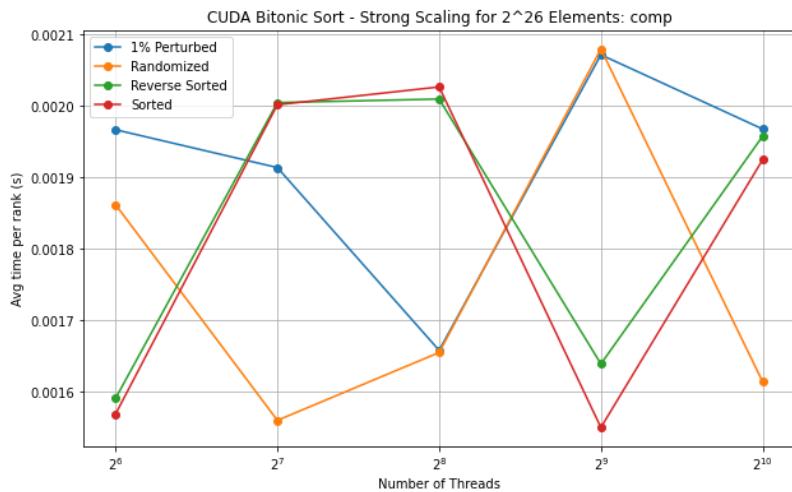
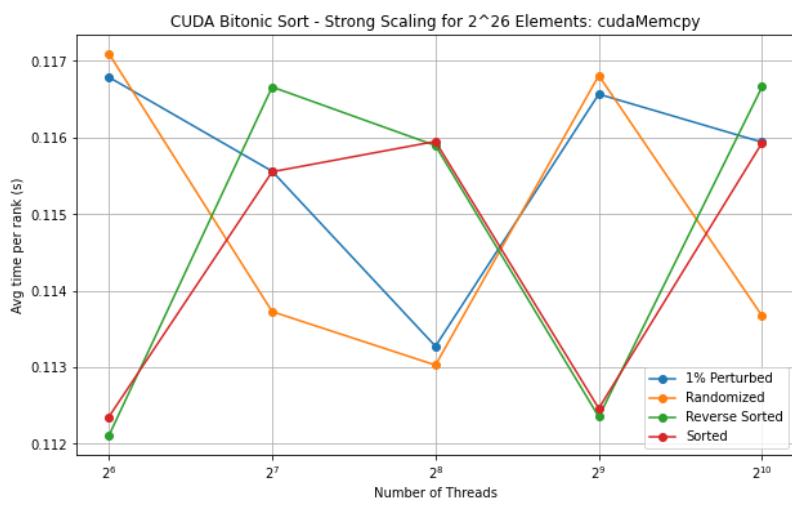


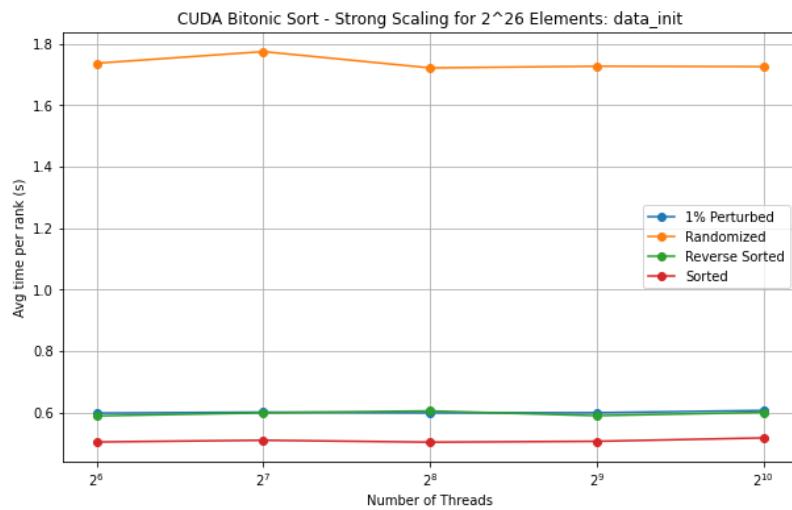
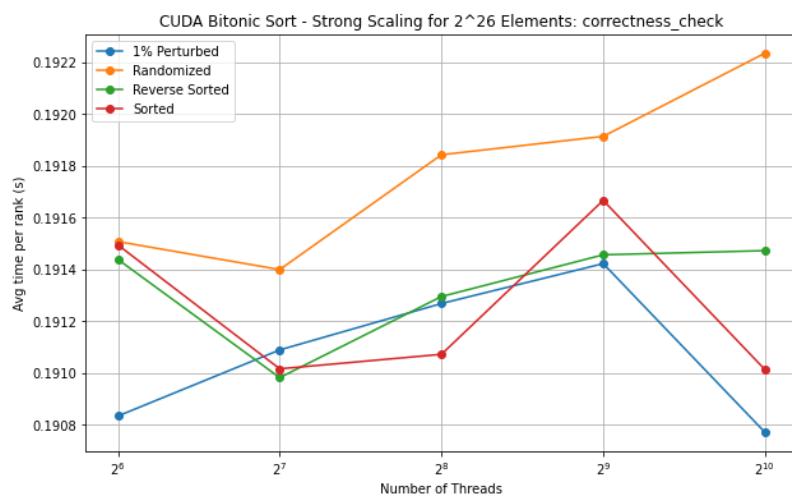




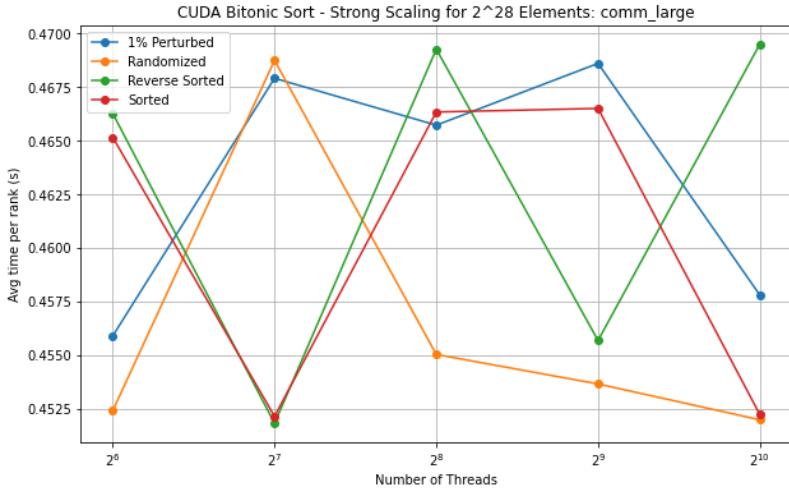
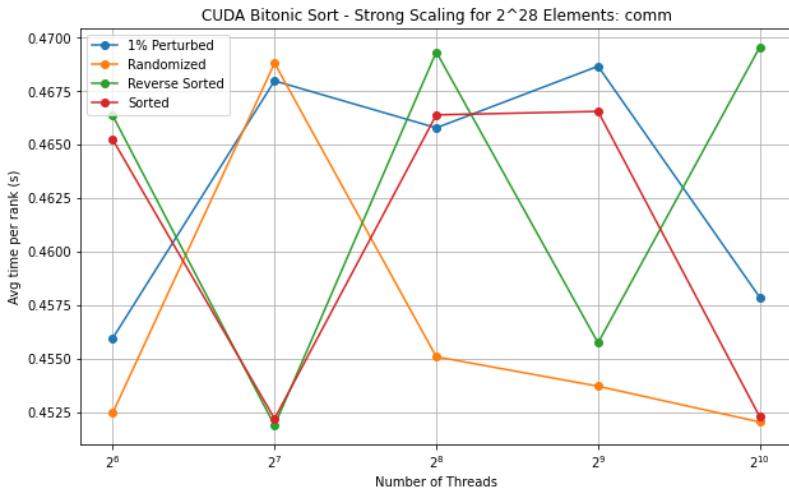
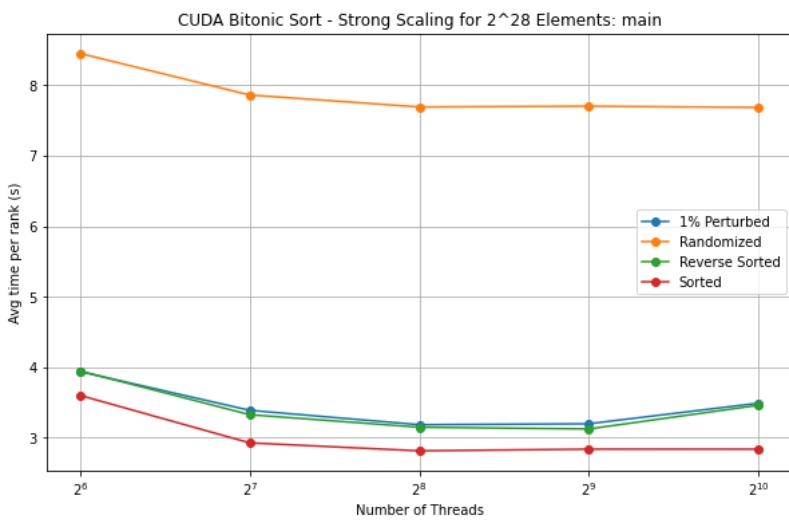
2^{26} Elements

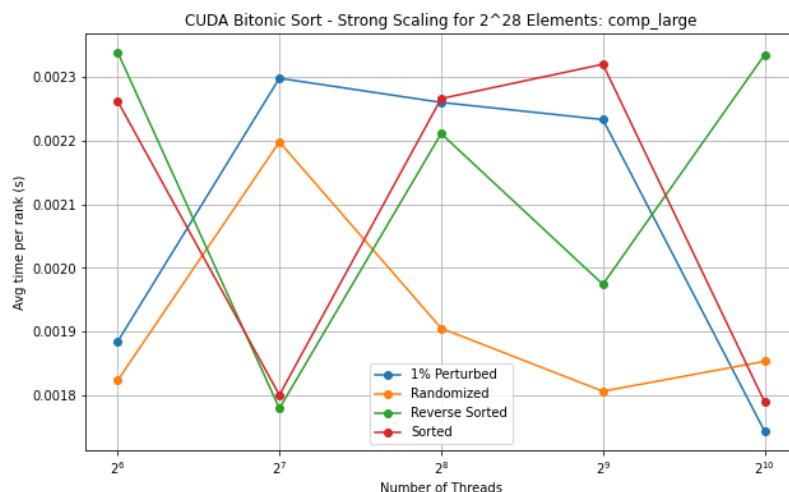
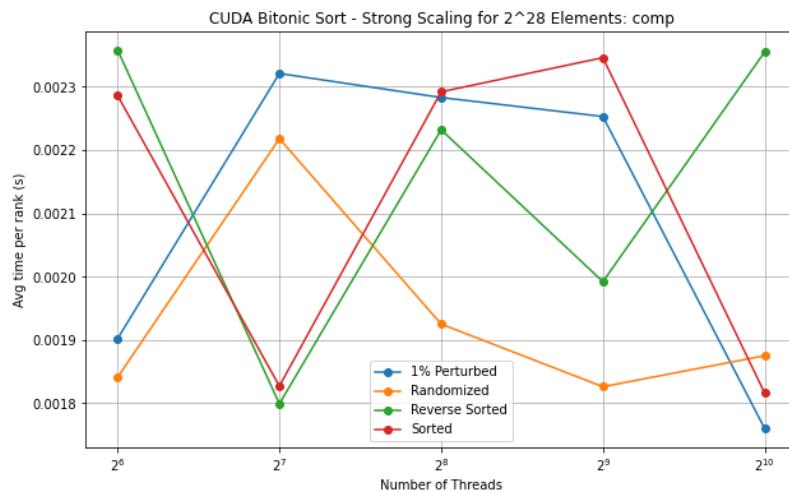
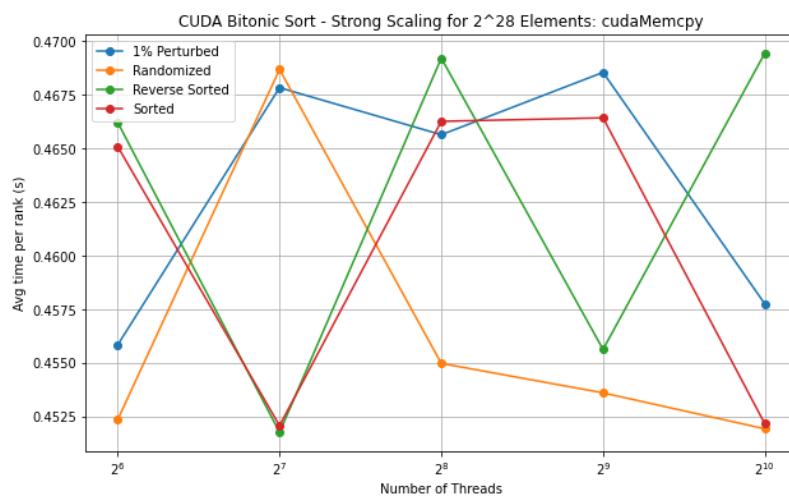


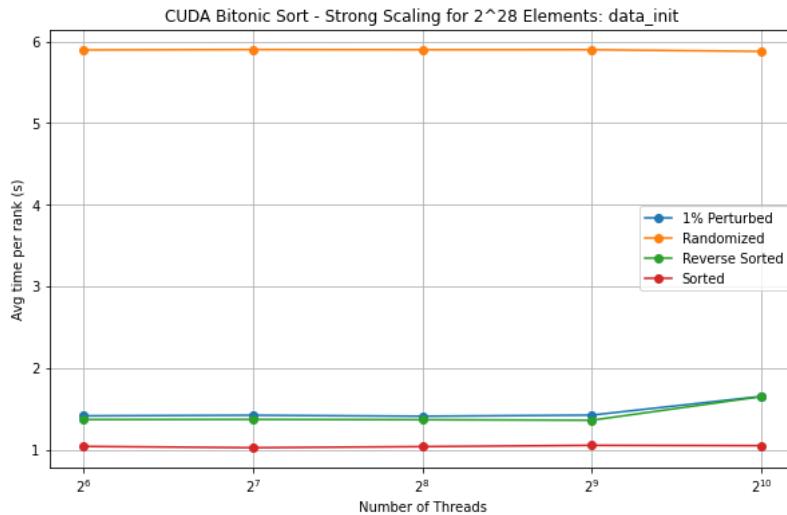
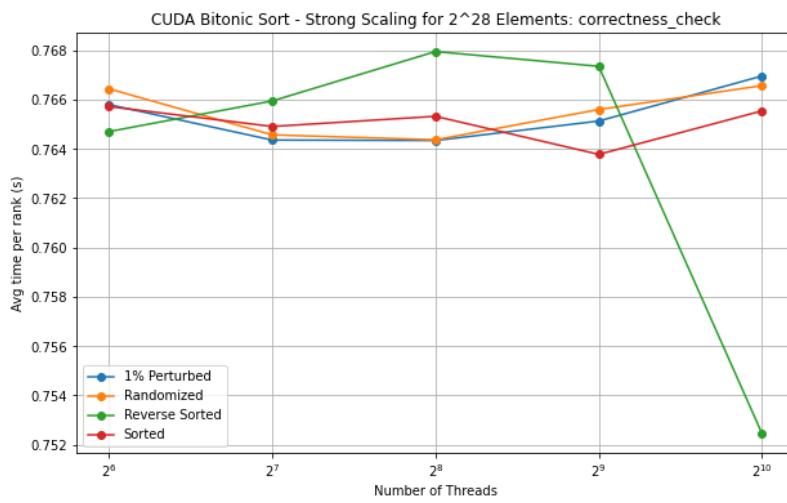




2^{28} Elements







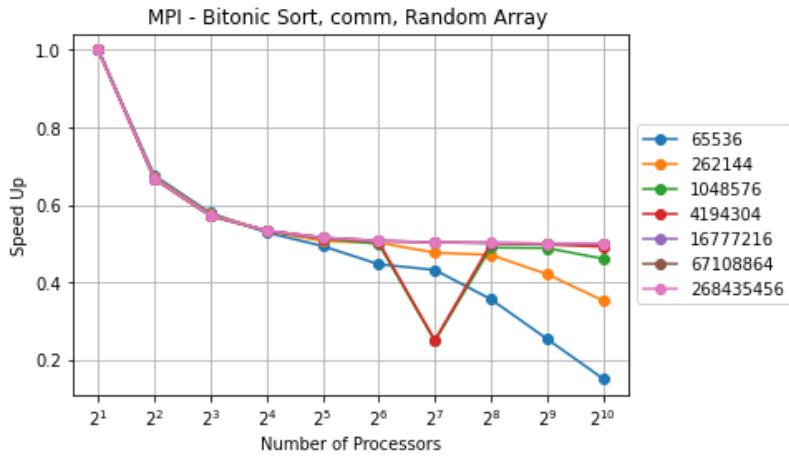
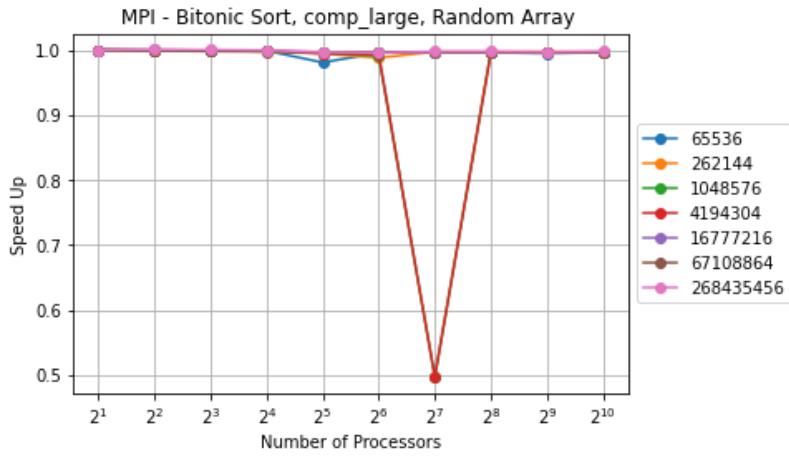
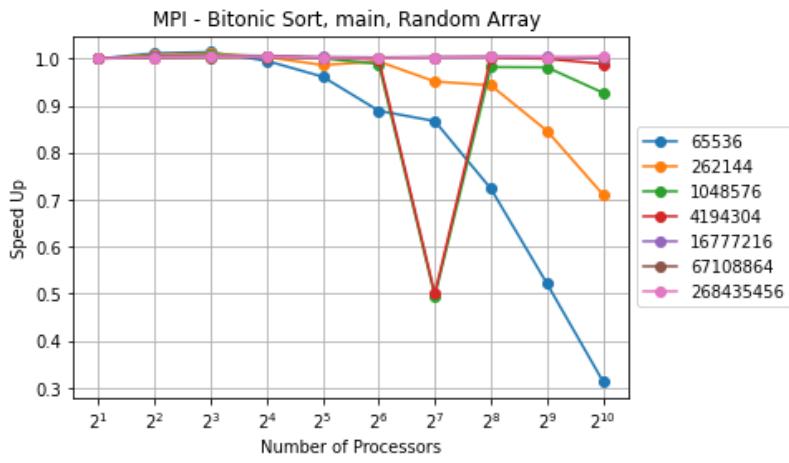
Speedup

MPI

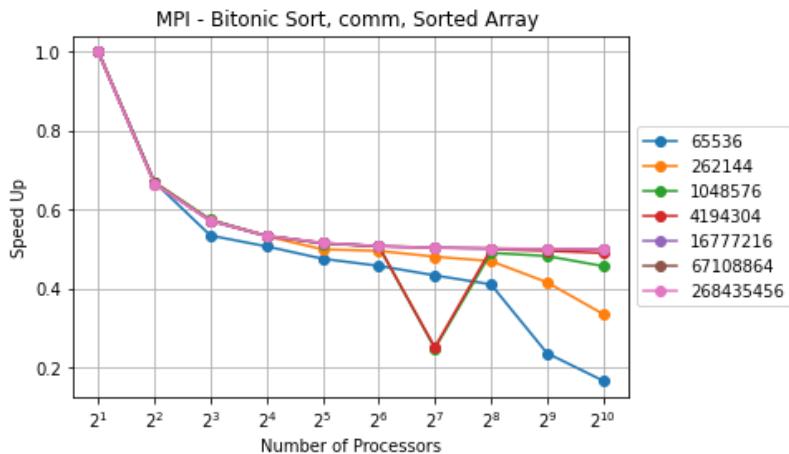
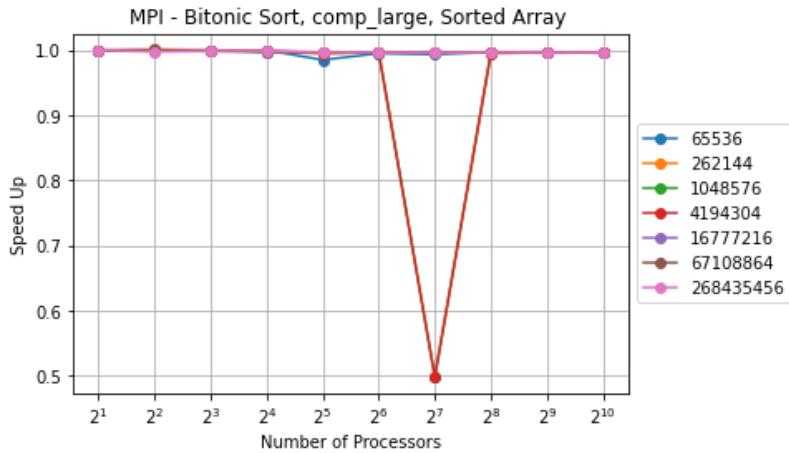
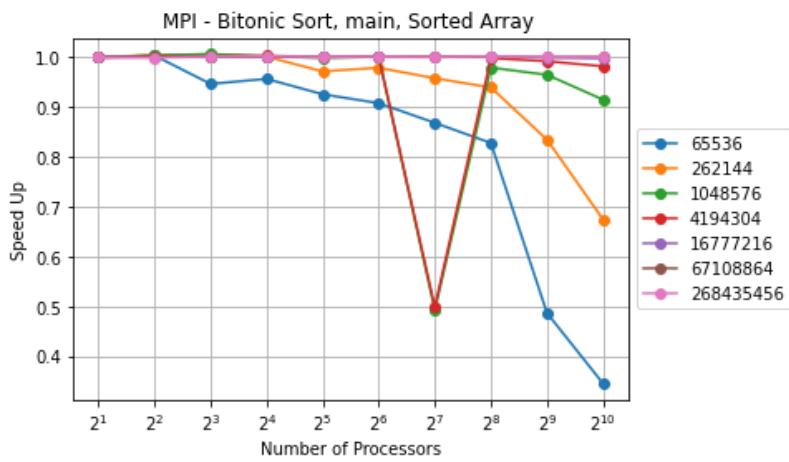
When looking at the computation speedups, we see a linear trend between the input sizes. This shows the importance of parallelism. For all sizes, as we increase the number of processors, the pure computation time gets better and better, this showing the higher speedup. We do see that speedup for the 2^{16} elements begins to taper off for the comp regions, indicating that the parallelism is good for very large sizes, but smaller sizes aren't as impacted by increasing processors. This makes sense because you would spend more percentage of your time increasing your processors and have more communication and face more overhead for the smaller values. If we look at the entire main region as a whole, we can see that the larger the input size, the more speedup there is. This further reinforces that there is a need for parallelism, and a definite benefit. For all of the input sizes, there seems to be a peak speedup at 256 processors before decreasing again. After 256 processors, there may be so much overhead that we lose overall performance. It is a little difficult to look at it easily for the smaller sizes, but the speedup graphs tend to all show that a bigger input size benefits from this parallelism. We also see that for communication, speedup decreases for more processors, again reiterating that we are likely facing some significant overhead.

FIX ME!

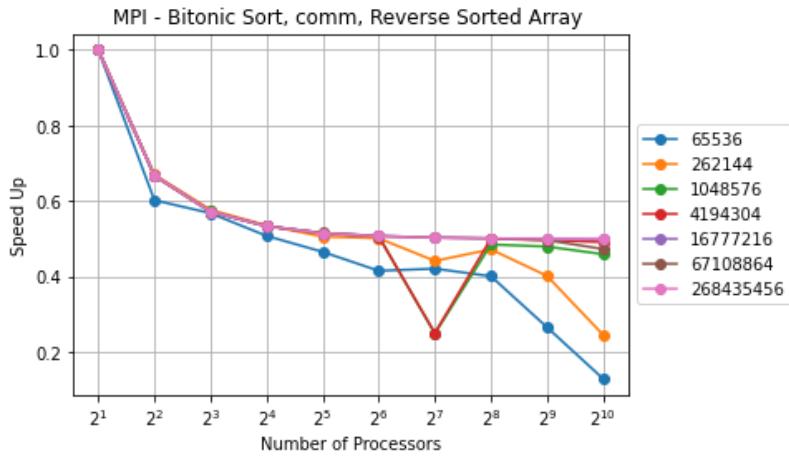
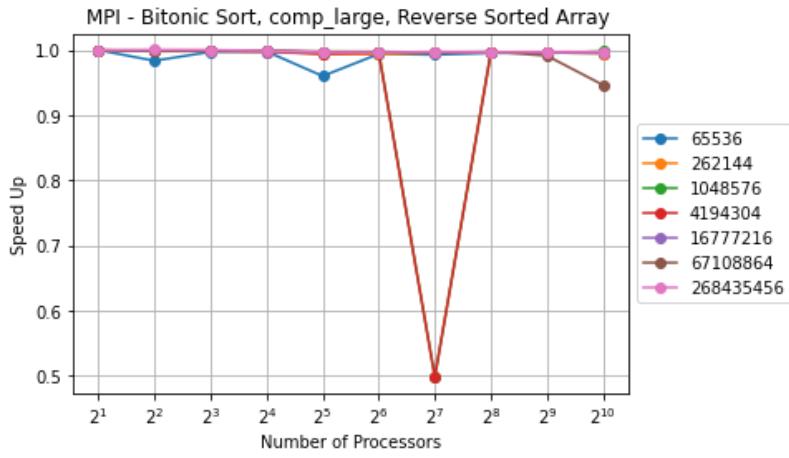
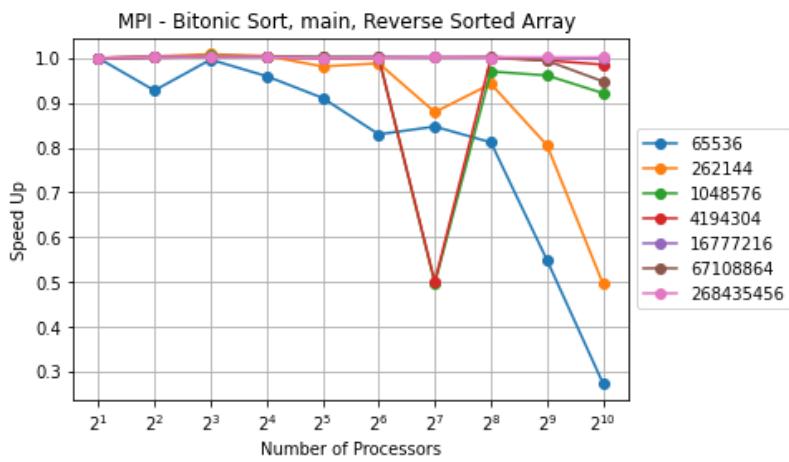
RANDOM INPUT ARRAY



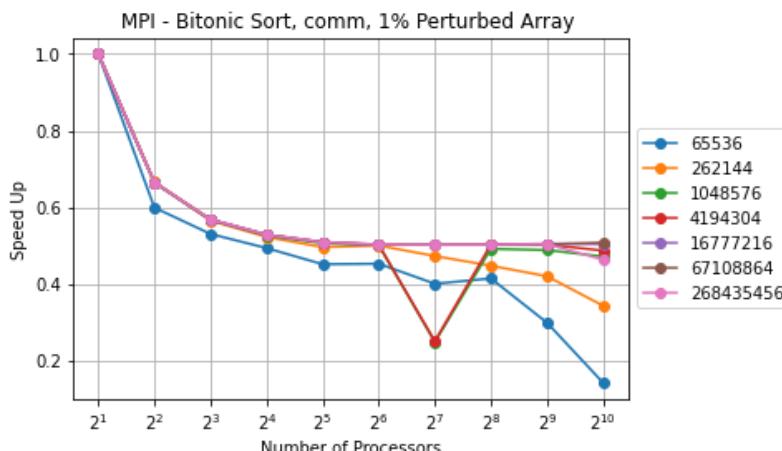
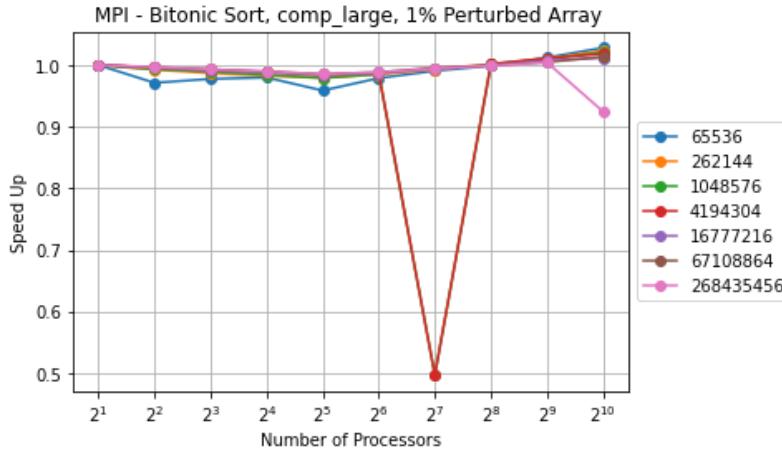
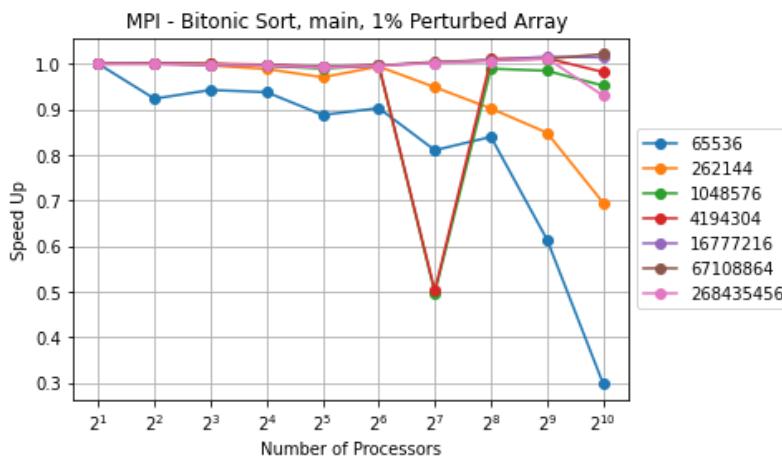
SORTED INPUT ARRAY



REVERSE SORTED INPUT ARRAY



1% PERTURBED INPUT ARRAY



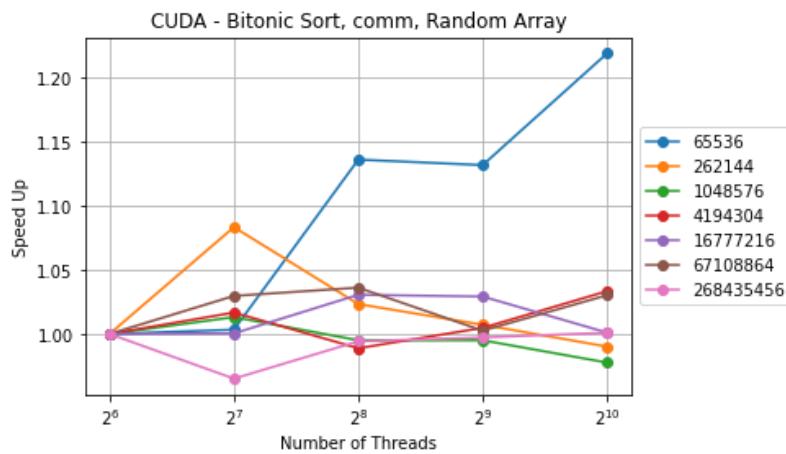
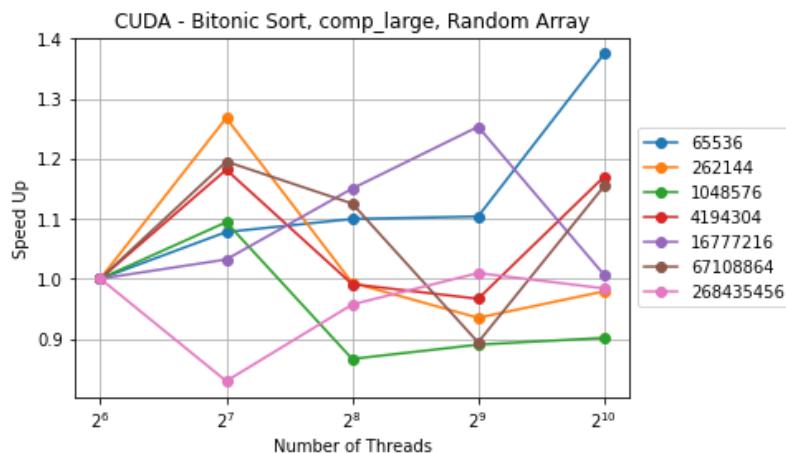
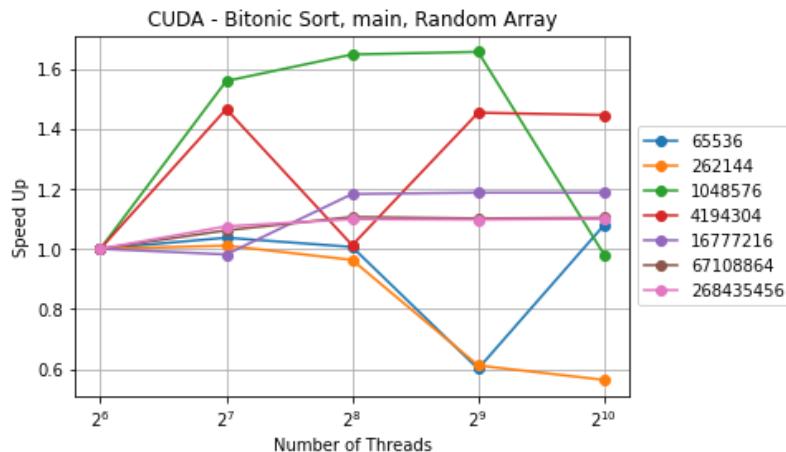
CUDA

In the CUDA speedup graphs, there doesn't seem to be any inherent benefit to parallelism. When we look at the main_function as a whole, we see that the speedup values tend to hover around 1.0. In fact, most of the lines tend to be below a value of 1, indicating that it is performing worse by being parallel. There also seems to be no pattern with the varying input size. We see that at the highest number of processors, the worst speedups are 2^{16} elements and 2^{28} elements, which are both the smallest and largest values respectively. It seems like as we increase processors, speedup for the communication is worst for the 2^{16} element, but the computation is worst for 2^{28} elements. Communication is hard to gauge completely because there is quite a bit of fluctuation here, where computation speedup hovers around 1.0 and steadily decreases for most of the input sizes. This could also be because there's no real inter-process

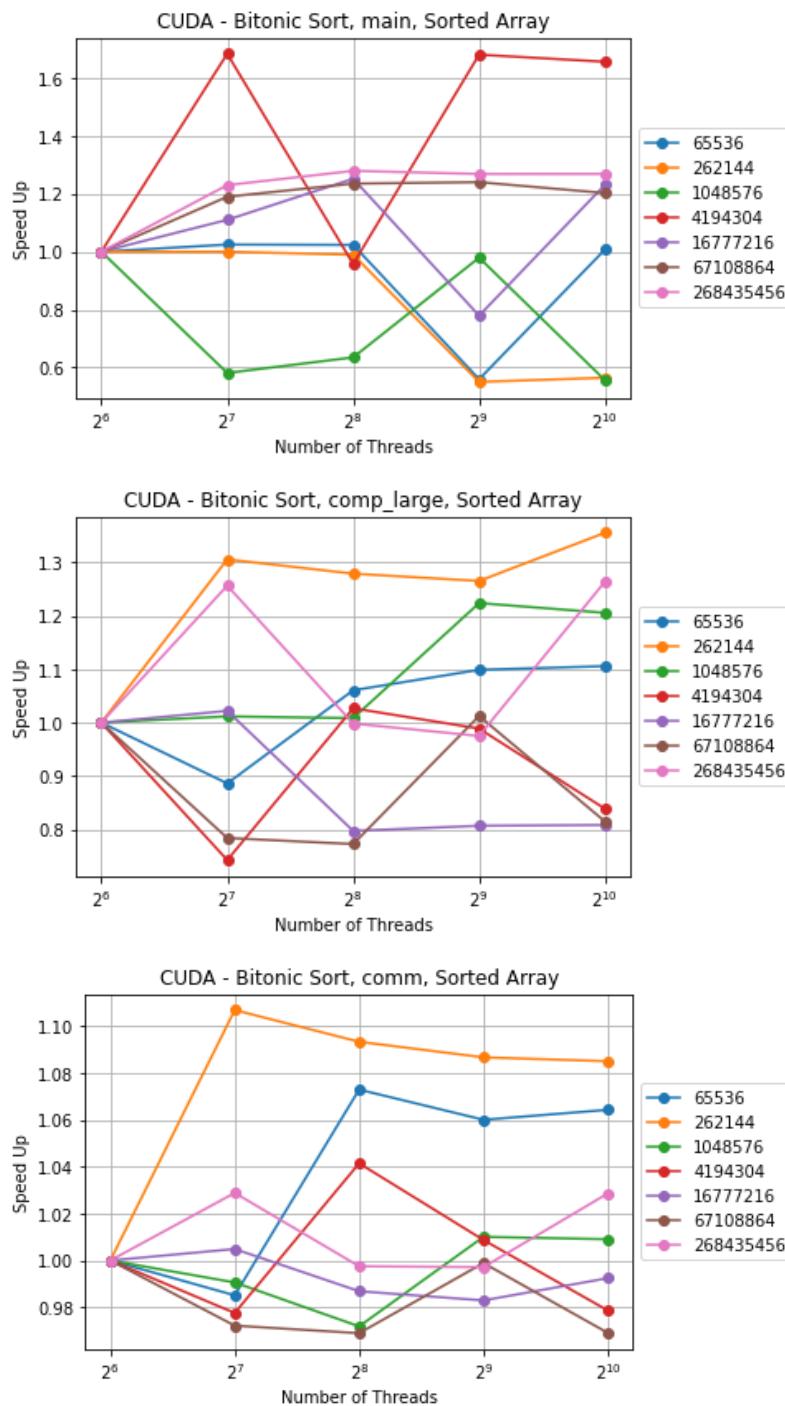
communication for CUDA like there is for MPI.

FIX ME!

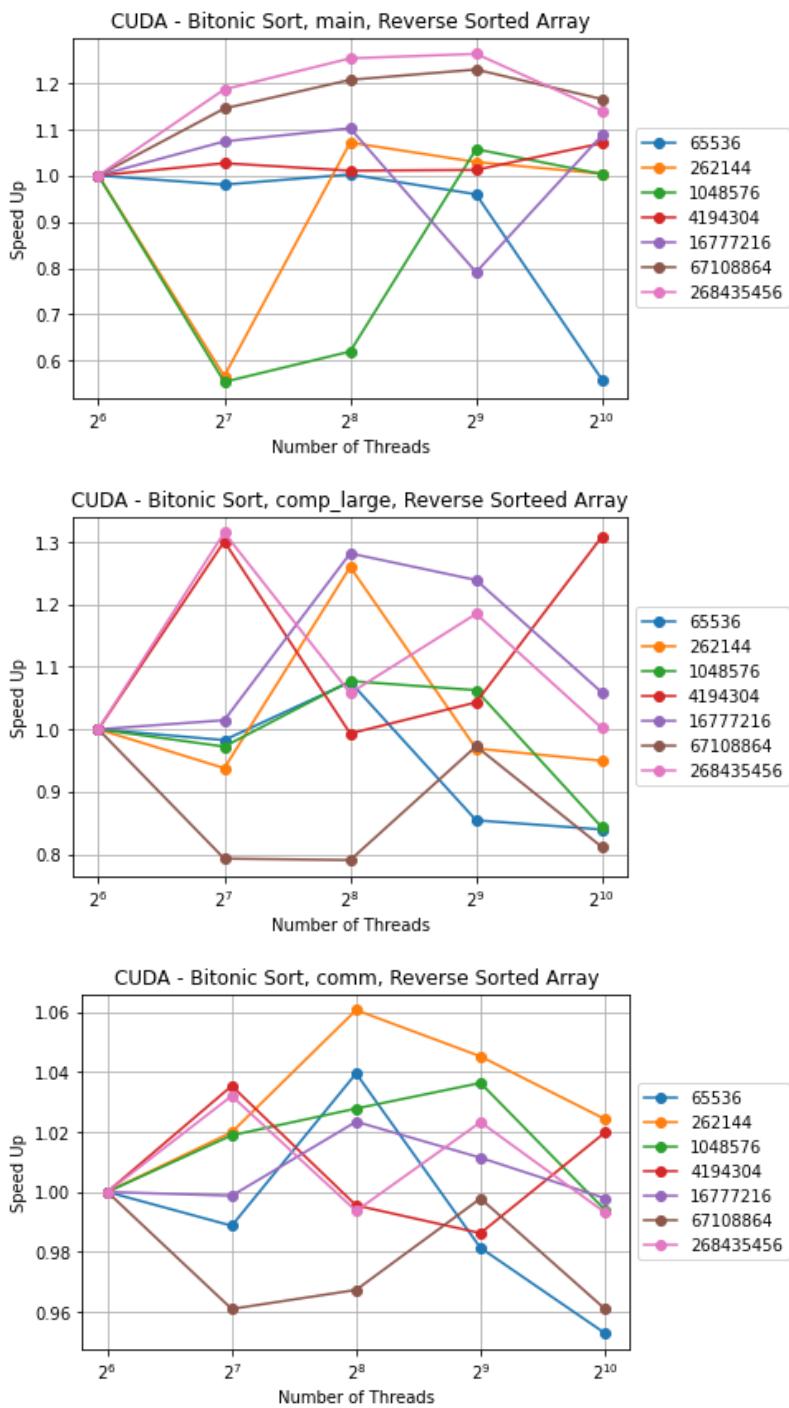
RANDOM INPUT ARRAY



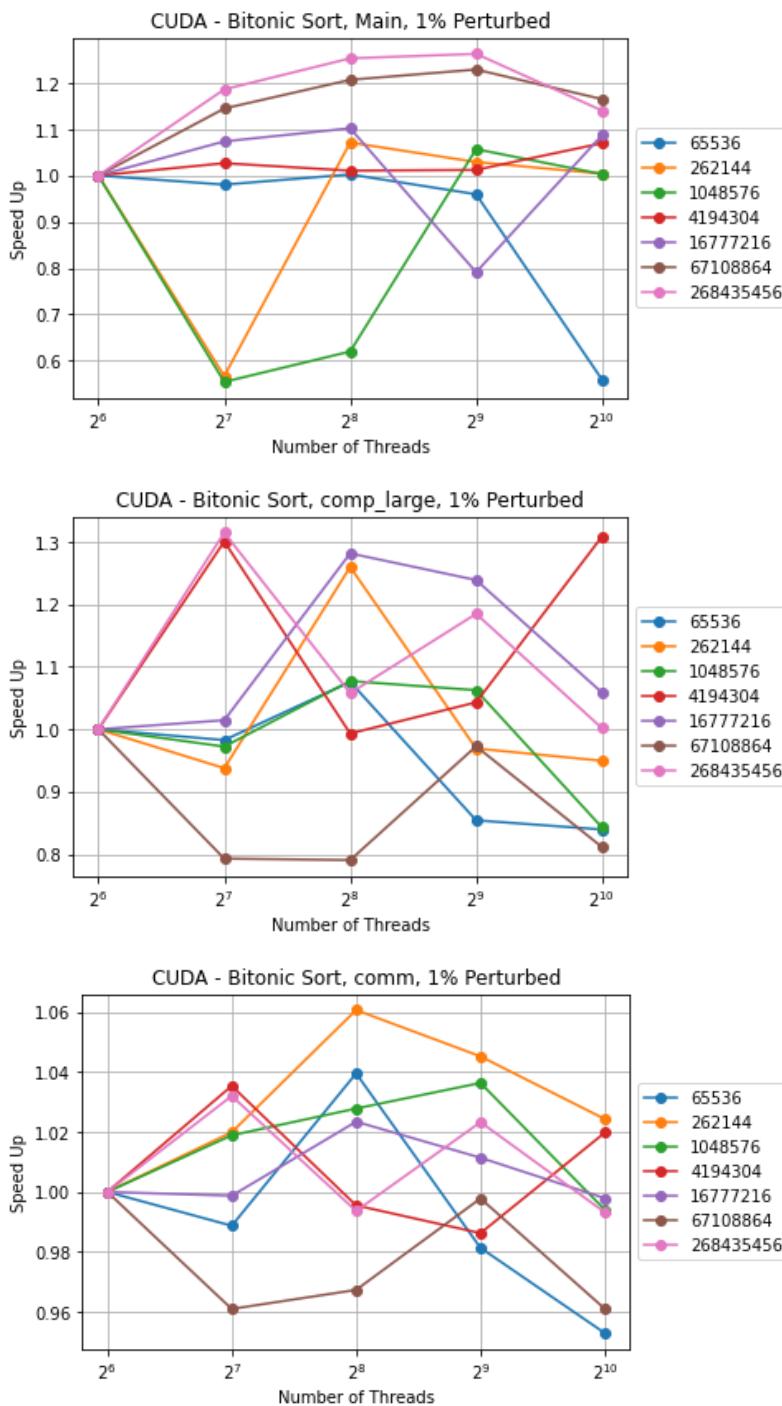
SORTED INPUT ARRAY



REVERSE SORTED INPUT ARRAY



1% PERTURBED INPUT ARRAY



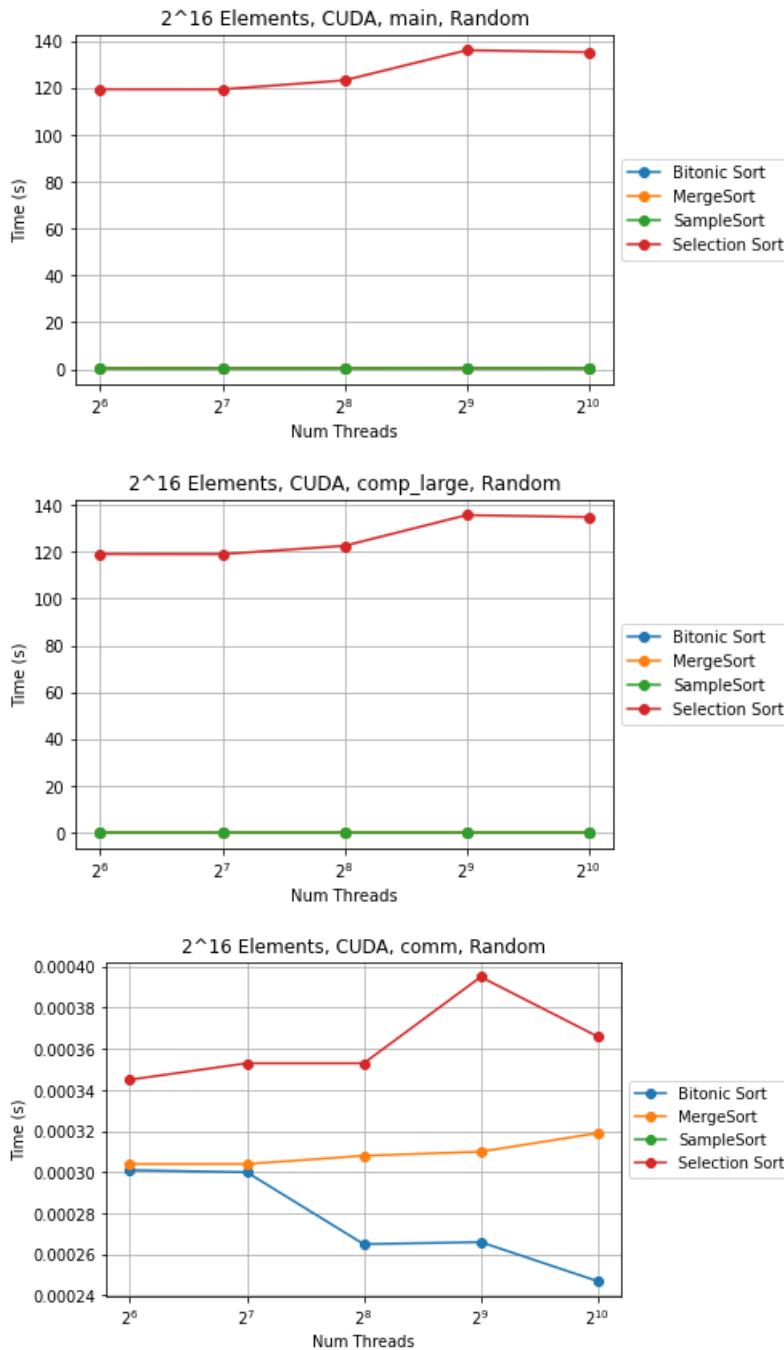
Comparison Analysis

For the comparison analysis, we compared the 2^{16} elements since selection sort was having issues at larger values.

When starting with CUDA, we noticed that selection sort was much, much slower in all aspects of total, communication, and computation time. To complete the analysis, we looked at both the graphs with all of the sorting algorithms, as well as graphs that excluded selection sort. This was needed because selection sort worked on such a large scale that it would wash out the lines for the other three sorts and it was difficult to tell which ones performed better. It is easiest to see in the communication step, since in most of the sorting algorithms, this was simply a Cudamemcpy which is a relatively quick step. Sample sort did not include a comm region, which is why it is not present on that graph. However, we can see that the

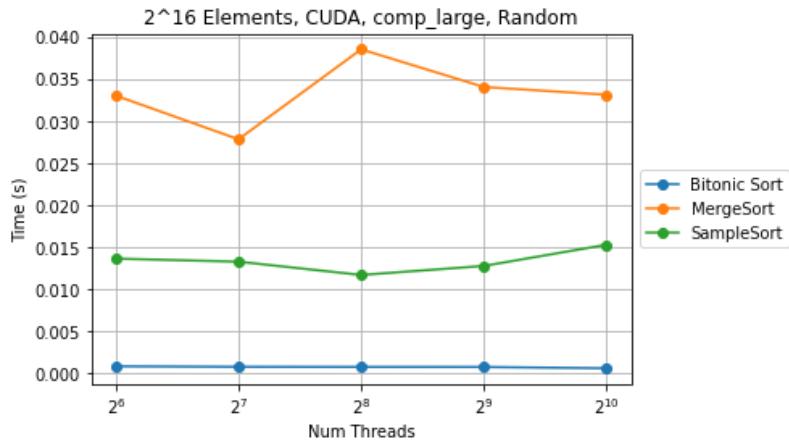
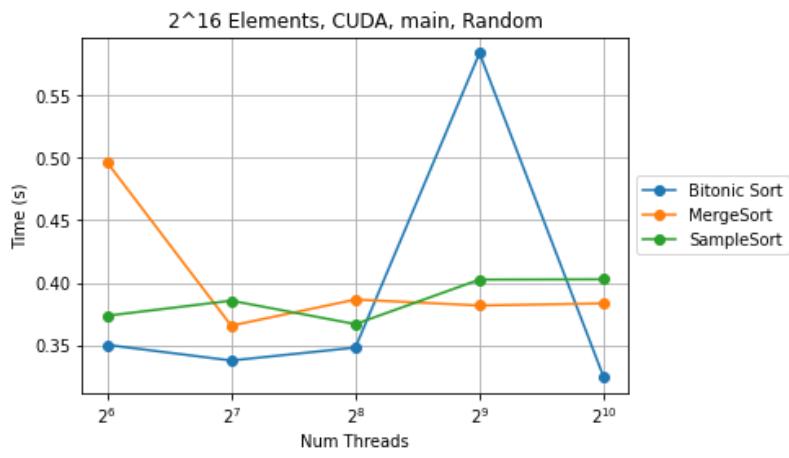
main and comp regions are completely overtaken by the selection sort graph. Since selection sort is not inherently a parallel algorithm, it performs the worst by far.

CUDA - All



CUDA - Bitonic, Merge, and Sample Sort

When looking at the main and computation regions without the selection sort line, we see that they are relatively constant, especially considering the fact they are at the scale of tenths of a second. We see a spike for bitonic sort at 512 threads in the main region, which could likely be an outlier looking at its other data points. When looking at just computation, we see that it is mostly constant for all three graphs. We can conclude here that for CUDA, the order from best to worst scaling sorting algorithms for us is bitonic, sample, merge, and then selection sort.



MPI

The main region shows mostly similar behavior outside of merge sort, with the other three sorts hovering around zero for most processors. Sample sort has a comm spike in time at higher processors, which explains the peak at the end. Merge sorts main function isn't explain by comp or comm, so other areas such as data init likely played a role. Looking at the computation region, we see that selection sort is by far the worst at lower processors, which indicates its a poor performing algorithm. However, it quickly converges towards 0 around 32 processors. It is difficult to see on the comp graph, but merge sort is the best performing sorting algorithm here. When looking at the comm graph, we see that most of the sorting algorithms are very slightly increasing, except for sample sort which shoots up at 1024 processors. In terms of communication time.

